arXiv:2610.20604v1 [cs.CV] 14 Oct 2026

地平线 Horizon Robotics

VITAL: From Visual Motion to Precise Action

Alan Liang[1,2], **Youquan Liu**[2,3], **Mingxiao Li**[1], **Wei Chow**[4], **Lingdong Kong**[4]
Yue Liao[4], **Dongyue Lu**[4], **Yu Yang**[5], **Wei Yin**[1], **Runnan Chen**[2]

[1]Horizon Robotics, [2]Shanghai Innovation Institute, [3]Fudan University
[4]National University of Singapore, [5]Zhejiang University

Project Page: https://worldbench.github.io/vital

GitHub: https://github.com/worldbench/VITAL

Figure 1 From world prediction to precise action. World-action models aim to translate visual predictions of the physical world into actions. This motivates deriving action targets from visual motion to reduce the visual–action representation gap. Yet even with recorded future images, our study recovers the dominant maneuver while leaving local metric errors. **VITAL** jointly addresses visual grounding and execution precision: visual motion defines the principal action target, and a paired residual encodes its remaining trajectory error.

Abstract

World-action models aim to turn visual anticipation into driving decisions, yet typically train action experts on trajectory representations defined independently of visual dynamics. This motivates deriving action targets directly from visual motion. We first examine a simpler setting in which trajectories are reconstructed from recorded future images. Even without the need for future prediction, reconstruction captures the dominant maneuver but leaves centimeter-scale errors. We therefore introduce **VITAL** (Vision-Induced Tokens for Action Learning), an action tokenizer that couples a visually derived motion scaffold with a compact residual encoding the trajectory error left by that scaffold. This representation makes visual motion the principal prediction target and gives small metric corrections a separate, normalized learning scale. Using flow matching, the action expert jointly predicts both components from the current image, ego state, and navigation command, without generating future images at inference. The same residual coordinates support bounded reinforcement post-training that refines execution while keeping the predicted visual scaffold fixed. With only one action denoising step, **VITAL** achieves 90.8 PDMS on NAVSIM v1 `navtest`, within 0.13 points of its ten-step score at 3.29× the end-to-end inference speed. Without retraining, it achieves 91.1 EPDMS on NAVSIM v2 `navtest` and 38.6 on `navhard`.

1 Introduction

Learning to anticipate a driving scene should also help a vehicle decide how to move. World-action models pursue this connection by coupling visual prediction with trajectory generation (Xia et al., 2026a; Yuan et al., 2026; Zhao et al., 2026). A common design couples video and action experts in a mixture of transformers, allowing action tokens to attend to visual features before decoding an ego trajectory (Li et al., 2026a; Yuan et al., 2026; Zhao et al., 2026). As illustrated in Fig. 1, the promise is to learn how to act from a model that anticipates how the world moves.

Yet exchanging information does not necessarily give the two experts a shared prediction target. Attention (Vaswani et al., 2017) connects their features, while the video expert learns to predict scene evolution and the action expert often learns a code defined only by trajectory coordinates. The two experts describe the same physical motion through separately constructed representations. We therefore ask whether action targets can be derived directly from observed visual motion.

Recent advances in robot control suggest that this connection can be learned at the tokenizer level. RepWAM (Wang et al., 2026b) learns latent actions from transitions between semantically aligned visual representations, and LingBot-VA 2.0 (Zhang et al., 2026a) adopts this strategy for joint pretraining of video and action. Their gains in control precision make visual motion a promising source of action targets. Motion-token and continuous-motion learning offer complementary evidence for connecting videos to control (Chen et al., 2025; Yang et al., 2026a). However, these tokenizers reconstruct transitions in visual feature space. It remains unclear how much of the metric detail needed for driving can be recovered from visual motion alone, even before a planner has to predict that motion from the observations available at deployment.

We therefore remove future prediction from the problem altogether. As shown in Appendix A.1, in a simpler task than world-action modeling, the model receives recorded current and future images and reconstructs the corresponding ego trajectory. Even with the future already in view, a precision gap remains. As shown in Fig. 2, reconstruction leaves 123.6 mm average displacement error (ADE) and 221.6 mm final displacement error (FDE) across all NAVSIM validation scenes (Dauner et al., 2024). Looking more closely at the recovered motion, we find that its HF fraction, the mean share of velocity energy at high frequencies across x and y, is only 0.1%. What emerges is a smooth motion scaffold. It captures the maneuver but still misses metric detail that matters near road boundaries, where small deviations can make an otherwise plausible trajectory noncompliant.

This finding suggests preserving the motion that vision captures while giving the missing detail a representation of its own. We introduce **VITAL**, short for Vision-Induced Tokens for Action Learning, an action tokenizer built around this separation. A visual motion code decodes into the trajectory scaffold. A second code, the paired metric residual, encodes the difference between that scaffold and the demonstrated trajectory. Normalization gives these small corrections their own learning scale alongside the dominant motion. Decoding both codes restores the missing trajectory detail. As shown in Fig. 2, supplying recorded images and the residual encoded from ground truth reduces reconstruction ADE and FDE to 0.3 mm and 0.5 mm and raises PDMS from 92.1 to 94.7. The residual's HF fraction reaches 61.4%, revealing faster local variations in the component that restores metric precision. The resulting representation retains both the visual motion scaffold and the finer corrections needed to reconstruct the demonstrated trajectory.

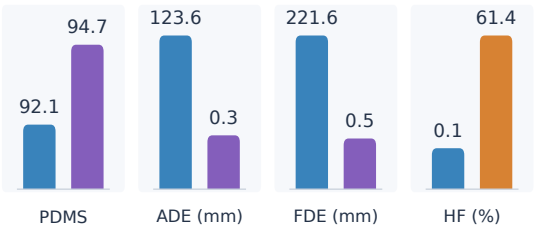

Figure 2 Visual reconstruction and residual precision. ■ vision only, ■ **VITAL**, ■ residual. HF denotes the mean x/y high-frequency velocity energy fraction.

Precise reconstruction is useful for planning only if the representation can be predicted before the future is observed, consistent with recent analyses of action tokenization beyond reconstruction fidelity (Chen et al., 2026). We train the action expert with flow matching to generate both codes jointly from the current image, ego state, and navigation command. One flow step then produces a trajectory without generating future images. Separating the metric correction also opens a route to policy improvement (Johannink et al., 2019). After imitation learning, we freeze the action generator and decoder and learn bounded updates in the residual coordinates. At deployment, a single deterministic correction refines local execution while keeping the predicted visual scaffold fixed.

The gains extend from reconstruction to planning. On NAVSIM, **VITAL** strengthens visual conditioning. With equal action training budgets, it gains 5.1 PDMS points in one step over an autoencoder trained only on trajectories. Residual refinement brings the final policy to 90.8 PDMS in one step, within 0.1 points of ten steps at 3.29× the inference speed, including image encoding.

Our contributions bridge visual motion, precise action learning, and efficient policy refinement.

- We propose **VITAL**, an action tokenizer that pairs visual motion with a normalized metric residual to preserve both visual grounding and fine trajectory precision for autonomous driving.
- We propose residual reinforcement learning that improves planning through a single correction, without additional denoising, candidate search, or reward evaluation at inference.
- We achieve 90.8 PDMS on NAVSIM v1 with a single flow evaluation. Without retraining, the same policy reaches 91.1 EPDMS on NAVSIM v2 and 38.6 EPDMS on reactive `navhard`.

2 Related Work

2.1 World-action models

The survey of 3D and 4D world modeling organizes video-, occupancy-, and LiDAR-based approaches and their evaluation protocols (Kong et al., 2025b). Quo Vadis, World Modeling? broadens this view to agent-centric proxies that provide dynamics, spatial, execution, memory, skill, and reward feedback (Yang et al., 2026b). Together, these perspectives distinguish scene prediction from the information a policy needs to select and improve executable actions.

World-action models connect visual prediction to control by using generated futures or their latent representations to inform action generation. Drive-WM evaluates candidate plans through generated multiview videos (Wang et al., 2024), while MaskGWM combines video generation with masked reconstruction to learn generalizable driving dynamics (Ni et al., 2025). DriveLaW injects video-generator latents into a trajectory planner (Xia et al., 2026a), and SUV lets the action expert attend to predicted appearance, semantics, relative depth, and instance dynamics (Yuan et al., 2026). Metis and SimWAM use attention masks to support joint video-action training while avoiding future-video generation during action inference (Li et al., 2026a; Zhao et al., 2026).

Latent world models also support driving without reconstructing full future videos. LAW predicts action-conditioned future features to improve scene representations (Li et al., 2025c), while World4Drive incorporates intention-aware physical latent modeling (Zheng et al., 2025). Latent-WAM combines geometry-informed scene compression with dynamic latent prediction (Wang et al., 2026d). A separate use of future prediction is to evaluate plans: WoTE forecasts BEV states for online trajectory evaluation (Li et al., 2025d), and DA-WAM predicts an action-specific future latent for each candidate and uses it for scoring (Zhong et al., 2026).

WorldLens evaluates driving world models across generation, reconstruction, action-following, downstream tasks, and human preference (Liang et al., 2026b). The associated workshop study, Is Your Driving World Model an All-Around Player?, examines six driving world models across 24 evaluation dimensions within the same WorldLens framework (Kong et al., 2026a). These evaluations show that visual realism alone does not establish geometric or behavioral fidelity. This motivates assessing the control interface separately: our recorded-future reconstruction study isolates retained action precision, while conditional prediction and NAVSIM scores test whether that representation supports planning from current observations.

Broader surveys place both directions within agentic world modeling and vision-language-action driving systems (Chu et al., 2026; Hu et al., 2025a). LCVN studies language-conditioned navigation through latent imagination or joint observation-action prediction, while UniWM couples future prediction with memory-augmented planning (Dong et al., 2026b,c). OneVL uses language and future-video supervision to learn compact reasoning tokens, discarding auxiliary decoders during planning inference (Lu et al., 2026a). Beyond video, DynamicCity generates dynamic occupancy, while LiDARCrafter and U4D model controllable or uncertainty-aware LiDAR sequences (Bian et al., 2025; Liang et al., 2026a; Xu et al., 2026b). OmniLiDAR extends diffusion-based scan generation across weather, sensor configurations, and acquisition platforms through a shared range-image representation and domain-conditioned modeling (Liu et al., 2026c). These scene-generation interfaces provide a broader context for our focus on visually grounded, metric action targets.

Related strategies appear in robotic manipulation. Prediction with Action couples visual and action prediction through joint denoising (Guo et al., 2024), and UWM uses independent video and action diffusion timesteps within one transformer to support policy and dynamics modeling (Zhu et al., 2025). WorldVLA studies autoregressive action-world modeling (Cen et al., 2025); VideoVLA and DreamZero leverage pretrained video generators to jointly predict actions and visual outcomes (Shen et al., 2025; Ye et al., 2026). Video Prediction Policy uses predictive visual representations with inverse dynamics (Hu et al., 2025b), while DreamVLA forecasts compact dynamic, spatial, and semantic knowledge for action prediction (Zhang et al., 2025d). V-JEPA 2-AC plans through an action-conditioned predictor in visual feature space (Assran et al., 2025). Gen2Act transfers information from generated human videos to manipulation (Bharadhwaj et al., 2025), and World Guidance predicts compact conditions derived from future observations alongside actions (Su et al., 2026).

VITAL constructs the action target from recorded visual motion and explicitly encodes the metric trajectory error left unresolved by that representation. Future images supply training supervision and target-construction inputs; deployment predicts the paired action codes from current context without generating future images or scoring trajectory candidates.

2.2 Representation alignment and latent actions

REPA aligns diffusion-model features with pretrained visual representations (Yu et al., 2025). For driving, Drive-JEPA learns predictive video representations for planning (Wang et al., 2026e), ReWorld supervises intermediate video and action representations through prediction, alignment, and geometric discrimination (Xia et al., 2026b), and RAP introduces 3D geometric supervision through rasterization (Feng et al., 2026). Cross-modal alignment also supports 3D perception: CLIP2Scene and Seal transfer pretrained visual semantics to point clouds (Chen et al., 2023; Liu et al., 2023b), while LargeAD scales cross-sensor pretraining and LaserMix++ exploits multimodal consistency under limited annotation (Kong et al., 2025a, 2026b). PointLoRA adapts point-cloud representations with low-rank updates and token selection (Wang et al., 2025a); Pi3DET and 3EED examine cross-platform detection and language-grounded 3D localization (Li et al., 2025b; Liang et al., 2025).

Representation design also shapes the fidelity of driving perception. RangeFormer addresses projection-induced semantic and geometric distortions in range-view LiDAR segmentation (Kong et al., 2023a), while LaserMix uses spatially structured beam mixing and prediction consistency to learn from sparsely labeled scans (Kong et al., 2023c). RoboBEV evaluates BEV perception under camera corruptions, sensor failures, and domain shifts (Xie et al., 2025b). These works concern perception accuracy, label efficiency, and robustness; our complementary question is how an action representation preserves metric detail and remains predictable from visual conditioning.

Genie and LAPA infer latent actions from visual transitions, making videos without action annotations useful for controllable generation or policy pretraining (Bruce et al., 2024; Ye et al., 2025). Moto learns motion tokens for video pretraining and subsequently connects them to robot control (Chen et al., 2025). UniVLA uses language-conditioned, task-centric latent actions to reduce task-irrelevant variation (Bu et al., 2025). CoMo instead learns continuous latent motion, addressing information loss in discrete representations and the gap to continuous actions (Yang et al., 2026a). Visual transitions can contain exogenous changes, motivating supervision that separates action-relevant motion from distractors (Nikulin et al., 2025; Zhang et al., 2025a).

RepWAM (Wang et al., 2026b) and LingBot-VA 2.0 (Zhang et al., 2026a) learn latent actions as transitions between semantically aligned visual states, with transport and residual components for reconstructing visual changes. RLA-WM learns residual latent actions from DINO feature differences and models their distribution through flow matching (Zhang et al., 2026b). The residual in **VITAL** has a different target: the metric difference between a visually decoded trajectory scaffold and the demonstrated trajectory. Trajectory supervision grounds the scaffold in physical coordinates, and a normalized compact residual code retains its remaining reconstruction errors. Both components are jointly predicted for planning, and the residual coordinates also support bounded refinement.

2.3 Action tokenization and reconstruction fidelity

FAST efficiently tokenizes labeled action sequences (Pertsch et al., 2025). VQ-BeT uses hierarchical vector quantization for multimodal behavior generation (Lee et al., 2024), while VQ-VLA scales vector-quantized

action tokenizers (Wang et al., 2025b). QueST learns temporal skill abstractions with a flexible quantized latent representation (Mete et al., 2024). OAT structures the token sequence for compression, total decodability, and ordered prediction, allowing prefix-based decoding to trade inference cost for action fidelity (Liu et al., 2026a). VITA additionally learns shared discrete codes that decode into future frames and robot actions (Ma et al., 2026). For autonomous driving, LaPla combines quantized trajectory representation learning with continuous latent prediction to address the discrete-to-continuous planning interface (Yao et al., 2026a).

VITAL instead assigns separate targets to visually grounded motion and the metric error remaining after scaffold reconstruction.

Reconstruction fidelity is only one requirement of a useful policy interface. BAAT aligns quantized representations of behaviorally corresponding action chunks jointly with reconstruction (Dong et al., 2026a). ProAct explicitly studies token predictability and robustness of decoded policy outputs beyond clean action reconstruction (Chen et al., 2026). Our analysis separates clean-code reconstruction from conditional code prediction, while downstream planning scores test usefulness for control; improved reconstruction alone does not establish improved planning.

2.4 Few-step action generation

Diffusion Policy models multimodal action sequences through iterative denoising (Chi et al., 2025), while flow matching and rectified flow learn velocity fields that transport noise to data (Lipman et al., 2023; Liu et al., 2023a).

Consistency Models learn direct noise-to-data mappings (Song et al., 2023), and Consistency Policy and One-Step Diffusion Policy distill iterative policies into faster samplers (Prasad et al., 2024; Wang et al., 2025c). Shortcut Models condition generation on step size and use self-consistency to support different sampling budgets (Frans et al., 2025); FlowPolicy applies consistency flow matching to one-step robot control (Zhang et al., 2025b). DanceOPD learns from capability-specific velocity fields queried at student-induced states, and DiffusionOPSD converts reward guidance into refreshed diffusion post-training targets (Zhou et al., 2026a,b). Both study on-policy target construction for generative models. MeanFlow learns average velocity over an interval (Geng et al., 2025), and conditional optimal-transport couplings encourage straighter flow trajectories for few-step control (Sochopoulos et al., 2025). FLUX applies rectified flow to efficient cross-embodiment navigation and follows a static-to-dynamic curriculum with reinforcement learning in social environments (Gong et al., 2026). Where this approach shapes the flow and training curriculum, our approach changes the paired action targets while retaining the standard flow-matching objective. Streaming Flow Policy permits action execution during integration (Jiang et al., 2025).

DRIFT, the residual flow adapter for VLMs, combines a coarse base predictor with iterative flow-based refinement of a localized residual distribution (Liu et al., 2026e). Its iterative refinement differs from our joint prediction of scaffold and metric residual codes. In driving, DiffusionDrive uses truncated denoising (Liao et al., 2025), while MeanFuser combines one-step generation of multiple proposals with adaptive reconstruction (Wang et al., 2026c). **VITAL** retains the standard flow-matching prediction objective and changes the action representation supplied to it. The paired scaffold and normalized metric residual support one-step prediction of a single executable trajectory without teacher distillation or candidate selection.

2.5 Residual policy refinement

Reward-based refinement addresses a limitation of imitation learning: matching demonstrations does not directly optimize deployment performance. DPPO optimizes diffusion policies through reinforcement learning (Ren et al., 2025), and ReinFlow adapts flow-matching policies with online RL (Zhang et al., 2025c). Flow Q-Learning combines flow-based behavior modeling with an expressive one-step policy optimized for value (Park et al., 2025).

For driving, AD-R1 learns an impartial world model using counterfactual collisions and off-road events, then uses predicted action outcomes as an internal critic for closed-loop policy refinement (Yan et al., 2026). It emphasizes the quality of learned feedback for optimization. Simulation-grounded video–trajectory alignment also supplies matched rendered-video supervision for reward-optimized driving trajectories (Jing et al., 2026). Our residual post-training instead uses the NAVSIM scorer during training and freezes the

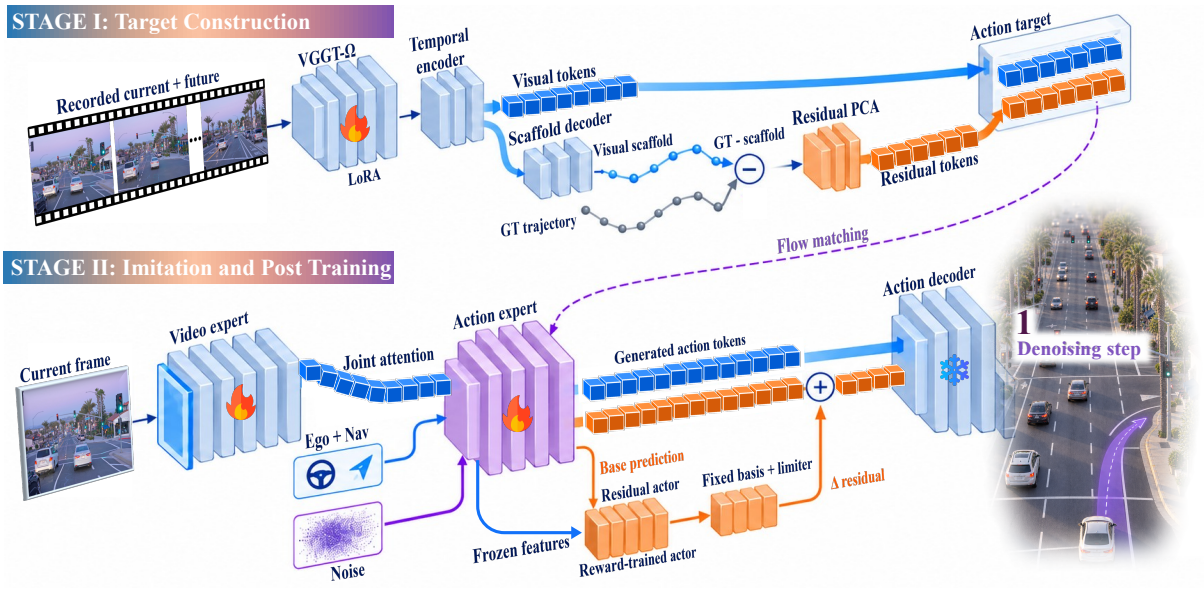

Figure 3 VITAL learns a shared representation for action prediction and refinement. The upper stage constructs visual motion tokens and paired metric residuals from recorded images and trajectories. The lower stage predicts the joint code in one flow step, applies a bounded residual correction, and decodes the executable trajectory. The diagram shows the driving implementation.

base policy and codec, confining updates to bounded metric corrections around the predicted scaffold; deployment requires neither a learned world-model critic nor reward-based candidate search.

Residual RL learns corrections around an existing controller (Johannink et al., 2019), while Policy Decorator and ResiP refine imitation policies through residual control (Ankile et al., 2025; Yuan et al., 2025). DSRL keeps the diffusion policy fixed and learns to steer its latent-noise inputs (Wagenmaker et al., 2025). VITAL also freezes the base action generator and decoder, but restricts intervention to the metric residual coordinates already defined by its tokenizer. A fixed, bounded correction space preserves the predicted visual scaffold and provides a shared physical decoding map for imitation learning and policy improvement. At deployment, the actor applies one deterministic correction, with no extra denoising or reward-based search. This construction bounds the scope of local refinement; it does not guarantee that every correction improves reward or satisfies task-level safety.

3 Learning Precise Actions from Visual Motion

VITAL combines visual motion scaffolds with paired metric residuals to connect target construction, action prediction, and reward-based refinement through one executable representation.

3.1 Problem Setup and Method Overview

We study visually conditioned action planning from context c comprising observations, agent state, and task instruction. Our driving implementation uses SimWAM's video/action architecture (Zhao et al., 2026). Training pairs c with recorded images V^+ and a demonstrated trajectory $\tau \in \mathbb{R}^{T \times d_a}$, with T points and d_a coordinates per point. Fig. 3 summarizes the method. Future images in V^+ provide training supervision, while deployment uses only c to predict trajectories. In driving, the target trajectory contains 40 ego-frame position and heading samples over four seconds. Training proceeds in three stages: fitting and freezing the visual–residual codec, learning a causal predictor of its paired targets, and optimizing a small residual actor around the frozen imitation policy. The target encoder and the deployed visual conditioning network have separate parameters, so recorded future observations shape the supervision without becoming an input to the deployed planner.

3.2 Visual Motion and Metric Residual Tokenization

The representation assigns different roles to visual motion and execution detail. A visual scaffold can recover the overall path while retaining a small local displacement from the demonstrated trajectory. We encode that displacement relative to the recovered scaffold, giving the planner an explicit prediction target for the metric trajectory errors that remain after visual reconstruction.

Learning a visual trajectory scaffold. The visual encoder E_v extracts motion from the recorded image sequence, and a learned decoder D_v maps its code into trajectory coordinates. The encoder receives images as its input and summarizes motion across the visual sequence. Trajectory supervision grounds the decoded trajectory in metric coordinates required for precise action execution.

$$z_v^\star = E_v(V^+), \qquad \tau_v^\star = D_v(z_v^\star). \tag{1}$$

The visual code $z_v^\star \in \mathbb{R}^{L\times d_v}$ contains L tokens of dimension d_v for successive temporal intervals. We fit the adaptable encoder parameters and the decoder through trajectory reconstruction with temporal and latent regularization, then freeze the complete visual tokenizer. Appendix A specifies the backbone, adaptation scheme, decoder architecture, and fitting objectives for driving. Our tokenizer adapts a pretrained visual backbone on five frames spanning the prediction horizon. Temporal aggregation produces four 64-dimensional motion tokens, while the decoder integrates predicted body-frame velocities and yaw rates onto the trajectory grid. This makes trajectory supervision responsible for metric scale, including when the image-derived motion features themselves provide only non-metric monocular translation.

Encoding the paired metric residual. The scaffold's remaining error defines $e^\star = \mathrm{vec}(\tau - \tau_v^\star)$. Residuals use a common local coordinate system for subtraction and additive decoding. We fit a mean μ, an orthonormal PCA basis U, and a positive diagonal coefficient scale S on training residuals. These statistics define the paired residual target and the common trajectory decoder.

$$\begin{aligned} z_r^\star &= S^{-1}U^\top(e^\star - \mu), \qquad M = US, \\ D(z_v, z_r) &= D_v(z_v) + \mathrm{unvec}(\mu + Mz_r). \end{aligned} \tag{2}$$

The residual code has dimension d_r, with $U, M \in \mathbb{R}^{Td_a\times d_r}$ and $S \in \mathbb{R}^{d_r\times d_r}$. Each residual target is paired with its visual scaffold, and all examples share the fitted affine residual decoder.

For the driving instantiation, each one-second interval contributes twelve residual coefficients alongside its visual token, forming a 4×76 action interface. A shared interval basis keeps the residual representation compact across the horizon. Heading differences are wrapped before fitting, and both residual statistics and the outer planner normalization are fitted exclusively on training data and held fixed during action learning.

The scale S standardizes the PCA coefficients, giving small physical corrections an explicit learning scale. These coordinates retain whichever errors the scaffold leaves, including both rapid velocity variations and slowly varying offsets. This division follows reconstruction error rather than a prescribed frequency cutoff. Appendix A details the shared interval basis and rank selection, and Appendix A.1 measures the amplitudes and temporal variation of both components.

3.3 Joint Action Learning and One-Step Inference

Joint flow matching. The action expert jointly predicts both token groups from $h(c)$, the video and task features available at deployment, with future-video tokens masked. Let $z^\star = \mathcal{N}(\mathrm{concat}(\mathrm{vec}(z_v^\star), z_r^\star))$ with training-fitted affine normalization $\mathcal{N}$. For Gaussian noise $\epsilon \sim \mathcal{N}(0, I)$ and $t \in [0, 1]$, we minimize the weighted joint flow-matching objective.

$$\begin{aligned} z_t &= (1-t)\epsilon + tz^\star, \qquad u = z^\star - \epsilon, \\ \mathcal{L}_{\mathrm{FM}}^a &= \mathbb{E}\left[w(t)\|v_\theta(z_t, t, h(c)) - u\|_W^2\right]. \end{aligned} \tag{3}$$

Here $w(t)$ weights time and the fixed positive diagonal matrix W balances the two token groups. The driving model assigns visual and residual coordinates relative weights of $1:2$, normalized across the action interface. Current-frame key/value adapters provide trainable visual conditioning while the remaining video foundation stays frozen. Isolated attention allows action queries to use the observed-frame features but prevents them from accessing future-video tokens, even when future-video supervision is included

during training.

Supervising decoded motion. Inverse normalization and splitting of the clean estimate $\hat{z}_{\text{clean}} = z_t + (1 - t)v_\theta(z_t, t, h(c))$ give $(\hat{z}_v, \hat{z}_r)$. The decoded loss uses the target reconstruction.

$$\mathcal{L}_{\text{IL}} = \mathcal{L}^a_{\text{FM}} + \lambda_v \mathcal{L}^v_{\text{FM}} + \lambda_d \, \ell_\tau(D(\hat{z}_v, \text{sg}[\hat{z}_r]), D(z^\star_v, z^\star_r)) \, . \tag{4}$$

$\mathcal{L}^v_{\text{FM}}$ supervises future video and ℓ_τ measures trajectory error. Stop-gradient sg blocks the decoded-loss gradient through $\hat{z}_r$ only. Flow matching trains both groups through shared parameters.

One-step execution. Let $\widetilde{D}$ apply inverse normalization, splitting, and decoding. One Euler update produces an executable trajectory with a single evaluation of the action expert.

$$\hat{z}^{(1)} = \epsilon + v_\theta(\epsilon, 0, h(c)), \qquad \tau_{\text{IL}} = \widetilde{D}(\hat{z}^{(1)}). \tag{5}$$

Appendix A.2 details the training conventions and iterative sampling of the same field. Changing the evaluation budget therefore changes numerical integration of the same conditional field, with observed-frame features reused across updates. It does not change the target codec or introduce an additional planning model. Decoding and residual arithmetic use FP32 to retain metric precision when converting the predicted latent coordinates into an executable plan.

3.4 Reconstruction and Prediction Analysis

The representation must both retain a precise trajectory and remain predictable from the current context. Work on behavior alignment and robust action tokenization likewise emphasizes the downstream policy interface beyond clean reconstruction (Chen et al., 2026; Dong et al., 2026a). We analyze these requirements separately. All affine identities use a common local coordinate chart for demonstrated and decoded trajectories. A fixed trajectory norm $\|a\|_Q = (\text{vec}(a)^\top Q \, \text{vec}(a))^{1/2}$, with $Q \succ 0$, assigns a consistent physical scale to the trajectory coordinates throughout the following analysis. Planning scores are evaluated separately.

Reconstruction capacity. Because $U^\top U = I$ and S is invertible, substituting the residual target into the decoder gives the reconstruction error of the complete representation. The remaining discrepancy is the component of the centered scaffold error outside the fitted residual subspace.

$$\text{vec}(\tau - D(z^\star_v, z^\star_r)) = (I - UU^\top)(e^\star - \mu). \tag{6}$$

The residual thus directly targets the scaffold's reconstruction error. Its retained span controls the trade-off between code size and reconstruction fidelity. This identity describes capacity when the recorded images and paired trajectory residual are supplied. Deployment additionally requires predicting both token groups from the current observation, which introduces further error.

Decoded prediction error. Assume D_v is locally L_v-Lipschitz from visual tokens with the Frobenius norm into the trajectory norm. Thus, small token changes produce bounded trajectory changes. Define $\varepsilon_{\text{codec}} = \|\tau - D(z^\star_v, z^\star_r)\|_Q$ and $\hat{\tau} = D(\hat{z}_v, \hat{z}_r)$. For predicted tokens in this local region, adding and subtracting the clean reconstruction yields the following trajectory error bound.

$$\|\tau - \hat{\tau}\|_Q \le \underbrace{\varepsilon_{\text{codec}}}_{\text{representation}} + \underbrace{L_v \|\hat{z}_v - z^\star_v\|_F}_{\text{visual prediction}} + \underbrace{\|M\|_{2\to Q} \|\hat{z}_r - z^\star_r\|_2}_{\text{residual prediction}} . \tag{7}$$

The three terms separate representation capacity from visual and residual prediction. This distinction also clarifies the role of the empirical diagnostics: clean-code reconstruction tests retained information, frozen-feature probes test conditional predictability, and sampling sweeps test the effect of numerical integration. A representation can perform well on one of these measures without automatically improving the others, particularly when token normalization and decoder sensitivity differ. The induced norm $\|M\|_{2\to Q}$ measures how residual token errors translate into physical trajectory errors. Consequently, low clean-code error must be accompanied by accurate conditional prediction. Appendix A.4 gives the derivation and the coordinate conventions for our driving implementation.

The physical metric $M^\top Q M$ guides residual refinement. Appendix A.4 analyzes conditional predictability and decoded integration error, evaluated in Sec. 4.3.2 and Appendix B.4.

3.5 Structure-Preserving Residual Policy Refinement

The residual decoder provides a common coordinate system for imitation learning and policy improvement (Johannink et al., 2019). Residual control and latent-noise steering provide related ways to improve imitation policies (Ankile et al., 2025; Wagenmaker et al., 2025); here the intervention is confined to the paired metric residual. After imitation learning, we freeze the action generator, its visual conditioning, and the complete codec. A residual actor reads the frozen policy's hidden states, current context, and base prediction, then outputs correction coordinates $q_\phi \in \mathbb{R}^k$. Reward optimization uses these coordinates to refine execution around the predicted scaffold.

A bounded correction space. We project a family of smooth trajectory perturbations into the residual decoder's span. Regularized whitening of the decoded Gram matrix balances the physical scales of these directions, giving a fixed matrix $B \in \mathbb{R}^{d_r \times k}$. A radial limiter $\Pi_{\mathcal{C}}$ rescales each proposed update to satisfy task-defined bounds on decoded motion and token magnitude.

In driving, six longitudinal and lateral perturbation templates define the correction family. Their projection and whitening are computed from the frozen decoder geometry, then calibrated using training observations. The limiter enforces maximum positional and heading displacements of 1 m and 0.2 rad, respectively, together with a bound on residual-token magnitude. For $\tau_{\mathrm{IL}} = D(\hat{z}_v, \hat{z}_r)$, the correction uses the same local coordinates as the fixed trajectory decoder.

$$\begin{aligned} \Delta z_r &= \Pi_{\mathcal{C}}(B q_\phi), \\ \tau_\phi &= D(\hat{z}_v, \hat{z}_r + \Delta z_r) = \tau_{\mathrm{IL}} + \mathrm{unvec}(M \Delta z_r). \end{aligned} \tag{8}$$

The update preserves the predicted visual tokens $\hat{z}_v$ and their decoded scaffold. The affine residual mean is already included in the base plan, so a zero correction recovers that plan exactly. The limiter bounds the physical intervention, while the fixed basis confines exploration to a k-dimensional subspace of the residual coordinates. These properties preserve the representation's structure without guaranteeing improved reward or task-level safety for every individual correction.

Reward optimization. We treat each training scenario as a contextual bandit, selecting one trajectory and receiving its task reward. The frozen base policy identifies a hard-context pool $\mathcal{H}$ and a disjoint set of high-scoring anchors $\mathcal{A}$. Hard contexts receive reward exploration, while anchors regularize the actor toward zero correction. Gaussian exploration replaces q_ϕ with $q_\phi + \sigma\xi$, where $\xi \sim \mathcal{N}(0, I)$, producing the perturbed plan τ_ϕ^ξ. The objective combines rewards from these perturbed plans with penalties on the actor's deterministic correction for hard contexts and anchors.

$$\begin{aligned} J(\phi) = \mathbb{E}_{s\sim\mathcal{H},\,\xi}[R_\lambda(s, \tau_\phi^\xi)] - \beta\, \mathbb{E}_{s\sim\mathcal{H}} \Omega(q_\phi(s)) \\ - \eta\, \mathbb{E}_{s\sim\mathcal{A}} \Omega(q_\phi(s)). \end{aligned} \tag{9}$$

Here R_λ combines the task reward with penalties for regressions in selected performance criteria.

The prior Ω penalizes the decoded mean correction before limiting. Paired positive and negative Gaussian perturbations estimate the gradient of the smoothed reward for actor updates. The scorer and limiter require no derivatives. The actor starts with a zero output layer, making its initial deterministic behavior identical to the base policy. Both hard contexts and anchors come only from the original training partition; held-out validation and test scenes are excluded. The actor is trained around a ten-step base prediction, then reused unchanged across reported inference budgets, including the one-step setting. Appendix A.3 details the basis, limiter, and reward estimator.

At deployment, the actor applies its deterministic mean correction once, using the same basis and normalization as during training. It requires no reward evaluation, candidate selection, or additional denoising. The complete policy therefore combines one action-expert evaluation with one bounded correction in the existing metric residual coordinates of the learned action representation.

4 Experiments

4.1 Experimental Setups

Datasets and benchmarks. NAVSIM evaluates driving trajectories in simulation (Cao et al., 2025; Dauner et al., 2024). We report PDMS on v1 `navtest`, corrected EPDMS on v2 `navtest`, and reactive two-stage

Table 1 NAVSIM v1 `navtest`. "Frame" counts include the current time step; "C/L" denote camera/LiDAR inputs, and "NFE/plans" denotes generator evaluations/candidate counts. Symbol † marks released-checkpoint results; all other baseline values are from the cited papers, respectively. Bold marks the best PDMS, and blue highlights **VITAL**@1, the proposed one-trajectory setting, against multi-step methods.

Method	Input	Frames	NFE/plans	NC↑	DAC↑	EP↑	TTC↑	C↑	**PDMS↑**
Regression / selection									
TransFuser (Chitta et al., 2023)	3C+L	1	– / 1	97.7	92.8	79.2	92.8	100.0	84.0
Hydra-MDP++ (Li et al., 2025a)	3C	2	– / many	97.6	96.0	80.4	93.1	100.0	86.6
DiffusionDrive (Liao et al., 2025)	3C+L	1	2 / 20	98.2	96.2	82.2	94.7	100.0	88.1
DiffusionDriveV2 (Zou et al., 2025)	3C+L	1	2 / 20	98.3	97.9	87.5	94.8	99.9	91.2
Vision–language									
ReCogDrive (Li et al., 2026d)	1C	1	5 / 1	97.9	97.3	87.3	94.9	100.0	90.8
SGDrive (Li et al., 2026b)	1C	4	5 / 1	98.6	97.8	85.8	96.2	100.0	91.1
Predictive / world-action									
DriveLaW (Xia et al., 2026a)	1C	4	– / 1	99.0	97.1	81.3	96.7	100.0	89.1
ReWorld (Xia et al., 2026b)	1C	4	5 / 1	99.1	98.2	82.0	97.7	99.8	90.4
EponaV2 (Xu et al., 2026a)	1C	4	– / 1	98.6	97.9	84.8	95.7	100.0	90.4
Metis (Top-6) (Li et al., 2026a)	1C	1	10 / 6	98.5	97.5	84.0	95.1	100.0	89.7
SimWAM (RL) (Zhao et al., 2026)	1C	1	10 / 1	98.4	98.7	86.4	95.5	100.0	**91.5**
SUV† (Yuan et al., 2026)	1C	4	1 / 1	99.0	96.8	84.1	95.9	100.0	89.7
SUV†	1C	4	10 / 1	99.1	97.8	84.5	96.8	100.0	90.9
SimWAM (IL)† (Zhao et al., 2026)	1C	1	1 / 1	97.3	91.2	79.1	81.8	20.5	66.6
SimWAM (IL)†	1C	1	10 / 1	98.7	98.0	83.9	95.7	100.0	90.1
One-pass, multi-proposal									
MeanFuser (Wang et al., 2026c)	3C	1	1 / multi	98.6	97.0	82.8	95.0	100.0	89.0
DRIFT (Xing et al., 2026)	1C	n.r.	1 / 48	98.6	97.6	84.0	94.6	100.0	89.6
VITAL *with one generated trajectory*									
VITAL w/o residual RL	1C	1	1 / 1	98.6	98.0	84.1	95.6	100.0	90.2
VITAL w/o residual RL	1C	1	10 / 1	98.6	98.1	84.1	95.6	100.0	90.3
VITAL	1C	1	**1 / 1**	98.7	98.5	84.8	95.8	100.0	90.8
VITAL	1C	1	10 / 1	98.6	98.7	85.0	95.7	100.0	91.0

EPDMS on `navhard`. These scores summarize collision avoidance, road compliance, progress, and comfort on a scale from 0 to 100. We also report average and final displacement errors, ADE and FDE, in meters. All three evaluations use the policy trained on v1. The diagnostic experiments use deterministic, matched evaluation panels: reconstruction uses all 18,179 validation scenes, action learning uses a log-balanced panel of 4,096 scenes, and frozen-feature probes use a nested 512-scene panel. All fitted codecs and normalization statistics use training data only. Reconstruction and probe errors use the dense 10 Hz grid; the action-learning comparison evaluates displacement at the eight nominal planning times, so errors on these different grids are interpreted within their respective protocols.

Training and inference. We train the main model on one server with eight NVIDIA H200 GPUs. At inference, **VITAL** predicts one four-second trajectory from the current front image, ego state, and navigation command. NFE counts action expert evaluations. Latency includes image encoding and residual correction at batch size one on an RTX 4090. See Appendices A and B for full protocols. Timing includes preprocessing, host-to-device transfer, visual conditioning, action integration, and trajectory construction, with scene-file I/O and initial weight loading excluded. We synchronize the GPU around the complete call and measure the held-out timing panel after warmup. Across the sampling sweep, the policy, inference seed, decoder, and residual actor remain fixed, isolating the cost and benefit of additional action-flow evaluations.

4.2 Main Benchmark Results

One-step planning quality and inference cost. As shown in Tab. 1, **VITAL** achieves 90.8 PDMS with one denoising step. Fig. 4a and Fig. 4b show that increasing NFE from one to ten adds only 0.1 PDMS points while latency rises from 149 to 490 ms. One evaluation therefore retains nearly all of the planning quality at 3.29× the inference speed. The measurement includes image encoding and residual correction, with visual features reused across action evaluations. MeanFuser (Wang et al., 2026c) and DRIFT (Xing et al., 2026) combine multiple proposals in one pass. NFE/plans distinguishes their proposal budgets from **VITAL**'s single generated trajectory. These published comparisons provide benchmark context, while conditioning inputs, pretraining, and proposal mechanisms vary across methods. We therefore use the

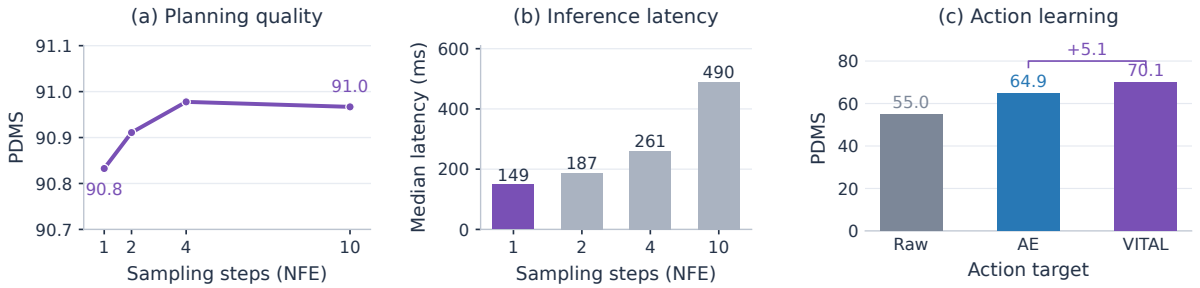

Figure 4 Planning quality, inference latency, and action learning. (a) and (b) report the final policy's PDMS and median latency on NAVSIM v1 `navtest` at different flow evaluation budgets. (c) compares one-step PDMS for raw trajectory, AE, and **VITAL** targets after one training pass with shared action-core initialization and batch order, before residual reinforcement learning.

matched target-learning experiment below to assess representation choices under a shared action-learning budget, rather than attributing every cross-model score difference to the action representation alone.

One-step capability before residual post-training. The imitation-trained **VITAL** policy scores 90.2 PDMS at one evaluation and 90.3 at ten, showing that the small sampling gap precedes residual post-training. The released SimWAM imitation checkpoint (Zhao et al., 2026) scores 66.6 and 90.1 at these budgets. SUV (Yuan et al., 2026) scores 89.7 and 90.9 using four front-camera frames. Section 4.3.3 further compares action learning from shared random initialization.

NAVSIM v2 under corrected scoring. As shown in Tab. 2, rescoring the deployed v1 trajectories under corrected v2 EPDMS gives **VITAL** a score of 91.1 with one evaluation, compared with 91.0 for ten-step SUV (Yuan et al., 2026) and 90.3 for top-6 Metis (Li et al., 2026a). Ten evaluations add only 0.1 EPDMS. This result shows that the small sampling gap persists under corrected scoring on the same `navtest` scenes, without further policy training or candidate selection. This evaluation changes the scoring protocol while retaining the deployed trajectories. It complements the reactive test below: the former checks whether the planning advantage survives corrected metrics, whereas the latter evaluates execution on continued scenes with reactive behavior.

NAVSIM v2 on reactive continuations. The reactive continuations in `navhard` test the same policy beyond rescoring recorded scenes. Without retraining, **VITAL** reaches 38.6 two-stage EPDMS with one evaluation, compared with 37.6 for ten-step SimWAM (Zhao et al., 2026), 36.9 for ten-step SUV (Yuan et al., 2026), 36.6 for GeoWAM (Lu et al., 2026b), and 34.5 for PerceptDrive (Liu et al., 2026d). Ten evaluations yield 38.1 for **VITAL**, so additional sampling offers no gain here. Appendices B.2 and C.3 examine the scope of local refinement under these reactive continuations.

4.3 Ablation Study

4.3.1 From Visual Motion to Metric Precision

To separate representation error from forecasting error, we reconstruct trajectories from recorded current and future images. As shown in Table 3, the visual scaffold captures the maneuver but leaves 0.12 m ADE. Adding the paired residual encoded from ground truth reduces ADE to 3.11×10^{-4} m. This measures the precision retained by the representation before action prediction. Shuffling the visual code raises ADE to 7.41 m, whereas shuffling only the residual gives 0.19 m. The contrast supports complementary roles for the visual scaffold in defining the primary path and the residual in restoring local detail. Because both reconstructions receive recorded future images, their difference isolates what the codec retains rather than the planner's ability to forecast an unseen future.

The residual is also computed from the paired demonstrated trajectory in this diagnostic; its near-exact reconstruction is a capacity result, while deployment must predict the residual from current observations. PDMS measures simulated execution and is assessed separately from pointwise reconstruction fidelity. As detailed in Appendix A.1, the residual also has a larger share of velocity energy at high frequencies, consistent with its role in capturing finer trajectory variations.

Table 2 NAVSIM v2 transfer. One-stage `navtest` is on the left and two-stage `navhard` on the right. NFE/plans denotes generator evaluations/candidate counts; dashes mark non-iterative predictions, n.r. unreported values, and AR autoregression. * marks extra refinement or guidance; bold marks the best EPDMS in each split, while blue highlights VITAL@1 throughout the table.

Method	NFE/plans	EPDMS↑
Regression / selection		
DiffusionDriveV2 (Zou et al., 2025)	2/20	87.5
DriveSuprim (Yao et al., 2026b)	–/256	87.1
SparseDriveV2 (Sun et al., 2026)	–/200	90.1
Vision–language		
DriveVLA-W0 (Li et al., 2026c)	AR/1	86.9
DriveWorld-VLA (Jia et al., 2026)	n.r.	86.8
SGDrive (Li et al., 2026b)	5/1	86.2
DriveFine (Dang et al., 2026)	13*/1	89.7
Predictive / world-action		
EponaV2 (Xu et al., 2026a)	n.r.	88.9
Metis (Li et al., 2026a)	10/1	89.5
Metis (Top-6)	10/6	90.3
SUV (Yuan et al., 2026)	1/1	89.8
SUV	10/1	91.0
VITAL w/o RL	1/1	90.7
VITAL w/o RL	10/1	90.7
VITAL	**1/1**	91.1
VITAL	10/1	**91.2**

Method	NFE/plans	EPDMS↑
Regression / selection		
LTF (Chitta et al., 2023)	–/1	25.1
GuideFlow (Liu et al., 2026b)	100*/256	27.1
Vision–language		
DriveVLA-W0 (Li et al., 2026c)	AR/1	24.4
ReCogDrive (Li et al., 2026d)	5/1	25.7
SGDrive (Li et al., 2026b)	5/1	25.5
Predictive / world-action		
EponaV2 (Xu et al., 2026a)	n.r.	36.1
Metis (Li et al., 2026a)	10/1	32.2
SUV (Yuan et al., 2026)	1/1	33.0
SUV	10/1	36.9
SimWAM (Zhao et al., 2026)	10/1	37.6
GeoWAM (Lu et al., 2026b)	–/1	36.6
PerceptDrive (Liu et al., 2026d)	25/1	34.5
VITAL w/o RL	1/1	38.5
VITAL w/o RL	10/1	**38.6**
VITAL	**1/1**	38.6
VITAL	10/1	38.1

Table 3 Ablation on recorded-image reconstruction. Residuals use ground truth.

Code pairing	ADE (m)↓	FDE (m)↓	Yaw (°)↓
Scaffold only	0.12	0.22	0.43
Zero residual	0.12	0.22	0.43
Matched pair	$\mathbf{3.11 \times 10^{-4}}$	$\mathbf{4.62 \times 10^{-4}}$	$\mathbf{3.38 \times 10^{-3}}$
Shuffle visual	7.41	15.10	14.70
Shuffle residual	0.19	0.34	0.65

4.3.2 Visual Conditioning and Target Predictability

Equal-capacity ridge probes predict VITAL or Trajectory AE tokens, both with 304 coordinates, from the same frozen features. Tab. 4 compares ego/navigation input E with input V that also includes current-image features. Each frozen codec decodes the predicted tokens, and MSE measures standardized token error with vision. This controls both output dimension and probe capacity.

With E alone, AE has lower ADE at 1.08 m versus 1.62 m. Matched vision reverses this ordering. VITAL reaches 0.72 m ADE and 79.4 PDMS, versus AE's 0.80 m and 73.2 PDMS. Shuffling vision raises VITAL's ADE to 2.20 m and lowers PDMS to 52.3, tying its advantage to scene correspondence. Appendix B.3 provides the full controls. The lower decoded ADE despite higher token MSE also shows why prediction error must be assessed in physical coordinates, as analyzed in Sec. 3.4. The shuffled-image control preserves feature dimensionality while breaking scene correspondence, helping distinguish useful visual conditioning from the mere presence of additional inputs. These linear probes characterize predictability within a fixed model class; the subsequent joint action-learning comparison checks whether the same representation is beneficial for a trained nonlinear planner.

4.3.3 Joint Alignment and Precision in Action Learning

We compare VITAL, raw trajectories, and Trajectory AE under matched training and evaluation. Trajectory AE matches VITAL's 4×76 token shape, while raw prediction retains the 8×3 target of SimWAM (Zhao et al., 2026). All three share the visual foundation, random action-core initialization, batch order, and 4,256 updates at batch size 20, without residual post-training.

As shown in Tab. 5, VITAL reaches 70.1 PDMS after one training pass, versus 64.9 for AE and 55.0 for raw prediction, with the lowest ADE and FDE at one evaluation. Its advantage over AE persists at ten evaluations, with 71.4 versus 66.3 PDMS. AE already has higher comfort scores at both budgets, while VITAL improves NC, DAC, EP, TTC, and DDC. The decomposition in Appendix B.3.2 attributes the

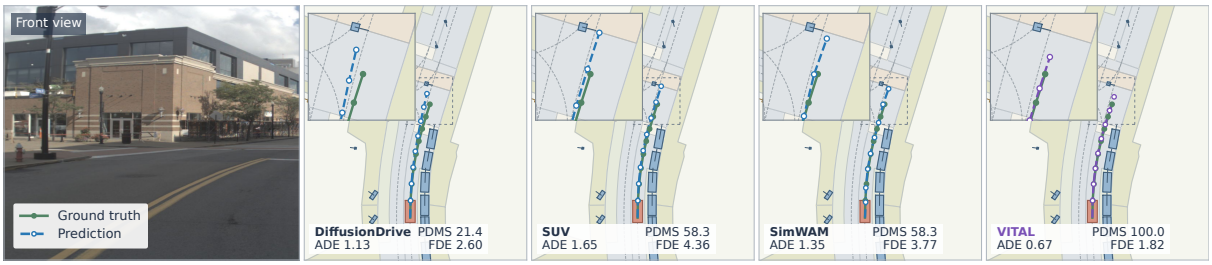

Figure 5 A curved-road scene with local trajectory insets. **VITAL** uses one sampling step, SUV and SimWAM-RL use ten, and DiffusionDrive uses two. ADE and FDE are in meters.

PDMS gain to a larger combined contribution from progress and TTC despite lower comfort. Together with the visual probes, these results support the planning value of combining visual alignment and metric precision in a shared action representation.

Table 4 Ablation on feature probes. Pairs add matched/shuffled vision to ego/navigation.

Target	ADE (m)↓	PDMS↑	MSE_V ↓
Trajectory AE	1.08 → 0.80	57.3 → 73.2	**0.12**
VITAL shuf.	1.62 → 2.20	56.4 → 52.3	–
VITAL	1.62 → **0.72**	56.4 → **79.4**	0.38

Table 5 Ablation on one-step action learning. All targets share action-core initialization.

Action target	PDMS↑	ADE (m)↓	FDE (m)↓
Raw trajectory	55.0	1.46	3.49
Trajectory AE	64.9	1.25	2.45
VITAL	**70.1**	**1.16**	**2.13**

4.3.4 Structure-Preserving Residual Post-Training

Residual post-training adds 0.6 PDMS at NFE= 1 and 0.7 at NFE= 10 through local corrections. As shown in Tab. A15, three training seeds yield positive gains across the tested reward pools and prior weights. Corrected v2 EPDMS also improves after refinement. Appendix B.2 shows that refinement improves Stage 1 while slightly lowering Stage 2 EPDMS. Local corrections help execution but cannot replace a change to the predicted maneuver when the continuation requires one. The refinement mechanism keeps the action generator and predicted visual scaffold fixed, so its gain reflects execution changes within the bounded residual space. Its behavior on reactive continuations also makes the scope of this intervention concrete: a local correction can adjust a path around the existing plan, but a substantially different continuation may require reconsidering the maneuver itself.

4.4 Qualitative Analysis

As shown in Fig. 5, one-step **VITAL** follows the bend more precisely than released DiffusionDrive (Liao et al., 2025), SUV (Yuan et al., 2026), and SimWAM-RL (Zhao et al., 2026). It achieves 0.67 m ADE and 1.82 m FDE, versus 1.35 m and 3.77 m for SimWAM-RL. This example illustrates how finer trajectory detail can be retained even with a single flow evaluation. Further qualitative comparisons and examples of residual refinement are provided in Appendices B.5 and C.

5 Conclusion

In this paper, we present **VITAL**, a framework for world-action planning that jointly addresses visual–action alignment and metric precision. Guided by our reconstruction study, we pair a visual motion scaffold with a normalized metric residual to preserve both visual grounding and trajectory precision. We further develop residual reinforcement post-training that refines planning through bounded corrections while preserving the predicted scaffold. With only one denoising step, **VITAL** achieves 90.8 PDMS on NAVSIM v1, 91.1 EPDMS on v2 and 38.6 EPDMS on `navhard`. The controlled reconstruction and action-learning studies indicate that visual grounding and metric fidelity are complementary properties of an action target. Recorded visual motion provides a useful scaffold, while its paired residual preserves execution detail and supplies a shared coordinate system for subsequent refinement. At deployment, both components are predicted from current observations, and reward evaluation is unnecessary.

Our experiments also delineate the scope of this design. Clean-code precision does not remove conditional prediction error, and bounded residual refinement cannot replace a different maneuver when reactive continuations demand one. Extending the approach to broader visual contexts and more interactive planning settings remains an important direction for future work.

Appendix

A Method Details and Analysis **15**
- A.1 Target Representation . . . 15
- A.2 Causal Planner Training . . . 16
- A.3 Residual Policy, Geometry and Reward Optimization . . . 17
- A.4 Derivations and Interpretation . . . 18

B Experimental Protocol and Additional Results **20**
- B.1 Data, Evaluation, and Comparison Protocol . . . 20
- B.2 NAVSIM v2 Transfer Details . . . 20
- B.3 Representation Controls . . . 22
  - B.3.1 Frozen-Feature Predictability Probes . . . 22
  - B.3.2 Learning with a Randomly Initialized Action Expert . . . 23
  - B.3.3 Decoder Sensitivity . . . 25
- B.4 Inference Efficiency and Solver Diagnostics . . . 25
- B.5 Residual-Policy Controls . . . 26

C Comparison Protocol and Limitations **28**
- C.1 Protocol for the Cross-Model Comparison . . . 28
- C.2 Additional Cross-Model Comparisons . . . 29
- C.3 Failure Cases and Scope of Residual Refinement . . . 29

D Potential Societal Impact & Limitations **29**
- D.1 Potential Societal Impact . . . 29
- D.2 Known Limitations . . . 30

E Public Resources Used **31**
- E.1 Public Datasets Used . . . 31
- E.2 Public Implementations Used . . . 31

A Method Details and Analysis

A.1 Target Representation

Driving instantiation. We instantiate the trajectory interface in Section 3.1 with $T = 40$ and $d_a = 3$. Each trajectory point contains planar position and heading in the current ego frame, sampled at 10 Hz over a four-second prediction horizon. The deployment context contains the current front-camera image, ego state, and navigation command. The tokenizer uses $L = 4$ visual tokens with $d_v = 64$, and the complete residual code has dimension $d_r = 48$. Pairing each visual token with twelve residual coefficients gives a 4×76 interface used by the action predictor and decoder.

Tokenizer architecture and fitting. The target encoder adapts VGGT-Ω (Wang et al., 2026a) through LoRA on five front-camera frames at $0, 1, 2, 3, 4$ seconds, keeping the pretrained base weights fixed. Its learned temporal encoder combines camera features, register summaries, inter-frame register correspondences and relative camera-motion features. Image-estimated monocular translation is treated as non-metric, with trajectory supervision supplying metric scale. Two temporal layers and two interval-mixing layers produce four 64-dimensional motion tokens. A width-256 decoder predicts body-frame forward/lateral velocities and yaw rates, bounded through its output parameterization and integrated on the 40-point, 10 Hz trajectory grid. The tokenizer is fitted using trajectory, increment and keyframe reconstruction, with auxiliary motion-smoothness and latent-conditioning terms. Tokenizer fitting also uses noisy-code reconstruction and an auxiliary flow task conditioned on a 256-dimensional feature from the jointly encoded five-frame sequence. Before the tokenizer's temporal mixer, this feature fuses the first-frame camera head hidden state and register mean, attention pooling over the first-frame registers, and relative camera-motion features. Latent-velocity supervision and a decoded loss on one-step noise-to-latent predictions update the tokenizer during fitting. The deployed action learner uses the causal current-frame features described in Appendix A.2. The visual tokenizer is selected by held-out reconstruction error. The encoder, LoRA adapters, and decoder are then frozen to fit the paired residual to their reconstruction errors.

For reproducibility, the tokenizer's principal reconstruction weights are 1 (position), 0.25 (increments), 0.75 (keyframes), and 0.1 (heading). The noisy-code standard deviation is 0.10, with reconstruction weight 0.15. The latent-flow and one-step decoded-probe weights are both 0.05, ramped over 3,000 updates. Other coefficients are 0.1 for body velocity, 0.02 for yaw rate, 0.002 for acceleration, $2 \cdot 10^{-5}$ for jerk, 0.01 for boundary continuity, and 0.01/0.02/0.05 for latent mean/variance/covariance. The planner losses use the separate weights specified in Equation 4.

Paired metric residual coordinates. For interval j, stack its ten (x, y, ψ) pose errors into $e_j \in \mathbb{R}^{30}$. The pose residual is computed in the ego reference frame, not as a body-frame increment residual. Positions are measured in meters and heading in radians. Heading residuals are wrapped using $\operatorname{atan2}(\sin\delta, \cos\delta)$ before fitting. We fit one shared interval mean $\mu_0 \in \mathbb{R}^{30}$, basis $U_0 \in \mathbb{R}^{30\times 12}$ and coefficient standard deviations S_0 on training residuals. Thus, with column-vector notation,

$$U = I_4 \otimes U_0, \quad S = I_4 \otimes S_0, \quad M = US \in \mathbb{R}^{120\times 48}, \quad \mu = \mathbf{1}_4 \otimes \mu_0. \tag{10}$$

PCA is fitted to $\tau - D_v(z_v^\star)$ in raw pose coordinates, capturing the dominant errors left by the paired visual scaffold. Its coefficient scale S^{-1} is distinct from the planner's outer token normalization $\mathcal{N}$. Both transformations are fitted on training data and frozen before action learning.

We swept per-interval ranks $\{2, 4, 6, 8, 9, 10, 11, 12, 16, 24, 30\}$ and selected 12 to balance compactness and reconstruction accuracy. Table A1 reports validation errors for ranks near the chosen operating point. Rank 16 further reduces reconstruction error at the cost of a larger action interface, while rank 12 retains submillimeter mean accuracy with fewer residual coordinates per interval.

Motion scale and temporal variation. On the training partition, we measure the decoded visual scaffolds and their paired residuals on the 40-point, 10 Hz grid. We compute orthonormal DCT-II spectra after removing each pose sequence's temporal mean. Velocity uses the 39 forward differences $(p_{i+1} - p_i)/(0.1\,\text{s})$ without demeaning. Bins $k \in [0, 4)$ define low-frequency energy, and the remaining bins define higher-frequency energy. Table A2 pools squared coefficients across scenes before computing these proportions. Figure 2 summarizes the high-frequency velocity fractions by their arithmetic mean over x and y. The scaffold carries predominantly smooth motion. Measured before mean removal, the residual's XY RMS is only 1.43% of the scaffold's, supporting its role as a small metric correction. Its velocity varies more rapidly, but its pose retains slower corrections.

Table A1 Residual-rank reconstruction. Errors are in millimeters. Bold marks the chosen operating rank.

Rank	8	10	11	**12**	16
Mean error↓	2.76	0.74	0.73	**0.31**	0.06
P99 error↓	10.20	1.48	1.47	**1.02**	0.19

Table A2 Temporal energy proportions (%). LF and HF denote low and high frequencies of decoded trajectories.

	Pose LF		Velocity HF	
Component	x	y	x	y
Visual scaffold	99.70	99.50	0.05	0.10
Conditional residual	91.90	92.50	52.90	69.90

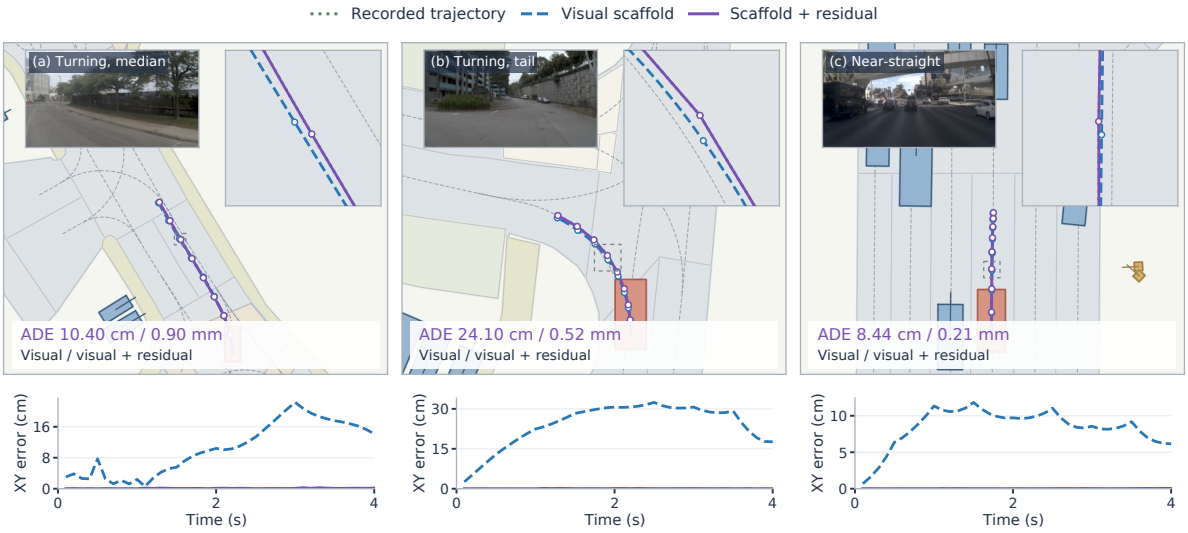

Figure A1 Clean-code trajectory reconstruction. Visual tokens recover the maneuver, while the paired residual corrects local path error in two turning cases and one near-straight case.

Trajectory reconstruction accuracy. The reconstruction metrics in Figure 2 use the full validation partition, as in Table 3. Pure visual reconstruction and the complete codec share the same recorded visual sequence and visual scaffold. The latter additionally receives the paired ground-truth trajectory residual. On the 40-point, 10 Hz grid, their ADE/FDE are 123.60/221.60 mm and 0.31/0.46 mm, respectively. Their PDMS on the same full validation set are 92.1 and 94.7.

PDMS evaluates simulated execution rather than pointwise agreement with the recorded trajectory. Progress uses the PDM reference, so exact trajectory reconstruction need not score 100.

The VGGT-Ω base feature path contains 1,111,483,401 parameters in its aggregator and camera head, excluding the unused dense prediction head and added LoRA adapters. The fitted tokenizer has 4,162,010 temporal encoder parameters, 219,523 motion decoder parameters, and 1,811,520 parameters in its training-only auxiliary predictor. These three learned modules total 6.19M parameters in addition to the VGGT LoRA adapters. Both reconstruction variants share the complete visual tokenizer. The paired residual adds 402 fitted scalars, comprising a 30-dimensional mean, a 30×12 basis, and 12 coefficient scales. These are training-data statistics held fixed after fitting.

Clean-code examples. Figure A1 separates the two parts of a clean target code. Across two turning scenes and one near-straight scene, the visual token recovers the route-level shape, while the paired residual removes the remaining local displacement. The enlarged views show how the combined code follows the recorded trajectory through both straight motion and changes in direction.

A.2 Causal Planner Training

Planner training and causal conditioning. The driving-pretrained video/action foundation and isolated attention follow SimWAM (Zhao et al., 2026). **VITAL** supplies a different action target and decoder. The target encoder above does not run at deployment. The video expert and visual target encoder have separate parameters. The trainable planner components are the action expert and rank-16 key/value adapters applied only to the first-frame token prefix in the last eight video blocks. Future-video and action queries can consume the adapted current-frame keys/values, while action queries cannot access future-video tokens. The remaining video foundation weights are frozen. All examples share a fixed text

prompt. At deployment the current-frame features are reused across action evaluations. Initialized from the released SimWAM imitation weights, the principal model inherits 130,000 updates of visual-token action pretraining: 74,000 with action-only supervision and 56,000 with joint video/action supervision. Its action expert and current-frame adapters are transferred to the visual–residual representation as in SimWAM, causal masking prevents action queries from accessing future-video tokens. We train the main model on one server with eight NVIDIA H200 GPUs. The final stage uses 26,000 updates with AdamW, a peak learning rate of 2×10^{-5}, and an effective batch size of 64. This budget excludes the preceding visual-token pretraining. The raw-trajectory, trajectory-AE, and VITAL action-learning comparison in Appendix B.3.2 randomly initializes only the action expert, retains the same pretrained video and proprioceptive foundation, and uses one complete pass through the training data.

Flow matching and sampling conventions. We use the data-to-noise fraction σ, with $x_\sigma = (1-\sigma)z^\star + \sigma\epsilon$ and target velocity $\epsilon - z^\star$. Training draws $\sigma = \varphi_5(u)$, $u \sim \mathcal{U}(0,1)$, where $\varphi_a(u) = au/[1+(a-1)u]$. It also applies the scheduler's positive interior time weight. Equation 3 uses $t = 1-\sigma$ and $v_\theta(z,t,h) = -f_\theta(z, 1000(1-t), h)$. The clean estimate and inverse-normalized token split are

$$\hat{z}_{\text{clean}} = z_t + (1-t)v_\theta(z_t, t, h(c)), \qquad (\hat{z}_v, \hat{z}_r) = \text{split}\big(\mathcal{N}^{-1}(\hat{z}_{\text{clean}})\big)\,. \tag{11}$$

Inference uses the reversed shifted grid $\varphi_5(1-k/K)$, with a complete noise-to-data interval. For $K = 1$, this is exactly one Euler update between the endpoints. Both sampling modes use the same learned field and reuse the observed-frame features.

Visual and residual dimensions have relative flow-matching weights $1:2$, normalized to mean one across the 76 coordinates. The visual prefix therefore accounts for $64/(64+24) = 72.7\%$ of the dimension-weight mass. The imitation objective combines action flow matching with video flow matching weighted by $\lambda_v = 0.25$. Both token groups are learned jointly through the action flow objective. Decoding and PCA arithmetic use FP32 to preserve numerical precision when converting the predicted tokens into metric trajectories.

A.3 Residual Policy, Geometry and Reward Optimization

A fixed, controllable correction family. We set the correction dimension to $k = 6$. The scene-independent correction basis contains six analytic templates. For normalized horizon time $\rho_i = i/40$, define longitudinal and lateral position perturbations proportional to ρ_i^p for $p \in \{2,3,4\}$. The reference axes and heading linearization use a fixed straight, 5 m/s nominal trajectory for every scene. For its velocity v_i and a template velocity perturbation δv_i, the heading component is $(v_{i,x}\delta v_{i,y} - v_{i,y}\delta v_{i,x})/\max(\|v_i\|_2^2, 0.25)$.

Collect these templates in $A \in \mathbb{R}^{120\times 6}$. Their coordinates in the existing residual decoder are obtained by a regularized weighted least-squares solve,

$$B_0 = (M^\top Q_b M + \lambda_b I)^{-1} M^\top Q_b A, \tag{12}$$

where Q_b uses a 5 m/rad heading scale and $\lambda_b = 10^{-6}\,\text{mean}(\text{diag}(M^\top Q_b M))$. The six decoded directions are subsequently whitened using their physical Gram matrix with a 1% diagonal ridge, and assigned a common scalar amplitude. This gives the fixed $B \in \mathbb{R}^{48\times 6}$ used in Equation 8. Amplitude calibration uses only frozen training observations and decoder geometry, with no reward queries. It matches initial executed XY RMS to a 0.2 m calibration target. Projection into the PCA span may change the ideal templates, so they are not exact polynomial controls after decoding.

Radial correction limiter. For a candidate $r = Bq$, let $d = \text{reshape}(Mr) \in \mathbb{R}^{40\times 3}$. We define

$$a(r) = \max\left\{1, \frac{\max_i \|d_{i,xy}\|_2}{1\,\text{m}}, \frac{\max_i |d_{i,\psi}|}{0.2\,\text{rad}}, \frac{\|r\|_2}{4\sqrt{48}}\right\}, \qquad \Pi_{\mathcal{C}}(r) = \frac{r}{a(r)}. \tag{13}$$

This radial rescaling satisfies all three bounds by construction. With unchanged $\hat{z}_v$ and fixed M, the resulting trajectory increment is exactly $M\Pi_{\mathcal{C}}(r)$ before heading wrapping. The affine mean μ is included in the base trajectory only once. Zero correction returns the original base trajectory.

Actor and prior. The actor concatenates the frozen action expert's final hidden states (4,096 values), ego/navigation vector (8), normalized predicted tokens (304), and its own decoded trajectory (120). Its $4528 \to 256 \to 256 \to 6$ network uses SiLU activations and has 1,226,758 trainable parameters. We keep training-fitted feature statistics fixed and clip normalized inputs to $[-10, 10]$. The output layer is initialized

to zero, so the initial deterministic policy equals the base exactly. For $d_{\text{raw}} = \text{reshape}(MBq)$, define

$$\Omega(q) = \frac{1}{120(0.15\,\text{m})^2} \sum_{i=1}^{40} \left(d_{\text{raw},i,x}^2 + d_{\text{raw},i,y}^2 + (5\,\text{m/rad})^2 d_{\text{raw},i,\psi}^2 \right). \tag{14}$$

The prior penalizes the actor's mean before limiting, rather than each sampled correction. Equation 9 uses $\beta = 0.01$ for hard scenes and $\eta = 0.1$ for the disjoint high-scoring anchors.

Training-only reward and component penalties. The hard pool contains all 17,013 scenes with base PDMS below 90 in the original training partition. From a deterministic subset selected before scoring, 6,598 high-scoring scenes form a disjoint no-op anchor pool. Post-training therefore uses 17,013 hard scenes for reward optimization and 6,598 anchors for consistency regularization, all drawn from the original training partition. Held-out validation logs and all `navtest` scenes and logs are excluded from both. Reward is evaluated by the official non-reactive NAVSIM v1 scorer (Dauner et al., 2024), with $R_{\text{PDM}} \in [0, 1]$. For j in no-at-fault collision, drivable-area compliance, time-to-collision and driving-direction compliance, define $c_j(s, \tau) = \max\{0, m_j(s, \tau_{\text{IL}}) - m_j(s, \tau)\}$. Then

$$R_\lambda(s, \tau) = R_{\text{PDM}}(s, \tau) - \sum_j \lambda_j c_j(s, \tau). \tag{15}$$

The four multipliers start at 0.5. After each actor update, an exponential average of deterministic-mean costs is updated with decay 0.95, and $\lambda \leftarrow \text{clip}_{[0,2]}[\lambda + 0.02(\bar{c} - 0.001)]$. Anchors contribute only the no-op prior, with no sampled reward rollouts.

Antithetic Gaussian policy gradient. For one training context, write $F_\lambda(q) = R_\lambda(s, \tau_{\text{IL}} + M\Pi_{\mathcal{C}}(Bq))$. The Gaussian-smoothed reward is $\bar{F}_\lambda(q) = \mathbb{E}_\xi F_\lambda(q + \sigma\xi)$. With fixed multipliers and integrable rewards, the Gaussian score identity gives

$$\nabla_q \bar{F}_\lambda(q) = \frac{1}{\sigma} \mathbb{E}_\xi[F_\lambda(q + \sigma\xi)\xi] = \frac{1}{2\sigma} \mathbb{E}_\xi[(F_\lambda(q + \sigma\xi) - F_\lambda(q - \sigma\xi))\xi]. \tag{16}$$

The second equality follows from Gaussian symmetry. The Monte Carlo estimator averages K independent antithetic pairs:

$$\hat{g}(q) = \frac{1}{2K\sigma} \sum_{k=1}^{K} [F_\lambda(q + \sigma\xi_k) - F_\lambda(q - \sigma\xi_k)]\xi_k. \tag{17}$$

It is unbiased for the gradient of the smoothed reward, not for the unsmoothed deterministic score. Equation 17 applies Gaussian smoothing to the six action coordinates rather than to all network parameters. The reward need not be differentiable through collision checks or the radial limiter. Backpropagating $-\langle q_\phi, \text{sg}[\hat{g}(q_\phi)]\rangle$ implements the corresponding policy gradient, while the two quadratic priors are differentiated directly.

Optimization and deployment. Each update uses 16 hard scenes, eight antithetic pairs per scene, $\sigma = 1$, and 64 no-op anchors. AdamW uses learning rate 10^{-4}, weight decay 10^{-3}, and gradient norm clipping at 1. The residual actor trains for 1,600 updates, giving 409,600 perturbed candidate reward queries in addition to deterministic-mean scoring. Hard scenes are traversed in shuffled uniform passes. The frozen base uses ten-step sampling during residual training. The same residual actor, feature normalization, basis and correction scale are used for every reported inference-step budget. Thus one-step evaluation changes the base sampling budget without retraining the actor. At deployment, $\xi = 0$ and the mean correction is applied once, with no scorer, candidate bank, ensemble or per-scene selection.

A.4 Derivations and Interpretation

Heading and units in the analysis. The decoder wraps final headings to $[-\pi, \pi]$. Affine decoder identities are exact for positions and for locally lifted headings before this final wrap. They are not identities between arbitrary wrapped real representatives. For a pose vector a, define $\|a\|_Q = (a^\top Q a)^{1/2}$ with fixed $Q \succ 0$ after choosing a meter-per-radian scale. The error bounds assume a consistent local chart containing the compared trajectories. Alternatively, shortest-arc heading error cannot exceed the corresponding lifted error. This is the driving instantiation of the trajectory norm in Section 3.4. Evaluation uses NAVSIM PDMS. The RL quadratic prior uses a heading scale of 5 m/rad to balance angular and positional corrections in the trajectory norm.

Reconstruction and prediction bounds. We prove Equations 6 and 7 under the assumptions in Section 3.4. Throughout this derivation, trajectories are identified with vectors in a common local coordinate chart. The PCA basis is orthonormal, its coefficient scales are invertible, and the visual decoder is locally L_v-Lipschitz into the fixed trajectory norm. Predicted and target trajectories use the same local coordinate chart.

Proof. Since $Mz_r^\star = USS^{-1}U^\top(e^\star - \mu) = UU^\top(e^\star - \mu)$, subtracting the reconstructed trajectory proves Equation 6. Next, add and subtract $D(z_v^\star, z_r^\star)$ to obtain

$$\tau - D(\hat{z}_v, \hat{z}_r) = [\tau - D(z_v^\star, z_r^\star)] + [D_v(z_v^\star) - D_v(\hat{z}_v)] + M(z_r^\star - \hat{z}_r). \tag{18}$$

The triangle inequality and the Lipschitz bound control the first two terms. The third uses the induced norm $\|M\|_{2\to Q} = \sqrt{\lambda_{\max}(M^\top Q M)}$, completing the prediction bound.

Conditional prediction is representation-dependent. Fix the conditioning map h. Let $u = z^\star - \epsilon$, $X = (z_t, t, h(c))$, and $m(X) = \mathbb{E}[u \mid X]$. Assume finite second moments and a fixed positive definite token weight W. For the flow-matching term alone,

$$\begin{aligned}\mathbb{E}[w(t)\|v_\theta(X) - u\|_W^2] = \mathbb{E}[w(t)\|v_\theta(X) - m(X)\|_W^2] \\ + \mathbb{E}[w(t)\,\mathrm{tr}(W\,\mathrm{Cov}[u \mid X])]\,.\end{aligned} \tag{19}$$

Proof. Expand $v_\theta - u = (v_\theta - m) + (m - u)$ and condition on X. The cross term vanishes because $\mathbb{E}[u - m \mid X] = 0$. The time weight is measurable with respect to X. Where $w(t) > 0$, the unconstrained population minimizer is therefore $m(X)$. This identity characterizes the action flow-matching term, while shared video supervision also contributes to the training objective. Whether a target is more predictable from visual conditioning remains empirical, and the finite linear probe measures predictability only within its model class. Changing normalization and decoder geometry changes the meaning of the same numerical token error. Let $\widetilde{D}$ denote the composition of inverse token normalization, splitting, and the trajectory decoder, as in Section 3.3. Locally, for differentiable $\widetilde{D}$ and small δ,

$$\|\widetilde{D}(z+\delta) - \widetilde{D}(z)\|_Q^2 = \delta^\top J_{\widetilde{D}}(z)^\top Q J_{\widetilde{D}}(z)\delta + o(\|\delta\|_2^2). \tag{20}$$

In raw residual coordinates, the exact physical metric is $M^\top Q M$. For normalized tokens, the Jacobian additionally includes the inverse normalization map. Equal standardized token errors can therefore imply different trajectory errors, so these metrics need not rank representations identically.

Decoded one-step integration error. Fix h and an initial noise sample z_0. Let z_t solve $\dot{z}_t = v_\theta(z_t, t, h)$ and define the Euler endpoint $z_{\mathrm{E}} = z_0 + v_\theta(z_0, 0, h)$. Assume the learned field is continuously differentiable along a solution on $[0,1]$, and $\widetilde{D}$ is L_D-Lipschitz between its exact endpoint and the one-step output. Define the material acceleration $a(t) = \partial_t v_\theta(z_t, t, h) + J_z v_\theta(z_t, t, h) v_\theta(z_t, t, h)$. Then

$$\begin{aligned} z_1 - z_{\mathrm{E}} &= \int_0^1 (1-t)a(t)\,dt, \\ \|\widetilde{D}(z_1) - \widetilde{D}(z_{\mathrm{E}})\|_Q &\le L_D \int_0^1 (1-t)\|a(t)\|_2\,dt \\ &\le \frac{L_D}{2} \sup_{t\in[0,1]} \|a(t)\|_2. \end{aligned} \tag{21}$$

Proof. The chain rule gives $d\,v_\theta(z_t, t, h)/dt = a(t)$. Thus, $z_1 - z_{\mathrm{E}} = \int_0^1 [v_\theta(z_t, t, h) - v_\theta(z_0, 0, h)]\,dt = \int_0^1 \int_0^t a(s)\,ds\,dt$. Interchanging the integrals yields the identity. The triangle inequality followed by decoder Lipschitz continuity gives the first bound in Equation 21. Taking the supremum of acceleration and using $\int_0^1 (1-t)\,dt = 1/2$ yields the final bound.

The one-step prediction error is bounded by $\|\tau - \widetilde{D}(z_1)\|_Q$ plus the first bound in Equation 21. Better prediction need not imply better integration. Appendix B.4 evaluates these errors against recorded trajectories and an 80-step numerical reference, respectively.

B Experimental Protocol and Additional Results

B.1 Data, Evaluation, and Comparison Protocol

Data partitions. Action learning uses 85,109 training scenes and a fixed, log-balanced validation panel of 4,096 scenes from 212 logs. Reconstruction uses all 18,179 validation scenes. Frozen-feature probes use a nested panel of 512 scenes, while latency and solver diagnostics use 64 scenes. The tokenizer, normalization statistics, residual basis, trajectory-only autoencoder, and residual actor are fitted only on training data. Panels are selected deterministically without rewards or model errors, and each comparison uses identical scenes.

Trajectory error in the action-learning experiment is evaluated at the eight nominal planning times $0.5, 1.0, \ldots, 4.0$ seconds. Clean-code reconstruction, ridge probes, and solver diagnostics instead use the 40-point, 10 Hz target. Errors from the two grids are not compared directly.

Metric components. NAVSIM v1 reports PDMS and five components comprising no at-fault collision (NC), drivable-area compliance (DAC), ego progress (EP), time-to-collision (TTC), and comfort (C). NAVSIM v2 reports EPDMS with nine components comprising NC, DAC, driving-direction compliance (DDC), traffic-light compliance (TLC), EP, TTC, lane keeping (LK), history comfort (HC), and extended comfort (EC). All reported scores are scaled from 0 to 100.

Planner configuration. The planner uses 384×672 front images and the video/action backbone of SimWAM (Zhao et al., 2026). VITAL replaces its action target with four [64 | 12] visual–residual tokens supplied by a separately trained, frozen tokenizer. The residual actor has 1.23M parameters and outputs six smooth correction coordinates. The imitation model, visual prefix, residual basis, decoder, and normalization remain frozen. Across the NFE sweep, only the number of action-flow evaluations changes. Each scene yields one trajectory from a fixed inference seed, followed by one deterministic residual correction, with no reward-based candidate ranking or trajectory ensemble.

Released-checkpoint evaluation. We evaluate the released SUV (Yuan et al., 2026) and SimWAM (Zhao et al., 2026) imitation checkpoints on NAVSIM v1 `navtest`, using their native preprocessing and conditioning. SUV consumes four front-camera observations, scene-dependent text, and 11 ego-state scalars. Its inference jointly generates future visual tokens and actions. SimWAM uses the current front image, a fixed text prompt, and eight ego/navigation scalars, with action-only inference. Their native image resolutions are 384×640 and 384×672, respectively. Each checkpoint is evaluated at one and ten sampling steps without retraining or candidate selection. Scores are computed independently for each predicted trajectory against the PDM reference, using the same v1 metric cache and scorer configuration as VITAL. Both use bfloat16 inference. Their native random seeds, 42 for SUV and 0 for SimWAM, are held fixed across sampling budgets. Table 1 marks these reproduced measurements with a dagger and also reports the imitation-trained VITAL before residual post-training.

Published comparisons and score conventions. Other results are taken from the cited papers. Inputs, foundation models, training data, and proposal-selection mechanisms differ across these systems. Camera counts refer to physical views, including those stitched into a panorama. Temporal counts include current and past camera observations used for inference, excluding future prediction targets and ego-state history. We check these settings against the cited papers and their public inference configurations. The n.r. label identifies frame counts that are not publicly reported.

MeanFuser combines multiple proposals (Wang et al., 2026c), and DRIFT aggregates 48 proposals in one pass (Xing et al., 2026). VITAL returns a single decoded trajectory at every NFE. The one-stage v2 comparison uses the corrected official EPDMS reported by SUV (Yuan et al., 2026). Published aggregates without component breakdowns are retained with missing entries marked by dashes. The v2 two-stage `navhard` aggregation follows the benchmark protocol (Cao et al., 2025). Tables report PDMS, EPDMS, and their component scores to one decimal place. Displacement errors use two decimal places, with scientific notation for nonzero values below 0.01. Latency is reported in whole milliseconds and counts remain integers. Gains, ratios, and confidence intervals are computed before rounding, so displayed values need not reproduce their exact differences.

B.2 NAVSIM v2 Transfer Details

Formal evaluation follows the NAVSIM v2.2 `navhard` mapping without changing the v1-trained model (Cao et al., 2025). The official mapping pairs original scenes with synthetic continuations, and both stages use

Table A3 NAVSIM v2 `navhard`. S1/S2 denote the reactive stages. The last column gives mapping-wise two-stage EPDMS per method. **VITAL** uses its v1-trained weights, and baseline scores come from the cited papers. Scores range from 0 to 100 and dashes mark missing values. Bold marks the best EPDMS, and blue highlights **VITAL**@1.

Method	Stage	NC↑	DAC↑	DDC↑	TLC↑	EP↑	TTC↑	LK↑	HC↑	EC↑	**EPDMS↑**
Regression / selection											
LTF (Chitta et al., 2023)	S1	96.2	79.6	99.1	99.6	84.1	95.1	94.2	97.6	79.1	25.1
	S2	77.8	70.2	84.3	98.1	85.1	75.6	45.4	95.7	76.0	
GuideFlow (Liu et al., 2026b)	S1	96.6	80.5	96.3	99.3	82.3	94.9	91.5	97.7	67.8	27.1
	S2	87.3	76.7	88.8	99.2	84.3	85.1	49.7	93.1	44.5	
Vision–language											
DriveVLA-W0 (Li et al., 2026c)	S1	96.8	83.3	99.0	99.6	84.6	95.3	96.4	97.6	78.2	24.4
	S2	76.8	64.3	79.9	98.3	89.2	75.0	46.8	95.8	53.1	
ReCogDrive (Li et al., 2026d)	S1	96.4	78.9	98.7	99.8	82.6	95.6	94.4	97.6	74.2	25.7
	S2	80.2	65.0	82.4	98.7	85.2	76.9	43.8	96.6	71.8	
SGDrive (Li et al., 2026b)	S1	95.8	87.6	97.8	99.8	84.4	94.7	92.9	97.8	28.9	25.5
	S2	79.4	65.4	79.1	98.9	88.9	75.3	42.7	96.4	29.6	
Predictive / world-action											
EponaV2 (Xu et al., 2026a)	S1	97.3	90.7	99.4	100.0	83.3	97.3	97.3	97.6	60.9	36.1
	S2	83.6	78.0	88.0	98.9	86.0	80.3	50.1	96.1	52.0	
Metis (Li et al., 2026a)	S1	96.6	87.8	99.0	99.3	84.5	95.6	97.8	97.8	77.8	32.2
	S2	79.6	73.3	84.9	97.8	85.8	76.6	47.7	95.4	75.3	
SUV@1 (Yuan et al., 2026)	S1	–	–	–	–	–	–	–	–	–	33.0
	S2	–	–	–	–	–	–	–	–	–	
SUV@10	S1	96.9	94.2	99.3	99.6	84.1	95.6	96.7	97.8	79.6	36.9
	S2	82.7	74.2	85.9	98.4	86.0	78.9	47.2	95.9	69.1	
SimWAM@10 (Zhao et al., 2026)	S1	98.0	92.0	99.7	99.6	83.8	96.2	97.3	97.8	71.6	37.6
	S2	81.8	78.6	87.3	98.4	86.3	78.6	49.5	96.3	69.9	
GeoWAM (Lu et al., 2026b)	S1	97.7	91.5	99.1	99.8	83.8	95.8	96.0	97.8	79.0	36.6
	S2	80.4	76.3	87.3	98.7	88.9	76.2	49.9	94.0	56.0	
PerceptDrive (Liu et al., 2026d)	S1	97.5	94.2	99.0	99.4	83.9	95.6	96.4	97.7	78.1	34.5
	S2	81.9	71.4	85.1	98.5	84.5	77.9	48.7	96.6	69.3	
VITAL *with one generated trajectory*											
VITAL@1	S1	97.6	94.4	99.6	99.6	84.0	96.9	98.9	97.8	77.3	**38.6**
	S2	82.2	79.5	87.4	98.3	86.8	78.3	48.8	95.8	71.6	
VITAL@10	S1	97.0	94.4	99.6	99.8	84.2	97.1	98.7	97.8	76.4	38.1
	S2	82.3	79.8	87.5	98.3	86.8	78.3	48.3	96.0	70.0	

reactive traffic. We aggregate scores within each mapping and then average the stage-wise products. The combined score is not the product of the two reported stage means. The GeoWAM and PerceptDrive comparisons use their published two-stage results, including the separate Stage 1 and Stage 2 components (Liu et al., 2026d; Lu et al., 2026b). Their combined EPDMS values are 36.6 for GeoWAM and 34.5 for PerceptDrive, under the same protocol.

Inference budgets. NFE counts iterative trajectory-generator evaluations, and plans counts candidate trajectories before final selection. A dash denotes non-iterative prediction. GeoWAM forecasts future geometry from three frames of eight camera views and uses a deterministic trajectory decoder. PerceptDrive uses four frames from one front camera, predicts future features once, and integrates its flow actor for 25 Euler steps to return one trajectory (Liu et al., 2026d; Lu et al., 2026b). The one-stage v2 comparisons retain DriveSuprim with 256 candidates and SparseDriveV2 with 200 candidates (Sun et al., 2026; Yao et al., 2026b). DiffusionDriveV2 uses two sampling steps over 20 proposals (Zou et al., 2025). SGDrive and ReCogDrive use the released five-step DDIM planners, each returning one trajectory. DriveFine's marked budget comprises 12 masked-token sampling passes and one refinement pass (Dang et al., 2026). GuideFlow's marked budget denotes 100 flow-integration steps over 256 anchors, with additional constraint guidance and energy-based refinement (Liu et al., 2026b). These counts extend beyond denoising. Missing budgets use n.r.

Residual post-training raises one-step Stage 1 EPDMS from 82.3 to 83.3 but lowers Stage 2 from 47.1 to 46.2, leaving the product nearly unchanged. Training on non-reactive v1 observations therefore improves initial-stage planning without a corresponding gain on reactive continuations.

The complementary one-stage diagnostic rescores deployed v1 trajectories with the v2 simulator and non-reactive traffic. It retains the full `navtest` split and evaluates adjacent-frame extended comfort where available. This diagnostic is reported separately from two-stage EPDMS.

Table A4 Corrected NAVSIM v2 `navtest`. VITAL retains its v1 weights. Baseline scores come from the cited papers. Scores range from 0 to 100. NFE/plans gives generator evaluations and candidate counts, as detailed in Appendix B.2. Dashes mark unreported components. Bold marks best EPDMS and blue highlights VITAL@1.

Method	NFE/plans	NC↑	DAC↑	DDC↑	TLC↑	EP↑	TTC↑	LK↑	HC↑	EC↑	**EPDMS↑**
Regression / selection											
DiffusionDriveV2 (Zou et al., 2025)	2/20	97.7	96.6	99.2	99.8	88.9	97.2	96.0	97.8	91.0	87.5
DriveSuprim (Yao et al., 2026b)	–/256	98.4	98.6	99.6	99.8	90.5	97.8	97.0	98.3	78.6	87.1
SparseDriveV2 (Sun et al., 2026)	–/200	98.1	98.1	99.6	99.8	91.1	97.3	96.9	98.2	78.4	90.1
Vision–language											
DriveVLA-W0 (Li et al., 2026c)	AR/1	98.4	95.2	99.4	99.9	86.6	97.9	97.8	98.3	82.7	86.9
DriveWorld-VLA (Jia et al., 2026)	n.r.	98.6	99.1	99.6	99.8	87.4	97.9	97.0	97.8	78.6	86.8
SGDrive (Li et al., 2026b)	5/1	98.6	94.3	99.5	99.8	86.0	97.9	96.1	98.3	85.9	86.2
DriveFine (Dang et al., 2026)	13*/1	98.7	97.3	99.5	99.8	88.7	97.8	97.7	98.4	83.8	89.7
Predictive / world-action											
EponaV2 (Xu et al., 2026a)	n.r.	98.5	97.4	99.5	99.9	87.9	98.1	97.7	98.2	77.4	88.9
Metis (Li et al., 2026a)	10/1	98.4	97.2	99.6	99.8	87.8	97.7	97.8	98.4	88.0	89.5
Metis (Top-6) (Li et al., 2026a)	10/6	98.5	97.5	99.6	99.8	87.9	97.8	98.0	98.4	90.0	90.3
SUV@1 (Yuan et al., 2026)	1/1	–	–	–	–	–	–	–	–	–	89.8
SUV@10	10/1	99.1	97.8	99.7	99.8	87.8	98.7	98.1	98.4	88.3	91.0
VITAL *with one generated trajectory*											
VITAL@1	**1/1**	98.7	98.5	99.7	99.8	87.8	98.3	97.8	98.4	87.8	91.1
VITAL@10	10/1	98.6	98.7	99.7	99.9	87.9	98.3	97.7	98.4	87.5	**91.2**

Table A5 Residual post-training under NAVSIM v2 transfer. Unchanged v1 policies are evaluated on reactive two-stage `navhard` (left) and rescored one-stage `navtest` (right). Stage columns report EPDMS within each stage.

Model	NFE	Stage 1↑	Stage 2↑	EPDMS↑
w/o residual RL	1	82.3	47.1	38.5
VITAL	1	83.3	46.2	38.6
w/o residual RL	10	82.3	47.3	38.6
VITAL	10	82.9	46.3	38.1

Model	NFE	EPDMS↑
w/o residual RL	1	90.7
VITAL	1	91.1
w/o residual RL	10	90.7
VITAL	10	91.2

B.3 Representation Controls

B.3.1 Frozen-Feature Predictability Probes

Trajectory-code control. Trajectory AE divides the 40-point, 10 Hz trajectory into four segments, each containing ten (x, y, θ) poses. A shared encoder maps the 30 coordinates of each segment through $30 \to 256 \to 256 \to 76$ layers; the decoder is symmetric, with SiLU activations. The complete AE has 186,474 parameters and is fitted on training trajectories. Its four 76-dimensional tokens match VITAL's action interface. Each code is overcomplete relative to its 30-dimensional input, providing a control for token shape. Held-out clean-code reconstruction ADE is 1.80 cm. We use this trajectory-only control in both the probes and the action-learning comparison below.

Table A6 Probes at varying training budgets. E uses ego/navigation input and V adds current-image features. E and V report ADE in meters. Standardized token MSE uses V.

Target	Train	E↓	V↓	MSE↓
Trajectory AE	819	1.09	0.94	0.18
VITAL	819	1.69	0.85	0.40
Trajectory AE	2,048	1.09	0.95	0.15
VITAL	2,048	1.65	0.84	0.47
Trajectory AE	8,192	1.08	0.80	0.12
VITAL	8,192	1.62	0.72	0.38

Inputs and fitting. The probes receive 6,144 channel-wise summary statistics from the frozen causal current-frame video tokens and eight ego/navigation scalars. Training budgets of 819, 2,048, and 8,192 scenes are nested within a deterministic subset of the original training split. Evaluation uses the same held-out panel for every fitting budget. Each target has 304 standardized coordinates. Ridge regression

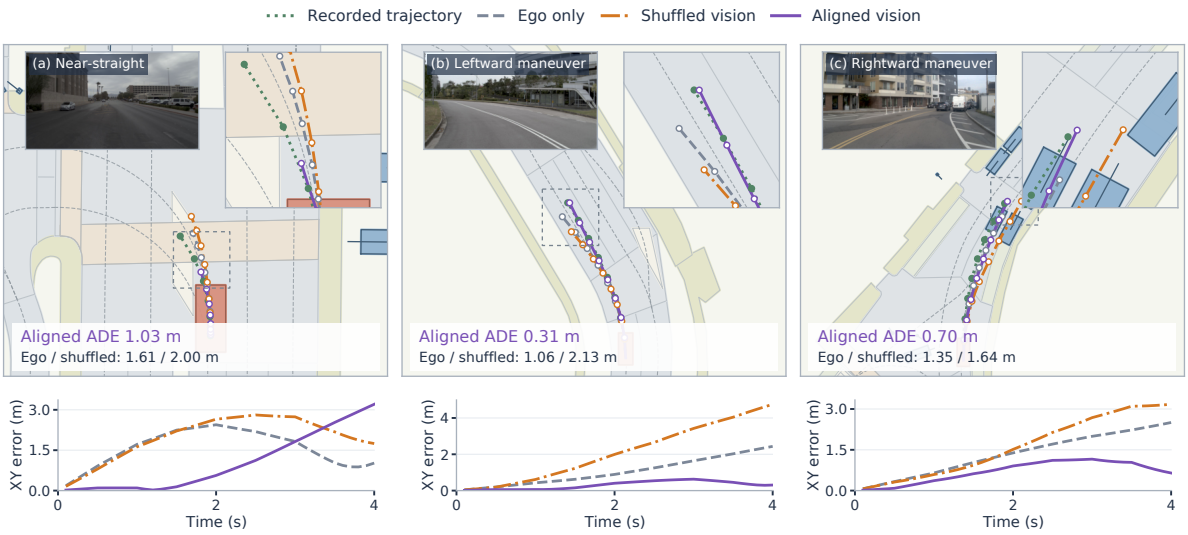

Figure A2 Aligned, ego-only, and shuffled-vision probe predictions. Lower plots show XY error.

uses an unpenalized intercept, penalty 0.01, and the same design matrix and 1,870,512 fitted parameters for both targets.

Planning performance. With 8,192 fitting scenes, adding matched visual features raises **VITAL** from 56.4 to 79.4 PDMS and the trajectory code from 57.3 to 73.2. **VITAL**'s stronger planning gain agrees with its lower decoded ADE despite its higher standardized token prediction MSE.

Scene correspondence. To test whether the gain comes from scene correspondence, we independently permute visual features within the training and validation partitions while preserving ego states, targets, and input size. Shuffling lowers **VITAL** from 79.4 to 52.3 PDMS and raises ADE from 0.72 to 2.20 m. The corresponding trajectory-code values change from 73.2 to 50.1 PDMS and from 0.80 to 1.38 m ADE. The matched-input advantage therefore depends on scene-specific visual information rather than feature dimensionality. ADE in Table A8 is measured in meters.

Table A7 Probe planning scores. E uses ego/navigation and V adds vision. Gains use unrounded PDMS.

Target	E↑	V↑	Gain↑
Trajectory AE	57.3	73.2	16.0
VITAL	56.4	79.4	23.0

Table A8 Visual correspondence control. Matched and shuffled features share the panel. ADE is in meters.

Target	Vision	ADE↓	PDMS↑
Trajectory AE	Matched	0.80	73.2
Trajectory AE	Shuffled	1.38	50.1
VITAL	Matched	0.72	79.4
VITAL	Shuffled	2.20	52.3

Visual correspondence in decoded predictions. Figure A2 shows decoded trajectories from the frozen-feature probes in Appendix B.3. All probes predict the same **VITAL** target with an 8,192-scene fitting budget. In the near-straight, leftward, and rightward maneuvers, matched visual features bring the decoded paths closer to the recorded motion than either ego state alone or shuffled visual features. Below each scene, the error curves trace position error over the prediction horizon.

B.3.2 Learning with a Randomly Initialized Action Expert

Initialization and training exposure. We compare raw trajectories, Trajectory AE, and **VITAL** under the same training and evaluation protocol. Raw prediction retains SimWAM's eight (x, y, θ) poses at 0.5-second intervals. AE and **VITAL** predict 4×76 tokens and use their frozen codecs for decoding. All three load the same pretrained video and proprioceptive foundation. Only the action expert is randomly initialized. Its shared core tensors have identical initial values, and target-specific projections map each trajectory or token interface into this common core. All methods optimize the action expert and the same current-frame adapters on identical batches, with no residual post-training. Table A9 summarizes the shared optimization budget and inference settings.

Table A9 Shared action-learning protocol. Raw, AE, and VITAL process one training pass with identical batch order. Presentations count samples after distributed-batch padding.

Setting	Value	Setting	Value
Initialization seed	42	Global batch size	20
Updates / presentations	4,256 / 85,120	GPUs per method	Two 48 GB RTX 4090s
Optimizer / peak LR	AdamW / 2×10^{-5}	Warmup / schedule	1,000 / cosine
Trainable modules	Action expert and current-frame adapters	Adapter scope	Last 8 blocks, rank 16
Evaluation panel	Fixed and log-balanced	Inference seed / NFE	0 / 1 and 10

Action objectives. All targets use normalized action flow matching and video flow matching weighted by 0.25, without auxiliary losses. The raw baseline predicts trajectory coordinates with unit action weight, while AE predicts a trajectory-only code. For VITAL, visual and residual flow terms have relative weights of one and two, normalized by their weighted coordinate count. AE and VITAL recover poses through their respective frozen decoders. Tokenizers and normalization statistics remain fixed throughout action learning, and the shared budget of 4,256 updates covers one complete pass through the training data after tokenizer and foundation pretraining.

Trajectory-only learning controls. Table A10 tracks AE and raw prediction through the shared training pass. AE's advantage emerges by 3,000 updates and reaches 9.9 one-step PDMS points at the final checkpoint. This establishes a strong learned-code baseline using trajectory supervision alone. As shown in Figure 4c, VITAL raises PDMS by a further 5.1 points over AE.

Table A10 Trajectory-only learning progress. Δ is AE minus raw PDMS; intervals use paired log-bootstrap differences.

Updates	Seen	Raw↑	AE↑	Δ↑	95% CI
1,000	20,000	45.6	47.6	+2.0	[−0.1, 4.0]
2,000	40,000	52.2	53.3	+1.1	[−0.9, 3.0]
3,000	60,000	55.2	64.5	+9.3	[7.2, 11.4]
4,000	80,000	55.0	65.0	+10.0	[8.0, 12.0]
4,256	85,120	55.0	64.9	+9.9	[8.0, 12.0]

Table A11 Trajectory-only score contributions. NC/DAC-gated terms sum to PDMS before rounding.

Target	NFE	EP↑	TTC↑	C↑	PDMS↑
Raw	1	24.4	25.7	4.9	55.0
Raw	10	26.7	32.3	11.6	70.6
AE	1	25.1	27.7	12.1	64.9
AE	10	25.2	28.7	12.5	66.3

Planning gains across sampling budgets. Table A12 evaluates the same final weights at one and ten steps. VITAL scores 70.1 and 71.4 PDMS, exceeding AE's 64.9 and 66.3 at both budgets. Raw prediction improves by 15.6 points with ten evaluations, while AE and VITAL each gain 1.4. Thus both learned token representations support small sampling gaps, while VITAL achieves the highest PDMS at each budget. Its 5.1-point advantage over AE persists with iterative refinement. VITAL also has the lowest one-step ADE and FDE. At ten steps, AE yields lower trajectory errors, while VITAL retains higher scores for collision avoidance, road compliance, and driving progress.

Planning value beyond comfort. The component scores distinguish the tokenizer's planning gains from comfort improvements. AE reaches 99.6/99.2 comfort at one/ten steps, compared with 51.6/57.4 for VITAL. Yet VITAL improves NC, DAC, EP, TTC, and DDC at both budgets. This advantage also survives the multiplicative gating in PDMS. For per-scene components in $[0, 1]$, write

$$\begin{aligned} \mathrm{PDMS} &= 100\,\mathbb{E}\left[\mathrm{NC}\,\mathrm{DAC}\,\frac{5\mathrm{EP} + 5\mathrm{TTC} + 2\mathrm{C}}{12}\right], \\ G_C &= \frac{100}{6}\mathbb{E}[\mathrm{NC}\,\mathrm{DAC}\,\mathrm{C}] \le \frac{\overline{C}}{6}, \end{aligned} \tag{22}$$

where $\overline{C} = 100\mathbb{E}[\mathrm{C}]$ is the reported comfort score. As shown in Table A11, comfort contributes 7.3 points of AE's 9.9-point one-step improvement over raw prediction. For VITAL, even the upper bounds on G_C, 8.6/9.6 points, are below AE's measured 12.1/12.5 points. VITAL's higher PDMS must therefore come from a larger combined gated progress and TTC contribution despite a smaller comfort contribution. Figure A3 summarizes this distinction. Combined with the visual correspondence probes and lower one-step trajectory errors, these results substantiate the benefit of jointly learning visually aligned and metrically precise action targets.

Table A12 Final-checkpoint planning at matched solver budgets. All targets use one training pass. DDC is reported separately from v1 PDMS. ADE/FDE use the same eight-point grid. Bold marks best PDMS at each NFE.

Target	NFE	NC↑	DAC↑	DDC↑	EP↑	TTC↑	C↑	PDMS↑	ADE (m)↓	FDE (m)↓
Raw trajectory	1	92.5	80.3	94.6	58.6	79.1	35.2	55.0	1.46	3.49
Trajectory AE	1	95.4	76.8	95.0	60.3	87.0	99.6	64.9	1.25	2.45
VITAL	1	97.1	90.0	97.3	71.2	87.7	51.6	**70.1**	1.16	2.13
Raw trajectory	10	95.2	87.2	97.4	64.2	89.4	84.0	70.6	0.99	2.82
Trajectory AE	10	95.9	78.8	96.2	60.7	87.7	99.2	66.3	1.10	2.22
VITAL	10	98.2	89.4	97.6	69.8	91.0	57.4	**71.4**	1.20	2.44

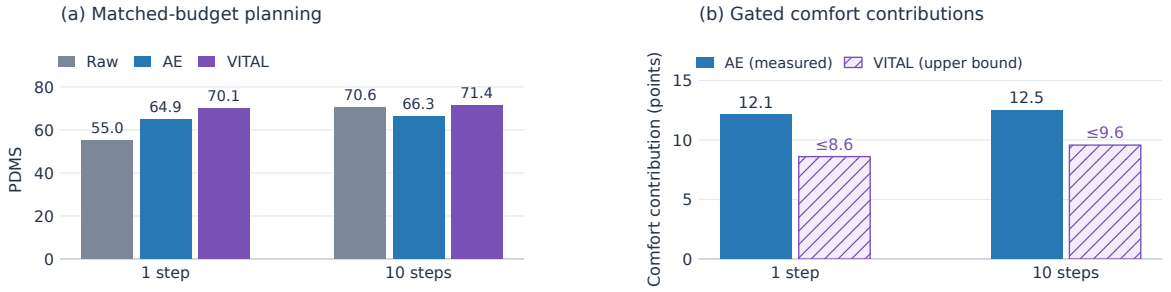

Figure A3 Planning gains beyond comfort. Left: final-checkpoint PDMS at one and ten evaluations. Right: measured AE comfort contributions and upper bounds for VITAL from Equation 22; hatched bars denote bounds.

Evaluation protocol. All evaluations use the same validation panel, with initial noise fixed across sampling budgets. Confidence intervals use 20,000 paired bootstrap replicates over validation logs, retaining all scenes within each sampled log and pairing the corresponding scores across the compared settings.

B.3.3 Decoder Sensitivity

We perturb one normalized token group with Gaussian noise while fixing the other at its clean value. As shown in Figure A4, visual-token noise causes larger deviations across the tested scales. At standard deviation 0.10, visual and residual perturbations yield 5.61 cm and 1.39 cm ADE, respectively, on the full validation partition. The comparison isolates local decoder sensitivity under the fixed training-set token normalization.

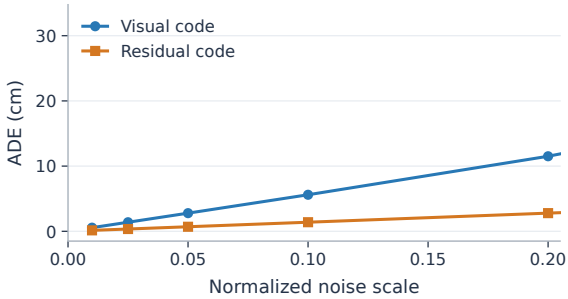

Figure A4 Sensitivity to normalized token noise.

B.4 Inference Efficiency and Solver Diagnostics

Timing protocol. Batch-one latency is measured on an exclusively reserved RTX 4090 after eight warmup calls. Inputs are front images and ego metadata resident in host memory. The measurement includes preprocessing, host-to-device transfer, image encoding, visual conditioning, action integration, and trajectory construction. Scene-file I/O and initial weight loading are excluded. Scene-dependent visual features are recomputed for every call. We synchronize CUDA around the complete call and report the median and 95th percentile over the held-out timing panel.

Residual post-training adds no measurable latency beyond run-to-run variation. At NFE= 1, the median is 149 ms with the actor and 150 ms without it. Increasing NFE from 1 to 10 raises the median to 490 ms

Table A13 Sampling quality and latency. PDMS uses v1 `navtest`. Median and P95 give batch-one latency (ms).

Model	NFE	PDMS↑	Median↓	P95↓
w/o residual RL	1	90.2	150	155
VITAL	1	90.8	149	157
w/o residual RL	2	90.2	184	193
VITAL	2	90.9	187	192
w/o residual RL	4	90.2	262	270
VITAL	4	91.0	261	270
w/o residual RL	10	90.3	487	495
VITAL	10	91.0	490	502

Table A14 Numerical solver refinement. Errors are in centimeters. The last column gives ADE to NFE 80.

NFE	GT ADE↓	GT FDE↓	To NFE 80↓
1	63.30	166.00	6.61
10	63.10	167.00	3.60
40	63.40	167.00	0.97
80	63.40	167.00	0.00

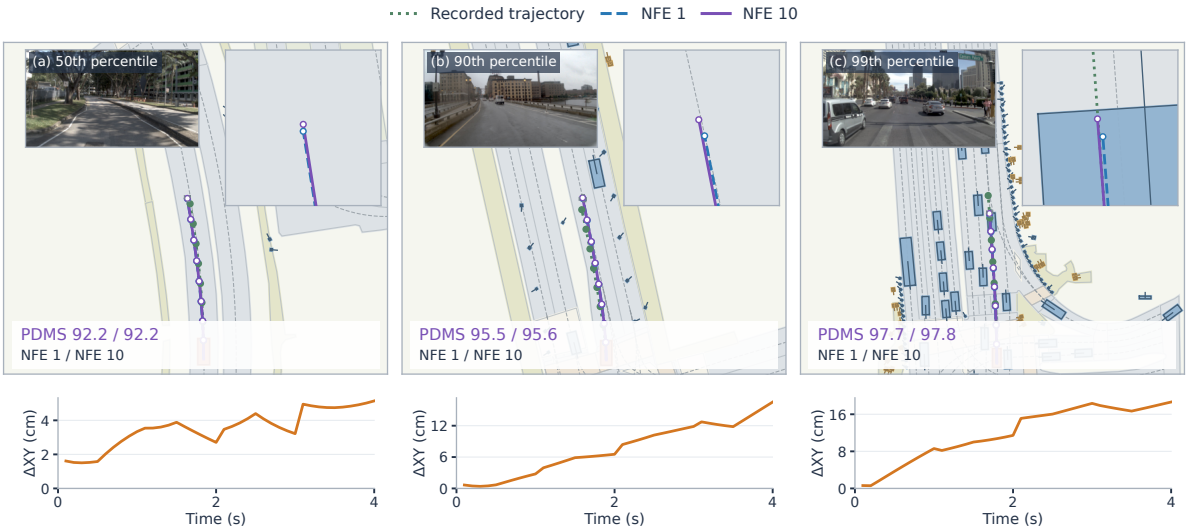

Figure A5 One/ten-step outputs at the 50th, 90th, and 99th percentiles of XY change in high-scoring scenes.

for only 0.1 additional PDMS, with little change in planning quality.

Numerical refinement. Using the held-out solver panel and identical initial noise, we compare NFE 1, 10, 40, and 80 in decoded trajectory space. The 80-step endpoint is a numerical reference, not ground truth. The one-to-80 discrepancy averages 6.61 cm, while the 40-to-80 discrepancy is 0.97 cm. Ground-truth ADE remains between 63.10 and 63.40 cm across all four budgets, so improved numerical agreement does not improve prediction accuracy on this panel.

One-step and ten-step examples. Figure A5 shows one-step and ten-step trajectories from the same trained policy and inference seed. From left to right, the enlarged views and displacement curves reveal increasingly visible local differences, while the paths retain the same maneuver. PDMS changes by at most 0.1 points in these examples. The additional flow evaluations mainly adjust local geometry in these examples, consistent with the small score changes in Table A13.

B.5 Residual-Policy Controls

The pool/prior factorial uses three seeds for each setting. As shown in Table A15 and Figure A6, expanding the reward-training pool from 13,105 to 17,013 hard scenes improves the mean gain by about 0.4 PDMS points, whereas reducing the prior weight has little effect. Within these settings, broader reward-pool coverage improves the residual policy more than the tested reduction in prior weight.

Local corrections in driving scenes. The one-step examples in Figure A7 illustrate three effects of residual correction. Panel a increases route progress while retaining the intended maneuver. Panel b restores road compliance with a maximum displacement of 7.9 cm, raising PDMS from 0.0 to 100.0. Panel c improves time-to-collision by adjusting the local path around another road user. These cases illustrate how small geometric corrections can improve the execution of a plausible maneuver. Appendices A.1 and B.3.1

Table A15 Reward-pool and prior controls. Mean PDMS gain and standard deviation (SD) use three seeds at 1,600 updates and NFE 10. Settings share training anchors, normalization, correction geometry, and query budget.

Pool	Scenes	Prior	ΔPDMS↑	SD
Partial	13,105	0.03	1.2	0.04
Partial	13,105	0.01	1.2	0.05
Full	17,013	0.03	1.6	0.15
Full	17,013	0.01	1.6	0.18

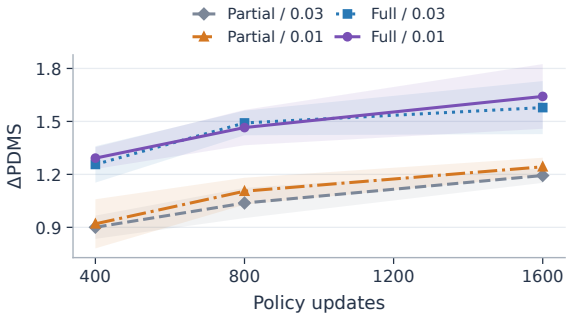

Figure A6 PDMS gains across policy updates. Bands indicate one standard deviation across three training seeds.

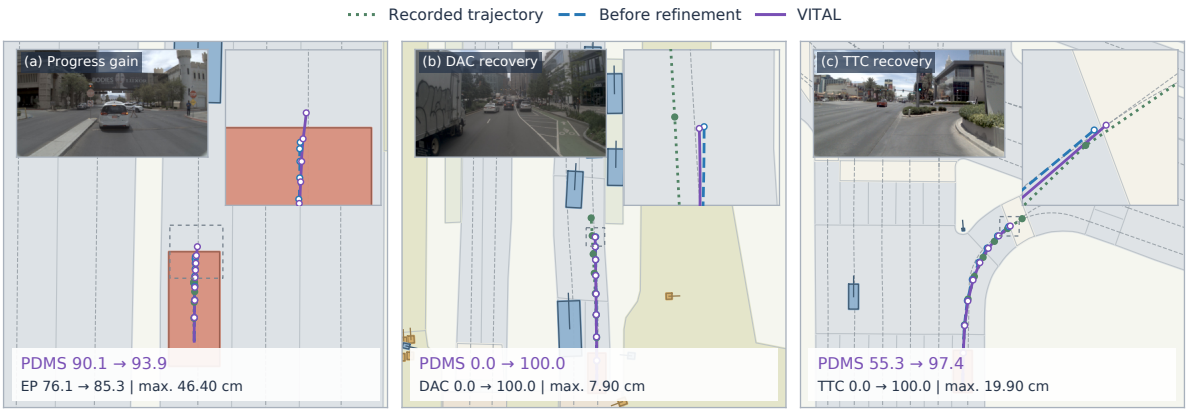

Figure A7 One-step trajectories before and after residual refinement. Panels a, b, and c show gains in progress, drivable-area compliance, and TTC, respectively. Insets show local corrections.

provide complementary examples of reconstruction precision and the dependence of decoded trajectories on correctly matched visual information.

Figure A8 Additional cross-model comparisons with shared local trajectory enlargements. ADE and FDE are in meters.

C Comparison Protocol and Limitations

C.1 Protocol for the Cross-Model Comparison

Figure 5 compares released DiffusionDrive (Liao et al., 2025), SUV (Yuan et al., 2026), and SimWAM-RL (Zhao et al., 2026) weights with **VITAL** on the same NAVSIM v1 scene. All models predict four seconds of motion from their native inputs. DiffusionDrive uses the current three front-facing cameras and LiDAR, with two DDIM steps over 20 proposals. SUV uses four front-camera frames, scene-dependent text, and an 11-dimensional ego-state vector, jointly generating future visual tokens and actions. SimWAM-RL and **VITAL** use the current front image and eight ego/navigation scalars. **VITAL** uses one sampling step, while SUV and SimWAM-RL use ten. All three use bfloat16 inference. DiffusionDrive uses its native float32 implementation. All planners use inference seed zero.

Each prediction is scored independently against the same PDM reference. ADE and FDE compare positions at eight times from 0.5 to 4.0 seconds. **VITAL**'s 10 Hz output contains these times exactly, so no interpolation is required. The BEV panels share the same metric viewport, and the insets magnify the same spatial region. Map and agent overlays use the recorded scene annotations.

As shown in Figure 5, **VITAL** follows the recorded path more closely and reduces endpoint overshoot relative to the three baselines. In this scene, **VITAL** scores 100.0 PDMS, while SUV and SimWAM score 58.3 with zero TTC compliance. DiffusionDrive scores 21.4 with both reduced NC and zero TTC compliance. The shared enlargements show where the predicted paths deviate from the recorded motion and how far they

extend beyond its endpoint near surrounding traffic.

C.2 Additional Cross-Model Comparisons

Figure A8 shows four additional cross-model trajectory comparisons. The local enlargements reveal smaller endpoint overshoot for **VITAL** than SimWAM in the right turn in panel a and the left bend in panel b. In panel b, SimWAM scores 100.0 PDMS with 1.71 m FDE, while **VITAL** scores 96.3 with 0.03 m FDE. This example illustrates the complementary information provided by geometric error and simulated driving quality when comparing trajectories through the same road layout.

C.3 Failure Cases and Scope of Residual Refinement

Safety regressions. As shown in Figure A9, an 11.40 cm maximum displacement moves a trajectory across the drivable-area boundary. DAC changes from one to zero and the multiplicative scorer drops PDMS from 98.7 to 0.0, even though NC remains one. The same multiplicative gate also zeros the progress contribution. Across the full paired one-step evaluation, 64 scenes regress on at least one of NC, DAC, TTC, or DDC, counting each affected scene once across the four component metrics.

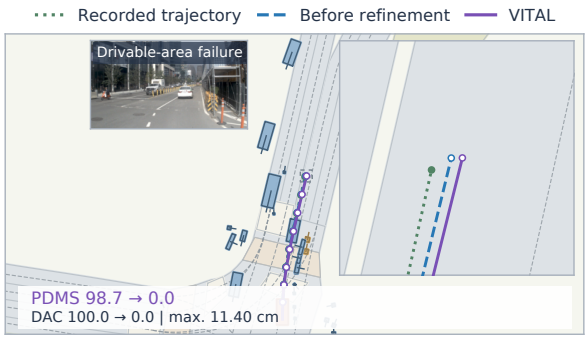

Figure A9 Drivable-area failure after refinement. The inset magnifies the displacement caused by the correction.

The residual space cannot replace the visual scaffold. The final actor changes six smooth coordinates while holding the predicted visual token fixed. This restriction keeps exploration local around the pretrained plan, so it cannot reliably repair a wrong high-level route choice or a large visual-prediction error. The Stage 2 regression on reactive `navhard` in Appendix B.2 further motivates refinement that can revise the scaffold when local corrections cannot adapt to a changed scene.

Perceptual grounding and environmental robustness. Our evaluation does not establish robustness to adverse weather, sensor corruptions, or unreliable visual-language reasoning. DriveBench shows that plausible driving answers need not be visually grounded, while Open-o3-Video explicitly localizes spatio-temporal evidence for video reasoning (Meng et al., 2026; Xie et al., 2025a). Robo3D benchmarks corruption sensitivity in 3D perception, and geometry-grounded weather synthesis provides a complementary way to study adverse conditions (Kong et al., 2023b; Qian et al., 2026). These findings motivate evaluating the visual scaffold under such perturbations: bounded metric refinement cannot compensate for arbitrary errors in the underlying scene interpretation.

D Potential Societal Impact & Limitations

D.1 Potential Societal Impact

VITAL connects visual motion to the metric precision needed to execute a driving maneuver. More accurate local trajectories could help a planner respect road boundaries and negotiate nearby traffic, while fewer flow evaluations could reduce the computation required for timely planning. The reported gains establish improvements in benchmark scores and measured inference latency. Their contribution to safer transportation depends on how the planner is integrated, validated, and monitored as part of an autonomous driving system that interacts with other road users.

Safety and distributional risks. A planner can improve average PDMS while making individual scenes less safe. As shown in Appendix C.3, a small residual correction can move a trajectory beyond the road boundary. Errors in visual interpretation, unusual road layouts, and interactions absent from the training distribution may therefore expose pedestrians, cyclists, and vehicle occupants to harm. Benchmark averages do not establish equal performance across locations, weather conditions, or road-user groups.

Deployment requires evaluation in the intended operating domain, tests of rare interactions, independent safety constraints, and a fallback for unreliable predictions.

Data and computational considerations. The experiments reuse released driving observations and annotations. These recordings may contain information about people and places, so access, visualization, and redistribution remain subject to the source datasets' privacy protections and use conditions. Public availability does not imply unrestricted reuse. Although one-step inference reduces measured latency, the study relies on substantial pretrained models and GPU training. Assessing environmental impact requires energy and emissions measurements across pretraining, tokenizer fitting, and planner optimization. These measurements are not reported in the present study.

D.2 Known Limitations

Simulation and transfer. The experiments assess planning within the NAVSIM benchmark. Generalization to unseen cities and robustness to sensor failures or prolonged interaction in real traffic remain open challenges. The stage-dependent behavior in Appendix B.2 motivates training that accounts for how surrounding traffic responds as the scene evolves beyond the initial observation.

Representation and control scope. The residual policy specializes in local refinement around a predicted maneuver. As shown in Appendix C.3, this scope leaves incorrect route choices and large visual prediction errors for future work. Broader control over the visual scaffold and explicit safety constraints could address these limitations, as bounded corrections can still cross road boundaries.

E Public Resources Used

We acknowledge the public resources supporting this work, including the datasets, pretrained model components, evaluation software, and tools used for analysis and visualization. The scenes in Figure 1 are taken from public videos[1].

E.1 Public Datasets Used

- nuPlan[2] ..... nuPlan Dataset Terms
- nuScenes[3] ..... CC BY-NC-SA 4.0
- OpenScene[4] ..... CC BY-NC-SA 4.0

E.2 Public Implementations Used

- NAVSIM v1/v2[5] ..... Apache License 2.0
- nuPlan-devkit[6] ..... Apache License 2.0
- VGGT-Ω[7] ..... FAIR Noncommercial Research License
- Depth Anything 3[8] ..... Apache License 2.0 for code
- LingBot-Map[9] ..... Apache License 2.0
- LingBot-VA / VA 2.0[10] ..... Apache License 2.0 for code
- StarVLA[11] ..... MIT with additional attribution terms
- SimWAM[12] ..... MIT License
- Wan2.2[13] ..... Apache License 2.0
- UMT5-XXL[14] ..... Apache License 2.0
- SUV[15] ..... MIT License
- DiffusionDrive[16] ..... MIT License
- DiffSynth-Studio[17] ..... Apache License 2.0
- PyTorch[18] ..... BSD 3-Clause License
- Transformers[19] ..... Apache License 2.0
- ROS 2[20] ..... Apache 2.0 / BSD 3-Clause
- NumPy[21] ..... BSD 3-Clause License

[1] https://www.youtube.com/watch?v=sUcWuiUpvF4&t=9083s
[2] https://www.nuscenes.org/nuplan
[3] https://www.nuscenes.org/nuscenes
[4] https://github.com/OpenDriveLab/OpenScene
[5] https://github.com/autonomousvision/navsim
[6] https://github.com/motional/nuplan-devkit
[7] https://github.com/facebookresearch/vggt-omega
[8] https://github.com/ByteDance-Seed/Depth-Anything-3
[9] https://github.com/Robbyant/lingbot-map
[10] https://github.com/Robbyant/lingbot-va
[11] https://github.com/starVLA/starVLA
[12] https://github.com/H-EmbodVis/SimWAM
[13] https://github.com/Wan-Video/Wan2.2
[14] https://huggingface.co/google/umt5-xxl
[15] https://github.com/ASH-2046/SUV
[16] https://github.com/hustvl/DiffusionDrive
[17] https://github.com/modelscope/DiffSynth-Studio
[18] https://github.com/pytorch/pytorch
[19] https://github.com/huggingface/transformers
[20] https://github.com/ros2/ros2
[21] https://github.com/numpy/numpy

References

Lars Ankile, Anthony Simeonov, Idan Shenfeld, Marcel Torne, and Pulkit Agrawal. From Imitation to Refinement – Residual RL for Precise Assembly. In *IEEE International Conference on Robotics and Automation*, pages 01–08, 2025. doi: 10.1109/icra55743.2025.11127442. URL https://arxiv.org/abs/2407.16677.

Mahmoud Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, et al. V-JEPA 2: Self-Supervised Video Models Enable Understanding, Prediction and Planning. *arXiv preprint arXiv:2506.09985*, 2025. URL https://arxiv.org/abs/2506.09985.

Homanga Bharadhwaj, Debidatta Dwibedi, Abhinav Gupta, Shubham Tulsiani, Carl Doersch, et al. Gen2Act: Human Video Generation in Novel Scenarios enables Generalizable Robot Manipulation. In *Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 3936–3951. PMLR, 2025. URL https://arxiv.org/abs/2409.16283.

Hengwei Bian, Lingdong Kong, Haozhe Xie, Liang Pan, Yu Qiao, et al. DynamicCity: Large-Scale 4D Occupancy Generation from Dynamic Scenes. In *International Conference on Learning Representations*, 2025. URL https://arxiv.org/abs/2410.18084.

Jake Bruce, Michael D Dennis, Ashley Edwards, Jack Parker-Holder, Yuge Shi, et al. Genie: Generative Interactive Environments. In *International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 4603–4623. PMLR, 2024. URL https://arxiv.org/abs/2402.15391.

Qingwen Bu, Yanting Yang, Jisong Cai, Shenyuan Gao, Guanghui Ren, et al. Learning to Act Anywhere with Task-centric Latent Actions. In *Robotics: Science and Systems*, 2025. doi: 10.15607/RSS.2025.XXI.014. URL https://arxiv.org/abs/2505.06111.

Wei Cao, Marcel Hallgarten, Tianyu Li, Daniel Dauner, Xunjiang Gu, et al. Pseudo-Simulation for Autonomous Driving. In *Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 4709–4722. PMLR, 2025. URL https://arxiv.org/abs/2506.04218.

Jun Cen, Chaohui Yu, Hangjie Yuan, Yuming Jiang, Siteng Huang, et al. WorldVLA: Towards Autoregressive Action World Model. *arXiv preprint arXiv:2506.21539*, 2025. URL https://arxiv.org/abs/2506.21539.

Haoran Chen, Jingtian Ji, Samuel Wheeler, Kaylene Caswell Stocking, and Matthew Walter. Beyond Reconstruction: What Matters in Action Tokenization for Robot Policies? *arXiv preprint arXiv:2610.09170*, 2026. URL https://arxiv.org/abs/2610.09170.

Runnan Chen, Youquan Liu, Lingdong Kong, Xinge Zhu, Yuexin Ma, et al. CLIP2Scene: Towards Label-Efficient 3D Scene Understanding by CLIP. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7020–7030, 2023. URL https://arxiv.org/abs/2301.04926.

Yi Chen, Yuying Ge, Weiliang Tang, Yizhuo Li, Yixiao Ge, et al. Moto: Latent Motion Token as the Bridging Language for Learning Robot Manipulation from Videos. In *IEEE/CVF International Conference on Computer Vision*, pages 19752–19763, 2025. URL https://arxiv.org/abs/2412.04445.

Cheng Chi, Zhenjia Xu, Siyuan Feng, Eric Cousineau, Yilun Du, et al. Diffusion policy: Visuomotor policy learning via action diffusion. *The International Journal of Robotics Research*, 44(10-11):1684–1704, 2025. doi: 10.1177/02783649241273668. URL https://arxiv.org/abs/2303.04137.

Kashyap Chitta, Aditya Prakash, Bernhard Jaeger, Zehao Yu, Katrin Renz, et al. TransFuser: Imitation With Transformer-Based Sensor Fusion for Autonomous Driving. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(11):12878–12895, 2023. doi: 10.1109/tpami.2022.3200245. URL https://arxiv.org/abs/2205.15997.

Meng Chu, Xuan Billy Zhang, Kevin Qinghong Lin, Lingdong Kong, Jize Zhang, et al. Agentic World Modeling: Foundations, Capabilities, Laws, and Beyond. *arXiv preprint arXiv:2604.22748*, 2026. URL https://arxiv.org/abs/2604.22748.

Chenxu Dang, Sining Ang, Yongkang Li, Haochen Tian, Jie Wang, et al. DriveFine: Refining-Augmented Masked Diffusion VLA for Precise and Robust Driving. *arXiv preprint arXiv:2602.14577*, 2026. URL https://arxiv.org/abs/2602.14577.

Daniel Dauner, Marcel Hallgarten, Tianyu Li, Xinshuo Weng, Zhiyu Huang, et al. NAVSIM: Data-Driven Non-Reactive Autonomous Vehicle Simulation and Benchmarking. In *Advances in Neural Information Processing Systems*, volume 37, pages 28706–28719, 2024. URL https://arxiv.org/abs/2406.15349.

Junbo Dong, Ze Chen, Zhendong Xie, Junjie Li, Lixin Xu, et al. Behavior-Aligned Action Tokenization for Robot Policy Learning. *arXiv preprint arXiv:2609.27513*, 2026a. URL https://arxiv.org/abs/2609.27513.

Yifei Dong, Fengyi Wu, Guangyu Chen, Lingdong Kong, Xu Zhu, et al. Towards Unified World Models for Visual Navigation via Memory-Augmented Planning and Foresight. In *European Conference on Computer Vision*, 2026b. URL https://arxiv.org/abs/2510.08713.

Yifei Dong, Fengyi Wu, Yilong Dai, Lingdong Kong, Guangyu Chen, et al. Language-Conditioned World

Modeling for Visual Navigation. In *Advances in Neural Information Processing Systems*, 2026c. URL https://arxiv.org/abs/2603.26741.

Lan Feng, Yang Gao, Eloi Zablocki, Quanyi Li, Wuyang Li, et al. RAP: 3D Rasterization Augmented End-to-End Planning. In *International Conference on Learning Representations*, 2026. URL https://arxiv.org/abs/2510.04333.

Kevin Frans, Danijar Hafner, Sergey Levine, and Pieter Abbeel. One Step Diffusion via Shortcut Models. In *International Conference on Learning Representations*, 2025. URL https://arxiv.org/abs/2410.12557.

Zhengyang Geng, Mingyang Deng, Xingjian Bai, J. Zico Kolter, and Kaiming He. Mean Flows for One-step Generative Modeling. In *Advances in Neural Information Processing Systems*, volume 38, pages 75460–75482, 2025. URL https://arxiv.org/abs/2505.13447.

Zeying Gong, Yangyi Zhong, Yiyi Ding, Tianshuai Hu, Guoyang Zhao, et al. FLUX: Accelerating Cross-Embodiment Generative Navigation Policies via Rectified Flow and Static-to-Dynamic Learning. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2026. URL https://arxiv.org/abs/2603.12806.

Yanjiang Guo, Yucheng Hu, Jianke Zhang, Yen-Jen Wang, Xiaoyu Chen, et al. Prediction with Action: Visual Policy Learning via Joint Denoising Process. In *Advances in Neural Information Processing Systems*, volume 37, pages 112386–112410, 2024. URL https://arxiv.org/abs/2411.18179.

Tianshuai Hu, Xiaolu Liu, Song Wang, Yiyao Zhu, Ao Liang, et al. Vision-Language-Action Models for Autonomous Driving: Past, Present, and Future. *arXiv preprint arXiv:2512.16760*, 2025a. URL https://arxiv.org/abs/2512.16760.

Yucheng Hu, Yanjiang Guo, Pengchao Wang, Xiaoyu Chen, Yen-Jen Wang, et al. Video Prediction Policy: A Generalist Robot Policy with Predictive Visual Representations. In *International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 24328–24346. PMLR, 2025b. URL https://arxiv.org/abs/2412.14803.

Feiyang Jia, Lin Liu, Ziying Song, Caiyan Jia, Hangjun Ye, et al. DriveWorld-VLA: Unified Latent-Space World Modeling with Vision-Language-Action for Autonomous Driving. In *International Conference on Machine Learning*, 2026. URL https://arxiv.org/abs/2602.06521.

Sunshine Jiang, Xiaolin Fang, Nicholas Roy, Tomás Lozano-Pérez, Leslie Pack Kaelbling, et al. Streaming Flow Policy: Simplifying diffusion/flow-matching policies by treating action trajectories as flow trajectories. In *Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 238–257. PMLR, 2025. URL https://arxiv.org/abs/2505.21851.

Changhui Jing, Haochen Tian, Hongchen Li, Tianyu Li, Caojun Wang, et al. SimVTA: Simulation-Grounded Video-Trajectory Alignment for Driving World Action Models. *arXiv preprint arXiv:2610.13781*, 2026. URL https://arxiv.org/abs/2610.13781.

Tobias Johannink, Shikhar Bahl, Ashvin Nair, Jianlan Luo, Avinash Kumar, et al. Residual Reinforcement Learning for Robot Control. In *IEEE International Conference on Robotics and Automation*, pages 6023–6029, 2019. doi: 10.1109/icra.2019.8794127. URL https://arxiv.org/abs/1812.03201.

Lingdong Kong, Youquan Liu, Runnan Chen, Yuexin Ma, Xinge Zhu, et al. Rethinking Range View Representation for LiDAR Segmentation. In *IEEE/CVF International Conference on Computer Vision*, pages 228–240, 2023a. URL https://arxiv.org/abs/2303.05367.

Lingdong Kong, Youquan Liu, Xin Li, Runnan Chen, Wenwei Zhang, et al. Robo3D: Towards Robust and Reliable 3D Perception against Corruptions. In *IEEE/CVF International Conference on Computer Vision*, pages 19994–20006, 2023b. URL https://arxiv.org/abs/2303.17597.

Lingdong Kong, Jiawei Ren, Liang Pan, and Ziwei Liu. LaserMix for Semi-Supervised LiDAR Semantic Segmentation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 21705–21715, 2023c. URL https://arxiv.org/abs/2207.00026.

Lingdong Kong, Xiang Xu, Jiawei Ren, Wenwei Zhang, Liang Pan, et al. Multi-Modal Data-Efficient 3D Scene Understanding for Autonomous Driving. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47 (5):3748–3765, 2025a. doi: 10.1109/tpami.2025.3535625. URL https://arxiv.org/abs/2405.05258.

Lingdong Kong, Yu Yang, Jianbiao Mei, Youquan Liu, Ao Liang, et al. 3D and 4D World Modeling: A Survey. *arXiv preprint arXiv:2509.07996*, 2025b. URL https://arxiv.org/abs/2509.07996.

Lingdong Kong, Ao Liang, Tianyi Yan, Hongsi Liu, Wesley Yang, et al. Is Your Driving World Model an All-Around Player? In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, 2026a. URL https://arxiv.org/abs/2605.10858. VideoWorldModel Workshop.

Lingdong Kong, Xiang Xu, Youquan Liu, Jun Cen, Runnan Chen, et al. LargeAD: Large-Scale Cross-Sensor Data Pretraining for Autonomous Driving. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 48(2): 1291–1308, 2026b. doi: 10.1109/tpami.2025.3617126. URL https://arxiv.org/abs/2501.04005.

Seungjae Lee, Yibin Wang, Haritheja Etukuru, H. Jin Kim, Nur Muhammad Mahi Shafiullah, et al. Behavior

Generation with Latent Actions. In *International Conference on Machine Learning*, volume 235, pages 26991–27008, 2024. URL https://arxiv.org/abs/2403.03181.

Jingyu Li, Zhe Liu, Dongnan Hu, Junjie Wu, Zipei Ma, et al. Metis: A Generalizable and Efficient World-Action Model for Autonomous Driving and Urban Navigation. *arXiv preprint arXiv:2606.15869*, 2026a. URL https://arxiv.org/abs/2606.15869.

Jingyu Li, Junjie Wu, Dongnan Hu, Xiangkai Huang, Bin Sun, et al. SGDrive: Scene-to-Goal Hierarchical World Cognition for Autonomous Driving. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4032–4042, 2026b. URL https://arxiv.org/abs/2601.05640.

Kailin Li, Zhenxin Li, Shiyi Lan, Yuan Xie, Zhizhong Zhang, et al. Hydra-MDP++: Advancing End-to-End Driving via Expert-Guided Hydra-Distillation. *arXiv preprint arXiv:2503.12820*, 2025a. URL https://arxiv.org/abs/2503.12820.

Rong Li, Yuhao Dong, Tianshuai Hu, Alan Liang, Youquan Liu, et al. 3EED: Ground Everything Everywhere in 3D. In *Advances in Neural Information Processing Systems*, 2025b. URL https://arxiv.org/abs/2511.01755.

Yingyan Li, Lue Fan, Jiawei He, Yuqi Wang, Yuntao Chen, et al. Enhancing End-to-End Autonomous Driving with Latent World Model. In *International Conference on Learning Representations*, 2025c. URL https://arxiv.org/abs/2406.08481.

Yingyan Li, Yuqi Wang, Yang Liu, Jiawei He, Lue Fan, et al. End-to-End Driving with Online Trajectory Evaluation via BEV World Model. In *IEEE/CVF International Conference on Computer Vision*, pages 27137–27146, 2025d. URL https://arxiv.org/abs/2504.01941.

Yingyan Li, Shuyao Shang, Weisong Liu, Bing Zhan, Haochen Wang, et al. DriveVLA-W0: World Models Amplify Data Scaling Law in Autonomous Driving. In *International Conference on Learning Representations*, 2026c. URL https://arxiv.org/abs/2510.12796.

Yongkang Li, Kaixin Xiong, Xiangyu Guo, Fang Li, Sixu Yan, et al. ReCogDrive: A Reinforced Cognitive Framework for End-to-End Autonomous Driving. In *International Conference on Learning Representations*, 2026d. URL https://arxiv.org/abs/2506.08052.

Alan Liang, Youquan Liu, Yu Yang, Dongyue Lu, Linfeng Li, et al. LiDARCrafter: Dynamic 4D World Modeling from LiDAR Sequences. In *AAAI Conference on Artificial Intelligence*, volume 40, pages 18406–18414, 2026a. doi: 10.1609/aaai.v40i22.38905. URL https://arxiv.org/abs/2508.03692.

Ao Liang, Lingdong Kong, Dongyue Lu, Youquan Liu, Jian Fang, et al. Perspective-Invariant 3D Object Detection. In *IEEE/CVF International Conference on Computer Vision*, pages 27725–27738, 2025. URL https://arxiv.org/abs/2507.17665.

Ao Liang, Lingdong Kong, Tianyi Yan, Hongsi Liu, Wesley Yang, et al. WorldLens: Full-Spectrum Evaluations of Driving World Models in Real World. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026b. URL https://arxiv.org/abs/2512.10958.

Bencheng Liao, Shaoyu Chen, Haoran Yin, Bo Jiang, Cheng Wang, et al. DiffusionDrive: Truncated Diffusion Model for End-to-End Autonomous Driving. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 12037–12047, 2025. URL https://arxiv.org/abs/2411.15139.

Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matthew Le. Flow Matching for Generative Modeling. In *International Conference on Learning Representations*, 2023. URL https://arxiv.org/abs/2210.02747.

Chaoqi Liu, Xiaoshen Han, Jiawei Gao, Yue Zhao, Haonan Chen, et al. OAT: Ordered Action Tokenization. In *Robotics: Science and Systems*, 2026a. doi: 10.15607/RSS.2026.XXII.075. URL https://arxiv.org/abs/2602.04215.

Lin Liu, Caiyan Jia, Guanyi Yu, Ziying Song, Junqiao Li, et al. GuideFlow: Constraint-Guided Flow Matching for Planning in End-to-End Autonomous Driving. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3719–3728, 2026b. URL https://arxiv.org/abs/2511.18729.

Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow Straight and Fast: Learning to Generate and Transfer Data with Rectified Flow. In *International Conference on Learning Representations*, 2023a. URL https://arxiv.org/abs/2209.03003.

Youquan Liu, Lingdong Kong, Jun Cen, Runnan Chen, Wenwei Zhang, et al. Segment Any Point Cloud Sequences by Distilling Vision Foundation Models. In *Advances in Neural Information Processing Systems*, volume 36, pages 37193–37229, 2023b. URL https://arxiv.org/abs/2306.09347.

Youquan Liu, Weidong Yang, Ao Liang, Xiang Xu, Lingdong Kong, et al. OmniLiDAR: A Unified Diffusion Framework for Multi-Domain 3D LiDAR Generation. *arXiv preprint arXiv:2605.13815*, 2026c. URL https://arxiv.org/abs/2605.13815.

Yushan Liu, Tianxiong Lv, Bohua Wang, Hangqi Fan, Chenxu Zhao, et al. PerceptDrive: Perception Prior World-Action Modeling with Adaptive Expert Routing for End-to-End Autonomous Driving. *arXiv preprint arXiv:2607.20175*, 2026d. URL https://arxiv.org/abs/2607.20175.

Zhuoming Liu, Jinhong Lin, Kwan Man Cheng, Lin Zhang, Shayok Bagchi, et al. DRIFT: A Residual Flow Adapter for Decoding Continuous Outputs in Vision-Language Models. *arXiv preprint arXiv:2606.05758*, 2026e. URL https://arxiv.org/abs/2606.05758.

Jinghui Lu, Jiayi Guan, Zhijian Huang, Jinlong Li, Guang Li, et al. OneVL: One-Step Latent Reasoning and Planning with Vision-Language Explanation. In *Advances in Neural Information Processing Systems*, 2026a. URL https://arxiv.org/abs/2604.18486.

Yiren Lu, Xin Ye, Jiaming Liu, Philip Jacobson, Jin Yao, et al. GeoWAM: Visual Geometry World Action Models for Autonomous Driving. *arXiv preprint arXiv:2608.23486*, 2026b. URL https://arxiv.org/abs/2608.23486.

Xiangkai Ma, Lekai Xing, Han Zhang, Wenzhong Li, and Sanglu Lu. Unifying Perception and Action: A Hybrid-Modality Pipeline with Implicit Visual Chain-of-Thought for Robotic Action Generation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22380–22390, 2026. URL https://arxiv.org/abs/2511.19859.

Jiahao Meng, Xiangtai Li, Haochen Wang, Yue Tan, Tao Zhang, et al. Open-o3-Video: Grounded Video Reasoning with Explicit Spatio-Temporal Evidence. In *International Conference on Machine Learning*, 2026. URL https://arxiv.org/abs/2510.20579.

Atharva Mete, Haotian Xue, Albert Wilcox, Yongxin Chen, and Animesh Garg. QueST: Self-Supervised Skill Abstractions for Learning Continuous Control. In *Advances in Neural Information Processing Systems*, volume 37, pages 4062–4089, 2024. doi: 10.52202/079017--0133. URL https://arxiv.org/abs/2407.15840.

Jingcheng Ni, Yuxin Guo, Yichen Liu, Rui Chen, Lewei Lu, et al. MaskGWM: A Generalizable Driving World Model with Video Mask Reconstruction. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22381–22391, 2025. URL https://arxiv.org/abs/2502.11663.

Alexander Nikulin, Ilya Zisman, Denis Tarasov, Nikita Lyubaykin, Andrei Polubarov, et al. Latent Action Learning Requires Supervision in the Presence of Distractors. In *International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 46427–46447. PMLR, 2025. URL https://arxiv.org/abs/2502.00379.

Seohong Park, Qiyang Li, and Sergey Levine. Flow Q-Learning. In *International Conference on Machine Learning*, volume 267, pages 48104–48127, 2025. URL https://arxiv.org/abs/2502.02538.

Karl Pertsch, Kyle Stachowicz, Brian Ichter, Danny Driess, Suraj Nair, et al. FAST: Efficient Action Tokenization for Vision-Language-Action Models. In *Robotics: Science and Systems*, 2025. URL https://arxiv.org/abs/2501.09747.

Aaditya Prasad, Kevin Lin, Jimmy Wu, Linqi Zhou, and Jeannette Bohg. Consistency Policy: Accelerated Visuomotor Policies via Consistency Distillation. In *Robotics: Science and Systems*, 2024. URL https://arxiv.org/abs/2405.07503.

Chenghao Qian, Nedko Savov, Lingdong Kong, Yeying Jin, Rui Song, et al. Semantic-Aware, Physics-Informed, Geometry-Grounded Weather Synthesis. In *European Conference on Computer Vision*, 2026. URL https://arxiv.org/abs/2606.29020.

Allen Z. Ren, Justin Lidard, Lars Lien Ankile, Anthony Simeonov, Pulkit Agrawal, et al. Diffusion Policy Policy Optimization. In *International Conference on Learning Representations*, 2025. URL https://arxiv.org/abs/2409.00588.

Yichao Shen, Fangyun Wei, Zhiying Du, Yaobo Liang, Yan Lu, et al. VideoVLA: Video Generators Can Be Generalizable Robot Manipulators. In *Advances in Neural Information Processing Systems*, volume 38, pages 95597–95621, 2025. doi: 10.52202/085713--3197. URL https://arxiv.org/abs/2512.06963.

Andreas Sochopoulos, Nikolay Malkin, Nikolaos Tsagkas, Joao Moura, Michael Gienger, et al. Fast Flow-based Visuomotor Policies via Conditional Optimal Transport Couplings. In *Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 3357–3377. PMLR, 2025. URL https://arxiv.org/abs/2505.01179.

Yang Song, Prafulla Dhariwal, Mark Chen, and Ilya Sutskever. Consistency Models. In *International Conference on Machine Learning*, volume 202, pages 32211–32252, 2023. URL https://arxiv.org/abs/2303.01469.

Yue Su, Sijin Chen, Haixin Shi, Mingyu Liu, Zhengshen Zhang, et al. World Guidance: World Modeling in Condition Space for Action Generation. In *International Conference on Machine Learning*, 2026. URL https://arxiv.org/abs/2602.22010.

Wenchao Sun, Xuewu Lin, Keyu Chen, Zixiang Pei, Xiang Li, et al. SparseDriveV2: Scoring is All You Need for End-to-End Autonomous Driving. *arXiv preprint arXiv:2603.29163*, 2026. URL https://arxiv.org/abs/2603.29163.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, et al. Attention is All you Need. In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL https://arxiv.org/abs/1706.03762.

Andrew Wagenmaker, Yunchu Zhang, Mitsuhiko Nakamoto, Seohong Park, Waleed Yagoub, et al. Steering Your Diffusion Policy with Latent Space Reinforcement Learning. In *Conference on Robot Learning*, volume 305, pages 258–282, 2025. URL https://arxiv.org/abs/2506.15799.

Jianyuan Wang, Minghao Chen, Shangzhan Zhang, Nikita Karaev, Johannes Schönberger, et al. VGGT-Ω. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 21486–21499, 2026a. URL https://arxiv.org/abs/2605.15195.

Junke Wang, Qihang Zhang, Shuai Yang, Yiming Luo, Yujun Shen, et al. RepWAM: World Action Modeling with Representation Visual-Action Tokenizers. *arXiv preprint arXiv:2606.13674*, 2026b. URL https://arxiv.org/abs/2606.13674.

Junli Wang, Yinan Zheng, Xueyi Liu, Zebin Xing, Pengfei Li, et al. MeanFuser: Fast One-Step Multi-Modal Trajectory Generation and Adaptive Reconstruction via MeanFlow for End-to-End Autonomous Driving. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17884–17893, 2026c. URL https://arxiv.org/abs/2602.20060.

Linbo Wang, Yupeng Zheng, Qiang Chen, Shiwei Li, Yichen Zhang, et al. Latent-WAM: Latent World Action Modeling for End-to-End Autonomous Driving. *arXiv preprint arXiv:2603.24581*, 2026d. URL https://arxiv.org/abs/2603.24581.

Linhan Wang, Zichong Yang, Chen Bai, Guoxiang Zhang, Xiaotong Liu, et al. Drive-JEPA: Video JEPA Meets Multimodal Trajectory Distillation for End-to-End Driving. *arXiv preprint arXiv:2601.22032*, 2026e. URL https://arxiv.org/abs/2601.22032.

Song Wang, Xiaolu Liu, Lingdong Kong, Jianyun Xu, Chunyong Hu, et al. PointLoRA: Low-Rank Adaptation with Token Selection for Point Cloud Learning. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6605–6615, 2025a. URL https://arxiv.org/abs/2504.16023.

Yating Wang, Haoyi Zhu, Mingyu Liu, Jiange Yang, Hao-Shu Fang, et al. VQ-VLA: Improving Vision-Language-Action Models via Scaling Vector-Quantized Action Tokenizers. In *IEEE/CVF International Conference on Computer Vision*, pages 11089–11099, 2025b. URL https://arxiv.org/abs/2507.01016.

Yuqi Wang, Jiawei He, Lue Fan, Hongxin Li, Yuntao Chen, et al. Driving into the Future: Multiview Visual Forecasting and Planning with World Model for Autonomous Driving. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14749–14759, 2024. URL https://arxiv.org/abs/2311.17918.

Zhendong Wang, Max Li, Ajay Mandlekar, Zhenjia Xu, Jiaojiao Fan, et al. One-Step Diffusion Policy: Fast Visuomotor Policies via Diffusion Distillation. In *International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 63399–63416. PMLR, 2025c. URL https://arxiv.org/abs/2410.21257.

Tianze Xia, Yongkang Li, Lijun Zhou, Jingfeng Yao, Kaixin Xiong, et al. DriveLaW: Unifying Planning and Video Generation in a Latent Driving World. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 39701–39712, 2026a. URL https://arxiv.org/abs/2512.23421.

Tianze Xia, Lijun Zhou, Kaixin Xiong, Jingfeng Yao, Zhenxin Zhu, et al. ReWorld: Representation Learning for World Action Models. *arXiv preprint arXiv:2606.27504*, 2026b. URL https://arxiv.org/abs/2606.27504.

Shaoyuan Xie, Lingdong Kong, Yuhao Dong, Chonghao Sima, Wenwei Zhang, et al. Are VLMs Ready for Autonomous Driving? An Empirical Study from the Reliability, Data and Metric Perspectives. In *IEEE/CVF International Conference on Computer Vision*, pages 6585–6597, 2025a. URL https://arxiv.org/abs/2501.04003.

Shaoyuan Xie, Lingdong Kong, Wenwei Zhang, Jiawei Ren, Liang Pan, et al. Benchmarking and Improving Bird's Eye View Perception Robustness in Autonomous Driving. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(5):3878–3894, 2025b. doi: 10.1109/TPAMI.2025.3535960. URL https://arxiv.org/abs/2405.17426.

Yining Xing, Zhiyuan Liu, Zehong Ke, Wenhao Yu, and Jianqiang Wang. DRIFT: Drift and Aggregation for Motion Planning. *arXiv preprint arXiv:2607.14507*, 2026. URL https://arxiv.org/abs/2607.14507.

Jiawei Xu, Zhizhou Zhong, Zhijian Shu, Mingkai Jia, Mingxiao Li, et al. EponaV2: Driving World Model with Comprehensive Future Reasoning. *arXiv preprint arXiv:2605.14696*, 2026a. URL https://arxiv.org/abs/2605.14696.

Xiang Xu, Alan Liang, Youquan Liu, Linfeng Li, Lingdong Kong, et al. U4D: Uncertainty-Aware 4D World Modeling from LiDAR Sequences. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10027–10039, 2026b. URL https://arxiv.org/abs/2512.02982.

Tianyi Yan, Tao Tang, Xingtai Gui, Yongkang Li, Jiasen Zheng, et al. AD-R1: Closed-Loop Reinforcement Learning for End-to-End Autonomous Driving with Impartial World Models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Findings*, 2026. URL https://arxiv.org/abs/2511.20325.

Jiange Yang, Yansong Shi, Haoyi Zhu, Mingyu Liu, Kaijing Ma, et al. CoMo: Learning Continuous Latent Motion from Internet Videos for Scalable Robot Learning. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 42352–42363, 2026a. URL https://arxiv.org/abs/2505.17006.

Yu Yang, Xuemeng Yang, Licheng Wen, Lingdong Kong, Xiaobin Hu, et al. Quo Vadis, World Modeling? *arXiv preprint arXiv:2608.02713*, 2026b. URL https://arxiv.org/abs/2608.02713.

Ruoyu Yao, Yusen Xie, Qingzhao Liu, Pei Liu, Zewei Yang, et al. Continuous Actions from Discrete Minds: Latent-Aligned Planning for End-to-End Autonomous Driving. *arXiv preprint arXiv:2609.04070*, 2026a. URL https://arxiv.org/abs/2609.04070.

Wenhao Yao, Zhenxin Li, Shiyi Lan, Zi Wang, Xinglong Sun, et al. DriveSuprim: Towards Precise Trajectory Selection for End-to-End Planning. In *AAAI Conference on Artificial Intelligence*, volume 40, pages 11910–11918, 2026b. doi: 10.1609/aaai.v40i14.38178. URL https://arxiv.org/abs/2506.06659.

Seonghyeon Ye, Joel Jang, Byeongguk Jeon, Se June Joo, Jianwei Yang, et al. Latent Action Pretraining from Videos. In *International Conference on Learning Representations*, 2025. URL https://arxiv.org/abs/2410.11758.

Seonghyeon Ye, Yunhao Ge, Kaiyuan Zheng, Shenyuan Gao, Sihyun Yu, et al. World Action Models are Zero-shot Policies. *arXiv preprint arXiv:2602.15922*, 2026. URL https://arxiv.org/abs/2602.15922.

Sihyun Yu, Sangkyung Kwak, Huiwon Jang, Jongheon Jeong, Jonathan Huang, et al. Representation Alignment for Generation: Training Diffusion Transformers Is Easier Than You Think. In *International Conference on Learning Representations*, 2025. URL https://arxiv.org/abs/2410.06940.

Xiu Yuan, Tongzhou Mu, Stone Tao, Yunhao Fang, Mengke Zhang, et al. Policy Decorator: Model-Agnostic Online Refinement for Large Policy Model. In *International Conference on Learning Representations*, 2025. URL https://arxiv.org/abs/2412.13630.

Yibo Yuan, Jiacheng Fu, Jiangtong Zhu, Yi Li, Jianhua Han, et al. SUV: Future Scene Understanding as Video Generation for End-to-End Driving. *arXiv preprint arXiv:2608.03084*, 2026. URL https://arxiv.org/abs/2608.03084.

Chuheng Zhang, Tim Pearce, Pushi Zhang, Kaixin Wang, Xiaoyu Chen, et al. What Do Latent Action Models Actually Learn? In *Advances in Neural Information Processing Systems*, volume 38, pages 146676–146697, 2025a. URL https://arxiv.org/abs/2506.15691.

Qihang Zhang, Lin Li, Luyao Zhang, Shuai Yang, Yiming Luo, et al. Native Video-Action Pretraining for Generalizable Robot Control. *arXiv preprint arXiv:2607.08639*, 2026a. URL https://arxiv.org/abs/2607.08639.

Qinglun Zhang, Zhen Liu, Haoqiang Fan, Guanghui Liu, Bing Zeng, et al. FlowPolicy: Enabling Fast and Robust 3D Flow-Based Policy via Consistency Flow Matching for Robot Manipulation. In *AAAI Conference on Artificial Intelligence*, pages 14754–14762, 2025b. doi: 10.1609/aaai.v39i14.33617. URL https://arxiv.org/abs/2412.04987.

Tonghe Zhang, Chao Yu, Sichang Su, and Yu Wang. ReinFlow: Fine-tuning Flow Matching Policy with Online Reinforcement Learning. In *Advances in Neural Information Processing Systems*, volume 38, pages 106282–106319, 2025c. doi: 10.52202/085713--3547. URL https://arxiv.org/abs/2505.22094.

Wenyao Zhang, Hongsi Liu, Zekun Qi, Yunnan Wang, XinQiang Yu, et al. DreamVLA: A Vision-Language-Action Model Dreamed with Comprehensive World Knowledge. In *Advances in Neural Information Processing Systems*, volume 38, pages 24195–24228, 2025d. doi: 10.52202/085713--0818. URL https://arxiv.org/abs/2507.04447.

Xinyu Zhang, Zhengtong Xu, Yutian Tao, Yeping Wang, Yu She, et al. Learning Visual Feature-Based World Models via Residual Latent Action. In *Advances in Neural Information Processing Systems*, 2026b. URL https://arxiv.org/abs/2605.07079.

Zongchuang Zhao, Xin Zhou, Tianyang Xu, Zhengyang Sun, Kaixuan Zhou, et al. SimWAM: A Simple World Action Model for End-to-End Autonomous Driving. *arXiv preprint arXiv:2608.07468*, 2026. URL https://arxiv.org/abs/2608.07468.

Yupeng Zheng, Pengxuan Yang, Zebin Xing, Qichao Zhang, Yuhang Zheng, et al. World4Drive: End-to-End Autonomous Driving via Intention-aware Physical Latent World Model. In *IEEE/CVF International Conference on Computer Vision*, pages 28632–28642, 2025. URL https://arxiv.org/abs/2507.00603.

Ruiguo Zhong, Benshan Ma, Xiaolong Chen, Lang Zhang, Mingyue Feng, et al. DA-WAM: Decision-Aligned Future Latents for Driving World Models. *arXiv preprint arXiv:2608.19085*, 2026. URL https://arxiv.org/abs/2608.19085.

Wei Zhou, Xiongwei Zhu, Lingdong Kong, Bo Chen, Lei Zhang, et al. On-Policy Self-Distillation in Diffusion Models. *arXiv preprint arXiv:2608.24646*, 2026a. URL https://arxiv.org/abs/2608.24646.

Wei Zhou, Xiongwei Zhu, Zelin Xu, Bo Dong, Lixue Gong, et al. DanceOPD: On-Policy Generative Field Distillation. *arXiv preprint arXiv:2606.27377*, 2026b. URL https://arxiv.org/abs/2606.27377.

Chuning Zhu, Raymond Yu, Siyuan Feng, Benjamin Burchfiel, Paarth Shah, et al. Unified World Models: Coupling Video and Action Diffusion for Pretraining on Large Robotic Datasets. In *Robotics: Science and Systems*, 2025. doi: 10.15607/RSS.2025.XXI.015. URL https://arxiv.org/abs/2504.02792.

Jialv Zou, Shaoyu Chen, Bencheng Liao, Zhiyu Zheng, Yuehao Song, et al. DiffusionDriveV2: Reinforcement Learning-Constrained Truncated Diffusion Modeling in End-to-End Autonomous Driving. *arXiv preprint arXiv:2512.07745*, 2025. URL https://arxiv.org/abs/2512.07745.