Through a PRISM, Clearly: Step-Sequence Reasoning for Mistake Recognition in Procedural Videos

Olga Loginova
University of Trento
Trento, Italy
olga.loginova@unitn.it

Oleksandr Bezrukov
Gran Sasso Science Institute
L'Aquila, Italy
oleksandr.bezrukov@gssi.it

Frank Keller
University of Edinburgh
Edinburgh, United Kingdom
keller@inf.ed.ac.uk

arXiv:2610.31117v1 [cs.CV] 22 Oct 2026

Abstract

Procedural mistake recognition requires identifying deviations from valid multi-step activities. We study post-completion recognition from videos with benchmark-provided step boundaries. We introduce PRISM (Procedural mistake Recognition via Interpretable Semantic Modeling), which grounds video segments in language and maps their descriptions to procedural states using a lightweight hidden Markov model trained on correct executions. Reference alignment and semantic checks identify errors in step content and order. To support evaluation across heterogeneous benchmarks, we map dataset-specific annotations to four edit-based error types: insertion, deletion, substitution, and transposition. Across three benchmarks, we compare gold and VLM-generated descriptions in raw, semantic-role, and action–object forms. Decoder controls show consistent gains from the trained local decoder over nearest-candidate matching, with additional, dataset-dependent benefits from sequence decoding. Semantic role labeling improves Any-F1 in all six matched grounded comparisons, although the gains vary across datasets and do not translate into uniform improvements under separate configuration selection. These results establish language grounding as an effective interface for explicit procedural mistake reasoning and characterize the contributions of representation and decoding.

1 Introduction

A visually flawless plate of zoodles reaches the table, only for the first bite to reveal overboiled vegetables. The procedure appears finished before the failure becomes clear.

Some procedural errors become apparent only after task completion (Byrne and Davis, 2006; Reason, 2000). For AI systems monitoring or assisting human activity, the key question is retrospective: given the full video of an apparently completed task and a reference procedure, did the execution succeed, and if not, where did it go wrong? We focus on post-completion mistake recognition, which identifies errors from the complete recorded execution. This setting matters for troubleshooting, quality assurance, training, and human–AI assistance in real procedural environments (Norman, 2002; Garfield and Franklin, 2016).

Despite progress in temporal action segmentation and recognition (Ding et al., 2023b; Ding and Wang, 2025; Lai et al., 2024), procedural mistake detection remains underexplored (Mazzamuto et al., 2025; Bacharidis and Argyros, 2025). Existing approaches are fragmented: they are often vision-first and use different error labels, annotation units, and evaluation protocols (Peddi et al., 2024; Haneji et al., 2025; Lee and Elhamifar, 2025). Recent work on procedural error taxonomies also highlights the need for shared categories across datasets (Loginova and Keller, 2026). These differences complicate comparisons across datasets and obscure the contributions of grounding, representation, and procedural reasoning.

Following the procedural video taxonomy of Bacharidis and Argyros (2025), we distinguish temporal localization from logical error identification after action completion. Temporal segmentation and action recognition determine what happened and when. We study the subsequent reasoning problem: given step-level evidence and a reference procedure, does the completed execution deviate from the protocol, and where?

We evaluate this reasoning problem on EgoPER (Lee et al., 2024), CaptainCook4D (Peddi et al., 2024), and EgoOops (Haneji et al., 2025). Related systems include EgoPED (Lee et al., 2024),

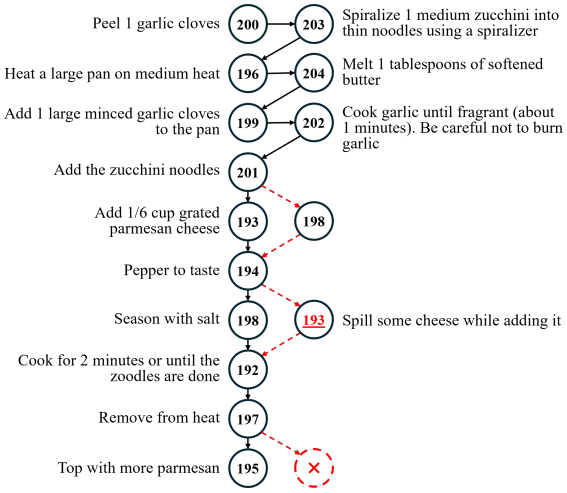

Figure 1: PRISM representation of the valid Zoodles procedure and an erroneous execution from CaptainCook4D. Circles contain Step IDs, adjacent text gives step descriptions, and solid arrows trace the reference order. Red marks show a transposition of Steps 193, 194, and 198, a substitution at Step 193, and the deletion of Step 195.

AMNAR (Huang et al., 2025), and AEM (Guo et al., 2026). Their temporal input and evaluation protocols are discussed in Section 4. We model this reasoning problem over discrete step semantics: procedural errors depend on what step was performed, with which arguments, and in what order.

This view motivates our language-grounded framework, because these constructs are naturally expressed in language. We propose a *video → language → reasoning* pipeline.

Like an optical prism refracting raw light into its constituent colors, our PRISM (Procedural mistake Recognition via Interpretable Semantic Modeling) makes procedural structure explicit through steps, semantic roles, and their order.[1] Benchmark-provided segment boundaries isolate mistake recognition from temporal segmentation. For each segment, a vision-language model (VLM) receives the clip and the task's candidate descriptions as context, then generates a free-form step description. PRISM encodes this description as raw text or as a semantic role representation produced through semantic role labeling (SRL). A hidden Markov model (HMM) maps the representation to a task-local step identifier (Step ID). Edit alignment identifies deviations from a valid reference sequence. We reconcile the benchmark-specific annotation schemes through a common taxonomy of four edit operations: insertion, deletion, substitution, and transposition. Figure 1 contrasts a valid procedure with an erroneous execution under this taxonomy. Circles denote HMM states labeled by Step IDs. The adjacent step descriptions are encoded as observations, and the ordered state sequence represents the procedure.

Our contributions are:

- **PRISM, a language-grounded framework** that combines a lightweight HMM trained on correct executions with reference alignment and semantic checks to identify procedural mistakes.
- **Cross-benchmark error analysis** through a shared edit representation of procedural deviations.
- **A controlled study of step representations and decoding**, comparing gold and VLM-generated descriptions, matched Raw–SRL configurations, action–object abstractions, and independent versus transition-aware inference.

2 Related work

Mistake-aware procedural video benchmarks. Despite progress in procedural video understanding, only a small subset of benchmarks explicitly annotates mistakes. Current benchmarks span cooking (Peddi et al., 2024; Lee et al., 2024), assembly (Sener et al., 2022; Schoonbeek et al., 2023; Jang et al., 2019; Wang et al., 2023b), laboratory tasks (Qian et al., 2022), and mixed scenarios (Haneji et al., 2025) but suffer from heterogeneous error definitions, annotation levels (video vs. step), and description styles (free-form vs. labels). Loginova and Keller (2026) similarly argue that procedural video datasets need shared categories for mistakes and corrections across tasks and domains. Our experiments span CaptainCook4D, EgoPER, and EgoOops: CaptainCook4D contains 384 videos across 24 cooking tasks; EgoPER concentrates 386 videos in five cooking tasks; and EgoOops adds 50 videos from five mixed-domain procedures. This combination tests VLM grounding and PRISM under broad task coverage, dense repetition within tasks, and limited data outside cooking.

Mistake recognition in procedural videos. Prior approaches are predominantly vision-centric,

[1] Code will be released upon acceptance.

modeling deviations from video features via recognition, anticipation, action-effect modeling, or structured prediction (Lee et al., 2024; Peddi et al., 2024; Lee and Elhamifar, 2025; Mazzamuto et al., 2025; Huang et al., 2025; Guo et al., 2026). Recent work also explores interpretability in error reasoning (Storks et al., 2025). Some work focuses on online detection and intervention forecasting using modified datasets (Flaborea et al., 2024; Plini et al., 2024; Wang et al., 2023b; Seminara et al., 2025). We target post-completion recognition, similarly to Huang et al. (2025); Narasimhan et al. (2023); Moriwaki et al. (2022); Ding et al. (2023a); Haneji et al. (2025); Li et al. (2025). We perform core reasoning in the language space by grounding video executions into textual step sequences.

Semantic roles and structured step representations. Semantic role labeling provides a structured interface for events through predicate–argument structures (Baker et al., 1998; Gildea and Jurafsky, 2002; Palmer et al., 2005; Chen et al., 2025). In multimodal settings, semantic structure supports video understanding (Sadhu et al., 2021), procedural instruction tracking (Batra et al., 2025), and fine-grained mistake attribution (Li et al., 2025).

We use SRLs as an intermediate normalization layer to bridge free-form step text and mistake reasoning.

3 Methodology

Figure 2 summarizes PRISM. Intuitively, our HMM captures two complementary aspects of a correct procedure: what a valid step looks like (emission) and how steps logically progress (transition), learning both exclusively from successful executions.

Each task provides candidate step descriptions, valid reference sequences, and correct training executions. Candidate steps define the task-local state inventory. We encode each segment using the representations in Section 3.4 and estimate HMM parameters from correct executions. At inference, the HMM decodes the executed step sequence. Reference alignment identifies structural edits, while emission, transition, and semantic similarity scores identify atypical observations and transitions. Appendix D reports the configuration-specific settings.

The grounding interface, SRL schema, text-encoding stage, HMM estimator and decoder, and edit-alignment procedure remain unchanged across tasks. Each procedure supplies its task-local candidate inventory and reference sequences, while its HMM parameters are estimated from correct executions.

3.1 HMM training

Each video is represented as a sequence of pre-segmented step clips. Section 3.4 defines the textual, SRL, action–object, visual, and multimodal representations used for each clip. For a video with T clips, let

$$\boldsymbol{x}_{1:T} = (\boldsymbol{x}_1, \ldots, \boldsymbol{x}_T), \qquad \boldsymbol{x}_t \in \mathbb{R}^{D_0}.$$

When specified by the configuration, we L2-normalize the embeddings and add Gaussian jitter to a random subset of training embeddings. We fit principal component analysis (PCA) on the training embeddings and obtain

$$\boldsymbol{z}_t = \mathrm{PCA}(\boldsymbol{x}_t) \in \mathbb{R}^D.$$

For task a, the hidden-state set $\mathcal{S}_a = \{1, \ldots, K_a\}$ contains its Step IDs. The projected embeddings $\boldsymbol{z}_{1:T}$ are the observations. The model parameters are

$$\theta_a = (\boldsymbol{\pi}, \boldsymbol{A}, \{\phi_s\}_{s \in \mathcal{S}_a}),$$

where $\boldsymbol{\pi}$ is the initial-state distribution, $\boldsymbol{A}$ is the transition matrix, and ϕ_s contains the emission parameters for state s. For Gaussian emissions, $\phi_s = (\boldsymbol{\mu}_s, \boldsymbol{\Sigma}_s)$. We estimate $\boldsymbol{\mu}_s$ from the state observations and $\boldsymbol{\Sigma}_s$ with Ledoit–Wolf shrinkage. For cosine emissions, $\boldsymbol{\mu}_s$ serves as the state prototype.

For the transition model, let $c_{ss'}$ denote the number of times the transition $s \to s'$ is observed in the training set. We then apply Laplace smoothing with parameter β:

$$A_{ss'} = \frac{c_{ss'} + \beta}{\sum_{k=1}^{K_a}(c_{sk} + \beta)}. \tag{1}$$

After Laplace smoothing, we apply local transition constraints based on the valid reference sequences by suppressing out-of-band transitions and renormalizing each row. We estimate the initial state distribution in the same way. Let m_s be the number of correct training videos whose first step is s. We then define $\pi_s = \frac{m_s+\beta}{\sum_{j=1}^{K_a}(m_j+\beta)}$. Thus, both the start distribution and the transition matrix are smoothed empirical estimates from correct executions.

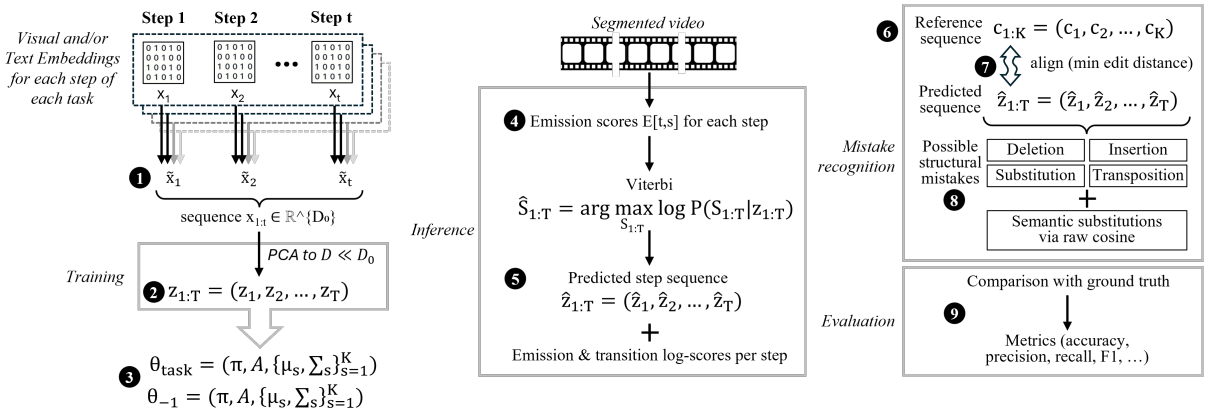

Figure 2: The PRISM pipeline for post-completion mistake recognition. (1) For each step of each training video, we obtain a textual, visual, or multimodal encoding, $\boldsymbol{x}_t$. (2) After PCA regularization, we obtain the representation of each step, $\boldsymbol{z}_t$. (3) During training, the model θ_{task} learns from sequences of $\boldsymbol{z}_t$ for each task. Additionally, when we put all tasks together, we obtain a global HMM θ_{-1}. (4) At inference time, we obtain embeddings for each video segment in the same way as during training and compute the emissions $[\boldsymbol{E}]_{t,s}$. (5) We then run Viterbi for sequence inference and keep the emission and transition scores for each step. (6) For each task, we also have a set of possible non-erroneous sequences. (7) We align the predicted sequence to the reference sequence using the selected sequence-alignment method. (8) Mismatches in the alignment indicate possible structural errors. Falling below the predefined thresholds for emission and transition scores indicates possible low-probability events. (9) During evaluation, we compare the predicted errors with the gold-standard error annotations of the dataset.

During training we record calibration statistics (mean, quantiles) of emission log-likelihoods and transition log-probabilities on correct videos. Let $\boldsymbol{\delta}_{t,s} = \boldsymbol{z}_t - \boldsymbol{\mu}_s$. The log emission likelihood is

$$\log p(\boldsymbol{z}_t \mid S_t = s) = -\frac{1}{2}\Big[\boldsymbol{\delta}_{t,s}^\top \boldsymbol{\Sigma}_s^{-1} \boldsymbol{\delta}_{t,s} + \log|\boldsymbol{\Sigma}_s| + D\log(2\pi)\Big]. \tag{2}$$

These standard log-likelihood statistics are stored for analysis and threshold tuning.

The term $\log|\boldsymbol{\Sigma}_s|$ penalizes the volume of the Gaussian, and $D\log(2\pi)$ is the normalization constant ensuring that the Gaussian density integrates to 1.

3.2 HMM inference

At test time, the model extracts, normalizes, and projects step-level features $\boldsymbol{x}_t$ into $\boldsymbol{z}_t$ identically to training.

We then compute an emission score matrix $\boldsymbol{E} \in \mathbb{R}^{T\times K}$, where $E_{t,k}$ is the emission score assigned to observation $\boldsymbol{z}_t$ under state k. For Gaussian emissions, we compute $E_{t,k} = \log \mathcal{N}(\boldsymbol{z}_t; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$.

We also evaluate cosine-based emissions for normalized embeddings. In that case, $E_{t,k} = T_{\cos}\cos(\boldsymbol{z}_t, \boldsymbol{\mu}_k)$, where $T_{\cos} > 0$ is a temperature parameter.

This cosine form matches the geometry of embedding spaces trained or commonly used with angular similarity, and can be interpreted as an unnormalized state score. We L2-normalize vectors before cosine to obtain a bounded, temperature-scaled score.

We also support a global emission scaling factor α, so that $[\tilde{\boldsymbol{E}}]_{t,k} = \alpha[\boldsymbol{E}]_{t,k}$. This adjusts the relative weight between emission and transition terms in the Viterbi objective.

We then use the Viterbi algorithm (Viterbi, 1967) to decode the highest-scoring step sequence $\hat{S}_{1:T} = \arg\max_{S_{1:T}} \text{Score}(S_{1:T}, \boldsymbol{z}_{1:T})$. The path score is

$$\text{Score}(S_{1:T}, \boldsymbol{z}_{1:T}) = \log \pi_{S_1} + \sum_{t=2}^{T} \log A_{S_{t-1},S_t} + \sum_{t=1}^{T} \tilde{E}_{t,S_t}. \tag{3}$$

For Gaussian emissions, this reduces to a sum of start, transition, and emission log-likelihood terms.

After predicting the sequence, we record for each time step t the emission score $\ell_t^{\text{em}} = \tilde{E}_{t,\hat{S}_t}$, and for each step $t > 1$ the transition score $\ell_t^{\text{tr}} = \log A_{\hat{S}_{t-1},\hat{S}_t}$. We refer to these as the HMM scores along the Viterbi path.

3.3 Step alignment and mistake recognition

Each procedure, regardless of its specific execution in the dataset, has one or more valid correct step sequences. This is an abstract representation of the

different ways to successfully complete the task. The pool of such reference sequences is formed either from the corresponding written instructions provided by annotation guidelines or from error-free videos. Let a reference sequence for a given procedure be $\boldsymbol{r}_{1:L} = (r_1, \ldots, r_L)$, where each r_ℓ is a step ID and L is the length of that reference sequence.

Next we compare the sequence predicted by the algorithm $\hat{S}_{1:T}$ with one of the reference sequences, namely, the closest reference variant under the selected alignment method. We align the predicted sequence to the chosen reference sequence using global sequence alignment.[2] This alignment identifies structural errors: **insertions** (extra steps), **deletions** (missing steps), **transpositions** (wrong order of correct steps)[3], and **substitutions** (mismatched steps at an aligned position).

In addition to structural errors, we use the Viterbi-path scores directly. Low values of ℓ_t^{em} indicate atypical observations under the assigned state, and low values of ℓ_t^{tr} indicate atypical state transitions. Given the configuration-specific thresholds τ_{em} and τ_{tr}, reported in Appendix D, we flag a LowProbEmission if $\ell_t^{\mathrm{em}} < \tau_{\mathrm{em}}$, and for $t > 1$ a LowProbTransition if $\ell_t^{\mathrm{tr}} < \tau_{\mathrm{tr}}$. These two classes provide score-based diagnostics that complement the edit labels: LowProbEmission flags an observation that is unlikely under its decoded step state, while LowProbTransition flags an unlikely adjacent step transition in the decoded path. They can overlap with edit-based errors, but they also identify low-confidence or locally implausible events that are not cleanly summarized by a single insertion, deletion, substitution, or transposition[4].

Structural alignment cannot expose an altered execution when the predicted and reference steps share the same Step ID. In this case, we compute the cosine similarity between the observation and the mean embedding of its assigned state. A similarity below $\tau_{\mathrm{sem}} = 0.85$ produces a semantic Substitution. This rule captures changes in quantities, tools, destinations, and other arguments hidden by a shared Step ID.

3.4 Language grounding and step representations

We use language as the interface between video and procedural reasoning. We call benchmark-provided or benchmark-derived step descriptions gold text and VLM-generated descriptions grounded text. Gold-text runs bypass video grounding and do not use video. Grounded runs use descriptions generated by a VLM. The HMM decodes the resulting embedding sequence into task-local Step IDs.

Gold step text. We start from the canonical step descriptions provided by each benchmark. For erroneous steps, we use the benchmark-specific observed or modified step descriptions when they are available. When a benchmark provides only a short error cue instead of a full observed step, we convert it into a complete modified description by combining the canonical step text with the error cue. We also apply minimal text cleanup to remove benchmark-specific artifacts that do not carry procedural meaning. Appendix A gives representative dataset-specific examples.

The resulting representation keeps the lexical detail of the original annotation, including quantities, destinations, surfaces, and local modifiers. That detail matters because many procedural mistakes differ from the canonical step in terms of the arguments, not only in terms of the main verb.

Semantic role representations. We convert each step description into a predicate–argument form that makes its action, participants, and attributes explicit. We use the form `PREDICATE(Role: value, Role: value, ...)`. For example, `PEEL(Agent: you, Object: onion(Size: medium, Quantity: 1))`. SRL representations preserve distinctions in object identity, quantity, destination, surface, and manner while reducing variation in surface form. These arguments often determine whether a step is correct. This structure is useful for CaptainCook4D, where near-duplicate steps can have different descriptions or IDs,[5] and

[2] We use either Needleman–Wunsch or difflib. Needleman–Wunsch is a classical dynamic-programming algorithm for global sequence alignment. It optimizes a well-defined objective and is preferable when we need full control over substitution and gap penalties. Python's difflib.SequenceMatcher is a heuristic matcher that focuses on maximizing the length of matching blocks.

[3] We identify a transposition by pairing a deletion and an insertion of the same step ID at different alignment positions.

[4] LowProbEmission flags a step whose decoded ID is correct but whose visual or textual content is atypical, such as using the wrong quantity, tool, or destination. For example, in EgoPER, the grounded text *Drop tortilla on floor* is decoded as the procedural state *Place tortilla on cutting board*, but receives a very low emission score. LowProbTransition can flag an unlikely local ordering even when both neighboring steps are individually valid, e.g., a CaptainCook4D transition from rinsing mushrooms directly to adding them to the pan before the expected drying, stemming, and slicing steps.

[5] For example, *Peel-Peel 1 garlic clove*, Step ID 14, and *Peel-Peel 1 garlic cloves*, Step ID 200.

for EgoOops, where long step descriptions are paired with short error cues.

We construct SRL representations with both a proprietary model, GPT-5.3 (OpenAI, 2025), and an open-weight model, Qwen3.6-27B (Qwen Team, 2026b). The two generators follow the same constrained predicate–argument format. We treat their outputs as two alternative SRL variants and compare how differences in role labeling and event decomposition affect downstream performance. Appendix C gives additional SRL details, including the prompt and generation settings for each model and systematic differences with examples.

Action–object representations. We also derive an action–object abstraction that retains the main predicate and a flat set of salient entities, including tools and contacted objects. It removes quantities, attributes, manner, and the semantic relations that distinguish roles such as Destination and Surface. This is used as an ablation to test how much procedural detail can be removed before mistake recognition drops.

Video-to-text grounding with VLMs. For each segment, the VLM receives the clip and the candidate descriptions for its task. The candidates provide task-specific closed-set guidance. The VLM returns a structured output containing `predicted_description`, `action_name`, and `objects`. Raw-text runs encode `predicted_description` directly. The action–object ablation uses the two auxiliary fields. Grounded SRL runs retrieve a cached representation keyed only by the normalized `predicted_description`. Annotated Step IDs, canonical descriptions, and modified descriptions do not enter this lookup.

4 Experiments and results

Benchmarks and protocol. We adopt a common post-completion protocol across CaptainCook4D (Peddi et al., 2024), EgoPER (Lee et al., 2024), and EgoOops (Haneji et al., 2025): the input is a fully observed, step-segmented procedural video, and the model must recover the executed step sequence and detect deviations from a valid reference procedure. We follow the text augmentation and cleanup protocol described in Section 3.4.

For CaptainCook4D and EgoPER, we use the benchmark-provided train, validation, and test splits and train each HMM on correct training videos. For Table 1, we select one complete configuration within each representation family by validation Any-F1; Type-F1, AUC, and Seq break ties in that order. Test data are used only for final reporting. Decision thresholds are fixed before test evaluation and reported in Appendix D. EgoOops has no official validation split. We train its HMMs on all 21 correct videos. For each representation family, we transfer the complete configuration with the highest mean validation Any-F1 across CaptainCook4D and EgoPER; Type-F1, AUC, and Seq again break ties. EgoOops error-video labels are not used for selection. The matched representation and decoder studies use the configurations reported in Table 15.

Under this step-level class imbalance (20.8%–39.2% errors across datasets), accuracy can obscure poor error recall. We prioritize F1 and AUC metrics.

Models and representations. For visual step embeddings, we use frozen VideoMAE-v2 (Wang et al., 2023a) and V-JEPA 2 (Assran et al., 2025) backbones. For text, we sweep frozen CLIP (Radford et al., 2021), all-MiniLM (Wang et al., 2020), all-MPNet (Song et al., 2020), SimCSE (Gao et al., 2021), Instructor-XL (Su et al., 2023), and OpenAI text-embedding-3-large (OpenAI, 2024). For video-to-text grounding, we compare Gemini 3 Flash (Doshi and the Gemini Team, 2025) and Qwen 3.5-27B (Qwen Team, 2026a). For *Qwen + GPT-SRL* and *Gemini + GPT-SRL*, we first generate grounded step descriptions with the VLM and then convert these generated descriptions into SRL representations[6] before text encoding and HMM inference. On the language side, we evaluate step text, semantic role representations, and action–object abstractions.

Metrics. We report four main metrics for the unified post-completion task. **Any-F1** is our primary metric. It treats a segment as erroneous if it carries any of the four edit-based labels: Insertion, Deletion, Substitution, or Transposition. **Type-F1** requires both the error location and its edit-based type to match the gold annotation. **Seq** (Sequence accuracy) aligns the decoded and annotated executed Step-ID sequences with Needleman–Wunsch

[6]The main-paper SRL results use representations generated by GPT-5.3 (OpenAI, 2025). We treat these as the primary SRL setting in the main comparisons, and report the corresponding Qwen3.6-27B (Qwen Team, 2026b) results in Appendix F.

	EgoOops				EgoPER				CC4D			
Family	**Seq**	**Any-F1**	**Type-F1**	**AUC**	**Seq**	**Any-F1**	**Type-F1**	**AUC**	**Seq**	**Any-F1**	**Type-F1**	**AUC**
Gold text	0.713	0.480	0.446	0.756	0.903	**0.741**	0.673	**0.918**	0.617	**0.631**	0.466	**0.761**
Vision	0.328	0.454	0.439	0.523	0.372	0.519	0.527	0.580	0.159	0.556	0.464	0.478
Gold + GPT-SRL	**0.742**	**0.510**	**0.533**	**0.775**	**0.913**	0.736	**0.675**	0.908	**0.638**	0.629	0.495	0.757
Qwen text	0.493	0.441	0.451	0.564	0.802	0.607	0.599	0.736	0.517	0.603	0.498	0.640
Qwen + GPT-SRL	0.499	0.409	0.407	0.623	0.816	0.601	0.594	0.742	0.485	0.592	0.505	0.648
Gemini text	0.557	0.408	0.361	0.610	0.845	0.642	0.589	0.748	0.577	0.603	**0.516**	0.656
Gemini + GPT-SRL	0.560	0.353	0.314	0.611	0.854	0.631	0.595	0.749	0.554	0.595	0.513	0.669

Table 1: Family-level leaderboard. Each row reports all metrics from one complete configuration. Encoder and HMM settings may differ across rows. Seq denotes sequence accuracy, and CC4D denotes CaptainCook4D. Appendix E.1 lists the corresponding configurations.

	Gold				Qwen				Gemini			
Dataset	**Raw**	**GPT-SRL**	**Δ**	**Act.+Obj.**	**Raw**	**GPT-SRL**	**Δ**	**Act.+Obj.**	**Raw**	**GPT-SRL**	**Δ**	**Act.+Obj.**
EgoOops	0.400	**0.434**	+0.034	0.318	0.321	0.356	+0.036	**0.444**	0.349	0.353	+0.003	**0.414**
EgoPER	0.681	**0.725**	+0.044	0.540	0.578	**0.601**	+0.024	0.550	0.602	**0.631**	+0.030	0.552
CaptainCook4D	**0.629**	**0.629**	+0.000	0.571	0.587	**0.592**	+0.005	0.566	0.595	**0.596**	+0.001	0.570

Table 2: Any-F1 by text source and representation. Each Raw–GPT-SRL pair uses identical downstream settings, and Δ is GPT-SRL minus Raw. Act.+Obj. reports the best configuration within that representation family. Bold marks the highest score within each dataset–source group. Differences are computed before rounding.

and reports the number of identical non-gap pairs divided by the number of evaluated segments. **AUC** uses the negative emission score as a continuous anomaly score and measures threshold-free discrimination over observed segments. A gold Deletion is represented by an annotated position without an observed clip; it contributes to Any-F1 and Type-F1 but has no emission score and is excluded from AUC.

For comparison with prior work, we report Error Detection Accuracy (EDA) and area under the ROC curve (AUC) on EgoPER and CaptainCook4D and Precision only for CaptainCook4D, following the original EgoPED (Lee et al., 2024), AMNAR (Huang et al., 2025) and AEM (Guo et al., 2026) evaluation protocols. We use these metrics only to compare with prior work.

For Tables 1 and 4, all metrics within each row come from a single configuration. The family-level leaderboard uses the validation-based selection protocol described above. The matched Raw–SRL and decoder comparisons use fixed downstream configurations and perform no independent selection between conditions. For Table 4, we select one configuration per family using validation EDA and report all benchmark-specific test metrics from that configuration.

Family-level results. Table 1 compares representation families under the configuration-selection protocol described above. Gold-text families achieve the highest Any-F1 on each dataset. Grounded text-based variants outperform visual embeddings on EgoPER and CaptainCook4D, while vision performs better on EgoOops. This contrast highlights the importance of the available step descriptions. Within the separately selected grounded configurations, Raw achieves higher Any-F1 than SRL in all six pairs, although SRL improves Seq in four. Table 2 therefore examines representation changes under identical downstream settings.

Matched Raw–SRL comparison. Under identical downstream settings, GPT-SRL improves Any-F1 in all six grounded comparisons by 0.001–0.036. The largest gain occurs for Qwen on EgoOops. Gains on CaptainCook4D are small for both VLMs. Gold-description comparisons show improvements on EgoOops and EgoPER and a tie on CaptainCook4D. Together with Table 1, these results indicate that the benefit of SRL depends on the downstream configuration and selection criterion. Paired video-level bootstrap intervals for the Any-F1 difference exclude zero in three of the six grounded comparisons (Appendix F.5).

GPT-SRL also exceeds the independently se-

Dataset	VLM	Nearest candidate		Greedy emission		Viterbi	
		Raw	SRL	Raw	SRL	Raw	SRL
EgoOops	Qwen	0.337	0.335	**0.339**	0.351	0.321	**0.356**
	Gemini	0.305	0.315	0.339	0.328	**0.349**	**0.353**
EgoPER	Qwen	0.547	0.524	**0.578**	**0.601**	**0.578**	**0.601**
	Gemini	0.571	0.557	**0.602**	**0.631**	**0.602**	**0.631**
CaptainCook4D	Qwen	0.567	0.559	0.581	0.578	**0.587**	**0.592**
	Gemini	0.538	0.528	0.564	0.561	**0.595**	**0.596**

Table 3: Decoder controls under the matched Raw–SRL configurations. Nearest-candidate matching uses cosine similarity and no HMM. Greedy decoding uses learned emissions without start or transition probabilities. Viterbi uses the complete HMM. Bold marks the highest Any-F1 within each dataset–VLM–representation condition.

Method	EgoPER		CaptainCook4D		
	EDA	AUC	Precision	EDA	AUC
AMNAR	0.644	0.685	0.668	0.723	0.602
AEM (pred. segments)	0.667	0.738	0.681	0.719	0.625
AEM (GT segments)	**0.747**	**0.867**	–	–	–
Ours (Qwen)	0.704	0.751	0.253	0.646	0.694
Ours (Qwen + GPT-SRL)	0.699	0.753	0.785	0.841	**0.900**
Ours (Gemini)	0.725	0.768	0.291	0.711	0.753
Ours (Gemini + GPT-SRL)	0.713	0.762	**0.809**	**0.843**	**0.900**

Score sources. AMNAR and AEM results come from their original papers (Huang et al., 2025; Guo et al., 2026). AEM (GT segments) is an inference-time ablation that replaces predicted segments with ground-truth segments.

Table 4: System-level results under benchmark-specific protocols. PRISM uses benchmark-provided step boundaries. AEM results are shown for predicted and ground-truth (GT) segmentation. Baseline scores are taken from the original papers; all metrics in each PRISM row come from one configuration.

lected Act.+Obj. configurations in seven of nine dataset–source settings; both exceptions occur with grounded EgoOops descriptions. This comparison complements the matched Raw–SRL analysis by showing the performance of a coarser representation. Appendix F.3 reports effects by grounding stratum.

Decoder controls. Table 3 separates independent candidate matching, the trained local decoder, and sequence decoding. Greedy emission decoding improves Any-F1 over nearest-candidate matching in all twelve conditions. Adding Viterbi improves seven conditions, matches four, and reduces one; it exceeds nearest-candidate matching in eleven conditions overall. EgoPER's reported Any-F1 is unchanged by transition-aware decoding. The only decrease relative to greedy decoding occurs for Qwen Raw on EgoOops, where transition estimation relies on 21 correct videos across five tasks. These results show consistent benefits from the trained local decoder and dataset-dependent benefits from sequence decoding.

Comparison with prior work. Table 4 contextualizes PRISM against published systems under their respective protocols. On CaptainCook4D, the SRL configurations obtain EDA of 0.841–0.843 and AUC of 0.900. On EgoPER, AEM with ground-truth segmentation achieves higher EDA and AUC than PRISM. Differences in segmentation and configuration selection matter when interpreting these system-level results. Tables 2 and 3 provide the controlled comparisons of representation and decoding within PRISM.

5 Conclusion

We introduced PRISM, a language-grounded framework for post-completion procedural mistake recognition that combines an HMM trained on correct executions with reference alignment and semantic checks. Across three benchmarks, the trained local decoder consistently improves over nearest-candidate matching, while sequence decoding provides additional, dataset-dependent gains. SRL improves Any-F1 in all six matched grounded comparisons, with benefits that depend on the dataset and downstream configuration. Our unified edit-based evaluation makes these effects comparable across benchmarks. Together, the re-

sults support language grounding as an effective interface for reasoning about procedural step content and order, making the underlying deviation explicit even when the failure appears only as a plate of overboiled zoodles.

Limitations

Our evaluation assumes benchmark-provided step boundaries, task-specific candidate descriptions, and valid reference procedures. Gold-text experiments bypass video grounding, whereas grounded experiments depend on the VLM preserving error-relevant details. The results therefore characterize mistake reasoning under these inputs. Deployment additionally requires reliable temporal segmentation and grounding, and adapting PRISM to a new procedure requires specifying its candidate steps and reference sequences and estimating its task-specific parameters.

Although the SRL generator comparison between GPT-5.3 and Qwen3.6-27B suggests that the pipeline is stable across SRL generators, the method still assumes consistent role normalization. Inconsistent semantic parses can propagate into HMM reasoning.

Our four label classes capture structural sequence edits, but they do not exhaust the full granularity of procedural mistake semantics. They align with recent unified mistake-taxonomy efforts, but for compatibility with current benchmarks we collapse Correction into Insertion and Wrong Execution into Substitution. This choice supports cross-dataset evaluation, at the cost of merging some semantically distinct cases. A natural next step is to refine the substitution class into full-step substitutions and role-level wrong executions, using SRL arguments to localize which part of a step changed and to support more fine-grained automated mistake attribution.

Ethical considerations

This work uses existing benchmark datasets and does not involve new data collection, human-subject experiments, or deployment in real-world decision-making settings. Nevertheless, procedural mistake detection may raise ethical concerns in workplace, educational, or assistive settings, where video-based monitoring could be used for surveillance or overly punitive evaluation. Our method also depends on VLM-based grounding, so mistakes in generated step descriptions may propagate into downstream reasoning and lead to incorrect or misleading error labels. These systems should therefore be used only as decision-support tools under human oversight, rather than as standalone grounds for sanctions or safety-critical decisions. Any real deployment should additionally account for privacy, consent, secure data handling, and appropriate safeguards for affected users. Upon acceptance, we will release our code and scripts for generating SRL representations, in accordance with the licenses of the original benchmarks.

References

Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Mojtaba Komeili, Matthew Muckley, Ammar Rizvi, Claire Roberts, Koustuv Sinha, Artem Zholus, Sergio Arnaud, Abha Gejji, Ada Martin, Francois Robert Hogan, Daniel Dugas, Piotr Bojanowski, Vasil Khalidov, Patrick Labatut, and 10 others. 2025. V-JEPA 2: Self-supervised video models enable understanding, prediction and planning. *Preprint*, arXiv:2506.09985.

Konstantinos Bacharidis and Antonis A. Argyros. 2025. Vision-based mistake analysis in procedural activities: A review of advances and challenges. *Preprint*, arXiv:2510.19292.

Collin F. Baker, Charles J. Fillmore, and John B. Lowe. 1998. The Berkeley FrameNet project. In *36th Annual Meeting of the Association for Computational Linguistics and 17th International Conference on Computational Linguistics, Volume 1*, pages 86–90, Montreal, Quebec, Canada. Association for Computational Linguistics.

Anil Batra, Laura Sevilla-Lara, Marcus Rohrbach, and Frank Keller. 2025. Predicting implicit arguments in procedural video instructions. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 30399–30419, Vienna, Austria. Association for Computational Linguistics.

Michael D. Byrne and Elizabeth M. Davis. 2006. Task structure and postcompletion error in the execution of a routine procedure. *Human Factors*, 48(4):627–638.

Huiyao Chen, Meishan Zhang, Jing Li, Min Zhang, Lilja Øvrelid, Jan Hajič, and Hao Fei. 2025. Semantic role labeling: A systematical survey. *Preprint*, arXiv:2502.08660.

Guodong Ding, Fadime Sener, Shugao Ma, and Angela Yao. 2023a. Every mistake counts in assembly. *ArXiv*, abs/2307.16453.

Guodong Ding, Fadime Sener, and Angela Yao. 2023b. Temporal action segmentation: An analysis of modern techniques. *Preprint*, arXiv:2210.10352.

Xi Ding and Lei Wang. 2025. The journey of action recognition. *Companion Proceedings of the ACM on Web Conference 2025*.

Tulsee Doshi and the Gemini Team. 2025. Gemini 3 Flash: Frontier intelligence built for speed. Official announcement.

Alessandro Flaborea, Guido Maria D'Amely di Melendugno, Leonardo Plini, Luca Scofano, Edoardo De Matteis, Antonino Furnari, G. Farinella, and Fabio Galasso. 2024. PREGO: Online mistake detection in procedural egocentric videos. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18483–18492.

Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. SimCSE: Simple contrastive learning of sentence embeddings. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6894–6910, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.

Sara Garfield and B.D. Franklin. 2016. Understanding models of error in clinical practice. *Pharmaceutical Journal*, 296:361–364.

Daniel Gildea and Daniel Jurafsky. 2002. Automatic labeling of semantic roles. *Computational Linguistics*, 28(3):245–288.

Wenliang Guo, Yujiang Pu, and Yu Kong. 2026. Procedural mistake detection via action effect modeling. *Preprint*, arXiv:2512.03474.

Yuto Haneji, Taichi Nishimura, Hirotaka Kameko, Keisuke Shirai, Tomoya Yoshida, Keiya Kajimura, Koki Yamamoto, Taiyu Cui, Tomohiro Nishimoto, and Shinsuke Mori. 2025. EgoOops: A dataset for mistake action detection from egocentric videos referring to procedural texts. *Preprint*, arXiv:2410.05343.

Wei-Jin Huang, Yuan-Ming Li, Zhi-Wei Xia, Yu-Ming Tang, Kun-Yu Lin, Jian-Fang Hu, and Wei-Shi Zheng. 2025. Modeling multiple normal action representations for error detection in procedural tasks. *Preprint*, arXiv:2503.22405.

Youngkyoon Jang, Brian T. Sullivan, Casimir J. H. Ludwig, Iain D. Gilchrist, Dima Damen, and W. Mayol-Cuevas. 2019. EPIC-Tent: An egocentric video dataset for camping tent assembly. *2019 IEEE/CVF International Conference on Computer Vision Workshop (ICCVW)*, pages 4461–4469.

Bolin Lai, Sam Toyer, Tushar Nagarajan, Rohit Girdhar, Shengxin Zha, James M. Rehg, Kris Kitani, Kristen Grauman, Ruta Desai, and Miao Liu. 2024. Human action anticipation: A survey. *Preprint*, arXiv:2410.14045.

Shih-Po Lee and Ehsan Elhamifar. 2025. Error recognition in procedural videos using generalized task graph. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 10009–10021.

Shih-Po Lee, Zijia Lu, Zekun Zhang, Minh Hoai, and Ehsan Elhamifar. 2024. Error detection in egocentric procedural task videos. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18655–18666.

Yayuan Li, Aadit Jain, Filippos Bellos, and Jason J. Corso. 2025. Mistake attribution: Fine-grained mistake understanding in egocentric videos. *Preprint*, arXiv:2511.20525.

Olga Loginova and Frank Keller. 2026. How to correctly make mistakes: A framework for constructing and benchmarking mistake aware egocentric procedural videos. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, pages 8843–8853.

Michele Mazzamuto, Antonino Furnari, Yoichi Sato, and Giovanni Maria Farinella. 2025. Gazing into missteps: Leveraging eye-gaze for unsupervised mistake detection in egocentric videos of skilled human activities. *Preprint*, arXiv:2406.08379.

Kosuke Moriwaki, Gaku Nakano, and Tetsuo Inoshita. 2022. The BRIO-TA dataset: Understanding anomalous assembly process in manufacturing. In *2022 IEEE International Conference on Image Processing (ICIP)*, pages 1991–1995.

Medhini Narasimhan, Licheng Yu, Sean Bell, Ning Zhang, and Trevor Darrell. 2023. Learning and verification of task structure in instructional videos. *Preprint*, arXiv:2303.13519.

Donald A. Norman. 2002. *The Design of Everyday Things*, with a new introduction by the author edition. Basic Books, New York, N.Y.

OpenAI. 2024. New embedding models and API updates. Introduces text-embedding-3-large.

OpenAI. 2025. GPT-5 system card. Technical report, OpenAI. Version August 13, 2025.

Martha Palmer, Daniel Gildea, and Paul Kingsbury. 2005. The Proposition Bank: An annotated corpus of semantic roles. *Computational Linguistics*, 31(1):71–106.

Rohith Peddi, Shivvrat Arya, Bharath Challa, Likhitha Pallapothula, Akshay Vyas, Bhavya Gouripeddi, Jikai Wang, Qifan Zhang, Vasundhara Komaragiri, Eric Ragan, Nicholas Ruozzi, Yu Xiang, and Vibhav Gogate. 2024. CaptainCook4D: A dataset for understanding errors in procedural activities. *Preprint*, arXiv:2312.14556.

Leonardo Plini, Luca Scofano, Edoardo De Matteis, Guido Maria D'Amely di Melendugno, Alessandro Flaborea, Andrea Sanchietti, G. Farinella, Fabio Galasso, and Antonino Furnari. 2024. TI-PREGO: Chain of thought and in-context learning for online mistake detection in procedural egocentric videos. *ArXiv*, abs/2411.02570.

Yicheng Qian, Weixin Luo, Dongze Lian, Xu Tang, Peilin Zhao, and Shenghua Gao. 2022. SVIP: Sequence verification for procedures in videos. *Preprint*, arXiv:2112.06447.

Qwen Team. 2026a. Qwen3.5: Towards native multimodal agents.

Qwen Team. 2026b. Qwen3.6-27B: Flagship-level coding in a 27B dense model.

Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. 2021. Learning transferable visual models from natural language supervision. *Preprint*, arXiv:2103.00020.

James Reason. 2000. Human error: models and management. *BMJ*, 320 7237:768–70.

Arka Sadhu, Tanmay Gupta, Mark Yatskar, Ram Nevatia, and Aniruddha Kembhavi. 2021. Visual semantic role labeling for video understanding. *Preprint*, arXiv:2104.00990.

Tim J. Schoonbeek, Tim Houben, Hans Onvlee, Peter H. N. de With, and Fons van der Sommen. 2023. IndustReal: A dataset for procedure step recognition handling execution errors in egocentric videos in an industrial-like setting. *Preprint*, arXiv:2310.17323.

Luigi Seminara, Giovanni Maria Farinella, and Antonino Furnari. 2025. Differentiable task graph learning: Procedural activity representation and online mistake detection from egocentric videos. *Preprint*, arXiv:2406.01486.

Fadime Sener, Dibyadip Chatterjee, Daniel Shelepov, Kun He, Dipika Singhania, Robert Wang, and Angela Yao. 2022. Assembly101: A large-scale multi-view video dataset for understanding procedural activities. *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 21064–21074.

Kaitao Song, Xu Tan, Tao Qin, Jianfeng Lu, and Tie-Yan Liu. 2020. MPNet: Masked and permuted pre-training for language understanding. *Preprint*, arXiv:2004.09297.

Shane Storks, Itamar Bar-Yossef, Yayuan Li, Zheyuan Zhang, Jason J. Corso, and Joyce Chai. 2025. Transparent and coherent procedural mistake detection. *Preprint*, arXiv:2412.11927.

Hongjin Su, Weijia Shi, Jungo Kasai, Yizhong Wang, Yushi Hu, Mari Ostendorf, Wen tau Yih, Noah A. Smith, Luke Zettlemoyer, and Tao Yu. 2023. One embedder, any task: Instruction-finetuned text embeddings. *Preprint*, arXiv:2212.09741.

A. Viterbi. 1967. Error bounds for convolutional codes and an asymptotically optimum decoding algorithm. *IEEE Transactions on Information Theory*, 13(2):260–269.

Limin Wang, Bingkun Huang, Zhiyu Zhao, Zhan Tong, Yinan He, Yi Wang, Yali Wang, and Yu Qiao. 2023a. VideoMAE V2: Scaling video masked autoencoders with dual masking. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 14549–14560.

Wenhui Wang, Furu Wei, Li Dong, Hangbo Bao, Nan Yang, and Ming Zhou. 2020. MiniLM: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. *Preprint*, arXiv:2002.10957.

Xin Wang, Taein Kwon, Mahdi Rad, Bowen Pan, Ishani Chakraborty, Sean Andrist, D. Bohus, Ashley Feniello, Bugra Tekin, F. Frujeri, Neel Joshi, and Marc Pollefeys. 2023b. HoloAssist: an egocentric human interaction dataset for interactive AI assistants in the real world. *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 20213–20224.

A Dataset-specific annotation profiles

The benchmarks differ in scale, domain, annotation style, and lexical variability. Table 5 summarizes CaptainCook4D (Peddi et al., 2024), EgoPER (Lee et al., 2024), and EgoOops (Haneji et al., 2025). CaptainCook4D spans 24 cooking tasks. EgoPER has a similar number of videos concentrated in five cooking tasks. EgoOops covers five non-cooking procedures. The following profiles describe the annotation differences that affect representation and evaluation.

Dataset	Videos	Tasks	Hours	Step instances
CaptainCook4D	384	24	94.5	5.3K
EgoPER	386	5	28.0	5,369
EgoOops	50	5	6.8	538

Table 5: Benchmark scale. Hours denote total video duration. The Step instances column reports annotated procedural units used for step-level training and evaluation. The benchmark papers use different names for these units.

CaptainCook4D. CaptainCook4D contains cooking procedures with relatively short step descriptions, but the surface forms are noisy. The benchmark includes duplicated predicate prefixes such as *Pour-Pour* and *Peel-Peel*, which we remove during text cleanup because they do not add procedural meaning. The dataset also contains near-duplicate step descriptions with different IDs, for example singular/plural variants such as *Peel 1 garlic clove* and *Peel 1 garlic cloves*. Modified descriptions often encode quantity, duration, temperature, instrument, or destination changes inside the step arguments. For example, *Microwave*

the plate, covered, on high for 5 minutes instead of 1.5 minutes differs from the canonical step mainly through duration. Similarly, the contrast between *Cut the English muffin into two pieces with a knife* and *Cut the English muffin into two pieces with hand* reflects a change in instrument, and the contrast between *Top cup with 1 tablespoon of salsa* and *Top cup with 1 teaspoon of salsa* reflects a change in quantity. These cases illustrate why argument-level information remains important for mistake recognition. SRL preserves these distinctions in an explicit predicate–argument form.

EgoPER. EgoPER uses short and fairly regular cooking instructions in a visually consistent environment. Its step descriptions are more atomic than in EgoOops and less lexically variable than in CaptainCook4D. Many errors are local substitutions of action, object, or destination, such as changing the target surface or replacing the intended manipulation with a visibly different one. For example, *Place tortilla on cutting board* becomes *Place tortilla on table*, and *Transfer water to kettle* becomes *Spill out water*. Because text descriptions and their deviations are often short, explicit, and lexically close, text descriptions and SRL remain close in performance. Action–object abstractions are still weaker because they remove distinctions that remain useful for mistake recognition, but the drop is more interpretable than in CaptainCook4D or EgoOops.

EgoOops. EgoOops has the most heterogeneous step language. Some canonical steps are long and contain multiple subevents, for example *Pour about 15mL of water into a cup, dip the tip of a red highlighter in the water and squeeze out a drop.* Its error annotations are also unusual: the benchmark often provides a short error cue instead of a full observed step description. We therefore construct complete modified descriptions by concatenating the canonical step with the error cue. For example, the canonical step *Turn on the switch S2.* plus the error cue *Blow up the propeller.* becomes *Turn on the switch S2. Blow up the propeller.* Other cases modify object identity or orientation inside the step, for example *Connect the battery box and switch S1 in series.* becoming *Connect the battery box and switch S2 instead of switch S1 in series.* EgoOops errors may also introduce an extra event inside a step, for example *Bridge two long thin cuboid blocks on the upright blocks from left to right in the order of yellow and red. Collapse blocks.* These patterns make the benchmark especially sensitive to role fillers, argument structure, and multi-event step descriptions. SRL makes these roles and appended deviations explicit.

Grounded descriptions. The same benchmark differences also affect video-to-text grounding. Gemini usually produces gold-aligned step descriptions more accurately than Qwen, especially on CaptainCook4D. On EgoPER, both VLMs are closer to the candidate-step vocabulary because the step language is shorter and more regular. On EgoOops, grounded descriptions remain harder because the benchmark steps can be long, compositional, and error cues may add an extra event to an otherwise canonical step.

B Unified label mapping

We map the benchmark-specific error annotations to a unified edit-based taxonomy with four labels: *Insertion*, *Deletion*, *Substitution*, and *Transposition*. Table 6 summarizes the dataset-specific mappings. This mapping defines the gold targets for Type-F1 and converts native annotations into the four edit labels predicted after sequence alignment. It does not affect HMM training or Step-ID decoding. Because the benchmarks use different error inventories and annotation conventions, the mapping rules are dataset-specific.

When a task admits several valid references, the evaluator selects the reference closest to the decoded sequence. Structural targets are instantiated relative to that reference, as detailed in Appendix F.4.

CaptainCook4D. CaptainCook4D provides explicit categories for erroneous steps. Additional executed actions that are not part of the reference sequence are treated as *Insertion*. In the original dataset annotations, omitted steps are represented with `start_time = -1.0` and `end_time = -1.0`; we keep this convention when constructing deletion labels.

EgoPER. EgoPER defines a psychologically motivated taxonomy of procedural errors and does not include a native *Transposition* category. We derive transpositions by aligning the annotated executed Step-ID sequence with the valid reference sequence and identifying steps that occur in a different position. As in CaptainCook4D, omitted steps are represented by missing aligned actions and receive

Dataset	Native label / description	Unified label
CC4D	Order Error	Transposition
	Missing Step	Deletion
	Preparation Error	Substitution
	Measurement Error	Substitution
	Timing Error	Substitution
	Technique Error	Substitution
	Temperature Error	Substitution
	Other	Substitution
	additional executed step outside the reference sequence	Insertion
EgoPER	step omission	Deletion
	addition	Insertion
	correction	Insertion
	modification	Substitution
	slip	Substitution
	step aligned as a deletion and insertion of the same Step ID	Transposition
EgoOops	working with wrong objects	Substitution
	grasping wrong objects and releasing them without using	Insertion
	correction of mistake actions	Insertion
	unintended actions	Substitution
	working in the wrong way or moving	Substitution
	others: mix all the liquid and shine a black light on it	Insertion
	others: not able to fold because there are no incisions	Substitution
	others: can't fold cardboard	Substitution
	others: after turning on the switch s1, turn off it	Substitution
	others: at the same time, also turn on the switch s1 then turn off it	Substitution
	others: place two magnesium plates on the last step and put one back	Substitution
	others: pour much water but should pour about 15 ml / knock over an empty cup	Substitution
	others: pour powdered detergent into a cup of yellow liquid but should mix in an empty cup	Substitution
	others: put two drops on a copper plate	Substitution
	others: scatter pieces of zinc / put three zinc plates on the right column, but should on the center column	Substitution
	others: put two drops on the metal plate in the center of the microplate	Substitution
	others: place blocks from left to right in the order of red and yellow but should put them in the order of yellow and red	Substitution

Table 6: Mapping from benchmark-specific error labels to the unified edit-based taxonomy used in this work. For EgoOops, the rare `others` cases are resolved individually from the textual error descriptions. CC4D is CaptainCook4D.

`start_time = -1.0` and `end_time = -1.0` in the derived sequence-level representation.

EgoOops. EgoOops often encodes the deviation as a short textual cue attached to an otherwise canonical step and does not contain *Deletion* or *Transposition*. The native category *others* is rare and is resolved case by case from the textual error description. Within this group, only *mix all the liquid and shine a black light on it* is mapped to *Insertion*; all remaining listed *others* cases are mapped to *Substitution*.

Dataset	Type	ID	Step description	SRL representation
CC4D	modified	20	Microwave the plate, covered, on high for 5 minutes instead of 1.5 minutes	`MICROWAVE(Agent: you, Object: plate(State: covered), Setting: high, Duration: minute(Quantity: 5), Manner: INSTEAD_OF(Duration: minute(Quantity: 1.5)))`
	modified	28	Measure 1/8 teaspoon of salt using dirty measuring spoon and add it to the mug	`ADD(Agent: you, Object: salt(Measurement: teaspoon(Quantity: 1/8)), Destination: mug, Instrument: dirty(measuring_spoon), Manner: after(MEASURE(Agent: you, Object: salt(Measurement: teaspoon(Quantity: 1/8)), Instrument: dirty(measuring_spoon))))`
	canonical	54	Put all the Vegetables in a microwave-safe bowl	`PUT(Agent: you, Object: vegetables(Quantity: all), Destination: bowl(Property: microwave_safe))`
EgoOops	canonical	211	Put chopsticks through the cutout in the lid to widen the hole.	`PUT(Agent: you, Object: chopsticks, Path: through(cutout(Location: in(lid))), Purpose: WIDEN(Agent: you, Object: hole))`
	modified	10	Pour about 15mL of water into a cup, dip the tip of a red highlighter in the water and squeeze out a drop. Knock over an empty cup.	`POUR(Agent: you, Object: water(Measurement: ml(Quantity: 15, Approx: about)), Destination: into(cup)); DIP(Agent: you, Object: tip(of(red_highlighter)), Destination: in(water(Location: in(cup)))); SQUEEZE_OUT(Agent: you, Object: drop); KNOCK_OVER(Agent: you, Object: empty_cup)`
EgoPER	canonical	57	Put banana slices on tortilla	`PUT(Agent: you, Object: banana_slices, Destination: on(tortilla))`
	canonical	37	Microwave for X seconds	`MICROWAVE(Agent: you, Temporal: for(second(Quantity: x)))`
	canonical	18	Drop tea bag	`DROP(Agent: you, Object: tea_bag)`

Table 7: Representative SRL examples from the three benchmarks. We include both canonical and modified step descriptions to show how the representation handles quantities, instruments, destinations, purposes, and multi-event modified steps. CC4D is CaptainCook4D.

Hyperparameter	Sym.	Tags	Meaning	Value effect	Default settings & rationale
Frames per step	n	t/i	*Number of frames* per step averaged into the image feature.	High: stronger denoising and object coverage, but more IO/latency and possible over-smoothing on very short steps; low: faster and crisper, but noisier and more sensitive to accidental frames.	**3**: a good robustness–latency trade-off; using more frames may introduce irrelevant content in short steps but can help with long, stable steps.
PCA dimension	d	t/i	Output dimensionality after PCA (fit on train; applied at inference) before Gaussian emissions.	High: preserves fine semantics, but makes covariance estimation harder and can overfit; low: better conditioning and sample-efficiency, but too low drops discriminative cues.	**128**: stable Ledoit–Wolf covariances with modest memory/compute; increase with ample per-class data.
Laplace smoothing	β	t	Add-β smoothing for transition $\boldsymbol{A}$ and start prior $\boldsymbol{\pi}$ (counts+β, renormalize).	High: flatter $\boldsymbol{A}, \boldsymbol{\pi}$, safer on rare or omitted transitions, and less brittle; low: sharper priors that mirror train paths, but more risk of overfitting ordering quirks.	**3.0**: applied before reference-based transition masking.
Banding (local moves)	band	t	For each valid reference sequence, let i and j be the positions of two states in that sequence. Mask $A_{ss'}$ unless $\lvert i-j\rvert \leq 1$ in at least one valid reference, then renormalize each row; encodes locally ordered procedures.	On: restricts non-local transitions and produces more locally ordered paths; off: allows global jumps for branched workflows, but raises transposition false positives.	**On** with $\lvert i-j\rvert \leq 1$: disabling recommended only if tasks truly need non-local transitions.
Jitter probability	p_{jit}	t	Probability to add Gaussian noise to a training vector before PCA.	High: more invariance to ambience and lighting, but too much blurs class boundaries; low: crisper clusters, but more brittle to domain shift.	**0.15**: mild diversity without collapsing adjacent steps; co-tuned with σ.

Table 8: Core HMM hyperparameters for sequence modeling and training-time regularization. Tags indicate whether a hyperparameter is used during training (**t**) or inference (**i**).

Hyperparameter	Sym.	Tags	Meaning	Value effect	Default settings & rationale
Jitter std	σ	t	Noise scale in $\mathcal{N}(0, \sigma^2 I)$ before PCA.	High: more robust, but can distort likelihoods and increase class overlap; low: minimal regularization.	**0.3** in configurations that use jitter.
Per-modality normalization	l2/None	t/i	L2-normalize each modality before fusion and HMM; stabilizes scales and PCA.	l2: avoids dominance by high-variance channels and yields steadier thresholds; none: preserves raw magnitudes, but can skew fusion and Gaussian likelihoods.	**l2**: matches CLIP best practice; lets w modulate content, not scale.
Emission family	Gauss or Cosine	i	Observation model at inference: Gaussian log-likelihoods vs. cosine similarity (with temperature).	Gauss: calibrated log-likelihoods and good use of PCA plus Ledoit–Wolf; cosine: magnitude-invariant, but needs a temperature and may require transition rebalancing.	**Gauss** default; switch to **Cosine** for strongly angular features or when skipping PCA.
Cosine temperature	T	i	Scale for cosine emissions ($E = T \cdot \cos$); controls emission/transition balance.	High: more emission-dominant, peak-following paths; low: more transition-dominant, smoother paths.	**N/A** for Gauss. For Cosine, tuned T to match the dynamic range of log-transitions.
Emission scaling	α	i	Global multiplier $\tilde{\boldsymbol{E}} = \alpha \boldsymbol{E}$ to retune emission vs. transition influence.	High: more emission-driven Viterbi, hence more sensitive to per-step scores; low: more transition-driven, smoother, and more prior-driven.	**1.0**: adjustment needed only if domain shift skews score ranges.
Semantic substitution threshold	τ_{sem}	i	Cosine-similarity threshold between an observation and the mean embedding of its decoded state. A lower score produces a semantic Substitution when structural alignment retains the same Step ID.	High: detects subtler argument-level changes but may increase false positives; low: is more conservative but may miss modified executions.	**0.85** in all reported experiments.

Table 9: Core HMM hyperparameters for normalization and emission control. Tags indicate whether a hyperparameter is used during training (**t**) or inference (**i**).

Hyperparameter	Sym.	Tags	Meaning	Value effect	Default settings & rationale
Emission anomaly threshold	τ_{em}	i	Flag if the selected-state emission score is below τ_{em}.	Higher thresholds flag more observations; lower thresholds are more conservative.	-10 for Gaussian emissions and 0 for cosine emissions.
Transition anomaly threshold	τ_{tr}	i	Flag if $\log A_{S_{t-1},S_t} < \tau_{\mathrm{tr}}$.	Higher thresholds flag more transitions; lower thresholds are more conservative.	-23.5 for CaptainCook4D and -5 for EgoPER and EgoOops.
Alignment mode	NW or difflib	i	Align prediction to a canonical variant to label Deletion, Insertion, Substitution, or Transposition.	NW: stable edit metric (use sub=2, gap=1 for reports). difflib: forgiving for narratives and qualitative inspection.	**difflib** is our default in the main experiments because it performs well across grounded settings; **NW** remains competitive and is preferred in some paper-facing or family-specific comparisons.
Global fallback HMM	task$= -1$	i	Task-agnostic model pooled over all tasks and available when task-specific parameters are missing.	Provides an implementation fallback when no task-specific model is available.	Implemented but not invoked in the reported experiments.
Vision weight (mm)	w	mm (t/i)	Fusion in joint space: $\boldsymbol{x} = w\,\boldsymbol{v}_{img} + (1-w)\,\boldsymbol{v}_{txt}$ (or scaled-concat in disjoint space).	High: vision-led. Low: text-led; better when phrasing disambiguates or clarifies visuals.	**Task-dependent**. Optimal $w\approx$0.2–0.4: for subtle manipulations; increase when distinctive visuals drive the step.
Use superstructure (mm)	on/off	mm (t/i)	Add structured phrases (objects/actions) as text before fusion.	On: boosts separability for visually similar steps; sensitive to noisy phrases. Off: simpler; may blur key discriminators.	**Off** for clean baselines; **On** in V + AO style variants with reliable extractors and prompt filtering.
Superstructure top-k	k	mm (t/i)	Keep top-k phrases by confidence to control context breadth.	High: broader context, more recall but more distractors. Low: focused cues; risk missing secondary hints.	$k = 3$–5: small slate gives gains without flooding the encoder.
Superstructure threshold	τ	mm (t/i)	Minimum confidence for phrases included in superstructure.	High: precise/sparse cues. Low: recall-oriented; may add noise and hurt calibration.	**Moderate** τ: to be tuned jointly with k for coverage vs. noise.

Table 10: Additional HMM hyperparameters for anomaly scoring, alignment, and multimodal variants. Tags indicate whether a hyperparameter is used during training (**t**), inference (**i**), or in a multimodal setting (**mm**); entries like **mm (t/i)** denote multimodal settings with separate training/inference use.

C Semantic representation construction

Each step is mapped to an SRL representation of the form `PREDICATE(Role: value, Role: value, ...)`. The representation makes the main event, its participants, and key attributes explicit. We use nested structures when the step meaning requires internal attachment, scope, or composition. Table 7 gives examples from all three benchmarks.

We generate SRL representations offline with two LLM generators: GPT-5.3 (OpenAI, 2025) and Qwen3.6-27B (Qwen Team, 2026b). Both generators use the same constrained few-shot prompt and the same output fields. This keeps the target SRL interface fixed across models. The model-specific differences are in decoding and output enforcement: GPT-5.3 is queried with strict JSON-schema decoding at temperature 0.1, while Qwen3.6-27B is run with thinking disabled, temperature 0.0, top-$p = 0.8$, and top-$k = 20$, after which we parse the returned JSON object.

The generation process covers three vocabularies: (1) step descriptions from correct videos, (2) modified descriptions for erroneous steps, and (3) VLM-predicted descriptions that are not already covered by the first two sets.

Both generators use the same predicate–argument interface, but they produce different formalisms. For example, GPT-5.3 tends to preserve the surface predicate, while Qwen3.6-27B more often rewrites the instruction into a normalized action representation.

The two generators also differ in role inventory: GPT-5.3 uses 98 distinct role labels, while Qwen3.6-27B uses 71. Several GPT-5.3 labels correspond to distinctions that Qwen3.6-27B encodes with more general roles or nested modifiers.

For example, GPT-5.3 sometimes uses `Setting`, `Surface`, `Position`, and `Container`, while Qwen3.6-27B usually uses `Location`, `Destination`, or nested entity structure. Similarly, GPT-5.3 separates labels such as `Duration`, `Range`, and `Approx`; Qwen3.6-27B usually places these distinctions inside `Temporal` or `Measurement`. Quantity-related labels such as `Fraction`, `Portion`, and `Amount` are also more often collapsed into `Quantity` or `Measurement` in Qwen3.6-27B. Table 11 gives more representative examples.

Manual inspection calibrates the prompt and role schema. We inspect preliminary outputs for role consistency, attachment, lexical form, and coordination scope. Recurring issues lead to a schema clarification or an additional few-shot example. This iterative process produced the final nine-example prompt.

After the final inspection, each generator regenerates the complete dictionaries automatically. No roles or argument structures in the final dictionaries are edited manually. Scripted checks verify valid JSON, required fields, and balanced delimiters.

Prompt calibration uses ontology descriptions without access to split membership, mistake labels, HMM predictions, or evaluation results. Each unique description is converted once, independent of its video occurrence or dataset split.

We apply the same calibrated prompt to all unique descriptions across the three datasets.

Table 13 summarizes the core prompt design and output schema.

D HMM hyperparameters and computational details

Tables 8–10 define the supported HMM parameters. Section E reports the settings used in each experiment.

Computational details. Experiments were run on an institutional compute node with one NVIDIA A100 (80 GB) GPU and an AMD EPYC 7282 CPU (8 cores and 56 GB RAM allocated per job). HMM training and inference are CPU-bound, and a per-task HMM trains in seconds; the cost of the experiments comes from configuration sweeps. The complete multimodal sweeps on CaptainCook4D, the largest benchmark, consumed $\approx$57 node-hours ($\approx$454 CPU-core-hours). GPU compute is required for the frozen encoders and the locally hosted open-weight models: video-to-text grounding with Qwen 3.5-27B took roughly 24 GPU-hours on the largest benchmark and less on the other two, and SRL generation with Qwen3.6-27B for all three benchmarks took 2.65 GPU-hours (batch size 1, reasoning mode disabled). Across all benchmarks, representation families, and configuration sweeps, the reported experiments consumed several hundred node-hours in total. Gemini 3 Flash and GPT-5.3 were accessed through their public APIs. We do not train or fine-tune any neural model; parameter counts of the open models are given in the cited papers and model names, and those of the proprietary models are not publicly disclosed.

GPT-5.3 label	Step description	GPT-5.3	Qwen3.6-27B
Setting	Microwave the ramekin cup uncovered on high for 10 seconds	`Setting: high`	`Manner: uncovered & on(high)`
Surface	Brush 2 slices of baguette with olive oil on both sides	`Surface: both_sides`	`Location: on(both_sides)`
Position	Put three copper plates on the left column of the microplate	`column(of(microplate), Position: left)`	`left_column( of(microplate))`
Container	In the mug, whisk two eggs	`eggs(Quantity: 2, Container: mug)`	`eggs(Quantity: 2), Location: in(mug)`
Duration	Microwave the plate, uncovered, on high for 5 minutes	`Duration: minute(Quantity: 5)`	`Temporal: for(minute(Quantity: 5))`
Range	Continue to microwave for 15–30 more seconds	`seconds(Range: 15-30)`	`seconds(Quantity: 15-30)`
Approx	Slice one tomato into about 1/2 inch thick slices	`Thickness: inch(Quantity: 0.5, Approx: about)`	`Thickness: about(1/2(inch))`
Portion	Slice 1/4 medium onion into pieces	`Portion: quarter(Quantity: 1)`	`Quantity: 1/4`
Fraction	Add 1/4 teaspoon cinnamon to the mug	`teaspoon(Fraction: 1/4)`	`teaspoon(Quantity: 1/4)`
Amount	Pour a small amount of milk into the filter	`Amount: small`	`Quantity: small_amount`
Alternative_to	Squeeze out a drop of a green highlighter into a cup and pour much water instead of about 15mL of water	`Alternative_to: water(Measurement: ml(Quantity: 15, Approx: about))`	`Manner: instead_of( about(15mL(of(water))))`
Contrast	Add tomatoes to the pan, but splash some tomatoes outside	`Contrast: BUT(SPLASH(...))`	`ADD(...) & SPLASH(...)`

Table 11: Examples of role-inventory differences between GPT-5.3 and Qwen3.6-27B. GPT-5.3 uses some low-frequency labels that Qwen3.6-27B often represents with broader roles, nested modifiers, or explicit event composition.

E Reported configurations

This section reports the configurations used for (1) the family-level leaderboard in Table 1, (2) the matched representation and decoder studies in Tables 2 and 3, and (3) the PRISM rows in the system-level comparison in Table 4.

E.1 Family-level configurations

Table 14 reports the configurations used for the family-level leaderboard. CaptainCook4D and EgoPER configurations were selected by validation Any-F1; Type-F1, AUC, and Seq break ties in that order. EgoOops uses the configuration with the highest mean validation Any-F1 across CaptainCook4D and EgoPER for each representation family. EgoOops error-video labels do not enter selection. All test metrics in a leaderboard row come from the same configuration.

E.2 Representation comparison configurations

Table 15 lists the fixed downstream configurations for the matched Raw–GPT-SRL comparisons in Table 2. Each configuration applies to both VLMs within one dataset.

E.3 Prior-work comparison configurations

Table 16 reports the configurations used for the PRISM rows in Table 4. CaptainCook4D and EgoPER configurations were selected by validation EDA. EgoPER's validation split contains only correct videos, so its validation EDA measures the proportion of correct segments accepted as normal and does not measure error recall. Test metrics were not used for selection. All benchmark-specific metrics within each row come from the same configuration.

F Additional results and diagnostics

F.1 Grounding accuracy

Table 17 reports Exact, the proportion of evaluated segments whose generated description matches the target description word for word. Seq instead measures Step-ID recovery after Needleman–Wunsch alignment, as defined in Section 4. A non-verbatim

Difference	Step description	GPT-5.3	Qwen3.6-27B
Predicate choice	Use a butter knife to scoop nut butter from the jar; Use scissors to cut the bag in half	`USE(Object: butter_knife, Purpose: SCOOP(Object: nut_butter, Origin: jar)); USE(Object: scissors, Purpose: CUT(Object: bag, Result: in_half))`	`SCOOP(Object: nut_butter, Instrument: butter_knife, Origin: jar); CUT(Object: bag, Instrument: scissors, Result: in(half))`
Alternative action	Chop or grate the cucumber	`CHOP_OR_GRATE(Object: cucumber)`	`CHOP(Object: cucumber) | GRATE(Object: cucumber)`
Event decomposition	Measure and add 2 tbsp water to the bowl	`ADD(Object: water(Measurement: tablespoon(Quantity: 2)), Destination: bowl, Manner: MEASURED)`	`MEASURE(Object: water(Measurement: tablespoon(Quantity: 2))) & ADD(Object: water(...), Destination: to(bowl))`
Predicate normalization	Slice mushrooms	`SLICE(Object: mushrooms)`	`CUT(Object: mushrooms, Result: slices)`
State change	Return to low heat	`RETURN(Destination: to(low_heat))`	`SET(Object: heat, Result: low)`
Error-event scope	Spill while cracking an egg in the bowl	`SPILL(Temporal: WHILE(CRACK(Object: egg, Destination: in(bowl))))`	`CRACK(Object: egg, Destination: bowl, Manner: spill)`
Negation	Can't fold along the incisions of the cardboard	`FOLD(Object: cardboard, Path: along(incisions(of(cardboard))), Polarity: negative, Ability: can_not)`	`NEG(CAN(Action: FOLD(Object: cardboard, Path: along(incisions(of(cardboard))))))`
Incorporated argument	Pepper to taste	`SEASON(Instrument: pepper, Manner: to_taste)`	`ADD(Object: pepper(Measurement: to_taste))`
Coordination	Extract and add contents of an egg to a mug	`ADD(Object: contents(of(egg)), Destination: mug, Purpose: EXTRACT(Object: contents(of(egg))))`	`EXTRACT(Object: contents(of(egg))) & ADD(Object: contents(of(egg)), Destination: to(mug))`

Table 12: Differences between SRL representations generated by GPT-5.3 and Qwen3.6-27B. The two models use the same predicate–argument syntax but differ in predicate choice, role assignment, event decomposition, and scope representation. Examples are shortened where needed.

description can therefore fail Exact while still producing the correct Step ID. Table 19 reports positionwise Step-ID accuracy within grounding strata, without sequence alignment.

F.2 Second SRL generator

Table 18 repeats the matched grounded comparison with Qwen3.6-27B SRLs. Qwen-SRL improves five of six Raw conditions. GPT-SRL gives higher Any-F1 in five conditions. The two generators tie after rounding on CaptainCook4D with Gemini grounding. Gains occur with both SRL generators, although their magnitude varies across conditions.

F.3 Stratified SRL analysis

Table 19 stratifies the matched Viterbi outputs to locate the Step-ID changes caused by GPT-SRL. For this diagnostic, each decoded ID is compared with the annotated ID of the same input segment before sequence alignment. Exact contains word-for-word target matches. Same ID contains non-verbatim descriptions associated with the annotated ID. Other ID contains descriptions associated with another ontology ID. Ambiguous/open contains strings associated with multiple IDs or absent from the ontology. Modified is an overlapping subset whose target annotation contains a modified description. The analysis uses annotations only to define diagnostic strata. This positionwise diagnostic differs from the alignment-aware Seq metric reported in Table 27.

F.4 Performance by error type

Table 20 reports type-specific F1 and support for all twelve matched grounded conditions in Table 2. The Support columns report $TP + FN$. A single value applies to both Raw and GPT-SRL. Raw/GPT-SRL values are shown when support differs after reference selection. EgoOops has no gold Deletion or Transposition instances. Their F1 is zero under the evaluator's zero-division convention.

For each decoded Step-ID sequence, the evaluator selects the valid reference with the lowest alignment cost; the number of ordering inversions breaks ties. The selected reference defines both the predicted structural edits and the structural targets for that condition. Different predictions may therefore select different valid references, so structural-error support may vary across conditions. Insertions and deletions correspond to unmatched executed and reference steps. A transposition pairs an insertion and deletion of the same Step ID at different alignment positions. Substitution targets come from the benchmark annotations after the mapping in Appendix B. Any-F1 uses the same reference-selection and error-construction procedure.

F.5 Bootstrap confidence intervals

We estimate uncertainty for the six matched Raw–SRL comparisons in Table 2 with a paired video-level bootstrap. We sample test videos with replacement and recompute Any-F1 from the pooled step-

Component	Content
System instruction	You are a linguistic semantic analyzer. For each procedural step description, generate the semantic representation in roles.
Prompt guidance	Predicates are uppercase. Concrete manipulated entities are usually encoded under `Object`. `Agent: you` is used unless another agent is explicitly specified. Entity names are normalized as lowercase_with_underscores. Nested structures are allowed when needed, e.g., `tip(of(red_highlighter))` or `water(Location: in(cup))`. The prompt also gives typical role names such as `Object`, `Coobject`, `Destination`, `Location`, `Instrument`, `Manner`, `Purpose`, `Temporal`, `Origin`, and `Result`.

Selected few-shot examples	`INSERT(Agent: you, Object: test_swab, Destination: into(nostril(of(her))));`
	`ADD(Agent: you, Object: coffee_grounds, Origin: bowl, Destination: filter(Location: in(french_press)));`
	`ADD(Agent: you, Object: cut(onions), Coobject: egg(Location: in(mixing_bowl)));`
	`TOP(Agent: you, Object: cup, Coobject: with(salsa(Measurement: tablespoon(Quantity: 1))));`
	`POUR(Agent: you, Object: egg(Quantity: 1), Destination: into(ramekin_cup)).`

Output schema	Strict JSON with one item per step, containing the fields `id`, `step_description`, and `semantic_representation`.
Prompt calibration and checks	Iterative manual inspection calibrates the role schema and nine-example prompt. Each generator then regenerates the complete dictionaries automatically. Scripted checks verify valid JSON, required fields, and balanced delimiters. No final role or argument structure is manually edited.
Covered vocabularies	Gold canonical descriptions, modified descriptions, and VLM-predicted descriptions not covered by the first two vocabularies. Each unique description is processed once. At grounded inference, the normalized VLM-generated description is the only SRL lookup key.

Table 13: Core design of the prompt and schema used to generate SRL representations.

level counts. Raw and GPT-SRL use the same sampled videos in each of 10,000 replicates. Table 21 reports percentile 95% confidence intervals for both conditions and their paired difference.

The interval for the SRL gain excludes zero for Qwen on EgoOops and for both VLMs on EgoPER. It includes zero for Gemini on EgoOops and for both CaptainCook4D comparisons. The point estimate remains positive in all six comparisons. These intervals quantify uncertainty within PRISM. They do not establish statistical superiority over the published systems in Table 4, whose per-video predictions are unavailable.

F.6 Multimodal fusion sweep

Table 22 reports the multimodal weight sweep. In the multimodal setting, we combine a visual step embedding with its paired text embedding by weighted interpolation, sweeping from text-dominant to vision-dominant mixtures over a fixed grid of weights to test how much visual signal can be added before HMM sequence recovery and mistake recognition degrade.

The result pattern is consistent with the family-level results. Text-heavy mixtures are strongest on EgoOops. On CaptainCook4D, increasing the visual weight gives a small Any-F1 gain but reduces Seq and AUC. On EgoPER, larger visual weight raises Seq and AUC but leaves Any-F1 nearly flat. We therefore do not present multimodal fusion as the universal solution. In this HMM setting, textual embeddings anchor the state space more sharply than frozen visual features.

F.7 Multimodal results with Qwen3.6-27B SRL

Table 23 compares multimodal runs that use GPT-SRL and Qwen-SRL. The runs are matched by dataset, VLM text source, visual encoder, feature dimension, fusion weights, emission model, and alignment method. The table reports mean deltas, computed as Qwen-SRL minus GPT-SRL.

Across the matched multimodal runs, the mean differences between GPT-SRL and Qwen-SRL are close to zero on all three datasets. Neither generator has a consistent average advantage under the evaluated configurations. Individual configurations may respond differently.

G Additional qualitative examples

This appendix illustrates the qualitative patterns behind the main quantitative results. We focus on three questions: when SRL helps over grounded text without SRL, what the action–object abstraction preserves and loses, and how the unified edit-based view applies across datasets.

Dataset	Family	Encoder	d	Norm	Emission	Temp	Align
EgoOops	Vision	V-JEPA	1024	L2	Gaussian	–	NW
	Gold text	CLIP	128	none	cosine	6	NW
	Gold + GPT-SRL	CLIP	128	none	cosine	12	difflib
	Qwen text	Instructor-XL	64	L2	Gaussian	–	difflib
	Qwen + GPT-SRL	Instructor-XL	64	none	Gaussian	–	difflib
	Gemini text	all-MPNet	768	L2	Gaussian	–	difflib
	Gemini + GPT-SRL	Instructor-XL	64	none	Gaussian	–	difflib
EgoPER	Vision	VideoMAE-v2	128	L2	cosine	6	difflib
	Gold text	CLIP	256	none	Gaussian	–	difflib
	Gold + GPT-SRL	RoBERTa-SimCSE	256	L2	Gaussian	–	NW
	Qwen text	CLIP	256	L2	Gaussian	–	NW
	Qwen + GPT-SRL	RoBERTa-SimCSE	256	none	Gaussian	–	NW
	Gemini text	CLIP	256	L2	Gaussian	–	NW
	Gemini + GPT-SRL	RoBERTa-SimCSE	128	L2	Gaussian	–	NW
CaptainCook4D	Vision	VideoMAE-v2	768	none	cosine	6	difflib
	Gold text	RoBERTa-SimCSE	256	none	cosine	10	NW
	Gold + GPT-SRL	RoBERTa-SimCSE	256	none	cosine	10	difflib
	Qwen text	all-MiniLM	256	L2	cosine	12	difflib
	Qwen + GPT-SRL	RoBERTa-SimCSE	256	L2	cosine	10	difflib
	Gemini text	RoBERTa-SimCSE	128	none	cosine	12	difflib
	Gemini + GPT-SRL	RoBERTa-SimCSE	256	none	cosine	10	difflib

Table 14: Configurations used for the family-level leaderboard in Table 1. CaptainCook4D and EgoPER configurations are selected by validation Any-F1. EgoOops configurations are transferred from the best mean validation result across these two datasets. EgoOops error-video labels do not enter selection.

Dataset	Encoder	Output d	PCA d	Norm	Emission	Temp.	Align
EgoOops	Instructor-XL	768	128	L2	cosine	8	NW
EgoPER	RoBERTa-SimCSE	1024	128	L2	Gaussian	–	NW
CaptainCook4D	RoBERTa-SimCSE	1024	256	L2	cosine	10	difflib

Table 15: Fixed downstream configurations for the matched Raw–GPT-SRL comparisons and decoder controls. Output d is the encoder output dimension; PCA d is the dimension after projection.

G.1 Error-relevant semantic distinctions

Table 24 lists error-relevant distinctions that SRL can represent explicitly, including relations, destinations, materials, and nested events.

G.2 What action–object preserves and loses

Table 25 shows what the action–object abstraction preserves and what it weakens. Our action–object representation retains the main action together with all visually involved objects, not only the semantic patient. Tools and contacted objects are therefore preserved, which makes the representation useful for some mistakes but too coarse for others.

EgoPER contributes a different qualitative pattern from EgoOops and CaptainCook4D. For EgoPER, the VLM receives the task-specific candidate-step vocabulary as grounding context. Its short and regular candidate descriptions limit lexical variation, so Raw and SRL remain close on this dataset. The more informative contrast in EgoPER may appear between richer textual representations and the compressed action–object abstraction, especially for manner-sensitive steps such as `Slowly pour the rest of water in circular motion`, `POUR(Agent: you, Object: rest(of(water)), Manner: slowly & in(circular_motion))`.

G.3 Examples of the unified error taxonomy

Table 26 shows how the unified taxonomy matches the procedural structure of the observed deviations across datasets.

H Decoder Ablations

We use the matched inputs and alignment methods from Table 15. The nearest–greedy comparison adds PCA and the trained emission model. The greedy–Viterbi comparison changes only the decoding rule.

Independent nearest-candidate matching (no HMM). Let x_t denote the description generated by the VLM for video segment t. Let $\mathcal{S}_a$ denote the candidate step set for task a, and let d_s denote

Dataset	VLM	Representation	Encoder	PCA d	Norm	Emission	Align
EgoPER	Qwen	Raw	all-MiniLM	256	L2	Gaussian	difflib
	Qwen	GPT-SRL	RoBERTa-SimCSE	64	none	Gaussian	difflib
	Gemini	Raw	all-MiniLM	256	L2	Gaussian	difflib
	Gemini	GPT-SRL	RoBERTa-SimCSE	64	none	Gaussian	NW
CaptainCook4D	Qwen	Raw	RoBERTa-SimCSE	256	none	Gaussian	difflib
	Qwen	GPT-SRL	RoBERTa-SimCSE	256	L2	Gaussian	difflib
	Gemini	Raw	RoBERTa-SimCSE	256	L2	Gaussian	difflib
	Gemini	GPT-SRL	RoBERTa-SimCSE	256	L2	Gaussian	difflib

Table 16: Configurations used for the PRISM rows in the system-level comparison in Table 4. PCA d is the projected dimensionality.

Dataset	Qwen Exact	Gemini Exact	Δ
EgoOops	0.425	**0.496**	+0.070
EgoPER	0.806	**0.864**	+0.057
CaptainCook4D	0.578	**0.695**	+0.117

Table 17: Word-for-word grounding accuracy before Step-ID decoding. Δ is Gemini minus Qwen; differences are computed before rounding.

the canonical description of step s. In the GPT-SRL condition, both x_t and d_s are replaced by their semantic role representations.

The text encoder e encodes the generated description and every candidate description. All embeddings are L2-normalized. Each segment is assigned independently to the candidate with the highest cosine similarity:

$$\hat{s}_t = \arg\max_{s \in \mathcal{S}_a} \cos(e(x_t), e(d_s)) \,.$$

This decoder uses no HMM checkpoint, start probabilities, transition probabilities, or trained emission parameters. It assigns each Step ID without information from adjacent segments. Candidate cosine similarity supplies the LowProbEmission score; LowProbTransition is disabled because this decoder has no transition model. The semantic-substitution check compares the observation with the assigned candidate embedding rather than a learned state mean. The predicted sequence $\hat{s}_{1:T}$ enters the same edit-alignment and error-detection modules as PRISM. Sequence information enters only at the alignment stage.

This baseline isolates the performance provided by the frozen language representation, the task-specific candidate vocabulary, and edit alignment, without any trained HMM component. The Raw–SRL difference measures the contribution of SRL without HMM training or sequence priors. The procedure is deterministic once the VLM descriptions and text encoder are fixed.

Greedy emission decoding (without transition modeling). PRISM obtains each observation vector as

$$z_t = P(N(e(x_t))) \,,$$

where e is the text encoder, N is the configured normalization, and P is the fitted PCA projection. It then computes an emission score $E_{t,s}$ for every segment t and state s. The complete HMM applies the Viterbi objective in Equation 3 to the emission matrix $\boldsymbol{E}$.

The greedy decoder uses the same observation vectors and the same trained emission model. It assigns each segment independently:

$$\hat{s}_t^{\text{greedy}} = \arg\max_{s \in \mathcal{S}_a} E_{t,s}.$$

In implementation, this operation applies a row-wise argmax to the emission matrix $E \in \mathbb{R}^{T \times |\mathcal{S}_a|}$. The decoder removes the start probabilities, transition probabilities, and Viterbi search. It retains the encoder, PCA projection, normalization, trained emission parameters, and all downstream settings. The resulting Step ID sequence is passed to the same edit-alignment module.

This ablation isolates the contribution of global sequence decoding. The comparison between nearest-candidate matching and greedy decoding measures the combined contribution of the fitted PCA projection and trained local emission model. The comparison between greedy and Viterbi decoding measures the contribution of the sequence prior.

Dataset	Grounding	Raw	GPT-SRL	Qwen-SRL
EgoOops	Qwen	0.321	**0.356**	0.341
	Gemini	0.349	**0.353**	0.341
EgoPER	Qwen	0.578	**0.601**	0.584
	Gemini	0.602	**0.631**	0.613
CaptainCook4D	Qwen	0.587	**0.592**	0.588
	Gemini	0.595	**0.596**	**0.596**

Table 18: Matched Any-F1 with GPT-5.3 and Qwen3.6-27B SRLs. Every row uses the fixed downstream configuration reported in Table 15.

Dataset	VLM	Exact	Same ID	Other ID	Ambiguous/open	Modified
EgoOops	Qwen	+0.021 (145)	+0.095 (21)	-0.006 (170)	-0.200 (5)	+0.034 (59)
EgoOops	Gemini	+0.000 (169)	+0.000 (19)	+0.020 (153)	–	+0.000 (59)
EgoPER	Qwen	+0.014 (1771)	-0.008 (121)	+0.000 (284)	+0.250 (20)	+0.054 (500)
EgoPER	Gemini	+0.006 (1897)	+0.000 (104)	-0.005 (182)	+0.077 (13)	+0.024 (500)
CaptainCook4D	Qwen	-0.004 (891)	-0.013 (80)	-0.013 (550)	+0.048 (21)	-0.006 (363)
CaptainCook4D	Gemini	+0.014 (1071)	+0.000 (130)	+0.013 (318)	-0.087 (23)	+0.025 (363)

Table 19: Change in positionwise decoded-ID accuracy from Raw to GPT-SRL within each grounding stratum, before sequence alignment. Support appears in parentheses. Modified overlaps the first four strata. A dash marks an empty stratum.

Configurations and evaluation. The nearest-candidate decoder uses the matched text encoder and alignment method. It operates on the native encoder embeddings with L2 normalization and raw cosine similarity. PCA dimension, HMM emission family, and emission temperature do not apply to this decoder.

The greedy decoder uses the complete matched PRISM configurations reported in Table 15. All settings remain fixed except the decoding rule.

Greedy decoding improves Any-F1 over nearest-candidate matching in all twelve conditions. Viterbi improves Any-F1 for both VLMs on CaptainCook4D and in three EgoOops conditions. It leaves all four EgoPER conditions unchanged: Gaussian emissions already select the same Step IDs before sequence decoding. It lowers Qwen Raw on EgoOops from 0.339 to 0.321, where transition estimation uses 21 correct videos across five tasks. On CaptainCook4D, Viterbi improves Any-F1 and lowers Seq. Seq measures Step-ID recovery; Any-F1 measures mistake recognition after alignment.

Dataset	Error type	Qwen grounding			Gemini grounding		
		Raw	GPT-SRL	Support	Raw	GPT-SRL	Support
EgoOops	Insertion	0.082	0.119	89	0.151	0.170	90
	Deletion	0.000	0.000	0	0.000	0.000	0
	Substitution	0.220	0.215	62	0.164	0.177	62
	Transposition	0.000	0.000	0	0.000	0.000	0
EgoPER	Insertion	0.203	0.189	199/195	0.185	0.208	209/207
	Deletion	0.742	0.720	307	0.778	0.777	306
	Substitution	0.349	0.430	354	0.384	0.442	354
	Transposition	0.093	0.117	66/64	0.100	0.133	53/57
CaptainCook4D	Insertion	0.000	0.000	4	0.000	0.000	4
	Deletion	0.167	0.218	71	0.208	0.257	71
	Substitution	0.137	0.146	238	0.146	0.175	238
	Transposition	0.279	0.281	366	0.298	0.306	366/367

Table 20: Type-specific F1 and evaluation support for all matched grounded conditions in Table 2. Support is $TP + FN$. Raw/GPT-SRL values are shown when support differs between conditions.

Dataset	VLM	Raw Any-F1 [95% CI]	GPT-SRL Any-F1 [95% CI]	Δ [95% CI]
EgoOops	Qwen	0.321 [0.252, 0.380]	0.356 [0.295, 0.407]	+0.036 [+0.005, +0.076]
	Gemini	0.349 [0.287, 0.414]	0.353 [0.279, 0.426]	+0.003 [−0.028, +0.035]
EgoPER	Qwen	0.578 [0.544, 0.611]	0.601 [0.565, 0.638]	+0.024 [+0.007, +0.041]
	Gemini	0.602 [0.569, 0.633]	0.631 [0.594, 0.668]	+0.030 [+0.011, +0.048]
CaptainCook4D	Qwen	0.587 [0.520, 0.643]	0.592 [0.525, 0.650]	+0.005 [−0.007, +0.017]
	Gemini	0.595 [0.530, 0.649]	0.596 [0.533, 0.649]	+0.001 [−0.013, +0.015]

Table 21: Paired video-level bootstrap 95% confidence intervals for Any-F1 in the six matched grounded comparisons. We use 10,000 replicates and percentile intervals. Δ is GPT-SRL minus Raw; each replicate uses identical sampled video IDs for both conditions. Differences are computed before rounding.

Weights $(w_{\text{img}}, w_{\text{txt}})$	EgoOops			EgoPER			CC4D		
	Seq	Any-F1	AUC	Seq	Any-F1	AUC	Seq	Any-F1	AUC
$(0.2, 0.8)$	0.651	0.437	0.558	0.235	0.478	0.441	0.598	0.547	0.604
$(0.4, 0.6)$	0.651	0.429	0.562	0.479	0.474	0.558	0.569	0.550	0.576
$(0.6, 0.4)$	0.651	0.429	0.563	0.629	0.463	0.597	0.538	0.554	0.561
$(0.8, 0.2)$	0.651	0.429	0.562	0.700	0.466	0.599	0.502	0.558	0.549

Table 22: Multimodal fusion sweep over interpolation weights between visual and textual step embeddings. For each dataset, we fix one visual–textual configuration and vary only the fusion weights.

Dataset	Matched runs	Δ Seq	Δ Any-F1	Δ Type-F1	Δ AUC
EgoOops	160	-0.0003	-0.0003	+0.0016	-0.0002
EgoPER	240	+0.0007	-0.0014	-0.0005	+0.0005
CaptainCook4D	240	+0.0007	+0.0002	-0.0006	-0.0004

Table 23: Matched multimodal comparison of GPT-SRL and Qwen-SRL. Deltas are computed as Qwen-SRL minus GPT-SRL. AUC denotes Step-AUC.

Dataset	Reference step	Observed / grounded variant	Semantic distinction
CC4D	Place the egg from the cup over the lettuce	Place the egg from the cup under the lettuce	Makes the relation contrast explicit: `over` vs. `under`.
	Invert the mug to release the cake onto a plate	Invert the mug to release the cake onto a bowl	Isolates destination and destination-property changes while preserving the same event structure.
		Invert the mug to release the cake onto an unclean plate	
	Invert the mug to release the cake onto a plate	Use a spoon to invert the mug to release the cake onto a plate	Preserves the added instrumental subevent as nested structure.
	Measure and add 1.5 tbsp sugar to the mixing bowl	Qwen: Measure and add 1.5 tbsp brown sugar to the mixing bowl	Preserves the `MEASURE & ADD` structure and highlights the ingredient mismatch.
EgoOops	Put three zinc plates on the center column of the microplate using the pair of tweezers	Qwen: Put three copper plates on the left column of the microplate using the pair of tweezers	Makes both material and location drift explicit.

Table 24: Examples of error-relevant distinctions that SRL can represent explicitly, including relations, destinations, properties, materials, and nested events. CC4D denotes CaptainCook4D.

Effect	Dataset	Reference step	Contrasting step	What action–object captures or loses
Preserves	EgoPER	Use knife to scoop jelly	Use spoon to scoop jelly	Preserves the tool distinction because tools remain in the object set.
	CC4D	Sprinkle 1 tablespoon of cheese on cup	Pour 1 tablespoon of cheese on cup	Preserves the predicate change.
Weakens	CC4D	Microwave just until cheese melts, about 10 seconds	Microwave for 30 seconds	Loses duration and completion-condition information.
	CC4D	Top cup with 1 tablespoon of salsa	Top cup with 1/2 tablespoon of salsa	Loses scalar quantity information.
	EgoPER	Slowly pour the rest of water in circular motion	–	Loses manner information.
	CC4D	Replace the top of the English muffin	Flip over the top muffin and replace it on top	Weakens multi-action structure.

Table 25: Qualitative examples of what the action–object abstraction preserves and what it weakens. It remains informative for predicate and tool changes, but loses detail when the mistake depends on duration, scalar quantity, manner, or multi-action structure. CC4D denotes CaptainCook4D.

Type	Dataset	Example	Why it belongs to this type
Insertion	EgoOops	Extra grasp actions such as `Grasp black light` or `Grasp a long blue thin cuboid block`	Extra steps outside the successful reference sequence.
	EgoPER	`Pour water from the second bowl to the first one`	Corrective action inserted after an earlier mistake.
	EgoOops	`Mix all the liquid and shine a black light on it`	Composite action inserted outside the reference sequence.
Deletion	EgoPER	`Put bananas` annotated as `Skipped this step`	Reference step missing from the observed sequence.
Substitution	CC4D	`Place the egg from the cup over the lettuce` vs. `... under the lettuce`	Same step slot, wrong relation.
	CC4D	`Invert the mug to release the cake onto a plate` vs. `... onto a bowl`	Same step slot, wrong destination.
	CC4D	`Measure and add 1.5 tbsp sugar to the mixing bowl` vs. `... 1.5 tbsp brown sugar ...`	Same step structure, wrong ingredient.
Transposition	CC4D	`Sprinkle 2 tablespoon of cheese on cup` is performed before `Microwave for 10 seconds`	Wrong local order of otherwise valid steps.

Table 26: Qualitative examples of the unified edit-based taxonomy across benchmarks. The same four types remain interpretable despite different dataset-specific error labels and annotation styles. CC4D denotes CaptainCook4D.

			Nearest		Greedy		Viterbi	
Dataset	**VLM**	**Repr.**	**Seq**	**Any-F1**	**Seq**	**Any-F1**	**Seq**	**Any-F1**
EgoOops	Qwen	Raw	0.501	0.337	0.501	0.339	0.575	0.321
	Qwen	GPT-SRL	0.499	0.335	0.501	0.351	0.587	0.356
	Gemini	Raw	0.560	0.305	0.572	0.339	0.619	0.349
	Gemini	GPT-SRL	0.560	0.315	0.569	0.328	0.613	0.353
EgoPER	Qwen	Raw	0.803	0.547	0.803	0.578	0.803	0.578
	Qwen	GPT-SRL	0.816	0.524	0.816	0.601	0.816	0.601
	Gemini	Raw	0.849	0.571	0.849	0.602	0.849	0.602
	Gemini	GPT-SRL	0.854	0.557	0.854	0.631	0.854	0.631
CaptainCook4D	Qwen	Raw	0.630	0.567	0.631	0.581	0.499	0.587
	Qwen	GPT-SRL	0.630	0.559	0.630	0.578	0.485	0.592
	Gemini	Raw	0.782	0.538	0.783	0.564	0.542	0.595
	Gemini	GPT-SRL	0.785	0.528	0.783	0.561	0.555	0.596

Table 27: Decoder ablations under the matched configurations. Nearest-candidate matching uses no HMM. Greedy decoding retains the trained local decoder. Viterbi adds start and transition probabilities and global sequence search.