arXiv:2610.37213v1 [physics.med-ph] 27 Oct 2026

RETROSPECTIVE MACROMOLECULAR POOL FRACTION MAPPING FROM CONVENTIONAL CLINICAL MRI: A PHYSICS-GUIDED SELF-SUPERVISED DEEP LEARNING FRAMEWORK

A PREPRINT

Jelmer van Lune[1], **Miha Fuderer**[1], **Oscar van der Heide**[1], **Matteo Maspero**[1], **Martin B. Schilder**[1], **Jan Willem Dankbaar**[2], **Stefano Mandija**[1], **and Alessandro Sbrizzi**[1]

[1]Computational Imaging Group for MRI Therapy & Diagnostics, University Medical Center Utrecht, Utrecht, The Netherlands

[2]Department of Radiology, University Medical Center Utrecht, Utrecht, The Netherlands

ABSTRACT

Macromolecular pool fraction (MPF) is a quantitative imaging biomarker sensitive to myelin integrity and other changes in tissue macromolecular content, but its dependence on specialized acquisitions and reconstruction methods limits its availability for large-scale studies. We present a physics-guided, self-supervised deep learning framework for retrospectively estimating MPF directly from conventional clinical brain MRI contrasts (T_1-weighted, T_2-weighted, and fluid-attenuated inversion recovery [FLAIR]), without paired quantitative reference scans. The framework embeds an analytical two-pool magnetization transfer signal model that relates sequence-averaged radiofrequency transmit power to apparent longitudinal relaxation. To accommodate heterogeneous clinical protocols, the network conditions intermediate feature representations on acquisition parameters via adaptive instance normalization. The framework was trained and evaluated on a multisite, multivendor, multiprotocol dataset comprising 7,705 MRI sessions (3 T) from 4,489 subjects across Philips and Siemens systems. Across 1,237 test sessions, mean MPF was 0.182 ± 0.007 in white matter and 0.144 ± 0.009 in gray matter, with limited variation across vendor-by-protocol groups (between-group coefficient of variation: 2.65% in white matter and 5.63% in gray matter). In 70 sessions with demyelinating-disease labels, MPF was 35.6% lower in segmented lesions than in normal-appearing white matter. As a technical demonstration, MPF maps from two different clinical protocols acquired in a healthy volunteer showed a mean voxel-wise absolute percentage difference of 7.7% and a concordance correlation coefficient of 0.888, indicating good cross-protocol consistency. These results suggest that magnetization transfer-informed self-supervised learning enables stable estimation of MPF maps from conventional clinical MRI, providing a scalable approach to investigate demyelinating pathology in both existing imaging archives and prospectively acquired conventional clinical scans without dedicated qMT acquisitions.

Keywords Magnetization transfer · Macromolecular pool fraction · Quantitative MRI · Self-supervised learning · Multiple sclerosis · Myelin mapping

1 Introduction

Magnetic resonance imaging (MRI) provides excellent soft-tissue contrast and is routinely used for clinical diagnosis. Signal intensity in conventional MR images depends on tissue properties such as longitudinal and transverse relaxation times (T_1, T_2, T_2^*) and proton-density (PD). In addition, magnetization transfer (MT) is a fundamental tissue- and sequence-dependent effect that influences the MR signal [1]. MT is commonly described using a two-pool model comprising a free-water pool and a semisolid pool [2]. The semisolid pool contains protons bound to macromolecules, including proteins and lipids, and has a transverse relaxation time on the order of microseconds. It is therefore not directly observable with conventional MRI. However, exchange of longitudinal magnetization between the free and semisolid pools alters the longitudinal relaxation of the observable free-water signal and produces biexponential

relaxation. MT is consequently a major contributor to longitudinal relaxation in biological tissue [3]. The relative size of the semisolid pool is quantified by the macromolecular pool fraction (MPF, denoted m_0^{s}), which quantifies the relative amount of protons bound to biological macromolecules that exchange magnetization with water protons [4].

Among quantitative MT tissue parameters, MPF has attracted particular interest as a marker of myelin integrity in multiple sclerosis (MS), the most common demyelinating disease of the central nervous system [5–7]. MPF is sensitive to myelin content and is markedly reduced in MS lesions [3, 4, 6, 8–10]. This sensitivity has enabled investigations of spatial heterogeneity in demyelination and of lesion evolution over time [4, 9, 11–13]. Moreover, post-mortem studies associated reduced MPF with histopathological evidence of demyelination [4, 14, 15]. MPF mapping could therefore provide a quantitative imaging readout for studying demyelinating pathology, monitoring disease progression, and assessing treatment response.

Despite this clinical potential, MPF mapping requires dedicated prospective quantitative MT (qMT) acquisitions with specialized pulse sequences and reconstruction pipelines that are generally not included in standard neuroimaging protocols. Several qMT methods have been developed over the past three decades, including multi-point off-resonance Z-spectroscopy [2, 7, 10], rapid single-point mapping [16, 17], on-resonance selective inversion recovery [18, 19], steady-state saturation [20], and transient hybrid-state methods [3, 21, 22]. These approaches require additional acquisition time and can only be applied prospectively. In addition, they cannot recover MPF from the large collections of conventional scans that already exist in clinical archives. These limitations restrict access to MPF maps and impede both large-scale biomarker studies and subsequent clinical translation.

The widespread availability of routine clinical MRI provides an opportunity to derive quantitative tissue parameters without dedicated prospective acquisitions. Initial efforts to exploit these routine scans have relied on mathematical combinations of conventional contrasts—such as $T_1\mathrm{w}/T_2\mathrm{w}$ and $T_1\mathrm{w}/\mathrm{FLAIR}$ ratios or empirical polynomial fits—as accessible proxies for myelin integrity and myelin water fraction (MWF) [23–25]. However, these approaches provide relative, intensity-based indices rather than quantitative physical parameters.

More recently, self-supervised, physics-guided deep learning frameworks have begun to address these limitations by retrospectively estimating quantitative maps from routine multi-contrast acquisitions ($T_1\mathrm{w}$, $T_2\mathrm{w}$, and FLAIR) without requiring paired reference scans [26–28]. These frameworks embed analytical Bloch equations-based signal models directly into the training objective, enabling scalable deployment across clinical archives. Existing frameworks, however, assume single-pool relaxation and remain restricted to relaxometry parameters (T_1 and T_2) alongside PD. Consequently, the retrospective estimation of quantitative MT parameters, including the MPF, from conventional clinical MRI remains unexplored.

In this work, we introduce a framework for estimating an MT-model-based MPF directly from routine $T_1\mathrm{w}$, $T_2\mathrm{w}$, and FLAIR images without paired qMT reference data. We extended a physics-guided, self-supervised deep learning framework by incorporating a simplified closed-form description of MT into its forward signal model. Following Teixeira et al. [29], the model relates the MT-induced apparent longitudinal relaxation time to the tissue-dependent MPF and the average RF transmit power, $B_{1,\mathrm{rms}}$, the latter being determined by the acquisition protocol. To reduce protocol-dependent variation in these estimated maps, we conditioned the network on the input scans' acquisition parameters. We trained and evaluated the framework using a clinically heterogeneous, multicenter, multivendor dataset comprising more than 7,700 MRI sessions (23,000+ image volumes), including sessions from patients with demyelinating lesions. As a technical demonstration, we assessed cross-protocol consistency using repeated acquisitions in a healthy volunteer.

2 Theory

2.1 Two-pool magnetization transfer model

MT in biological tissue is generally modeled using a two-pool framework comprising a liquid free-water pool (superscript f) and a semisolid macromolecular pool (superscript s) [2]. Each pool is characterized by its own longitudinal and transverse relaxation rates. However, the transverse relaxation time of the semisolid pool, T_2^{s}, is on the order of microseconds, preventing its direct observation with conventional MRI. Longitudinal magnetization is exchanged between the two pools with first-order rate constants k^{f} (free to semisolid) and k^{s} (semisolid to free) [29].

In this work, we define the macromolecular pool fraction (MPF) as the fractional equilibrium magnetization of the semisolid pool relative to the total equilibrium magnetization. Here, M_0^{f} and M_0^{s} denote the equilibrium longitudinal magnetizations of the free and semisolid pools, respectively. The MPF is therefore defined as $m_0^{\mathrm{s}} = M_0^{\mathrm{s}}/(M_0^{\mathrm{f}} + M_0^{\mathrm{s}})$, such that $m_0^{\mathrm{f}} + m_0^{\mathrm{s}} = 1$. Note that this definition follows recent quantitative MT literature [3, 16, 22, 30]. In contrast, earlier qMT studies [2, 7, 18] frequently report the pool size ratio relative to the free pool, $F = M_0^{\mathrm{s}}/M_0^{\mathrm{f}}$. The two definitions are related by

$$m_0^{\mathrm{s}} = \frac{F}{1+F} \quad \Longleftrightarrow \quad F = \frac{m_0^{\mathrm{s}}}{1-m_0^{\mathrm{s}}}. \tag{1}$$

Following Assländer et al. [3], we express both rate constants through a single exchange rate R_x as $k^{\mathrm{f}} = R_x m_0^{\mathrm{s}}$ and $k^{\mathrm{s}} = R_x(1-m_0^{\mathrm{s}})$.

2.2 Sequence-averaged RF saturation approximation

Following the continuous-wave-equivalent formulation of Teixeira et al. [29], the semisolid-pool saturation rate can be expressed as

$$\langle \overline{W} \rangle = \pi \left(\gamma B_{1,\mathrm{rms}}\right)^2 G(\Delta), \tag{2}$$

where γ is the gyromagnetic ratio, $B_{1,\mathrm{rms}}$ is the root-mean-square transmit RF amplitude, and $G(\Delta)$ is the absorption lineshape of the semisolid pool at frequency offset Δ.

Teixeira et al. [29] evaluated Eq. (2) for short-TR steady-state acquisitions, defining $B_{1,\mathrm{rms}}$ over a single repetition period. To apply the same analytical framework to conventional clinical MRI, we summarized the cumulative RF power of all excitation, refocusing, and inversion pulses by a sequence-averaged $B_{1,\mathrm{rms}}$ computed over the entire acquisition. This extension assumes that the combined RF pulse rate is high relative to the longitudinal relaxation, exchange, and semisolid-pool saturation rates. Under this assumption, the discrete RF pulse train can be approximated by continuous irradiation characterized by the sequence-averaged RF power. Appendix A.1 provides further justification.

Assuming a (i) Gaussian absorption lineshape for the semisolid pool [2, 31], $G(\Delta) = (T_2^{\mathrm{s}}/\sqrt{2\pi})\exp[-(2\pi\Delta T_2^{\mathrm{s}})^2/2]$, and (ii) RF irradiation near the water resonance, the lineshape can be evaluated at its peak value, $G(\Delta) \approx G(0) = T_2^{\mathrm{s}}/\sqrt{2\pi}$. Eq. (2) then reduces to

$$\langle \overline{W} \rangle = \sqrt{\pi/2}\,\gamma^2 T_2^{\mathrm{s}} B_{1,\mathrm{rms}}^2. \tag{3}$$

Appendix A.2 further discusses the validity of the near-resonance assumption for the clinical acquisitions considered in this work.

2.3 MT-informed forward signal modeling

Under the continuous-wave-equivalent approximation of Teixeira et al. [29], the two-pool longitudinal exchange reduces to an effective single-pool relaxation process. Expressing the exchange rate constants in terms of R_x and m_0^{s} yields the apparent longitudinal relaxation rate:

$$R_1^{\mathrm{app}} = R_1^{\mathrm{f}} + R_x m_0^{\mathrm{s}} \left(1 - \frac{R_x(1-m_0^{\mathrm{s}})}{R_1^{\mathrm{s}} + R_x(1-m_0^{\mathrm{s}}) + \langle \overline{W} \rangle}\right), \tag{4}$$

where $\langle \overline{W} \rangle$ (Eq. (3)) couples sequence-specific RF power ($B_{1,\mathrm{rms}}$) directly to the observed longitudinal relaxation rate. Figure 1 illustrates the theoretical dependence of $T_1^{\mathrm{app}} = 1/R_1^{\mathrm{app}}$ on $B_{1,\mathrm{rms}}$ across representative biological macromolecular pool fractions (m_0^{s}).

By substituting R_1^{app} into established Bloch-based signal models [32], conventional acquisitions can be modeled in closed form while accounting for sequence-specific MT effects. For a given voxel, the expected signal intensity S is expressed as a function of tissue parameters (PD, $R_1^{\mathrm{f}} = 1/T_1^{\mathrm{f}}$, T_2^{f}, m_0^{s}) and acquisition parameters (repetition time (TR), echo time (TE), inversion time (TI), flip angle (FA), and sequence-averaged $B_{1,\mathrm{rms}}$):

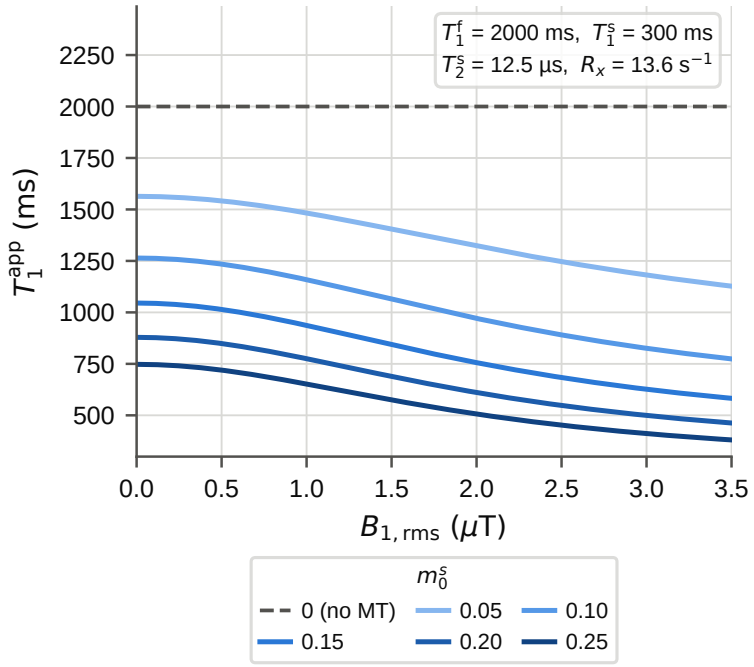

Figure 1: Theoretical behavior of apparent longitudinal relaxation time ($T_1^{\text{app}} = 1/R_1^{\text{app}}$) as a function of sequence-averaged RF amplitude ($B_{1,\text{rms}}$) across varying macromolecular pool fractions ($m_0^{\text{s}} \in [0, 0.25]$) according to Eq. (4). Fixed representative white matter parameters: $T_1^{\text{f}} = 2000$ ms, $T_1^{\text{s}} = 300$ ms, $T_2^{\text{s}} = 12.5$ μs, and $R_x = 13.6$ s^{-1}. The dashed line indicates the free-water longitudinal relaxation time (T_1^{f}) in the absence of MT ($m_0^{\text{s}} = 0$).

$$S_{\text{T1w,GRE}} = \text{PD} \cdot \sin(\text{FA}) \cdot \left[\frac{1 - \exp(-R_1^{\text{app}}\text{TR})}{1 - \cos(\text{FA})\exp(-R_1^{\text{app}}\text{TR})}\right] \cdot \exp(-\text{TE}/T_2^{\text{f}}), \tag{5}$$

$$S_{\text{T1w,MPRAGE}} = \left|\text{PD} \cdot \sin(\text{FA}) \cdot \left[\frac{1 - 2\exp(-R_1^{\text{app}}\text{TI}) + \exp(-R_1^{\text{app}}\text{TR})}{1 + \cos(\text{FA})\exp(-R_1^{\text{app}}\text{TR})}\right]\right| \cdot \exp(-\text{TE}/T_2^{\text{f}}), \tag{6}$$

$$S_{\text{T2w}} = \text{PD} \cdot [1 - \exp(-R_1^{\text{app}}\text{TR})] \cdot \exp(-\text{TE}/T_2^{\text{f}}), \tag{7}$$

$$S_{\text{FLAIR}} = |\text{PD} \cdot [1 - 2\exp(-R_1^{\text{app}}\text{TI}) + \exp(-R_1^{\text{app}}\text{TR})]| \cdot \exp(-\text{TE}/T_2^{\text{f}}). \tag{8}$$

Here, R_1^{app} is computed individually for each sequence via Eq. (4) using its respective $B_{1,\text{rms}}$, with fixed literature values assumed for the semisolid relaxation and exchange parameters (R_1^{s}, R_2^{s}, and R_x; detailed in Section 3.2.2). For the gradient-echo-based acquisitions, T_2^{f} was used as an approximation of $T_2^{*\text{f}}$.

Although the two-pool reduction formally also yields a sequence-dependent apparent equilibrium magnetization (M_0^{app}) [29], we represented its role in the signal models using a single sequence-independent relative PD map. Under this simplification, sequence-dependent MT effects are modeled only through R_1^{app}. The rationale for this simplification is discussed in Appendix A.2.

3 Methods

3.1 Dataset

3.1.1 Clinical dataset

We compiled a multicenter clinical dataset comprising 7,705 non-contrast-enhanced 3 T brain MRI sessions from 4,489 unique subjects from two different sources. The dataset combined (1) retrospectively collected examinations

from the University Medical Center Utrecht (UMCU) with (2) examinations from the publicly available MR-RATE dataset acquired at Istanbul Medipol University[1]. Each session included a 3D T_1-weighted (T_1w) acquisition, an axial 2D T_2-weighted (T_2w) acquisition, and a 3D T_2-weighted fluid-attenuated inversion recovery (FLAIR) acquisition. Gadolinium-enhanced T_1w acquisitions were excluded. The two data sources were combined to broaden the distribution of scanner systems, pulse-sequence implementations, and acquisition parameters, including $B_{1,\mathrm{rms}}$. Table 1 summarizes the acquisition characteristics of both datasets.

The UMCU dataset expanded upon the clinical dataset used in our previous work [26] by including additional examinations acquired with a broader range of protocols. The collected data were acquired between 2016 and 2023 and processed in accordance with institutional regulations and ethical guidelines. The resulting dataset comprised 5,312 MRI sessions from 2,192 unique subjects (age range, 6–91 years; mean $\pm$ standard deviation, 51.5 ± 17.6 years; 51.1% male and 48.9% female). The dataset included patients with a broad range of clinical conditions, including oncological, neurodegenerative, epileptic, and vascular disorders. Examinations were performed on four different Philips 3 T MRI systems: Achieva (33.4% of sessions), Ingenia (28.4%), Ingenia CX (20.6%), and Ingenia Elition X (17.6%). The T_1w acquisitions used either a 3D spoiled gradient-echo sequence (Spoiled GRE: 78.5%) or 3D magnetization-prepared rapid gradient-echo (MPRAGE; 21.5%). All T_2w and FLAIR acquisitions used 2D and 3D turbo spin-echo (TSE) sequences, respectively.

A subset of the MR-RATE data was selected based on the availability of the sequence parameters required for the forward signal model: TE, TR, inversion time TI, FA, and root-mean-square RF transmit field strength over the sequence ($B_{1,\mathrm{rms}}$). These criteria yielded 2,393 MRI sessions from 2,297 unique subjects, acquired between 2023 and 2026 (age range, 13–93 years; mean $\pm$ standard deviation, 42.0 ± 17.7 years; 43.2% male and 56.8% female). All selected examinations were acquired on a Siemens Magnetom Vida 3 T system using 3D MPRAGE for T_1w, 2D TSE for T_2w, and 3D TSE for FLAIR.

For each dataset, unique subjects were randomly assigned to the training, validation, and test sets using proportions of 80%, 5%, and 15%, respectively. All MRI sessions from the same subject were assigned to the same partition to prevent subject-level data leakage. We then combined the corresponding partitions from both datasets for model training, validation, and testing. The final training set comprised 6,084 sessions from 3,595 subjects, the validation set comprised 384 sessions from 225 subjects, and the test set comprised 1,237 sessions from 669 subjects. Table 2 summarizes the partition sizes.

MR-RATE provides pathology labels derived automatically from the accompanying radiology reports. Within the MR-RATE test set, 70 sessions from 65 subjects carried the label "demyelinating disease of the central nervous system" (mean age, 34.4 ± 11.0 years; 35.4% male and 64.6% female). Because diagnostic labels were not available for the UMCU dataset, the radiologist in our team (J.W.D.) performed a targeted review and identified four test-set subjects with clinically confirmed MS. This targeted review was not intended to identify all subjects with MS in the UMCU test set. Supplementary Table S1.1 provides the demographic and acquisition characteristics of these four subjects.

3.1.2 Data preprocessing

The UMCU MRIs were converted from DICOM to NIfTI format using dcm2niix [33]. We extracted sequence parameters from the DICOM headers. The MR-RATE dataset was already in NIfTI format, and we obtained the corresponding sequence parameters from the accompanying metadata files.

All MRI sessions from the combined dataset underwent a series of preprocessing steps. First, all T_2w acquisitions were resampled to an in-plane resolution of 1×1 mm while preserving their native slice thickness. Across the combined dataset, the native T_2w slice thickness was either 3.0 or 4.0 mm. Second, the 3D T_1w and FLAIR images were rigidly registered and resampled onto the voxel grid of the corresponding resampled 2D T_2w acquisition using ANTsPy [34]. Third, a brain mask was generated from the T_1w image using HD-BET [35], and all images were skull stripped using the resulting mask. Fourth, N4 bias-field correction [36] was applied to the skull-stripped T_1w, T_2w, and FLAIR images. Fifth, axial slices were center-cropped or zero-padded to a matrix size of 224×224 voxels. Finally, superior and inferior slices with low signal-to-noise ratio were excluded to reduce artifacts. Before training and inference of our proposed framework, each input image was independently normalized by dividing all voxel intensities by the histogram mode of the corresponding whole-brain volume.

To obtain tissue masks, we applied SynthSeg [37] to the preprocessed T_1w images from each MRI session. White matter (WM) and gray matter (GM) masks were extracted from the resulting segmentations. To reduce partial-volume contamination from cerebrospinal fluid (CSF), we excluded WM and GM voxels located within 2.0 mm of CSF.

[1] `https://huggingface.co/datasets/Forithmus/MR-RATE`

Table 1: Acquisition parameters for the UMCU and MR-RATE clinical datasets (training, validation, and test sets combined). Values denote median [min–max]; single values indicate dataset-wide constants. GRE: Spoiled Gradient-Echo; MPRAGE: Magnetization-Prepared Rapid Gradient-Echo; TSE: Turbo Spin Echo; IR: Inversion Recovery; NA: Not Applicable; FA: Flip Angle; ETL: Echo Train Length; BW: Pixel Bandwidth.

UMCU (Philips)

	T_1w		T_2w	FLAIR
Sequence	3D Spoiled GRE	3D MPRAGE	2D TSE	3D TSE (IR)
Sessions (n)	4,168	1,144	5,312	5,312
TR [ms]	5.26 [5.03–5.40]	2500	4063 [2779–4918]	4800
TE [ms]	2.38 [2.25–2.49]	3.94 [3.53–4.09]	80	314.5 [275.0–370.3]
TI [ms]	NA	1004	NA	1650
FA [°]	10	8	90	90
$B_{1,\mathrm{rms}}$ **[μT]**	0.93 [0.92–0.95]	0.73 [0.73–0.91]	2.34 [2.26–2.34]	1.00 [1.00–1.00]
ETL	NA	NA	15 [15–18]	199
BW [Hz/px]	539 [539–543]	190 [189–212]	195 [177–548]	1201 [863–1617]
Matrix (in-plane)	480^2 [480^2–560^2]	240^2 [224^2–288^2]	512^2 [448^2–560^2]	240^2 [224^2–288^2]
Slices (n)	360 [306–480]	160 [148–240]	37 [18–46]	267 [229–333]

MR-RATE (Siemens)

	T_1w	T_2w	FLAIR
Sequence	3D MPRAGE	2D TSE	3D TSE (IR)
Sessions (n)	2,393	2,393	2,393
TR [ms]	2000 [1780–2080]	5000 [3000–10350]	7000 [5200–7000]
TE [ms]	2.37 [2.20–4.71]	95 [89–106]	388 [386–393]
TI [ms]	920 [920–1050]	NA	2050
FA [°]	9 [8–9]	150 [119–150]	120
$B_{1,\mathrm{rms}}$ **[μT]**	0.78 [0.76–0.92]	1.96 [1.61–2.64]	1.11 [1.10–1.28]
ETL	NA	17 [13–20]	220 [218–220]
BW [Hz/px]	248 [199–248]	220 [216–252]	751 [651–751]
Matrix (in-plane)	320^2 [256^2–640^2]	512^2 [320^2–1024^2]	512^2 [512×432–568×512]
Slices (n)	176 [144–320]	33 [16–58]	176 [144–208]

Table 2: Subject-level partitioning of the UMCU, MR-RATE, and combined clinical datasets. Each source dataset was independently partitioned using an 80/5/15% training, validation, and test split.

Dataset	**Level**	**Training (80%)**	**Validation (5%)**	**Test (15%)**	**Total**
UMCU (Philips)	Sessions	4,173	265	874	5,312
	Subjects	1,755	110	327	2,192
MR-RATE (Siemens)	Sessions	1,911	119	363	2,393
	Subjects	1,840	115	342	2,297
Combined	Sessions	6,084	384	1,237	7,705
	Subjects	3,595	225	669	4,489

For the demyelinating lesion analyses, we segmented WM lesions using LST-AI, a deep learning model developed and trained for MS lesion segmentation on paired T_1w and FLAIR images [38]. Segmentation was performed for the four manually identified UMCU subjects with MS and the 70 MR-RATE MRI sessions with the demyelinating disease label. LST-AI was applied to the native-resolution, full-head T_1w and FLAIR acquisitions selected from each MRI session before any preprocessing steps. The resulting lesion masks were aligned with the preprocessed images using the same rigid transformation and resampling procedure applied during preprocessing. Lesion voxels were excluded from the WM masks to define normal-appearing white matter (NAWM) and gray matter (NAGM) for these MRI sessions.

3.2 Model

We extended our previously proposed physics-guided self-supervised framework for retrospective qMRI mapping from conventional clinical scans [26] by incorporating the MT-informed apparent longitudinal relaxation model into the differentiable forward signal model (Eqs. (4)–(8)). This extension enables the multi-contrast convolutional neural network (CNN) to estimate voxel-wise maps of the free-pool longitudinal relaxation time (T_1^{f}), free-pool transverse relaxation time (T_2^{f}), PD, and macromolecular pool fraction (m_0^{s}; MPF) directly from routine T_1w, T_2w, and FLAIR acquisitions. Training remains self-supervised and does not require reference quantitative maps, which are not available.

3.2.1 Network architecture

The network retains the three-level 2D U-Net architecture described in [26]. A shared-weight encoder processes each input contrast of an MRI session separately, and attention-based contrast-fusion modules combine the resulting feature representations at each resolution level. The fused features are passed to a common decoder with skip connections from the corresponding encoder levels. The network operates on 2D axial slices of 224×224 voxels.

We modified the original architecture in two respects. First, we expanded the output layer from three to four channels to generate T_1^{f}, T_2^{f}, PD, and m_0^{s} maps. A Sigmoid activation function constrained the outputs to $[0, 1]$, after which they were linearly scaled to physical ranges of $[0, 5]$ s for T_1^{f}, $[0, 3]$ s for T_2^{f}, $[0, 1]$ arbitrary units for PD, and $[0, 0.4]$ for m_0^{s}. We otherwise kept the decoder and attention-based contrast-fusion modules unchanged.

Second, we replaced the standard instance normalization in the encoder and bottleneck with acquisition-conditioned adaptive instance normalization (AdaIN) [39–41]. Conditioning neural networks on sequence parameters has previously been shown to mitigate protocol-dependent variability across heterogeneous protocols [42–44]. Here, AdaIN enables the shared encoder to dynamically modulate its intermediate feature representations based on contrast type and sequence timing parameters.

To condition the network on the acquisition settings, each input contrast $m \in \{T_1\mathrm{w}, T_2\mathrm{w}, \mathrm{FLAIR}\}$ is paired with a nine-dimensional acquisition parameter vector $\mathbf{c}_m \in \mathbb{R}^9$:

$$\mathbf{c}_m = \left[c_{T_1\mathrm{w}},\, c_{T_2\mathrm{w}},\, c_{\mathrm{FLAIR}},\, \ln(\mathrm{TR}/1\,\mathrm{ms}),\, \ln(\mathrm{TE}/1\,\mathrm{ms}),\, \sin(\mathrm{FA}),\, B_{1,\mathrm{rms}},\, c_{\mathrm{TI}},\, \ln(\mathrm{TI}/1\,\mathrm{ms})\right]^\top. \tag{9}$$

The first three terms are binary indicator variables ($c_{T_1\mathrm{w}}, c_{T_2\mathrm{w}}, c_{\mathrm{FLAIR}} \in \{0, 1\}$), with exactly one set to 1 to identify the contrast type. The continuous timing parameters TR, TE, and TI are log-transformed to compress their numerical ranges. The flip angle FA is transformed via $\sin(\mathrm{FA})$ and $B_{1,\mathrm{rms}}$ is the sequence-averaged root-mean-square RF transmit amplitude in μT. The binary indicator $c_{\mathrm{TI}} \in \{0, 1\}$ denotes whether an inversion preparation pulse was applied ($c_{\mathrm{TI}} = 1$ for MPRAGE and FLAIR; $c_{\mathrm{TI}} = 0$ for non-inversion GRE and TSE acquisitions). When $c_{\mathrm{TI}} = 0$, the corresponding logarithmic inversion time entry is set to zero.

A two-layer multilayer perceptron (MLP) with exponential linear unit (ELU) activations [45] projects $\mathbf{c}_m$ to a 128-dimensional acquisition embedding $\mathbf{w}_m$. For an intermediate feature map $\mathbf{x} \in \mathbb{R}^{C\times H\times W}$ with channel-wise mean $\mu_c(\mathbf{x})$ and standard deviation $\sigma_c(\mathbf{x})$, AdaIN normalizes each channel and applies sequence-specific affine modulation:

$$\mathrm{AdaIN}(\mathbf{x}, \mathbf{w}_m)_c = \gamma_c(\mathbf{w}_m)\left[\frac{\mathbf{x}_c - \mu_c(\mathbf{x})}{\sigma_c(\mathbf{x})}\right] + \beta_c(\mathbf{w}_m), \tag{10}$$

where the channel-wise scale and shift vectors are generated via learned linear projections:

$$\boldsymbol{\gamma}(\mathbf{w}_m) = \mathbf{W}_\gamma \mathbf{w}_m + \mathbf{b}_\gamma + \mathbf{1}, \qquad \boldsymbol{\beta}(\mathbf{w}_m) = \mathbf{W}_\beta \mathbf{w}_m + \mathbf{b}_\beta. \tag{11}$$

Both the projection weights and biases are initialized to zero. This ensures that each AdaIN block acts as standard instance normalization at the onset of training and progressively learns acquisition-dependent feature modulation. The conditioning MLP and AdaIN layers added approximately 0.50M trainable parameters, corresponding to 5.6% of the 9.06M trainable parameters of the full network.

3.2.2 Self-supervised physics-guided objective

The framework was trained using a physics-guided self-supervised objective that enforced consistency between the acquired input clinical scans and synthetic images generated by the forward signal model. For each MRI session, the network predicted voxel-wise maps of PD, T_1^{f}, T_2^{f}, and m_0^{s}. The MT-informed forward signal models described in Section 2.3 then combined these maps with the acquisition parameters of each input scan, including TR, TE, TI, FA,

and $B_{1,\mathrm{rms}}$, to synthesize images $\hat{S}_m$ for contrasts $m \in \{T_1\mathrm{w}, T_2\mathrm{w}, \mathrm{FLAIR}\}$ using Eqs. (5)–(8). Semisolid-pool and exchange parameters that were not estimated by the network were fixed to values from recent qMT literature for WM at 3 T: $\gamma = 2.675 \times 10^8\,\mathrm{rad\,s^{-1}T^{-1}}$, $R_x = 13.6\,\mathrm{s}^{-1}$, $T_1^{\mathrm{s}} = 300\,\mathrm{ms}$, and $T_2^{\mathrm{s}} = 12.5\,\mu\mathrm{s}$ [3, 30, 46].

For the TSE-based $T_2\mathrm{w}$ and FLAIR acquisitions, the forward signal model required an additional correction for effective T_2 weighting. Specifically, the analytical Bloch-based equations are less accurate for TSE acquisitions because these sequences use long echo trains with variable refocusing flip angles, so the effective T_2 weighting differs from that predicted by the nominal echo time. To account for this difference, we calculated TE correction factors following Busse et al. [47] and applied them to the nominal TE values in the forward signal model. The factors were 0.90 and 0.42 for the Philips $T_2\mathrm{w}$ and FLAIR acquisitions, respectively, and 0.86–0.96 and approximately 0.51 for the corresponding Siemens acquisitions. Supplementary Section S2 details the calculations.

Conventional MRI signal intensities are expressed in arbitrary, non-standardized units and therefore contain an unknown scan-dependent global scaling factor. Before computing the training loss, each synthesized image $\hat{S}_m$ was globally rescaled to the corresponding normalized input image S_m using the calibration procedure described in our previous work [26]. Because the whole-brain histogram mode used for input normalization corresponds to WM intensity, the scaling factor was constrained to within 20% of a reference value derived from representative WM tissue parameters: $T_1^{\mathrm{f}} = 2000\,\mathrm{ms}$, $T_2^{\mathrm{f}} = 70\,\mathrm{ms}$, $\mathrm{PD} = 0.70$, and $m_0^{\mathrm{s}} = 0.21$ [3, 30, 46, 48].

In the Bloch-based forward signal models, T_1^{f} and m_0^{s} both affect the synthesized signal intensity through R_1^{app} (Eq. (4)), resulting in an identifiability problem. For a fixed $B_{1,\mathrm{rms}}$, multiple combinations of T_1^{f} and m_0^{s} can therefore produce identical image contrast. Variation in $B_{1,\mathrm{rms}}$ across the dataset helps constrain this ambiguity. To further separate the two parameters, we imposed a prior on T_1^{f} within WM only, with $\mu_{T_1} = 2000\,\mathrm{ms}$ and $\sigma_{T_1} = 400\,\mathrm{ms}$. This prior mean reflects the free-water relaxation time at 3 T reported in recent qMT literature [3, 46], which is substantially longer than the apparent tissue T_1 ($\sim$ 700–1100 ms [48]). No direct prior or regularization term was imposed on m_0^{s}.

As in our previous work [26], a soft lower bound was imposed on PD to constrain the solution space of the ill-posed inverse problem. This constraint was implemented as a hinge penalty with a threshold $\tau = 0.60$. The total training loss combined a data-consistency image-reconstruction term with the PD constraint and the white-matter T_1^{f} prior regularization terms:

$$\mathcal{L} = \sum_m \left\| \tilde{S}_m - S_m \right\|_1 + \lambda_{\text{PD-constraint}} \left\langle \max\left(0,\, \tau - \mathrm{PD}\right) \right\rangle + \lambda_{T_1\text{-prior}} \left\langle \frac{\left(T_1^{\mathrm{f}} - \mu_{T_1}\right)^2}{2\sigma_{T_1}^2} \right\rangle_{\mathrm{WM}}, \tag{12}$$

where $\tilde{S}_m$ and S_m denote the rescaled synthesized and normalized acquired images for contrast $m \in \{T_1\mathrm{w}, T_2\mathrm{w}, \mathrm{FLAIR}\}$, respectively. The angle brackets $\langle\cdot\rangle$ represent spatial voxel averaging, with $\langle\cdot\rangle_{\mathrm{WM}}$ restricted to a white matter mask. The hyperparameters $\lambda_{\text{PD-constraint}}$ and $\lambda_{T_1\text{-prior}}$ control the relative contributions of the two regularization terms.

3.2.3 Training procedure and implementation details

The network was trained on the full combined dataset for 20 epochs using the Adam optimizer [49] with momentum parameters $\beta_1 = 0.9$ and $\beta_2 = 0.999$. Network weights were initialized using Kaiming initialization [50]. Each training batch consisted of slices randomly sampled from the entire training dataset. To identify the optimal model configuration, hyperparameter optimization was performed using a grid search. Table S3.1 in the Supplementary Materials summarizes the search space and explored values. Furthermore, the necessity of the two regularization terms (PD-constraint and T_1-prior) was evaluated through an ablation study, detailed in Supplementary Materials S7. Model selection was based on minimizing validation loss while ensuring quantitative outputs matched physiologically expected ranges for human brain qMT literature values in WM and GM [3, 46, 48].

The final model was trained using a learning rate of 1×10^{-5}, a batch size of 16 slices, and regularization weights $\lambda_{T_1\text{-prior}} = 0.0025$ and $\lambda_{\text{PD-constraint}} = 0.5$. We developed the implementation in Python 3.11.9 using PyTorch 2.1.2 [51], with data loading, preprocessing, and slice sampling handled by TorchIO 0.20.6 [52]. Training was performed on a single NVIDIA RTX 4500 Ada Generation GPU and took approximately 12 hours. The complete PyTorch implementation and pretrained network weights are available on GitHub: `https://github.com/JelmervanL/MPF-mapping-from-conventional-MRI`.

3.3 Evaluation

We evaluated the biological plausibility and protocol robustness of the generated MPF (m_0^{s}) maps, their sensitivity to demyelinating pathology, their cross-protocol reproducibility, and the contribution of acquisition-conditioned AdaIN.

3.3.1 Macromolecular pool fraction in white matter and gray matter

Estimated MPF values were evaluated in WM and GM across the full clinical test set ($n = 1{,}237$). To assess robustness against scanner and protocol variation, sessions were stratified into three vendor-by-protocol groups based on the $T_1\mathrm{w}$ acquisition: Philips Spoiled GRE, Philips MPRAGE, and Siemens MPRAGE. Stratification was based on the $T_1\mathrm{w}$ sequence because it exhibited the largest protocol variation in the dataset, including differences in pulse-sequence type. In contrast, $T_2\mathrm{w}$ and FLAIR acquisitions consistently employed turbo spin-echo sequences across the entire dataset, differing primarily in timing parameters (TR, TE, and TI). We quantified protocol stability using the between-group coefficient of variation (CV) of tissue-specific mean MPF values. The CV was calculated as the standard deviation of the three group means, each obtained by averaging the per-session means within that group, divided by the unweighted mean of the three group means.

3.3.2 Demyelinating pathology

Sensitivity to demyelinating pathology was evaluated using the 70 MR-RATE test sessions with report-derived demyelinating-disease labels and the four UMCU test sessions with confirmed MS. Mean MPF values within segmented WM lesions were compared with NAWM and NAGM using relative percentage differences.

Furthermore, because T_1^{f} and m_0^{s} jointly determine R_1^{app} in the forward signal model (Eq. (4)), we also assessed whether their estimates covaried across individual lesions. Following the approach of Assländer et al. [3], individual lesions were defined as distinct three-dimensional connected components of the WM lesion masks. Components containing fewer than 10 voxels were excluded from the analysis. Each lesion was represented by its median m_0^{s} and free-pool longitudinal relaxation rate ($R_1^{\mathrm{f}} = 1/T_1^{\mathrm{f}}$). Principal component analysis (PCA) of these paired lesion-level estimates quantified the dominant direction of joint variation and the remaining variation orthogonal to this direction.

3.3.3 Prospective volunteer reproducibility

As a technical demonstration, prospective cross-protocol reproducibility of the generated MPF maps was evaluated in a 26-year-old male volunteer scanned on a 3 T Philips Ingenia system with written informed consent and institutional review board approval. The volunteer was considered healthy based on the absence of visible lesions on the acquired images. Two clinical protocols, each comprising a separate set of $T_1\mathrm{w}$, $T_2\mathrm{w}$, and FLAIR acquisitions, were performed consecutively within the same scanning session. The protocols differed primarily in the $T_1\mathrm{w}$ sequence, which used Spoiled GRE in one protocol and MPRAGE in the other; detailed acquisition characteristics are provided in Supplementary Table S4.1. The prospective data were preprocessed using the same pipeline as the clinical dataset. Voxel-wise agreement between the generated MPF maps was assessed across brain parenchyma (WM + GM) using the mean absolute percentage difference, Pearson correlation coefficient (r), and concordance correlation coefficient (CCC).

3.3.4 Impact of acquisition-conditioned AdaIN

Finally, the contribution of acquisition-conditioned AdaIN was assessed through an ablation experiment in which all AdaIN layers were replaced with standard instance normalization, while all other model and training settings were kept unchanged. Performance was compared using test-set L1 reconstruction error, the between-group CV of tissue MPF estimates across vendor-by-protocol groups, and cross-protocol agreement in the prospective volunteer.

4 Results

4.1 Macromolecular pool fraction in white matter and gray matter

Across the clinical test set ($n = 1{,}237$), estimated macromolecular pool fraction (m_0^{s}; MPF) values were consistently higher in WM (0.182 ± 0.007) than in GM (0.144 ± 0.009). Estimates varied little across vendor-by-protocol groups, with between-group CVs of 2.65% for WM and 5.63% for GM and substantial overlap of the corresponding within-tissue distributions (Table 3; Figure 2). Results for all estimated quantitative parameters are provided in Supplementary Table S5.1.

Table 3: Mean estimated macromolecular pool fraction (m_0^s; MPF) in white matter (WM) and gray matter (GM) across the full clinical test set and the three vendor-by-protocol groups. Values represent mean ± standard deviation across MRI sessions.

		Philips (UMCU)		Siemens (MR-RATE)
Tissue	**Total** ($n = 1,237$)	**Spoiled GRE** ($n = 653$)	**MPRAGE** ($n = 221$)	**MPRAGE** ($n = 363$)
WM	0.182 ± 0.007	0.186 ± 0.006	0.179 ± 0.004	0.177 ± 0.006
GM	0.144 ± 0.009	0.149 ± 0.005	0.147 ± 0.005	0.134 ± 0.006

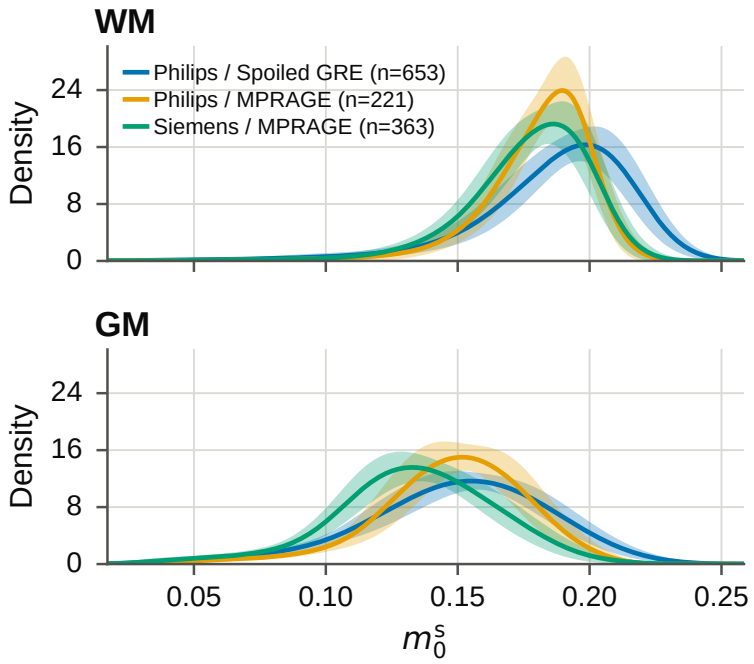

Figure 2: Distribution of estimated macromolecular pool fraction (m_0^s; MPF) values in white matter (WM) and gray matter (GM) for the three vendor-by-protocol groups in the clinical test set. Solid lines show the mean density across MRI sessions and shaded regions indicate one standard deviation.

4.2 Demyelinating pathology

In the 70 MR-RATE sessions with a report-derived demyelinating-disease label, mean MPF was 0.113 ± 0.015 in segmented WM lesions. Lesion MPF was 35.6% lower than in NAWM (0.176 ± 0.005) and 15.2% lower than in NAGM (0.133 ± 0.005). The corresponding results within lesions for T_1^f, T_2^f, and PD are also provided in Supplementary Table S5.2. Comparable MPF reductions were observed in the four UMCU subjects with MS (Supplementary Table S5.3). Representative examples from both subsets show that FLAIR-hyperintense lesions corresponded to reductions in MPF relative to the surrounding tissue across scanner vendors and T_1w acquisition protocols (Figures 3 and 4).

To investigate the relationship between free-pool longitudinal relaxation and MPF within demyelinating lesions, we performed PCA on lesion-wise R_1^f and MPF estimates from the MR-RATE subset (Figure 5). The first and second principal components explained 91% and 9% of the total variance, respectively.

4.3 Prospective volunteer reproducibility

In the controlled prospective cross-protocol evaluation, the generated m_0^s maps showed similar tissue values across the two acquisition protocols. For the Spoiled GRE protocol, the mean m_0^s was 0.171 ± 0.026 in WM and 0.135 ± 0.027 in GM. The corresponding values for the MPRAGE protocol were 0.171 ± 0.021 and 0.134 ± 0.023, respectively.

Figure 6 shows representative MPF maps from both protocols. The voxel-wise percentage-difference map showed only limited regional variation between protocols. Across brain parenchyma (WM+GM), the mean absolute percentage

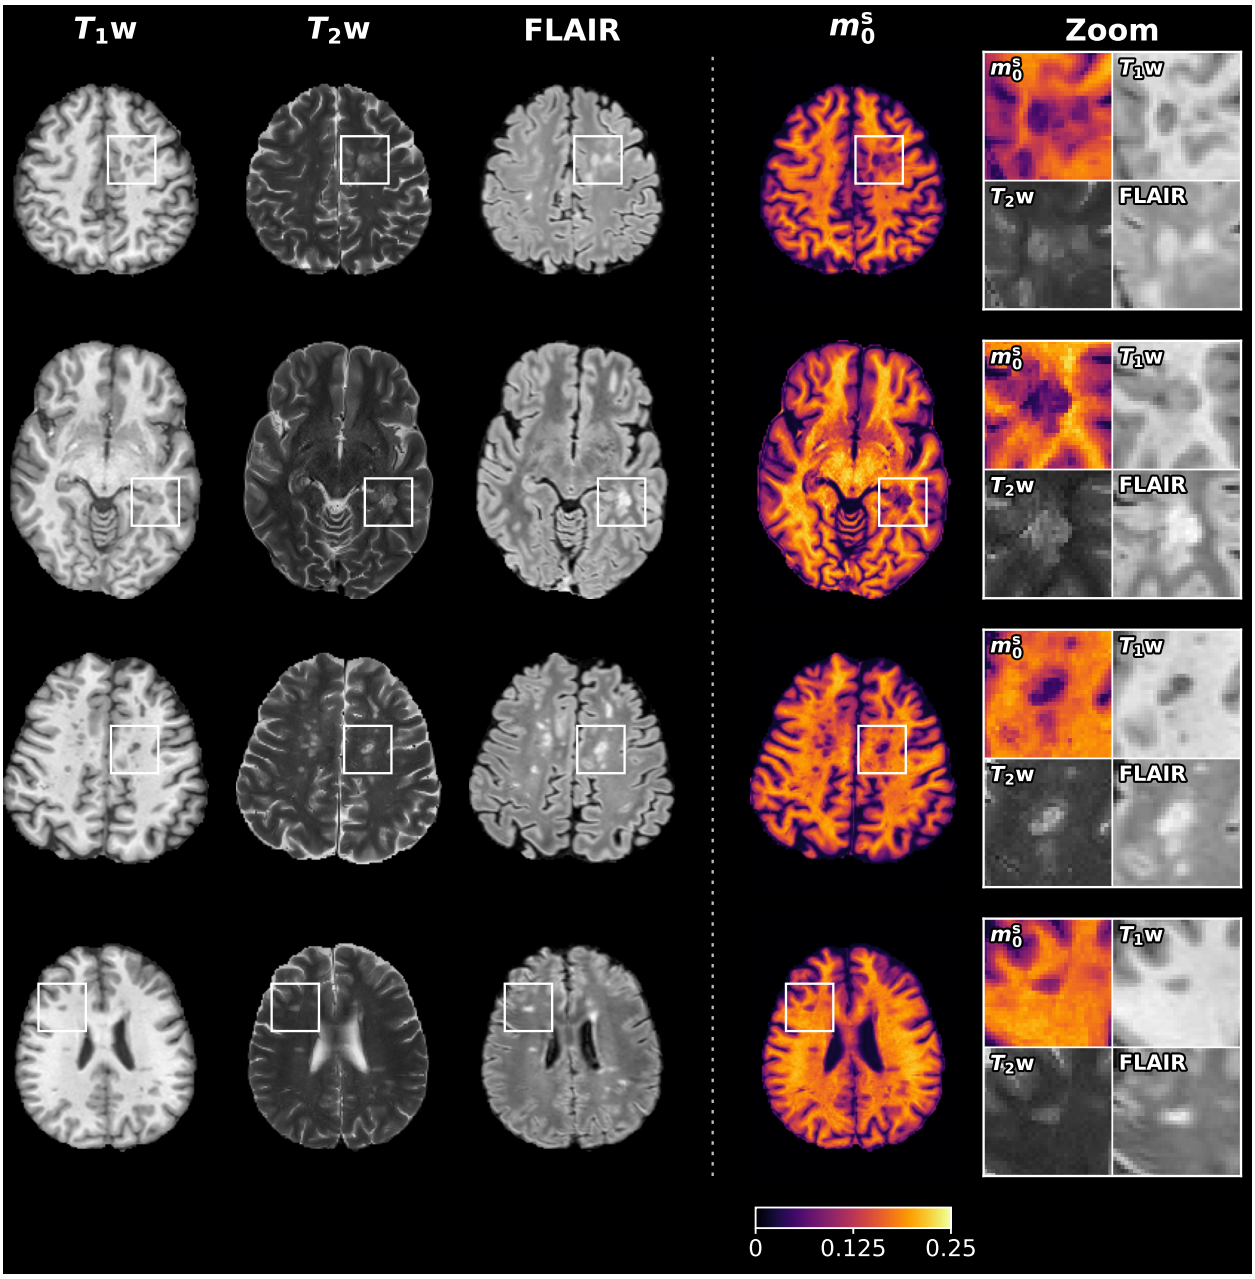

Figure 3: Representative MPF maps from the MR-RATE test set (Siemens MPRAGE). Each of the four rows shows a distinct, randomly selected subject with a report-derived demyelinating-disease diagnosis. Columns show the input T_1-weighted, T_2-weighted, and FLAIR images together with the generated m_0^s map. White boxes indicate regions enlarged in the zoom panels.

difference between the two generated MPF maps was 7.7%. The strong agreement between protocols was also evident in the voxel-wise density scatter plot (Figure 7), yielding a Pearson correlation coefficient of $r = 0.895$ and a CCC of 0.888. Similar reproducibility analyses for the generated T_1^f, T_2^f, and PD maps are presented in the Supplementary Figures S5.1 and S5.2.

4.4 Impact of acquisition-conditioned AdaIN

Replacing acquisition-conditioned AdaIN with standard instance normalization increased the between-group CV of estimated m_0^s across the three vendor-by-protocol groups. In WM, the CV increased from 2.65% with AdaIN to

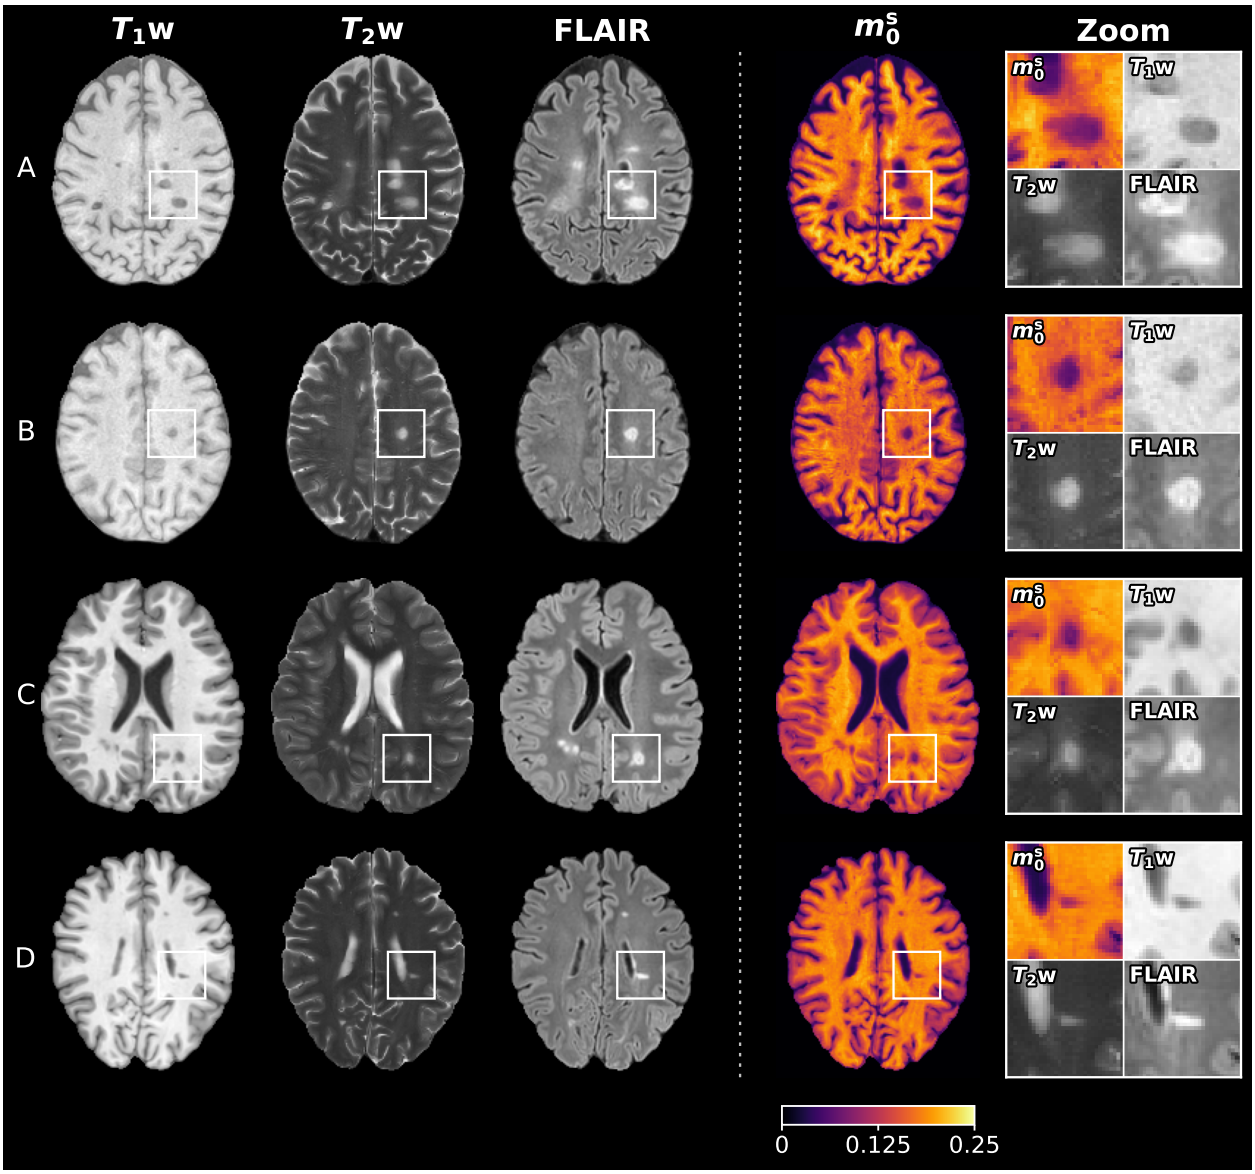

Figure 4: Representative MPF maps from the four subjects with MS identified in the UMCU test set. Rows A and B correspond to the Philips Spoiled GRE vendor-by-protocol group, whereas rows C and D correspond to the Philips MPRAGE vendor-by-protocol group. Columns show the input T_1-weighted, T_2-weighted, and FLAIR images together with the generated m_0^s map. White boxes indicate regions enlarged in the zoom panels.

10.27% without AdaIN; in GM, it increased from 5.63% to 10.37%, respectively (additional results are provided in Supplementary Table S6.1). The mean total L1 reconstruction error across the clinical test set ($n = 1{,}237$ sessions) increased by 22.3%, from 0.0466 ± 0.0177 with AdaIN to 0.0570 ± 0.0171 without AdaIN ($p < 0.001$, Wilcoxon signed-rank test).

In the prospective volunteer experiment, the cross-protocol mean absolute percentage difference between the estimated m_0^s maps increased from 7.7% with AdaIN to 15.3% without AdaIN. The voxel-wise CCC decreased from 0.888 to 0.672. The corresponding cross-protocol scatter plots are shown in Supplementary Figure S6.1.

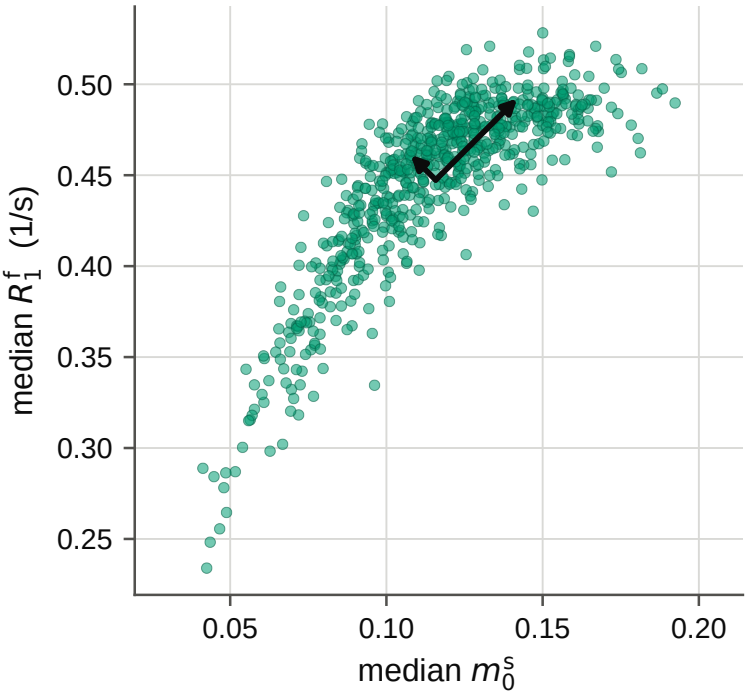

Figure 5: Relationship between lesion-wise free-pool longitudinal relaxation rate (R_1^{f}) and macromolecular pool fraction (m_0^{s}; MPF) in the MR-RATE demyelinating-disease subset ($n = 70$ MRI sessions, 673 unique lesions). Each point represents the median R_1^{f} and m_0^{s} value within a single demyelinating lesion. Black arrows indicate the principal-component eigenvectors and are scaled to one standard deviation along each component.

5 Discussion

In this work, we introduced a physics-guided, self-supervised deep learning framework for retrospectively estimating MPF from conventional $T_1\mathrm{w}$, $T_2\mathrm{w}$, and FLAIR brain MRI without paired qMT data. The framework extends retrospective quantitative MRI mapping by incorporating an analytical two-pool MT model that relates apparent longitudinal relaxation to sequence-averaged RF power, together with acquisition-conditioned AdaIN to account for differences among acquisition protocols. We trained and evaluated the framework using a large multicenter, multivendor clinical dataset and examined tissue-level MPF estimates, sensitivity to demyelinating pathology, vendor/protocol robustness, and the contribution of AdaIN acquisition conditioning.

We first assessed whether the retrospectively estimated MPF maps showed expected values for WM and GM. Across the clinical test set, estimated MPF was consistently higher in WM (0.182 ± 0.007) than in GM (0.144 ± 0.009), in agreement with the higher macromolecular and myelin content of WM [4, 53]. These tissue mean values fell within the ranges reported by qMT studies that model distinct longitudinal relaxation rates for the free and semisolid pools ($m_0^{\mathrm{s}} \in [0.17, 0.27]$ in WM and $m_0^{\mathrm{s}} \in [0.08, 0.18]$ across cortical and subcortical GM) [3, 30, 46, 54, 55]. However, absolute MPF values should be compared cautiously across studies, as they depend on the fitting strategy and the assumptions and constraints of the two-pool model. In particular, our estimates were higher than values commonly obtained using models that use the constraint $T_1^{\mathrm{s}} = T_1^{\mathrm{f}}$ [16, 17, 56]. Previous studies have shown that this constraint can substantially underestimate the semisolid pool fraction [3, 46, 54]. In contrast, our model fixes T_1^{s} to the substantially shorter value of 300 ms [3, 30, 46, 54] and estimates T_1^{f} voxel-wise. This difference in model assumptions likely contributes to the higher MPF estimates obtained with our framework.

Despite the tissue means falling within reported ranges, the contrast between WM and GM was smaller than expected. The GM-to-WM MPF ratio was approximately 0.79, whereas prospective qMT studies using both constrained and unconstrained two-pool models have generally reported ratios closer to 0.5 [3, 16, 56]. Our generated maps preserved the expected ordering of WM and GM values but showed reduced MPF contrast between these tissues. This reduced tissue contrast may stem from model simplifications, globally fixed semisolid-pool and exchange parameters, the WM-specific T_1^{f} prior, and partial-volume effects, as discussed further in Section 5.1.

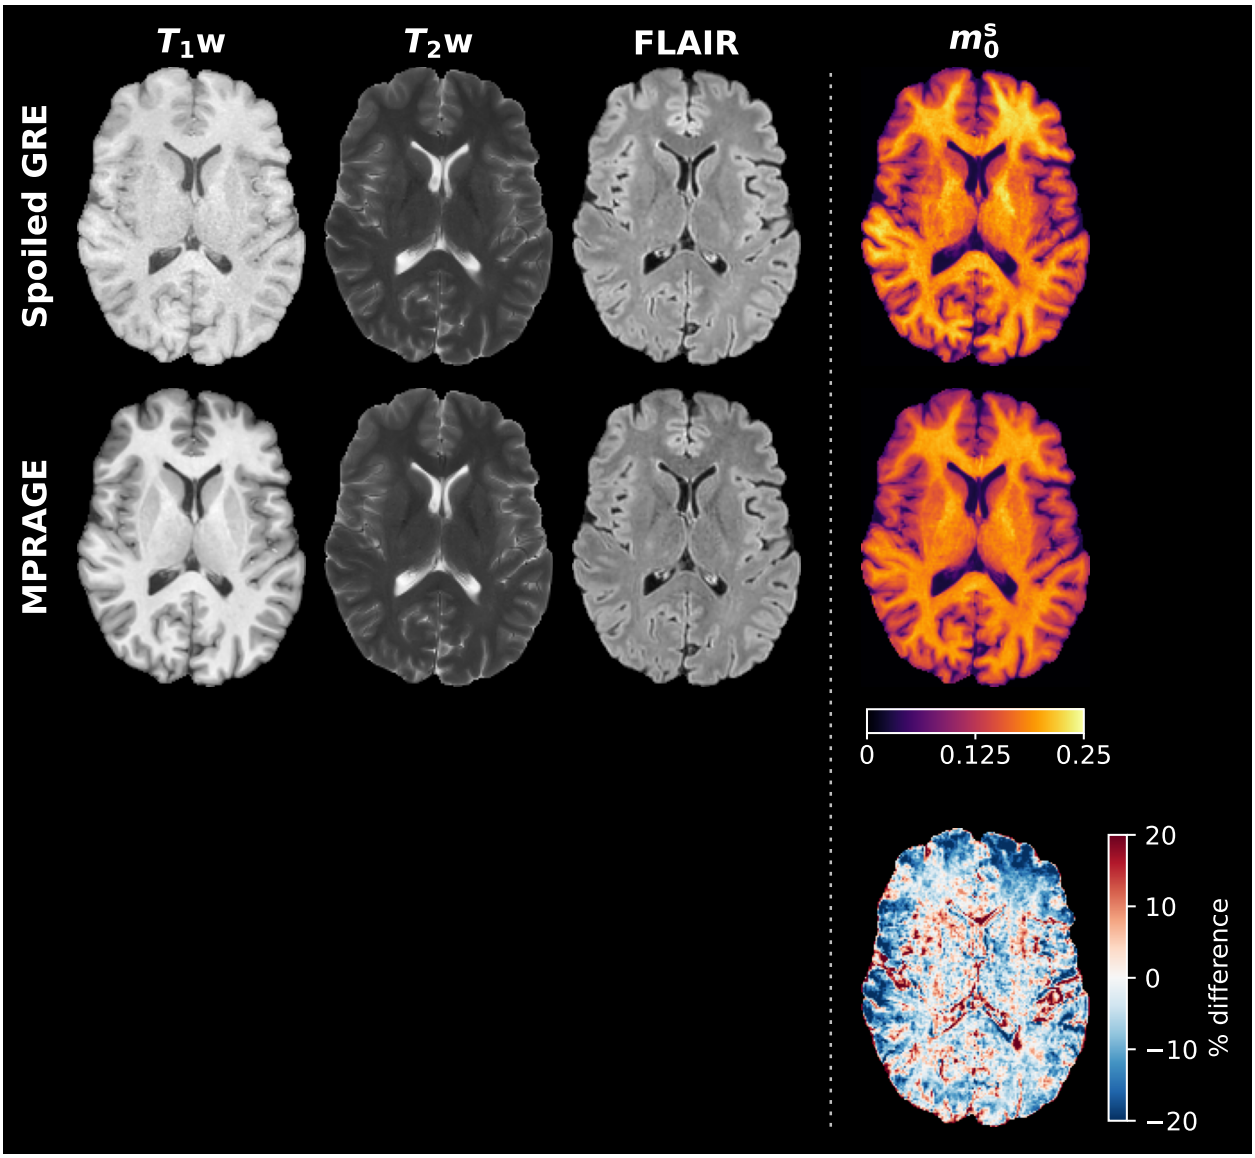

Figure 6: Prospective within-subject reproducibility in a healthy volunteer scanned with two different clinical protocols. Top row: estimated macromolecular pool fraction (m_0^s; MPF) map from the Spoiled GRE protocol. Middle row: corresponding MPF map from the MPRAGE protocol. Bottom row: voxel-wise percentage difference maps between the two generated MPF maps.

In segmented lesions, MPF was 35.6% lower than in NAWM in the MR-RATE demyelinating-disease subset, consistent with reductions of 23.5–34.2% in the four UMCU subjects with confirmed MS. These decreases align with prospective quantitative MT literature reporting substantial macromolecular pool reductions in MS lesions [3, 4], where relative decreases typically range from 30% to 45% compared to NAWM [9, 57–59]. This agreement supports the sensitivity of the generated MPF maps to tissue alterations associated with demyelinating pathology. To assess the extent to which lesion-wise MPF and R_1^f estimates covaried, we performed PCA on their paired values. The first principal component explained 91% of their joint variance (Figure 5). This proportion was similar to the 93% reported by Assländer et al. [3] for dedicated prospective qMT under the constraint of equal longitudinal relaxation times for the free and semisolid pools ($T_1^s = T_1^f$). Thus, our retrospectively estimated quantitative maps showed a similar degree of joint variation between lesion-wise R_1^f and MPF to that observed with prospective qMT. Section 5.1 discusses limitations related to the report-derived diagnoses, lesion segmentations, and clinical characterization of the data.

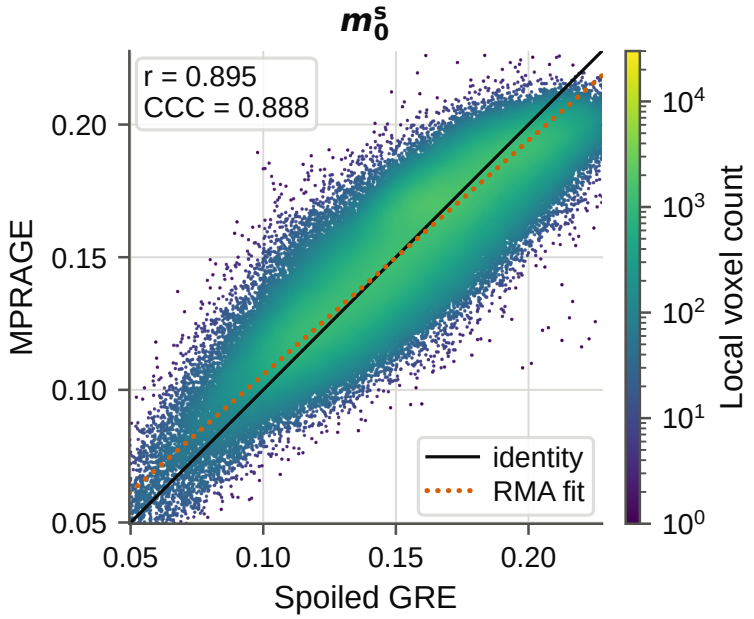

Figure 7: Voxel-wise agreement between macromolecular pool fraction (m_0^s; MPF) maps estimated from the Spoiled GRE and MPRAGE protocols in a healthy volunteer. Each point represents a brain parenchymal voxel, colored by local voxel density (log scale). The solid line denotes identity and the dotted line the reduced major axis (RMA) fit. Pearson's r and the concordance correlation coefficient (CCC) are shown in the top left.

For retrospective application to heterogeneous clinical MRI archives, MPF estimates must remain stable across acquisition protocols and scanner systems. Across the multicenter test cohort, mean MPF varied little among the three vendor-by-protocol groups, with between-group CVs of 2.65% in WM and 5.63% in GM, despite differences in scanner vendor, sequence implementation, timing parameters, and RF power. The prospective within-subject experiment further supported this protocol robustness: MPF maps estimated from consecutive Spoiled GRE and MPRAGE acquisitions showed good voxel-wise agreement (CCC $= 0.888$) and a mean absolute percentage difference of 7.7% (Figures 6 and 7).

Acquisition-conditioned AdaIN contributed substantially to this stability. Replacing AdaIN with standard instance normalization substantially increased the between-group CV of MPF, and reduced cross-protocol agreement in the prospective volunteer, with the CCC decreasing from 0.888 to 0.672 (Supplementary Materials S6). This finding is consistent with other quantitative mapping frameworks showing that conditioning neural networks on sequence parameters improves performance across heterogeneous acquisition settings [42, 43]. However, protocol hold-out experiments indicated limited generalization to sequence types absent from the training data (Supplementary Materials S8). Robust deployment across clinical archives therefore requires representative coverage of the target sequence types during training.

The observed protocol stability is particularly relevant given the ill-posed nature of the parameter-estimation problem. The framework estimates four quantitative maps (m_0^s, T_1^f, T_2^f, and PD) from only three conventional image contrasts (T_1w, T_2w, and FLAIR). Moreover, as discussed in the Methods (Section 3.2.2), T_1^f and m_0^s exhibit a specific identifiability problem. To constrain the solution space, the model incorporates two regularizers: a WM-specific T_1^f prior and a soft lower bound on PD. The importance of these regularizers was evaluated by the ablation study (Supplementary Materials S7). Removing the WM T_1^f prior caused shifts in T_1^f and MPF, increased differences between vendor-by-protocol groups, and reduced cross-protocol agreement in the prospective volunteer experiment. Removing the PD constraint primarily affected the estimated PD and T_2^f maps and increased vendor-by-protocol group variation. These findings suggest that the regularization improves the stability and reproducibility of the generated quantitative maps. This is consistent with previous physics-guided self-supervised approaches that used similar regularizers to derive quantitative maps from conventional MRI [26, 27]. Beyond these regularization constraints, protocol-dependent variation in sequence-averaged RF power ($B_{1,\mathrm{rms}}$) provides additional information that may help distinguish the contributions of T_1^f and m_0^s to the observed image intensities. Combining the UMCU and MR-RATE datasets broadened the range of acquisition protocols and corresponding $B_{1,\mathrm{rms}}$ values represented during training (Table 1).

5.1 Limitations

Several limitations should be considered when interpreting our findings. First, the framework relies on simplified descriptions of MT and conventional MRI signal formation. The analytical MT model adopted from Teixeira et al. [29] incorporates several assumptions, as described in the Theory section (Section 2). In addition, the Bloch-based signal equations do not account for factors such as B_1 inhomogeneity or detailed echo-train dynamics. More comprehensive approaches, including the MT models proposed by Graham and Henkelman [60], and Sled and Pike [7], EPG-X [61], and the generalized Bloch model [22], may provide more accurate descriptions of these effects. However, these methods generally require detailed sequence information that is often unavailable in retrospective clinical MRI archives.

The model also relies on fixed tissue parameters and explicit regularization to constrain parameter estimation. We fixed the exchange and semisolid relaxation parameters to literature values reported for healthy WM and applied them to all tissues and subjects. Because these parameters vary across tissue types and may also change with pathology [3], applying fixed healthy-WM values to all tissues and subjects may bias the estimated MPF. Moreover, although the WM-specific T_1^{f} prior and soft lower bound on PD improved model stability, they constrained the range of admissible solutions and made the estimated parameter maps dependent on the selected regularization settings.

In addition to these modeling assumptions, the heterogeneous acquisition geometry of clinical MRI imposed spatial limitations, as previously discussed in [26]. We registered all MRI contrasts to the 2D $T_2\mathrm{w}$ acquisition geometry and resampled them to an in-plane resolution of 1×1 mm. This procedure reduced the through-plane resolution of the original 3D $T_1\mathrm{w}$ and FLAIR images and removed volumetric detail. Registration errors and partial-volume effects may also propagate into the estimated parameter maps. Moreover, slice-wise processing does not use spatial information from adjacent slices.

The evaluation design also has several limitations. Establishing the absolute accuracy of the MPF maps remains challenging because no universally accepted qMT reference standard exists [4, 62]. Available qMT methods differ in sequence design, fitting procedures, and the constraints imposed on the two-pool model, which can lead to systematic differences in the estimated MPF [4]. Comparison with a dedicated prospective qMT implementation would therefore establish agreement with that specific methodology, rather than establishing absolute biological accuracy. In the absence of paired qMT data, we evaluated the generated MPF maps through complementary measures of biological plausibility, sensitivity to demyelinating pathology, stability across vendor-by-protocol groups, and prospective cross-protocol reproducibility. Furthermore, these analyses relied on tissue segmentations derived from clinical MRI, which may be inaccurate in the presence of lesions, anatomical abnormalities, or limited image quality. Errors in the WM and GM masks may affect the reported tissue-level MPF estimates, while errors in the automated WM lesion segmentations may affect the lesion analysis. Moreover, we relied primarily on report-derived diagnostic labels to evaluate demyelinating pathology, and FLAIR-hyperintense white-matter lesions are not specific to demyelination. The four clinically confirmed MS UMCU cases supported the findings but did not constitute a systematic validation cohort. These caveats should be considered when interpreting the observed reduction in lesion MPF. Future studies should evaluate the framework in clinically characterized cohorts with expert tissue and lesion annotations.

While the large and heterogeneous clinical dataset enabled evaluation across multiple vendors and protocols, the vendor-by-protocol groups differed in patient demographics. As a result, biological variability across the populations confounded the group comparisons, preventing an independent assessment of the effects of vendor and protocol on MPF estimates. Moreover, all MRIs were acquired at 3 T, and protocol diversity was concentrated mainly in the $T_1\mathrm{w}$ acquisitions; other common acquisitions, including 3D $T_2\mathrm{w}$ and 2D FLAIR, were not represented. Generalizability beyond the scanner systems and acquisition protocols represented in this study remains to be established. In particular, protocol hold-out experiments showed substantial parameter shifts for $T_1\mathrm{w}$ sequence types absent from training (Supplementary Materials S8); the framework should therefore not be applied to unseen sequence types without retraining or dedicated validation. Finally, the prospective reproducibility experiment included only a single healthy volunteer and should therefore be considered a technical demonstration; larger, dedicated cohorts are required to prospectively assess cross-protocol and cross-vendor reproducibility more comprehensively.

5.2 Outlook

By deriving MPF maps directly from routine $T_1\mathrm{w}$, $T_2\mathrm{w}$, and FLAIR scans, the framework enables quantitative MT analysis without dedicated qMT acquisitions. The framework can generate MPF maps from existing multicenter conventional clinical imaging archives, supporting large-scale retrospective biomarker research. Applied prospectively, it can provide MPF as a quantitative biomarker without increasing scan time and may support longitudinal analysis of lesions in patients undergoing repeated clinical MRI. Although the present study focuses on MS, the proposed MPF mapping approach may also enable investigation of macromolecular tissue changes in other neurological conditions.

Potential applications include Alzheimer's disease [63, 64], Parkinson's disease [65], Huntington's disease [66], and myelin damage and repair after stroke [67].

6 Conclusion

Physics-guided self-supervised learning enables retrospective estimation of MPF maps from conventional clinical brain MRI without requiring paired qMT reference data. Across heterogeneous protocols, the resulting maps showed good consistency and marked MPF reductions in lesions of patients with demyelinating disease. This highlights the framework's potential for large-scale biomarker research using existing clinical imaging archives. It may also enable prospective MPF mapping from conventional MRI without requiring a dedicated qMT acquisition or additional scan time.

Acknowledgments

The Hanarth Fonds for AI in Oncology supported this research.

Data availability

The UMCU data cannot be shared because of privacy and institutional restrictions. The MR-RATE dataset is available through Hugging Face at `https://huggingface.co/datasets/Forithmus/MR-RATE` (accessed September 8, 2026), subject to its access conditions and terms of use. The sessions included in this study were selected according to the eligibility criteria described in the Methods.

Code and model weights are available on GitHub: `https://github.com/JelmervanL/MPF-mapping-from-conventional-MRI`.

A MT signal model assumptions

The following sections provide further justification for the sequence-averaged RF saturation and near-resonance approximations introduced in the Theory section (Section 2), and explain the omission of the sequence-dependent apparent equilibrium magnetization from the forward signal model.

A.1 Sequence-averaged RF saturation

Teixeira et al. [29] applied the continuous-wave-equivalent formulation to short-TR steady-state acquisitions and defined $B_{1,\mathrm{rms}}$ over a single repetition period. We extend this formulation to conventional clinical acquisitions by averaging the cumulative RF power of all excitation, inversion, and refocusing pulses over the entire acquisition. This extension assumes that the combined RF pulse rate is high relative to the free- and semisolid-pool longitudinal relaxation rates (R_1^{f} and R_1^{s}), the exchange rate (R_x), and the sequence-averaged semisolid-pool saturation rate ($\langle\overline{W}\rangle$).

The combined RF pulse rate includes pulses applied to the imaged volume and slice-selective pulses applied to other slices, whether on- or off-resonance. Consequently, even acquisitions with a long TR can contain RF pulses at intervals on the order of milliseconds, primarily because of frequent TSE refocusing pulses. These intervals correspond to combined RF pulse rates on the order of $10^2\ \mathrm{s}^{-1}$, which exceed the longitudinal relaxation and exchange rates used in the model (R_1^{f}, R_1^{s}, and $R_x < 15\ \mathrm{s}^{-1}$). The modeled semisolid-pool saturation rate is also substantially lower than the combined RF pulse rate. For example, Eq. (3) gives $\langle\overline{W}\rangle \approx 6\ \mathrm{s}^{-1}$ for $B_{1,\mathrm{rms}} = 2.3\ \mu\mathrm{T}$ and $T_2^{\mathrm{s}} = 12.5\ \mu\mathrm{s}$. The semisolid longitudinal magnetization is therefore expected to change little between successive pulses, allowing its cumulative saturation to be approximated using the sequence-averaged RF power.

A.2 Near-resonance assumption and apparent equilibrium magnetization

In the continuous-wave-equivalent formulation of Teixeira et al. [29], the semisolid-pool saturation rate depends on the absorption lineshape $G(\Delta)$ evaluated at the RF frequency offset Δ. In this work, we assume a Gaussian lineshape:

$$G(\Delta) = (T_2^{\mathrm{s}}/\sqrt{2\pi})\exp[-(2\pi\Delta T_2^{\mathrm{s}})^2/2], \tag{A.1}$$

where Δ denotes the frequency offset from the water resonance. We assume near-resonance irradiation and therefore approximate the lineshape by its peak value, $G(\Delta) \approx G(0) = T_2^{\mathrm{s}}/\sqrt{2\pi}$, as used in Eq. (3). This approximation requires $2\pi|\Delta|T_2^{\mathrm{s}} \ll 1$, or equivalently,

$$|\Delta| \ll 1/(2\pi T_2^{\mathrm{s}}) \approx 12.7\ \mathrm{kHz} \qquad (\text{for } T_2^{\mathrm{s}} = 12.5\ \mu\mathrm{s}). \tag{A.2}$$

This near-resonance assumption is generally satisfied for the 3D acquisitions considered in this study, as all RF pulses are applied at or near the water resonance. The assumption is less valid for the multislice 2D T_2w acquisitions in our dataset. An RF pulse applied to one slice is off-resonant for spins in the other slices. The corresponding frequency offsets increase with interslice distance and may no longer be small relative to $1/(2\pi T_2^{\mathrm{s}})$, as required by Eq. (A.2). Explicitly modeling these contributions would require information about pulse timing, waveforms, and slice selection that is typically unavailable in retrospective clinical MRI archives. We therefore applied the assumption $G(\Delta) \approx G(0)$ to all acquisitions, yielding an analytical saturation model that can be used efficiently during self-supervised network training.

Under the continuous-wave-equivalent approximation, Teixeira et al. [29] showed that the longitudinal dynamics of the two-pool MT system can be represented by an effective single-pool model. This model is characterized by an apparent longitudinal relaxation rate, R_1^{app} (Eq. (4)), and a sequence-dependent apparent equilibrium magnetization:

$$M_0^{\mathrm{app}} = M_0^{\mathrm{f}} \frac{R_1^{\mathrm{f}} + R_x m_0^{\mathrm{s}} \left(1 - \frac{R_x(1-m_0^{\mathrm{s}}) + \langle\overline{W}\rangle}{R_1^{\mathrm{s}} + R_x(1-m_0^{\mathrm{s}}) + \langle\overline{W}\rangle}\right)}{R_1^{\mathrm{app}}}. \tag{A.3}$$

In our forward signal models (Eqs. (5)–(8)), we represented M_0^{app} by a single sequence-independent relative PD map. This simplification neglects the sequence-dependent modulation of the equilibrium-magnetization contribution by the semisolid-pool saturation rate $\langle\overline{W}\rangle$, as described by M_0^{app}.

This simplification was motivated primarily by the multislice 2D T_2w acquisitions. These acquisitions have the highest $B_{1,\mathrm{rms}}$ values in the dataset (Table 1) and would therefore show the largest changes in signal amplitude through

M_0^{app}. However, their modeled semisolid-pool saturation rates are also the most uncertain because the near-resonance assumption has limited validity for multislice 2D imaging (as discussed above). Explicitly modeling M_0^{app} would therefore propagate uncertainty in $\langle\overline{W}\rangle$ (Eq. (3)) directly into the modeled $T_2\mathrm{w}$ signal amplitude and potentially affect the estimated MPF maps.

By instead approximating the equilibrium signal amplitude using a sequence-independent PD, the influence of $\langle\overline{W}\rangle$ on the modeled $T_2\mathrm{w}$ signal is restricted to longitudinal recovery through R_1^{app}. Because the 2D $T_2\mathrm{w}$ acquisitions have long repetition times ($\mathrm{TR} \geq 2779\ \mathrm{ms}$; Table 1), longitudinal recovery is nearly complete over the relevant range of R_1^{app}. Uncertainty in $\langle\overline{W}\rangle$ therefore has only a limited effect on the modeled $T_2\mathrm{w}$ signal through this pathway. The estimated PD should therefore be interpreted as an effective, sequence-independent proton-density parameter rather than as the free-pool equilibrium magnetization M_0^{f} or the sequence-dependent apparent equilibrium magnetization M_0^{app}.

References

[1] Steven D. Wolff and Robert S. Balaban. Magnetization transfer contrast (MTC) and tissue water proton relaxation in vivo. *Magnetic Resonance in Medicine*, 10:135–144, 4 1989. ISSN 1522-2594. doi:10.1002/MRM.1910100113.

[2] R. Mark Henkelman, Xuemei Huang, Qing-San -S Xiang, G. J. Stanisz, Scott D. Swanson, and Michael J. Bronskill. Quantitative interpretation of magnetization transfer. *Magnetic Resonance in Medicine*, 29:759–766, 6 1993. ISSN 1522-2594. doi:10.1002/MRM.1910290607.

[3] Jakob Assländer, Andrew Mao, Elisa Marchetto, Erin S. Beck, Francesco La Rosa, Robert W. Charlson, Timothy M. Shepherd, and Sebastian Flassbeck. Unconstrained quantitative magnetization transfer imaging: Disentangling T1 of the free and semi-solid spin pools. *Imaging Neuroscience*, 2:2024, 5 2024. doi:10.1162/imag_a_00177.

[4] Alena A. Kisel, Anna V. Naumova, and Vasily L. Yarnykh. Macromolecular proton fraction as a myelin biomarker: Principles, validation, and applications. *Frontiers in Neuroscience*, 16:819912, 2 2022. ISSN 1662453X. doi:10.3389/fnins.2022.819912.

[5] Steven D. Wolff and Robert S. Balaban. Magnetization transfer imaging: Practical aspects and clinical applications. *Radiology*, 192:593–599, 9 1994. ISSN 00338419. doi:10.1148/radiology.192.3.8058919.

[6] R. M. Henkelman, G. J. Stanisz, and S. J. Graham. Magnetization transfer in MRI: a review. *NMR in Biomedicine*, 14:57–64, 4 2001. ISSN 1099-1492. doi:10.1002/NBM.683.

[7] John G. Sled and G. Bruce Pike. Quantitative imaging of magnetization transfer exchange and relaxation properties in vivo using MRI. *Magnetic Resonance in Medicine*, 46:923–931, 11 2001. ISSN 1522-2594. doi:10.1002/MRM.1278.

[8] John G. Sled. Modelling and interpretation of magnetization transfer imaging in the brain. *NeuroImage*, 182: 128–135, 11 2018. ISSN 1053-8119. doi:10.1016/J.NEUROIMAGE.2017.11.065.

[9] Ives R. Levesque, Paul S. Giacomini, Sridar Narayanan, Luciana T. Ribeiro, John G. Sled, Doug L. Arnold, and G. Bruce Pike. Quantitative magnetization transfer and myelin water imaging of the evolution of acute multiple sclerosis lesions†. *Magnetic Resonance in Medicine*, 63:633–640, 3 2010. ISSN 1522-2594. doi:10.1002/MRM.22244.

[10] Vasily L. Yarnykh. Pulsed Z-spectroscopic imaging of cross-relaxation parameters in tissues for human MRI: Theory and clinical applications. *Magnetic Resonance in Medicine*, 47:929–939, 5 2002. ISSN 1522-2594. doi:10.1002/MRM.10120.

[11] Ives Levesque, John G. Sled, Sridar Narayanan, A. Carlos Santos, Steven D. Brass, Simon J. Francis, Douglas L. Arnold, and G. Bruce Pike. The role of edema and demyelination in chronic T1 black holes: A quantitative magnetization transfer study†. *Journal of Magnetic Resonance Imaging*, 21:103–110, 2 2005. ISSN 1522-2586. doi:10.1002/JMRI.20231.

[12] Margareta A. Clarke, Dhairya A. Lakhani, Sijin Wen, Si Gao, Seth A. Smith, Richard Dortch, Junzhong Xu, Francesca Bagnato, and Ipek Oguz. Perilesional neurodegenerative injury in multiple sclerosis: Relation to focal lesions and impact on disability. *Multiple Sclerosis and Related Disorders*, 49:102738, 4 2021. ISSN 2211-0348. doi:10.1016/J.MSARD.2021.102738.

[13] Paul S. Giacomini, Ives R. Levesque, Luciana Ribeiro, Sridar Narayanan, Simon J. Francis, G. Bruce Pike, and Douglas L. Arnold. Measuring demyelination and remyelination in acute multiple sclerosis lesion voxels. *Archives of Neurology*, 66:375–381, 3 2009. ISSN 0003-9942. doi:10.1001/ARCHNEUROL.2008.578.

[14] Klaus Schmierer, Daniel J. Tozer, Francesco Scaravilli, Daniel R. Altmann, Gareth J. Barker, Paul S. Tofts, and David H. Miller. Quantitative magnetization transfer imaging in postmortem multiple sclerosis brain. *Journal of Magnetic Resonance Imaging*, 26:41–51, 7 2007. ISSN 1522-2586. doi:10.1002/JMRI.20984.

[15] Francesca Bagnato, Simon Hametner, Giulia Franco, Siddharama Pawate, Subramaniam Sriram, Hans Lassmann, John Gore, Seth E. Smith, and Richard Dortch. Selective inversion recovery quantitative magnetization transfer brain MRI at 7T: Clinical and postmortem validation in multiple sclerosis. *Journal of Neuroimaging*, 28:380–388, 7 2018. ISSN 1552-6569. doi:10.1111/JON.12511.

[16] Vasily L. Yarnykh. Fast macromolecular proton fraction mapping from a single off-resonance magnetization transfer measurement. *Magnetic Resonance in Medicine*, 68:166–178, 7 2012. ISSN 1522-2594. doi:10.1002/MRM.23224.

[17] Vasily L. Yarnykh. Time-efficient, high-resolution, whole brain three-dimensional macromolecular proton fraction mapping. *Magnetic Resonance in Medicine*, 75:2100–2106, 5 2016. ISSN 1522-2594. doi:10.1002/MRM.25811.

[18] Daniel F. Gochberg and John C. Gore. Quantitative imaging of magnetization transfer using an inversion recovery sequence. *Magnetic Resonance in Medicine*, 49:501–505, 3 2003. ISSN 1522-2594. doi:10.1002/MRM.10386.

[19] Richard D. Dortch, Ke Li, Daniel F. Gochberg, E. Brian Welch, Adrienne N. Dula, Ashish A. Tamhane, John C. Gore, and Seth A. Smith. Quantitative magnetization transfer imaging in human brain at 3 T via selective inversion recovery. *Magnetic Resonance in Medicine*, 66:1346–1352, 11 2011. ISSN 1522-2594. doi:10.1002/MRM.22928.

[20] Monika Gloor, K. Scheffler, and O. Bieri. Quantitative magnetization transfer imaging using balanced SSFP. *Magnetic Resonance in Medicine*, 60:691–700, 9 2008. ISSN 1522-2594. doi:10.1002/MRM.21705.

[21] Jakob Assländer, Dmitry S. Novikov, Riccardo Lattanzi, Daniel K. Sodickson, and Martijn A. Cloos. Hybrid-state free precession in nuclear magnetic resonance. *Communications Physics 2019 2:1*, 2:73, 6 2019. ISSN 2399-3650. doi:10.1038/s42005-019-0174-0.

[22] Jakob Assländer, Cem Gultekin, Andrew Mao, Xiaoxia Zhang, Quentin Duchemin, Kangning Liu, Robert W. Charlson, Timothy M. Shepherd, Carlos Fernandez-Granda, and Sebastian Flassbeck. Rapid quantitative magnetization transfer imaging: Utilizing the hybrid state and the generalized Bloch model. *Magnetic Resonance in Medicine*, 91:1478–1497, 4 2024. ISSN 1522-2594. doi:10.1002/MRM.29951.

[23] Marco Ganzetti, Nicole Wenderoth, and Dante Mantini. Whole brain myelin mapping using T1- and T2-weighted MR imaging data. *Frontiers in Human Neuroscience*, Volume 8 - 2014, 2014. ISSN 1662-5161. doi:10.3389/fnhum.2014.00671.

[24] S. Cappelle, D. Pareto, S. Sunaert, I. Smets, A. Laenen, B. Dubois, and Ph. Demaerel. T1w/FLAIR ratio standardization as a myelin marker in MS patients. *NeuroImage: Clinical*, 36:103248, 2022. ISSN 2213-1582. doi:https://doi.org/10.1016/j.nicl.2022.103248.

[25] Gerhard S. Drenthen, Walter H. Backes, Julie Staals, Maud van Dinther, Paulien H.M. Voorter, Jongho Lee, Sebastian Köhler, Willemijn J. Jansen, and Jacobus F.A. Jansen. An improved combination of T1-weighted and FLAIR contrasts for myelin water imaging. *NeuroImage*, 339:122182, 2026. ISSN 1053-8119. doi:https://doi.org/10.1016/j.neuroimage.2026.122182.

[26] Jelmer van Lune, Stefano Mandija, Oscar van der Heide, Matteo Maspero, Martin B. Schilder, Jan Willem Dankbaar, Cornelis A.T. van den Berg, and Alessandro Sbrizzi. Quantitative mapping from conventional MRI using self-supervised physics-guided deep learning: Applications to a large-scale, clinically heterogeneous dataset. *Medical Image Analysis*, 115:104295, 2027. ISSN 1361-8415. doi:https://doi.org/10.1016/j.media.2026.104295.

[27] Shihan Qiu, Lixia Wang, Pascal Sati, Anthony G Christodoulou, Yibin Xie, and Debiao Li. Physics-guided self-supervised learning for retrospective T1 and T2 mapping from conventional weighted brain MRI: Technical developments and initial validation in glioblastoma. *Magnetic Resonance in Medicine*, 92:2683–2695, 2024. ISSN 1522-2594. doi:10.1002/mrm.30226.

[28] Elisa Moya-Sáez, Rodrigo De Luis-García, and Carlos Alberola-López. A self-supervised deep learning approach to synthesize weighted images and T1, T2, and PD parametric maps based on MR physics priors. In *Proceedings of the 29th Annual Meeting of the International Society for Magnetic Resonance in Medicine (ISMRM)*, 2021. Abstract 2169.

[29] Rui Pedro A.G. Teixeira, Shaihan J. Malik, and Joseph V. Hajnal. Fast quantitative MRI using controlled saturation magnetization transfer. *Magnetic Resonance in Medicine*, 81(2):907–920, 2019. doi:10.1002/mrm.27442.

[30] Jakob Assländer and Sebastian Flassbeck. Magnetization transfer explains most of the T1 variability in the MRI literature. *Magnetic Resonance in Medicine*, 94:293–301, 7 2025. ISSN 1522-2594. doi:10.1002/MRM.30451.

[31] A. Ramani, C. Dalton, D.H. Miller, P.S. Tofts, and G.J. Barker. Precise estimate of fundamental in-vivo MT parameters in human brain in clinically feasible times. *Magnetic Resonance Imaging*, 20(10):721–731, 2002. ISSN 0730-725X. doi:https://doi.org/10.1016/S0730-725X(02)00598-2.

[32] J. B.M. Warntjes, O. Dahlqvist Leinhard, J. West, and P. Lundberg. Rapid magnetic resonance quantification on the brain: Optimization for clinical usage. *Magnetic Resonance in Medicine*, 60:320–329, 8 2008. ISSN 1522-2594. doi:10.1002/MRM.21635.

[33] Xiangrui Li, Paul S. Morgan, John Ashburner, Jolinda Smith, and Christopher Rorden. The first step for neuroimaging data analysis: DICOM to NIfTI conversion. *Journal of Neuroscience Methods*, 264:47–56, 2016. ISSN 1872-678X. doi:10.1016/j.jneumeth.2016.03.001.

[34] Brian B Avants, Nick Tustison, Gang Song, et al. Advanced normalization tools (ants). *Insight j*, 2(365):1–35, 2009.

[35] Fabian Isensee, Marianne Schell, Irada Pflueger, Gianluca Brugnara, David Bonekamp, Ulf Neuberger, Antje Wick, Heinz-Peter Schlemmer, Sabine Heiland, Wolfgang Wick, Martin Bendszus, Klaus H Maier-Hein, and Philipp Kickingereder. Automated brain extraction of multisequence MRI using artificial neural networks. *Human Brain Mapping*, 40:4952–4964, 2019. ISSN 1097-0193. doi:10.1002/hbm.24750.

[36] Nicholas J. Tustison, Brian B. Avants, Philip A. Cook, Yuanjie Zheng, Alexander Egan, Paul A. Yushkevich, and James C. Gee. N4ITK: improved N3 bias correction. *IEEE Transactions on Medical Imaging*, 29:1310–1320, 2010.

[37] Benjamin Billot, Douglas N. Greve, Oula Puonti, Axel Thielscher, Koen Van Leemput, Bruce Fischl, Adrian V. Dalca, and Juan Eugenio Iglesias. SynthSeg: Segmentation of brain MRI scans of any contrast and resolution without retraining. *Medical Image Analysis*, 86:102789, 5 2023. ISSN 1361-8415. doi:10.1016/J.MEDIA.2023.102789.

[38] Tun Wiltgen, Julian McGinnis, Sarah Schlaeger, Florian Kofler, CuiCi Voon, Achim Berthele, Daria Bischl, Lioba Grundl, Nikolaus Will, Marie Metz, David Schinz, Dominik Sepp, Philipp Prucker, Benita Schmitz-Koep, Claus Zimmer, Bjoern Menze, Daniel Rueckert, Bernhard Hemmer, Jan Kirschke, Mark Mühlau, and Benedikt Wiestler. LST-AI: A deep learning ensemble for accurate MS lesion segmentation. *NeuroImage: Clinical*, 42:103611, 2024. ISSN 2213-1582. doi:https://doi.org/10.1016/j.nicl.2024.103611. URL `https://www.sciencedirect.com/science/article/pii/S2213158224000500`.

[39] Vincent Dumoulin, Jonathon Shlens, and Manjunath Kudlur. A learned representation for artistic style. In *International Conference on Learning Representations (ICLR)*, 2017.

[40] Xun Huang and Serge Belongie. Arbitrary style transfer in real-time with adaptive instance normalization. *Proceedings of the IEEE International Conference on Computer Vision*, 2017-October:1510–1519, 12 2017. ISSN 15505499. doi:10.1109/ICCV.2017.167.

[41] Tero Karras, Samuli Laine, and Timo Aila. A style-based generator architecture for generative adversarial networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 43(12):4217–4228, 2021. doi:10.1109/TPAMI.2020.2970919.

[42] Lingjing Chen, Chengxiu Zhang, Yinqiao Yi, Yida Wang, Yang Song, Xu Yan, Shengfang Xu, Dalin Zhu, Mengqiu Cao, Yan Zhou, Chenglong Wang, and Guang Yang. A physics-driven neural network with parameter embedding for generating quantitative MR maps from weighted images. *Medical Physics*, 53:e70394, 3 2026. ISSN 2473-4209. doi:10.1002/MP.70394.

[43] Shihan Qiu, Yibin Xie, Anthony G. Christodoulou, Pascal Sati, Marcel Maya, Nancy L. Sicotte, and Debiao Li. Protocol-aware unsupervised retrospective T1 and T2 mapping with diverse imaging parameters. In *Proceedings of the 32th Annual Meeting of the International Society for Magnetic Resonance in Medicine (ISMRM)*, 2024. Abstract 1192.

[44] Pedro Borges, Virginia Fernandez, Petru Daniel Tudosiu, Parashkev Nachev, Sebastien Ourselin, and M. Jorge Cardoso. Unsupervised heteromodal physics-informed representation of MRI data: Tackling data harmonisation, imputation and domain shift. In Jelmer M. Wolterink, David Svoboda, Can Zhao, and Virginia Fernandez, editors, *Simulation and Synthesis in Medical Imaging*, pages 53–63, Cham, 2023. Springer Nature Switzerland. ISBN 978-3-031-44689-4. doi:10.1007/978-3-031-44689-4_6.

[45] Djork Arné Clevert, Thomas Unterthiner, and Sepp Hochreiter. Fast and accurate deep network learning by exponential linear units (ELUs). *4th International Conference on Learning Representations, ICLR 2016 - Conference Track Proceedings*, 11 2015.

[46] Peter van Gelderen, Xu Jiang, and Jeff H. Duyn. Effects of magnetization transfer on T1 contrast in human brain white matter. *NeuroImage*, 128:85–95, 3 2016. ISSN 1053-8119. doi:10.1016/J.NEUROIMAGE.2015.12.032.

[47] Reed F. Busse, Hari Hariharan, Anthony Vu, and Jean H. Brittain. Fast spin echo sequences with very long echo trains: Design of variable refocusing flip angle schedules and generation of clinical T2 contrast. *Magnetic Resonance in Medicine*, 55:1030–1037, 5 2006. ISSN 1522-2594. doi:10.1002/MRM.20863.

[48] Jorge Zavala Bojorquez, Stéphanie Bricq, Clement Acquitter, François Brunotte, Paul M. Walker, and Alain Lalande. What are normal relaxation times of tissues at 3 T? *Magnetic Resonance Imaging*, 35:69–80, 1 2017. ISSN 0730-725X. doi:10.1016/J.MRI.2016.08.021.

[49] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. 2017. URL `https://arxiv.org/abs/1412.6980`.

[50] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 1026–1034, 2015.

[51] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance deep learning library. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019.

[52] Fernando Pérez-García, Rachel Sparks, and Sébastien Ourselin. TorchIO: A Python library for efficient loading, preprocessing, augmentation and patch-based sampling of medical images in deep learning. *Computer Methods and Programs in Biomedicine*, 208:106236, 2021. ISSN 0169-2607. doi:10.1016/j.cmpb.2021.106236.

[53] S. H. Koenig, R. D. Brown, M. Spiller, and N. Lundbom. Relaxometry of brain: Why white matter appears bright in MRI. *Magnetic Resonance in Medicine*, 14:482–495, 6 1990. ISSN 1522-2594. doi:10.1002/MRM.1910140306.

[54] Gunther Helms and Gisela E Hagberg. In vivo quantification of the bound pool T1 in human white matter using the binary spin–bath model of progressive magnetization transfer saturation. *Physics in Medicine & Biology*, 54 (23):N529, nov 2009. doi:10.1088/0031-9155/54/23/N01.

[55] Peter van Gelderen, Xu Jiang, and Jeff H. Duyn. Rapid measurement of brain macromolecular proton fraction with transient saturation transfer MRI. *Magnetic Resonance in Medicine*, 77:2174–2185, 6 2017. ISSN 15222594. doi:10.1002/MRM.26304.

[56] Greg J. Stanisz, Ewa E. Odrobina, Joseph Pun, Michael Escaravage, Simon J. Graham, Michael J. Bronskill, and R. Mark Henkelman. T1, T2 relaxation and magnetization transfer in tissue at 3T. *Magnetic Resonance in Medicine*, 54(3):507–512, 2005. doi:https://doi.org/10.1002/mrm.20605.

[57] Davies GR, Tozer DJ, Cercignani M, Ramani A, Dalton CM, Thompson AJ, Barker GJ, Tofts PS, and Miller DH. Estimation of the macromolecular proton fraction and bound pool T2 in multiple sclerosis. *Multiple Sclerosis Journal*, 10(6):607–613, 2004. doi:10.1191/1352458504ms1105oa.

[58] M. Cercignani, B. Basile, B. Spanò, G. Comanducci, F. Fasano, C. Caltagirone, U. Nocentini, and M. Bozzali. Investigation of quantitative magnetisation transfer parameters of lesions and normal appearing white matter in multiple sclerosis. *NMR in Biomedicine*, 22(6):646–653, 2009. doi:https://doi.org/10.1002/nbm.1379.

[59] Vasily L. Yarnykh, James D. Bowen, Alexey Samsonov, Pavle Repovic, Angeli Mayadev, Peiqing Qian, Beena Gangadharan, Bart P. Keogh, Kenneth R. Maravilla, and Lily K.Jung Henson. Fast whole-brain three-dimensional macromolecular proton fraction mapping in multiple sclerosis. *Radiology*, 274:210–220, 1 2015. ISSN 15271315. doi:10.1148/radiol.14140528.

[60] S. J. Graham and R. Mark Henkelman. Understanding pulsed magnetization transfer. *Journal of Magnetic Resonance Imaging*, 7:903–912, 9 1997. ISSN 1522-2586. doi:10.1002/JMRI.1880070520.

[61] Shaihan J. Malik, Rui Pedro A.G. Teixeira, and Joseph V. Hajnal. Extended phase graph formalism for systems with magnetization transfer and exchange. *Magnetic Resonance in Medicine*, 80:767–779, 8 2018. ISSN 1522-2594. doi:10.1002/MRM.27040.

[62] Kathryn E. Keenan, Joshua R. Biller, Jana G. Delfino, Michael A. Boss, Mark D. Does, Jeffrey L. Evelhoch, Mark A. Griswold, Jeffrey L. Gunter, R. Scott Hinks, Stuart W. Hoffman, Geena Kim, Riccardo Lattanzi, Xiaojuan Li, Luca Marinelli, Gregory J. Metzger, Pratik Mukherjee, Robert J. Nordstrom, Adele P. Peskin, Elena Perez, Stephen E. Russek, Berkman Sahiner, Natalie Serkova, Amita Shukla-Dave, Michael Steckner, Karl F. Stupic, Lisa J. Wilmes, Holden H. Wu, Huiming Zhang, Edward F. Jackson, and Daniel C. Sullivan. Recommendations towards standards for quantitative MRI (qMRI) and outstanding needs. *Journal of Magnetic Resonance Imaging: JMRI*, 49:e26–e39, 2019. ISSN 1522-2586. doi:10.1002/jmri.26598.

[63] Claus Kiefer, Lisa Brockhaus, Katja Cattapan-Ludewig, Pietro Ballinari, Yuliya Burren, Gerhard Schroth, and Roland Wiest. Multi-parametric classification of Alzheimer's disease and mild cognitive impairment: The impact of quantitative magnetization transfer MR imaging. *NeuroImage*, 48:657–667, 12 2009. ISSN 1053-8119. doi:10.1016/J.NEUROIMAGE.2009.07.005.

[64] Jilu P. Mole, Fabrizio Fasano, John Evans, Rebecca Sims, Emma Kidd, John P. Aggleton, and Claudia Metzler-Baddeley. APOE-ε4-related differences in left thalamic microstructure in cognitively healthy adults. *Scientific Reports 2020 10:1*, 10:19787, 11 2020. ISSN 2045-2322. doi:10.1038/s41598-020-75992-9.

[65] Paula Trujillo, Paul E. Summers, Alex K. Smith, Seth A. Smith, Luca T. Mainardi, Sergio Cerutti, Daniel O. Claassen, and Antonella Costa. Pool size ratio of the substantia nigra in Parkinson's disease derived from two different quantitative magnetization transfer approaches. *Neuroradiology*, 59:1251–1263, 12 2017. ISSN 14321920. doi:10.1007/S00234-017-1911-2.

[66] José Bourbon-Teles, Sonya Bells, Derek K. Jones, Elizabeth Coulthard, Anne Rosser, and Claudia Metzler-Baddeley. Myelin breakdown in human Huntington's disease: Multi-modal evidence from diffusion MRI and quantitative magnetization transfer. *Neuroscience*, 403:79–92, 4 2019. ISSN 0306-4522. doi:10.1016/J.NEUROSCIENCE.2017.05.042.

[67] Marina Yu Khodanovich, Ilya L. Gubskiy, Marina S. Kudabaeva, Darya D. Namestnikova, Alena A. Kisel, Tatyana V. Anan'ina, Yana A. Tumentceva, Lilia R. Mustafina, and Vasily L. Yarnykh. Long-term monitoring of chronic demyelination and remyelination in a rat ischemic stroke model using macromolecular proton fraction mapping. *Journal of Cerebral Blood Flow and Metabolism*, 41:2856–2869, 11 2021. ISSN 15597016. doi:10.1177/0271678X211020860.

SUPPLEMENTARY MATERIALS OF: "RETROSPECTIVE MACROMOLECULAR POOL FRACTION MAPPING FROM CONVENTIONAL CLINICAL MRI: A PHYSICS-GUIDED SELF-SUPERVISED DEEP LEARNING FRAMEWORK"

A PREPRINT

Jelmer van Lune[1], **Miha Fuderer**[1], **Oscar van der Heide**[1], **Matteo Maspero**[1], **Martin B. Schilder**[1], **Jan Willem Dankbaar**[2], **Stefano Mandija**[1], **and Alessandro Sbrizzi**[1]

[1]Computational Imaging Group for MRI Therapy & Diagnostics, University Medical Center Utrecht, Utrecht, The Netherlands

[2]Department of Radiology, University Medical Center Utrecht, Utrecht, The Netherlands

S1 Characteristics of the UMCU multiple sclerosis subset

Table S1.1: Demographic and acquisition parameters for the four subjects with clinically confirmed MS in the UMCU test set. Subjects A–D correspond to Rows A–D in Figure 4 in the main manuscript. All examinations were performed on Philips 3 T systems. GRE: Gradient-Echo; TSE: Turbo Spin Echo; IR: Inversion Recovery; MPRAGE: Magnetization-Prepared Rapid Gradient-Echo; NA: Not Applicable; BW: Pixel Bandwidth; FA: Flip Angle; ETL: Echo Train Length.

	Spoiled GRE Protocol		**MPRAGE Protocol**	
Parameter	**Subject A**	**Subject B**	**Subject C**	**Subject D**
Demographics & Hardware				
Age [years]	35	51	13	49
Sex	Male	Male	Male	Female
Scanner model	Ingenia	Ingenia Elition X	Ingenia	Ingenia CX
T_1**w Sequence**				
Sequence	3D Spoiled GRE	3D Spoiled GRE	3D MPRAGE	3D MPRAGE
TR [ms]	5.30	5.30	2500	2500
TE [ms]	2.40	2.41	3.86	4.02
TI [ms]	NA	NA	1004	1004
FA [°]	10	10	8	8
$B_{1,\mathrm{rms}}$ **[μT]**	0.93	0.93	0.73	0.73
ETL	NA	NA	NA	NA
BW [Hz/px]	539	539	190	190
Rec. matrix (in-plane)	480×480	480×480	240×240	240×240
Rec. matrix (slices)	320	320	190	160
T_2**w Sequence**				
Sequence	2D TSE	2D TSE	2D TSE	2D TSE
TR [ms]	3906	4123	4340	3849
TE [ms]	80	80	80	80
TI [ms]	NA	NA	NA	NA
FA [°]	90	90	90	90
$B_{1,\mathrm{rms}}$ **[μT]**	2.34	2.34	2.34	2.34
ETL	15	15	15	15
BW [Hz/px]	195	195	195	194
Rec. matrix (in-plane)	512×512	512×512	512×512	512×512
Rec. matrix (slices)	35	37	39	35
FLAIR Sequence				
Sequence	3D TSE (IR)	3D TSE (IR)	3D TSE (IR)	3D TSE (IR)
TR [ms]	4800	4800	4800	4800
TE [ms]	355.5	355.5	338.3	309.0
TI [ms]	1650	1650	1650	1650
FA [°]	90	90	90	90
$B_{1,\mathrm{rms}}$ **[μT]**	1.00	1.00	1.00	1.00
ETL	199	199	199	199
BW [Hz/px]	1075	1075	1075	1250
Rec. matrix (in-plane)	240×240	240×240	240×240	240×240
Rec. matrix (slices)	267	267	317	267

S2 TSE echo time correction

Turbo spin-echo (TSE) acquisitions use echo trains that can produce stimulated-echo contributions, causing the effective T_2 weighting to differ from that predicted by the nominal echo time. In the forward signal models for the T_2w and FLAIR acquisitions, we accounted for this difference by replacing the nominal echo time with a contrast-equivalent effective echo time, $\mathrm{TE}_{\mathrm{eff}} = c_{\mathrm{TE}}\,\mathrm{TE}_{\mathrm{nom}}$, where c_{TE} denotes the echo-time correction factor [47].

For the Philips acquisitions, console-reported contrast-equivalent echo times yielded $c_{\mathrm{TE}} = 0.90$ for 2D T_2w and $c_{\mathrm{TE}} = 0.42$ for 3D FLAIR. These factors corresponded to an effective echo time of 72 ms for the nominal T_2w echo time of 80 ms, and a median effective echo time of approximately 132 ms for the median nominal FLAIR echo time of 314.5 ms.

Equivalent console values were unavailable for the Siemens acquisitions. Correction factors were therefore estimated using extended phase graph (EPG) simulations with representative tissue relaxation times of $T_1 = 1000$ ms and $T_2 = 100$ ms. Because the echo spacing and center-echo index were not provided in the metadata, the center of k-space was assumed to occur halfway through the echo train. For Siemens 2D T_2w, the constant refocusing flip angles ($119°$–$150°$), echo-train lengths (13–20), and nominal echo times reported in the metadata were simulated directly, yielding correction factors between 0.860 and 0.958 ($\mathrm{TE}_{\mathrm{eff}} = 76.5$–$100.1$ ms).

The Siemens 3D FLAIR acquisitions used a SPACE readout with a proprietary variable refocusing flip-angle schedule that was not exported in the metadata. We approximated this schedule using the pseudo-steady-state target-signal formulation of Busse et al. [47], constrained to a maximum refocusing angle of $120°$. Calibrating the normalized target plateau amplitude to $S_{\mathrm{target}} = 0.292$ reproduced the Philips 3D FLAIR correction factor ($c_{\mathrm{TE}} = 0.42$) under representative sequence parameters ($\mathrm{TE} = 314.5$ ms, $\mathrm{ETL} = 199$). Applying this calibrated schedule to the Siemens acquisitions yielded correction factors of 0.502–0.509, corresponding to a median effective echo time of approximately 196 ms.

S3 Hyperparameter tuning

Table S3.1 outlines the hyperparameter grid search space and the specific values explored during the optimization process to identify the optimal model configuration.

Table S3.1: Hyperparameter search space and range of values evaluated during the optimization process.

Hyperparameter	**Explored values**
Learning rate	[10^{-6}, 10^{-5}, 10^{-4}, 10^{-3}, 10^{-2}]
$\lambda_{\text{PD-constraint}}$	[0.1, 0.2, 0.25, 0.50, 0.75, 1.0, 2.0]
$\lambda_{T_1\text{-prior}}$	[0.0001, 0.001, 0.002, 0.0025, 0.005, 0.01]

S4 Prospective volunteer acquisition parameters

Table S4.1: Acquisition parameters for the two clinical protocols evaluated in the prospective healthy volunteer on a 3 T Philips Ingenia system. Protocols are designated by their respective T_1w sequence (Spoiled GRE protocol vs. MPRAGE protocol). GRE: Spoiled Gradient-Echo; TSE: Turbo Spin Echo; IR: Inversion Recovery; NA: Not Applicable; BW: Pixel Bandwidth; FA: Flip Angle; ETL: Echo Train Length.

	Spoiled GRE Protocol			**MPRAGE Protocol**		
Parameter	$\boldsymbol{T_1}$**w**	$\boldsymbol{T_2}$**w**	**FLAIR**	$\boldsymbol{T_1}$**w**	$\boldsymbol{T_2}$**w**	**FLAIR**
Sequence	3D Spoiled GRE	2D TSE	3D TSE (IR)	3D MPRAGE	2D TSE	3D TSE (IR)
TR [ms]	5.07	4340	4800	2500	4990	4800
TE [ms]	2.27	80	336	3.79	80	337
TI [ms]	NA	NA	1650	1004	NA	1650
FA [°]	10	90	90	8	90	90
$\boldsymbol{B_{1,\mathrm{rms}}}$ **[μT]**	0.95	2.34	1.00	0.73	2.34	1.00
ETL	NA	15	199	NA	15	199
BW [Hz/px]	542	195	1047	191	198	1033
Rec. matrix (in-plane)	560×560	512×512	240×240	256×256	560×560	256×256
Rec. matrix (slices)	380	39	317	180	45	300

S5 Quantitative results for all estimated maps

Table S5.1: Quantitative parameter estimates in white matter (WM) and gray matter (GM) across the clinical test set ($n = 1{,}237$ MRI sessions), stratified by the three vendor-by-protocol groups. Values represent mean ± standard deviation across MRI sessions. The between-group coefficient of variation (CV) was calculated from the three vendor-by-protocol group means (Philips Spoiled GRE, Philips MPRAGE, and Siemens MPRAGE).

		Philips (UMCU)		**Siemens (MR-RATE)**	
Metric	**Total** ($n = 1{,}237$)	**Spoiled GRE** ($n = 653$)	**MPRAGE** ($n = 221$)	**MPRAGE** ($n = 363$)	**Between-group CV (%)**
White Matter (WM)					
m_0^{s}	0.182 ± 0.007	0.186 ± 0.006	0.179 ± 0.004	0.177 ± 0.006	2.65
T_1^{f} [ms]	1982 ± 58	1967 ± 60	2017 ± 64	1987 ± 40	1.26
T_2^{f} [ms]	76.6 ± 13.1	73.8 ± 13.7	80.9 ± 18.9	78.8 ± 1.2	4.68
PD [a.u.]	0.671 ± 0.014	0.673 ± 0.010	0.664 ± 0.012	0.672 ± 0.019	0.74
Gray Matter (GM)					
m_0^{s}	0.144 ± 0.009	0.149 ± 0.005	0.147 ± 0.005	0.134 ± 0.006	5.63
T_1^{f} [ms]	2240 ± 74	2277 ± 60	2226 ± 61	2183 ± 66	2.12
T_2^{f} [ms]	87.4 ± 13.6	86.2 ± 14.1	91.9 ± 19.8	87.0 ± 3.8	3.52
PD [a.u.]	0.716 ± 0.020	0.711 ± 0.012	0.703 ± 0.017	0.733 ± 0.021	2.21

Table S5.2: Quantitative parameter estimates in demyelinating WM lesions, normal-appearing white matter (NAWM), and normal-appearing gray matter (NAGM) for the MR-RATE demyelinating-disease test subset ($n = 70$ MRI sessions; Siemens Magnetom Vida 3 T). Values represent mean ± standard deviation across MRI sessions. Relative percentage differences were calculated as $\Delta = (\mu_{\mathrm{lesion}} - \mu_{\mathrm{tissue}})/\mu_{\mathrm{tissue}} \times 100\%$, where the reference tissue was either NAWM or NAGM.

Metric	**NAWM**	**NAGM**	**Lesion**	**Δ vs. NAWM (%)**	**Δ vs. NAGM (%)**
m_0^{s}	0.176 ± 0.005	0.133 ± 0.005	0.113 ± 0.015	-35.6%	-15.2%
T_1^{f} [ms]	1981 ± 38	2182 ± 59	2369 ± 209	$+19.6\%$	$+8.6\%$
T_2^{f} [ms]	78.5 ± 0.8	86.5 ± 2.1	105.8 ± 16.2	$+34.8\%$	$+22.3\%$
PD [a.u.]	0.678 ± 0.015	0.736 ± 0.018	0.842 ± 0.045	$+24.2\%$	$+14.4\%$

Table S5.3: Subject-level quantitative parameter estimates in normal-appearing white matter (NAWM), normal-appearing gray matter (NAGM), and demyelinating WM lesions for the four clinically confirmed multiple sclerosis subjects in the UMCU test set acquired on Philips 3 T systems. Values represent voxel-wise mean ± standard deviation within each tissue compartment for each subject. Subjects A–D correspond to Rows A–D in Figure 4 and Supplementary Table S1.1. Relative differences were calculated separately for each subject as $\Delta = (\mu_{\text{lesion}} - \mu_{\text{tissue}})/\mu_{\text{tissue}} \times 100\%$, where the reference tissue was either NAWM or NAGM.

Subject	**Protocol**	**Compartment**	m_0^s	T_1^f [ms]	T_2^f [ms]	PD [a.u.]
		NAWM	0.193 ± 0.026	1922 ± 191	69.0 ± 6.6	0.656 ± 0.051
		NAGM	0.150 ± 0.036	2245 ± 430	82.8 ± 21.5	0.703 ± 0.090
Subject A	Spoiled GRE	Lesion	0.127 ± 0.036	2447 ± 580	114.6 ± 57.6	0.769 ± 0.090
		Δ vs. NAWM (%)	-34.2%	$+27.3\%$	$+66.1\%$	$+17.2\%$
		Δ vs. NAGM (%)	-15.3%	$+9.0\%$	$+38.4\%$	$+9.4\%$
		NAWM	0.182 ± 0.023	1949 ± 180	68.1 ± 4.5	0.671 ± 0.048
		NAGM	0.143 ± 0.031	2262 ± 354	78.3 ± 10.1	0.711 ± 0.081
Subject B	Spoiled GRE	Lesion	0.121 ± 0.038	2520 ± 599	108.4 ± 39.1	0.787 ± 0.100
		Δ vs. NAWM (%)	-33.5%	$+29.3\%$	$+59.2\%$	$+17.3\%$
		Δ vs. NAGM (%)	-15.4%	$+11.4\%$	$+38.4\%$	$+10.7\%$
		NAWM	0.170 ± 0.021	1988 ± 126	71.2 ± 5.9	0.701 ± 0.068
		NAGM	0.131 ± 0.021	2253 ± 224	83.6 ± 7.9	0.756 ± 0.092
Subject C	MPRAGE	Lesion	0.127 ± 0.029	2332 ± 473	115.4 ± 141.5	0.821 ± 0.102
		Δ vs. NAWM (%)	-25.3%	$+17.3\%$	$+62.1\%$	$+17.1\%$
		Δ vs. NAGM (%)	-3.1%	$+3.5\%$	$+38.0\%$	$+8.6\%$
		NAWM	0.179 ± 0.017	1959 ± 87	70.4 ± 4.7	0.658 ± 0.050
		NAGM	0.142 ± 0.027	2188 ± 264	81.7 ± 11.5	0.716 ± 0.097
Subject D	MPRAGE	Lesion	0.137 ± 0.034	2299 ± 520	107.1 ± 34.1	0.783 ± 0.123
		Δ vs. NAWM (%)	-23.5%	$+17.4\%$	$+52.1\%$	$+19.0\%$
		Δ vs. NAGM (%)	-3.5%	$+5.1\%$	$+31.1\%$	$+9.4\%$

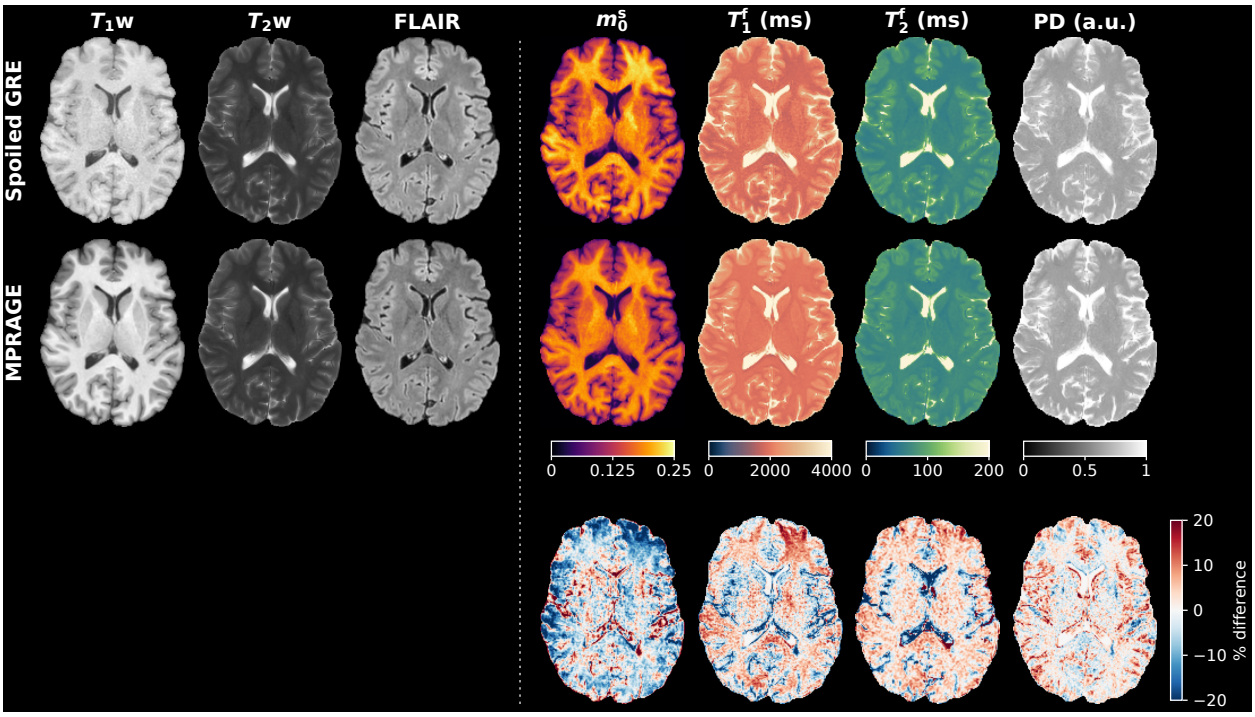

Figure S5.1: Prospective within-subject reproducibility in a healthy volunteer scanned with two different clinical protocols. Columns show the generated macromolecular pool fraction (m_0^s; MPF), free-pool longitudinal relaxation time (T_1^f), free-pool transverse relaxation time (T_2^f), and relative proton density (PD) maps. The top row corresponds to the Spoiled GRE protocol, the middle row to the MPRAGE protocol, and the bottom row shows voxel-wise percentage difference maps between the two generated quantitative maps. Mean absolute percentage differences across brain parenchyma (WM+GM) were 7.7% for m_0^s, 7.7% for T_1^f, 4.8% for T_2^f, and 4.8% for PD.

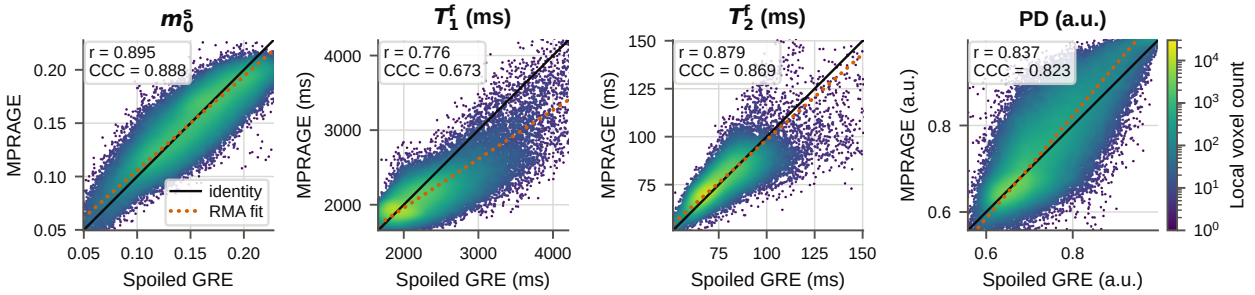

Figure S5.2: Voxel-wise agreement between quantitative maps estimated from the Spoiled GRE and MPRAGE protocols in a healthy volunteer. From left to right, panels show the macromolecular pool fraction (m_0^s; MPF), free-pool longitudinal relaxation time (T_1^f), free-pool transverse relaxation time (T_2^f), and relative proton density (PD). Each point represents a brain parenchymal voxel (WM+GM), colored by local voxel density on a logarithmic scale shared by all panels. The solid line denotes identity and the dotted line the reduced major axis (RMA) fit. Pearson correlation coefficients (r) and concordance correlation coefficients (CCC) are reported for each quantitative map to assess cross-protocol reproducibility.

S6 Impact of acquisition conditioning

Table S6.1: Ablation study evaluating the impact of acquisition-conditioned adaptive instance normalization (AdaIN) on quantitative parameter estimates across the multicenter clinical test set ($n = 1{,}237$ MRI sessions). Values represent mean ± standard deviation across MRI sessions for white matter (WM) and gray matter (GM). Results are shown for the proposed acquisition-conditioned model (AdaIN) and an otherwise identical unconditioned baseline model (No AdaIN). Group-specific estimates are reported for the three vendor-by-protocol groups (Philips Spoiled GRE, Philips MPRAGE, and Siemens MPRAGE). The between-group coefficient of variation (CV) was calculated from the corresponding groups.

			Philips (UMCU)		**Siemens (MR-RATE)**	
Parameter	**Conditioning**	**Total** ($n = 1{,}237$)	**Spoiled GRE** ($n = 653$)	**MPRAGE** ($n = 221$)	**MPRAGE** ($n = 363$)	**Between-group CV (%)**
White Matter (WM)						
m_0^{s}	AdaIN	0.182 ± 0.007	0.186 ± 0.006	0.179 ± 0.004	0.177 ± 0.006	2.65
	No AdaIN	0.189 ± 0.018	0.205 ± 0.008	0.174 ± 0.004	0.171 ± 0.004	10.27
T_1^{f} [ms]	AdaIN	1982 ± 58	1967 ± 60	2017 ± 64	1987 ± 40	1.26
	No AdaIN	2012 ± 62	2059 ± 40	1971 ± 45	1951 ± 20	2.87
T_2^{f} [ms]	AdaIN	76.6 ± 13.1	73.8 ± 13.7	80.9 ± 18.9	78.8 ± 1.2	4.68
	No AdaIN	75.0 ± 10.7	72.0 ± 12.1	76.2 ± 12.1	79.6 ± 1.2	5.03
PD [a.u.]	AdaIN	0.671 ± 0.014	0.673 ± 0.010	0.664 ± 0.012	0.672 ± 0.019	0.74
	No AdaIN	0.679 ± 0.013	0.685 ± 0.009	0.670 ± 0.010	0.673 ± 0.013	1.19
Gray Matter (GM)						
m_0^{s}	AdaIN	0.144 ± 0.009	0.149 ± 0.005	0.147 ± 0.005	0.134 ± 0.006	5.63
	No AdaIN	0.146 ± 0.014	0.158 ± 0.006	0.138 ± 0.005	0.129 ± 0.006	10.37
T_1^{f} [ms]	AdaIN	2240 ± 74	2277 ± 60	2226 ± 61	2183 ± 66	2.12
	No AdaIN	2176 ± 68	2215 ± 55	2135 ± 47	2129 ± 53	2.21
T_2^{f} [ms]	AdaIN	87.4 ± 13.6	86.2 ± 14.1	91.9 ± 19.8	87.0 ± 3.8	3.52
	No AdaIN	85.1 ± 11.3	82.9 ± 13.1	86.5 ± 12.3	88.1 ± 4.4	3.12
PD [a.u.]	AdaIN	0.716 ± 0.020	0.711 ± 0.012	0.703 ± 0.017	0.733 ± 0.021	2.21
	No AdaIN	0.725 ± 0.016	0.728 ± 0.012	0.706 ± 0.012	0.731 ± 0.016	1.86

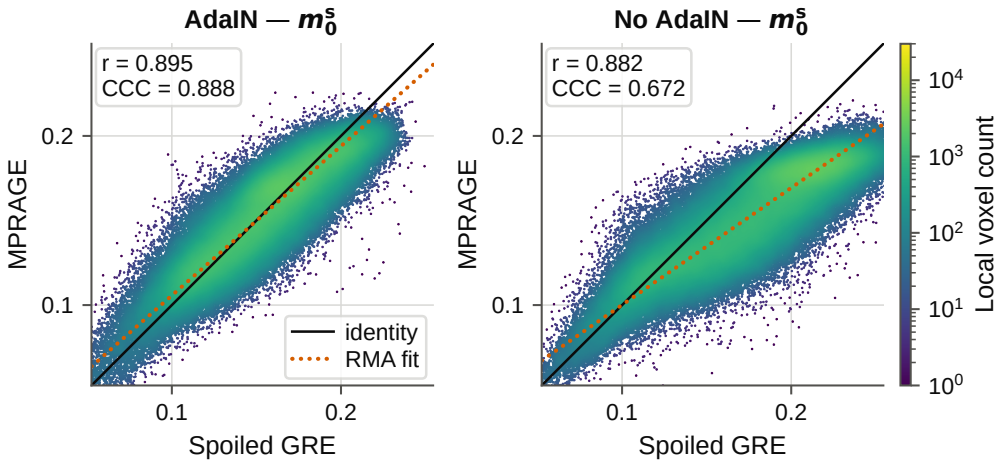

Figure S6.1: Impact of acquisition conditioning on cross-protocol reproducibility of the generated macromolecular pool fraction (m_0^s; MPF) maps in the prospective healthy volunteer. Voxel-wise agreement between MPF maps estimated from the Spoiled GRE and MPRAGE protocols is shown for the proposed acquisition-conditioned model (AdaIN; left) and an otherwise identical unconditioned baseline model (No AdaIN; right). Each point represents a brain parenchymal voxel (WM+GM), colored by local voxel density on a logarithmic scale; both panels share the same axis ranges and color scale. The solid line denotes identity and the dotted line the reduced major axis (RMA) fit. Pearson correlation coefficients (r) and concordance correlation coefficients (CCC) quantify cross-protocol agreement. Removal of acquisition conditioning reduced reproducibility, resulting in a lower CCC.

S7 Ablation of the regularization terms

The training objective in Eq. (12) combines a data-consistency loss with two regularization terms: a soft lower bound on PD and a prior on white-matter T_1^{f}. To evaluate their individual contributions, we trained four models differing only in the corresponding loss weights ($\lambda_{\text{PD-constraint}} \in \{0.5, 0\}$ and $\lambda_{T_1\text{-prior}} \in \{0.0025, 0\}$). Each model was evaluated across the clinical test set using the between-group coefficient of variation (CV) and in the prospective healthy volunteer using mean absolute percentage difference (MPD), concordance correlation coefficient (CCC), and Pearson correlation (r) (Supplementary Tables S7.1 and S7.2).

Removing the white-matter T_1^{f} prior primarily affected the estimated T_1^{f} and m_0^{s} maps. Compared with the full model, these configurations exhibited larger differences between vendor-by-protocol groups and reduced cross-protocol agreement in the prospective volunteer experiment. When both regularization terms were omitted, the estimated T_1^{f} values were substantially lower than those obtained with the full model.

Removing the PD constraint primarily affected PD and T_2^{f}, producing larger between-group variation and poorer prospective agreement, as reflected by higher MPD and generally lower CCC. It also shifted m_0^{s} to higher values, particularly in WM, although its cross-protocol agreement changed less markedly. Similar trends were observed when both regularization terms were omitted simultaneously.

Table S7.1: Regularization ablation: quantitative parameter estimates in white matter (WM) and gray matter (GM) across the clinical test set ($n = 1{,}237$ MRI sessions) for models trained with and without the PD constraint ($\lambda_{\text{PD-constraint}} \in \{0.5, 0\}$) and white-matter T_1^{f} prior ($\lambda_{T_1\text{-prior}} \in \{0.0025, 0\}$). Values denote mean ± standard deviation across sessions. Between-group coefficient of variation (CV) quantifies variation across the three vendor-by-protocol group means (Philips Spoiled GRE, Philips MPRAGE, and Siemens MPRAGE).

			Philips (UMCU)		**Siemens (MR-RATE)**	
Parameter	**Configuration**	**Total** ($n = 1{,}237$)	**Spoiled GRE** ($n = 653$)	**MPRAGE** ($n = 221$)	**MPRAGE** ($n = 363$)	**Between-group CV (%)**
White Matter (WM)						
m_0^{s}	Full model	0.182 ± 0.007	0.186 ± 0.006	0.179 ± 0.004	0.177 ± 0.006	2.65
	No T_1^{f} prior	0.177 ± 0.013	0.165 ± 0.005	0.187 ± 0.003	0.192 ± 0.004	7.96
	No PD constraint	0.236 ± 0.012	0.230 ± 0.011	0.239 ± 0.009	0.245 ± 0.009	3.26
	Neither	0.169 ± 0.011	0.159 ± 0.005	0.179 ± 0.004	0.180 ± 0.004	6.89
T_1^{f} [ms]	Full model	1982 ± 58	1967 ± 60	2017 ± 64	1987 ± 40	1.26
	No T_1^{f} prior	1866 ± 284	1611 ± 94	2110 ± 92	2177 ± 65	15.71
	No PD constraint	1935 ± 52	1935 ± 51	1966 ± 52	1915 ± 44	1.33
	Neither	1301 ± 115	1293 ± 111	1381 ± 100	1267 ± 107	4.54
T_2^{f} [ms]	Full model	76.6 ± 13.1	73.8 ± 13.7	80.9 ± 18.9	78.8 ± 1.2	4.68
	No T_1^{f} prior	77.5 ± 12.4	75.0 ± 13.1	80.7 ± 17.6	80.0 ± 1.2	3.93
	No PD constraint	86.0 ± 13.6	82.2 ± 14.0	96.6 ± 17.3	86.2 ± 1.1	8.41
	Neither	86.2 ± 13.8	82.3 ± 14.3	95.9 ± 17.6	87.5 ± 1.2	7.78
PD [a.u.]	Full model	0.671 ± 0.014	0.673 ± 0.010	0.664 ± 0.012	0.672 ± 0.019	0.74
	No T_1^{f} prior	0.664 ± 0.013	0.667 ± 0.011	0.660 ± 0.013	0.661 ± 0.017	0.52
	No PD constraint	0.541 ± 0.045	0.580 ± 0.016	0.503 ± 0.018	0.495 ± 0.019	8.94
	Neither	0.537 ± 0.045	0.576 ± 0.018	0.506 ± 0.017	0.487 ± 0.020	8.91
Gray Matter (GM)						
m_0^{s}	Full model	0.144 ± 0.009	0.149 ± 0.005	0.147 ± 0.005	0.134 ± 0.006	5.63
	No T_1^{f} prior	0.153 ± 0.009	0.147 ± 0.005	0.164 ± 0.004	0.159 ± 0.005	5.73
	No PD constraint	0.177 ± 0.011	0.176 ± 0.010	0.185 ± 0.010	0.173 ± 0.011	3.52
	Neither	0.151 ± 0.008	0.145 ± 0.005	0.163 ± 0.003	0.156 ± 0.004	5.64
T_1^{f} [ms]	Full model	2240 ± 74	2277 ± 60	2226 ± 61	2183 ± 66	2.12
	No T_1^{f} prior	2414 ± 233	2214 ± 97	2577 ± 86	2674 ± 90	9.76
	No PD constraint	2176 ± 80	2219 ± 62	2192 ± 48	2090 ± 53	3.14
	Neither	1934 ± 147	1916 ± 137	1985 ± 125	1936 ± 167	1.83
T_2^{f} [ms]	Full model	87.4 ± 13.6	86.2 ± 14.1	91.9 ± 19.8	87.0 ± 3.8	3.52
	No T_1^{f} prior	87.9 ± 12.6	87.0 ± 13.3	90.9 ± 18.3	87.7 ± 3.4	2.35
	No PD constraint	96.2 ± 14.1	94.6 ± 14.1	108.4 ± 17.8	91.5 ± 2.6	9.14
	Neither	96.0 ± 13.9	94.3 ± 14.1	106.8 ± 18.3	92.5 ± 2.7	7.99
PD [a.u.]	Full model	0.716 ± 0.020	0.711 ± 0.012	0.703 ± 0.017	0.733 ± 0.021	2.21
	No T_1^{f} prior	0.710 ± 0.018	0.708 ± 0.012	0.699 ± 0.018	0.722 ± 0.019	1.65
	No PD constraint	0.603 ± 0.035	0.629 ± 0.018	0.561 ± 0.022	0.581 ± 0.025	5.94
	Neither	0.601 ± 0.035	0.626 ± 0.019	0.566 ± 0.023	0.577 ± 0.027	5.43

Table S7.2: Regularization ablation: prospective cross-protocol reproducibility of quantitative maps in a healthy volunteer scanned with the Spoiled GRE and MPRAGE protocols. Metrics evaluate voxel-wise agreement across brain parenchyma (WM+GM). MPD: mean absolute percentage difference; CCC: concordance correlation coefficient; r: Pearson correlation coefficient.

Parameter	**Configuration**	**MPD (%)**	**CCC**	r
m_0^{s}	Full model	7.67	0.888	0.895
	No T_1^{f} prior	14.60	0.473	0.793
	No PD constraint	9.56	0.881	0.892
	Neither	12.14	0.419	0.805
T_1^{f} [ms]	Full model	7.70	0.673	0.776
	No T_1^{f} prior	19.57	0.722	0.898
	No PD constraint	6.67	0.711	0.796
	Neither	12.04	0.901	0.907
T_2^{f} [ms]	Full model	4.83	0.869	0.879
	No T_1^{f} prior	4.77	0.864	0.881
	No PD constraint	8.65	0.661	0.777
	Neither	10.10	0.587	0.754
PD [a.u.]	Full model	4.83	0.823	0.837
	No T_1^{f} prior	4.81	0.830	0.841
	No PD constraint	11.31	0.697	0.842
	Neither	11.49	0.696	0.851

S8 Generalization to unseen $T_1\mathrm{w}$ protocols

To evaluate whether the framework generalizes to vendor-by-protocol groups not represented during training, we performed two protocol hold-out experiments based on the $T_1\mathrm{w}$ sequence used to define these groups in the main analysis. We trained one model exclusively on sessions from the Philips and Siemens MPRAGE groups, withholding the Philips Spoiled GRE group. We trained a second model exclusively on sessions from the Philips Spoiled GRE group, withholding both MPRAGE groups. Both protocol-restricted models used the same network architecture, loss weights, and optimization schedule as the pooled model (presented in the main paper) and differed only in the training data provided to them. We evaluated all models on the complete clinical test set ($n = 1{,}237$ sessions). Table S8.1 summarizes the corresponding training subsets.

Table S8.1: Overview of training subsets evaluated in the protocol hold-out experiments. All models were evaluated on the complete clinical test set ($n = 1{,}237$ sessions).

Configuration	**Withheld protocol**	**Included protocol(s)**	**Training sessions**
Pooled (full model)	None	Spoiled GRE + MPRAGE	6,084
MPRAGE-only	Spoiled GRE	MPRAGE	2,773
Spoiled GRE-only	MPRAGE	Spoiled GRE	3,311

Withholding a protocol during training substantially shifted the parameter estimates for the corresponding out-of-domain vendor-by-protocol groups, indicated by † in Table S8.2. When applied to the unseen Philips Spoiled GRE group, the MPRAGE-only model produced a mean white-matter MPF of 0.126, compared with approximately 0.18 in the two in-domain MPRAGE groups. White-matter T_1^{f} and T_2^{f} also shifted to markedly higher values. Accordingly, the between-group CV of white-matter MPF increased from 2.65% for the pooled model to 19.71%.

Conversely, when applied to the unseen MPRAGE groups, the Spoiled-GRE-only model produced lower white-matter MPF values and markedly higher white-matter T_1^{f} and T_2^{f} values than in the in-domain Spoiled GRE group. Mean white-matter MPF decreased from 0.177 in the Spoiled GRE group to approximately 0.15 in both MPRAGE groups, increasing the between-group CV to 11.06%.

Despite differences in scanner vendor, the MPRAGE-only model produced similar estimates for the Philips and Siemens MPRAGE groups. In contrast, both protocol-restricted models produced substantial parameter shifts when applied to the withheld sequence regime. These findings suggest that stable estimation requires the relevant acquisition protocols to be represented during training.

Table S8.2: Quantitative parameter estimates in white matter (WM) and gray matter (GM) across the clinical test set ($n = 1{,}237$ MRI sessions) for models trained on protocol-restricted subsets. Values denote mean ± standard deviation across sessions. † indicates an out-of-domain evaluation where the corresponding protocol was withheld from the training set. Between-group coefficient of variation (CV) was calculated across the three vendor-by-protocol group means.

		Philips (UMCU)		**Siemens (MR-RATE)**	
Parameter	**Training data**	**Spoiled GRE** ($n = 653$)	**MPRAGE** ($n = 221$)	**MPRAGE** ($n = 363$)	**Between-group CV (%)**
White Matter (WM)					
m_0^{s}	Pooled (full model)	0.186 ± 0.006	0.179 ± 0.004	0.177 ± 0.006	2.65
	MPRAGE only	$0.126 \pm 0.004^{\dagger}$	0.182 ± 0.005	0.183 ± 0.006	19.71
	Spoiled GRE only	0.177 ± 0.007	$0.146 \pm 0.007^{\dagger}$	$0.148 \pm 0.008^{\dagger}$	11.06
T_1^{f} [ms]	Pooled (full model)	1967 ± 60	2017 ± 64	1987 ± 40	1.26
	MPRAGE only	$2240 \pm 48^{\dagger}$	2025 ± 72	2008 ± 52	6.19
	Spoiled GRE only	1996 ± 55	$3579 \pm 87^{\dagger}$	$3614 \pm 82^{\dagger}$	30.18
T_2^{f} [ms]	Pooled (full model)	73.8 ± 13.7	80.9 ± 18.9	78.8 ± 1.2	4.68
	MPRAGE only	$120.6 \pm 7.1^{\dagger}$	78.2 ± 14.8	79.2 ± 1.0	26.09
	Spoiled GRE only	77.2 ± 13.8	$118.4 \pm 20.0^{\dagger}$	$115.1 \pm 11.4^{\dagger}$	22.09
PD [a.u.]	Pooled (full model)	0.673 ± 0.010	0.664 ± 0.012	0.672 ± 0.019	0.74
	MPRAGE only	$0.680 \pm 0.018^{\dagger}$	0.660 ± 0.013	0.657 ± 0.015	1.84
	Spoiled GRE only	0.666 ± 0.010	$0.679 \pm 0.033^{\dagger}$	$0.682 \pm 0.034^{\dagger}$	1.29
Gray Matter (GM)					
m_0^{s}	Pooled (full model)	0.149 ± 0.005	0.147 ± 0.005	0.134 ± 0.006	5.63
	MPRAGE only	$0.141 \pm 0.004^{\dagger}$	0.152 ± 0.006	0.144 ± 0.006	3.83
	Spoiled GRE only	0.143 ± 0.006	$0.143 \pm 0.006^{\dagger}$	$0.147 \pm 0.009^{\dagger}$	1.77
T_1^{f} [ms]	Pooled (full model)	2277 ± 60	2226 ± 61	2183 ± 66	2.12
	MPRAGE only	$2254 \pm 57^{\dagger}$	2273 ± 71	2279 ± 58	0.58
	Spoiled GRE only	2319 ± 56	$3149 \pm 135^{\dagger}$	$3121 \pm 152^{\dagger}$	16.46
T_2^{f} [ms]	Pooled (full model)	86.2 ± 14.1	91.9 ± 19.8	87.0 ± 3.8	3.52
	MPRAGE only	$125.0 \pm 4.7^{\dagger}$	88.7 ± 15.4	86.8 ± 2.4	21.48
	Spoiled GRE only	91.5 ± 15.0	$130.8 \pm 21.2^{\dagger}$	$142.1 \pm 23.1^{\dagger}$	21.88
PD [a.u.]	Pooled (full model)	0.711 ± 0.012	0.703 ± 0.017	0.733 ± 0.021	2.21
	MPRAGE only	$0.652 \pm 0.013^{\dagger}$	0.693 ± 0.020	0.712 ± 0.019	4.44
	Spoiled GRE only	0.698 ± 0.012	$0.741 \pm 0.025^{\dagger}$	$0.767 \pm 0.022^{\dagger}$	4.75