arXiv:2611.04894v1 [cs.CV] 3 Nov 2026

VisSearch-Bench: Evaluating Long Horizon Visual Search in Multimodal Agents

ZhengXian Wu[1,2,*,‡], **Qianhao Nie**[1,2,*], **Zhuohong Chen**[1,*], **Junjie Gao**[2], **Yuyao Zhang**[2,3], **Yunyao Yu**[1], **Hangrui Xu**[1], **Cheng Cheng**[2], **Guiyu Ma**[2], **Jin Zhang**[2], **Xingzhong Xu**[2], **Kai Yang**[2,†], **Haoqian Wang**[1,†]

[1]Tsinghua University [2]Ant Group [3]Renmin University of China

[*]Equal contribution, [†]Corresponding author, [‡]Work done during an internship at Ant Group

Abstract

Multimodal agents use web and image tools to answer questions beyond their internal knowledge. Some search tasks use images mainly to identify an initial entity, then rely on text retrieval. These tasks offer limited insight into how agents connect information across images, objects, and search steps. We therefore focus on **visual context management**. This involves organizing and using visual information and linking it to other findings during search. Visually intensive tasks require agents to combine object appearance, local details, or spatial relations with external information. We introduce **VisSearch-Bench**, with 204 human-written tasks across 7 main domains and 25 subdomains. It covers Set-wise Visual Search, Sequential Visual Search, Hybrid Wide–Deep Search, Competitive Visual Disambiguation, and Visual-Structure-Guided Search. Each task provides a reference answer, traceable evidence, and scoring rules. **Accuracy** measures complete task success, while **Process Score** measures the correctness and completeness of individual results and subquestion answers. Across 16 multimodal models, the best system fully solves only 26.47% of tasks, even with search tools. These results highlight the challenge of linking visual evidence with search results to produce complete answers.

Website: Project page
Code: GitHub repository
Hugging Face: Dataset

Tsinghua University ANT GROUP

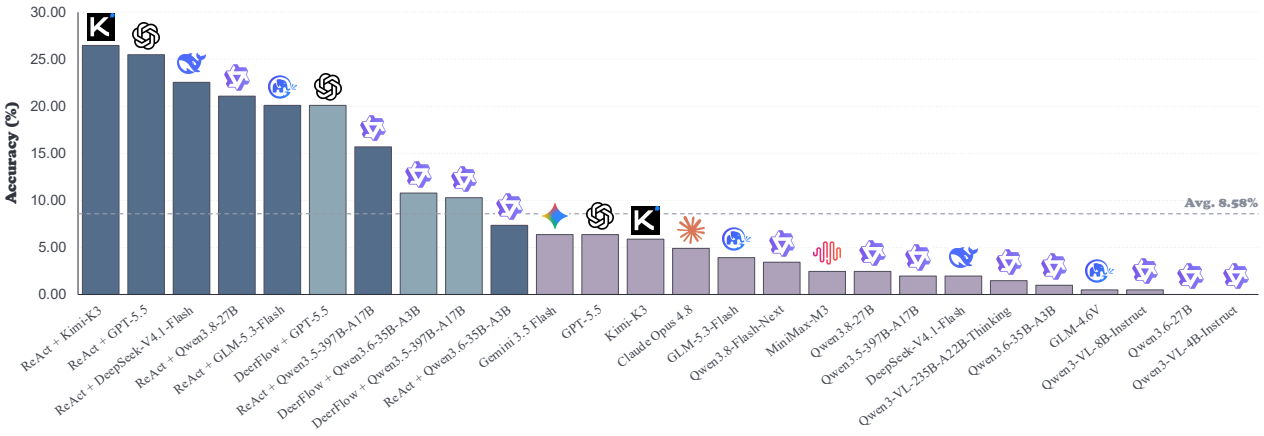

Figure 1 Accuracy of models and agents on **VisSearch-Bench**, comparing direct answering with ReAct and DeerFlow. The dashed line marks the average accuracy across the systems shown.

1 Introduction

Multimodal large language models (MLLMs) are evolving from one-shot visual question answering systems into agents that can plan, use tools, and interact with the open web over multiple rounds (Luo et al., 2026; Wu et al., 2026b). With text search, image search, webpage browsing, and region-level image inspection, these systems can gather evidence beyond the initial prompt and solve complex information-seeking tasks (Chng et al., 2025; Zhang et al., 2026b). This paradigm is closer to real-world information seeking than visual question answering based only on parametric knowledge.

At the same time, multimodal benchmarks are placing stronger demands on visual information. Models must not only understand images in the prompt, but also combine visual and textual evidence found during search (Hou et al., 2026). MMSearch-Plus (Tao et al., 2026) emphasizes the extraction and propagation of fine-grained visual cues. MM-BrowseComp (Li et al., 2025) and VisBrowse-Bench (Zhang et al., 2026c) require agents to actively obtain visual evidence while browsing. Together, these benchmarks make visual evidence an active component of search.

However, existing multimodal search benchmarks often focus on whether visual evidence is used at some point in a task. How images are used can lead to different task demands. In some existing tasks, images mainly serve to identify an initial entity (Wang et al., 2023; Chen et al., 2026). Once the agent obtains the name of a person, object, or place, it can rely mainly on text search. Such tasks provide limited insight into how agents handle multiple images or connect visual information across search steps. More recent benchmarks involve more image searches and tool calls. However, higher counts alone do not imply more complex visual context. We also need to examine how these visual operations are connected and how they support later searches and the final answer.

Real-world information seeking is more than a series of separate image recognition tasks (Zhang et al., 2026b). On the open web, agents need to update their understanding as they gather new evidence (Yao et al., 2022). Early evidence may suggest an initial hypothesis. Later evidence may help compare candidates, rule some out, or identify new ones. New information may also change how an earlier image is understood. An agent therefore needs to connect visual evidence across search steps and use it to guide further search (Hu et al., 2023). For example, finding all works in a catalog that meet a visual condition requires several steps. The agent must first establish the full set of entries, then inspect their images and combine the results. This raises a basic question: **Can agents reliably organize and use visual evidence as it accumulates during open-web search?**

We use **visual context management** to describe how agents organize and use visual information and its connections to findings during search. We study this ability through tasks that involve multiple images, multiple objects, and linked search steps. We examine whether agents can connect visual observations with external information to compare objects, filter candidates, and follow clues. We also examine whether they can combine results across steps into a complete and consistent answer.

To evaluate this ability, we introduce **VisSearch-Bench**, with five task categories that reflect different demands of long-horizon visual search. **Set-wise Visual Search** requires agents to collect and check visual information for a set of objects, then filter, compare, or combine the results. **Sequential Visual Search** uses visual clues from earlier steps to guide later searches. **Hybrid Wide–Deep Search** combines step-by-step search with search over a set of objects. Results from one stage determine the search scope or targets of another. **Competitive Visual Disambiguation** focuses on using fine visual details to distinguish similar-looking candidates. **Visual-Structure-Guided Search** combines object order, spatial positions, or connections in an image with information found through external search. Agents use this information to locate objects, check conditions, and combine results.

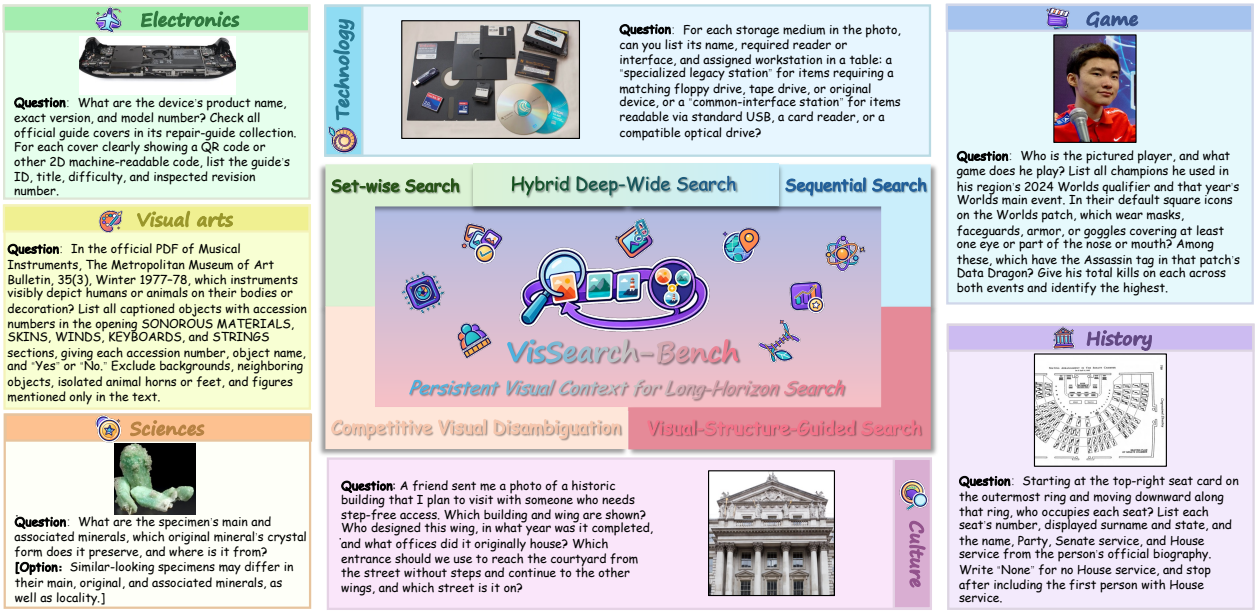

Figure 2 Overview of VisSearch-Bench. The center presents five task families that require long-horizon visual search, while the surrounding panels show selected examples from the dataset across diverse content domains.

Together, these tasks examine how agents handle multiple images, connect search results, and reason across steps. Figure 2 presents the five categories and representative tasks from different domains.

VisSearch-Bench contains 204 human-written tasks across 7 main domains and 25 subdomains. Dataset construction is led by human annotators and supported by automated checks. Annotators build each task from public, searchable sources. They gather the required images, entity relationships, and facts, then write the question and reference answer. They also document the evidence supporting the solution. Automated checks and human review examine source accessibility, answer consistency, and possible shortcuts that bypass the intended visual requirements. For evaluation, we use **Accuracy** to measure whether the full answer meets the task requirements. We also use a task-specific **Process Score** to assess the correctness and completeness of individual results and answers to subquestions. Our main contributions are as follows:

1. We introduce a task-based perspective for evaluating **visual context management** in open-web search. It focuses on how agents organize information from multiple images, connect different objects, and use visual findings to guide later searches.
2. We construct **VisSearch-Bench**, with 204 human-written tasks across 7 main domains and 25 subdomains. Its five categories cover Set-wise Visual Search, Sequential Visual Search, Hybrid Wide–Deep Search, Competitive Visual Disambiguation, and Visual-Structure-Guided Search.
3. We develop a human-led construction process supported by automated checks. Each task includes evidence sources, a reference answer, and scoring rules. Our evaluation combines full-answer accuracy with scores for individual results and subquestions. This measures overall task completion while also capturing which parts of an answer are correct or missing.

2 Related Work

2.1 Vision-Language Models

Vision-language models (VLMs) combine visual information with language understanding (Wu et al., 2026c; Yu et al., 2026). Their capabilities have expanded from single-image recognition

Benchmarks	Open-Web	Image Search	Multi-image	Multi-hop	Set Answers	Evidence Links	Partial Credit
BrowseComp	✓	—	—	✓	—	—	—
MMSearch	✓	✓	◐	◐	—	✓	◐
BrowseComp-VL	✓	✓	—	✓	—	—	—
MM-BrowseComp	✓	✓	✓	✓	—	✓	✓
MMSearch-Plus	✓	✓	✓	✓	—	◐	—
VisBrowse-Bench	✓	✓	✓	✓	—	—	—
BrowseComp-V3	✓	✓	✓	✓	—	✓	✓
InterLV-Search	✓	✓	✓	✓	—	—	—
VDR-Bench	✓	✓	—	✓	—	—	—
VisSearch-Bench (Ours)	✓	✓	✓	✓	✓	✓	✓

Table 1 Comparison of representative open-web and multimodal search benchmarks. ✓ Full; ◐ Partial; — Not covered. InterLV-Search refers to its open-web Level 3.

and question answering to document reading, multi-image understanding, video analysis, and visual reasoning. InternVL3 (Zhu et al., 2025) learns language and visual capabilities through joint pre-training and supports longer multimodal contexts. Kimi-VL (Du et al., 2025) focuses on high-resolution images, multi-image understanding, and long-context reasoning. Qwen3-VL (Bai et al., 2025) supports interleaved text, image, and video inputs, with improvements in multimodal reasoning and spatial understanding. These advances provide a foundation for handling more varied and complex visual inputs.

Recent models further combine visual understanding with general reasoning, tool use, and multi-step task execution. Qwen3.5 (Qwen Team, 2026a) uses unified vision-language training and reinforcement learning to improve reasoning and agent capabilities. Later models, including Qwen3.6-35B-A3B (Qwen Team, 2026c) and Qwen3.8-27B (Qwen Team, 2026d), support image and video understanding. They also improve code generation and multi-step task execution. Visual understanding is thus becoming an important part of general-purpose agent models. It supports tasks such as interpreting charts, reading documents, and using information from user interfaces.

Other model families follow a similar direction. GLM-5.3-Flash (GLM-5-Team et al., 2026) introduces native multimodal capabilities to the GLM-5 series. It supports browser and computer use based on visual information. DeepSeek-V4.1-Flash (Xu et al., 2026) supports both visual inputs and long-context reasoning. Kimi K3 (Team et al., 2026a) also provides native image and video understanding. Its one-million-token context window supports extended reasoning and tool interactions.

However, longer context windows do not guarantee effective visual-state management during search (Wu et al., 2026b). In open-web search, agents also need to correctly match objects across images and connect visual observations with external facts. They must use earlier findings to determine later search targets. VisSearch-Bench therefore focuses on visual context management in tasks involving multiple images, multiple objects, and linked search steps. It evaluates whether agents can use these connections to verify objects, compare candidates, filter results, and combine findings across steps.

2.2 Agentic Information Seeking

Existing benchmarks evaluate agentic information seeking through reasoning, tool use, and information gathering. GAIA focuses on multi-step reasoning and tool use in realistic tasks (Mialon et al., 2023). BrowseComp requires agents to persistently search the open web for facts that are difficult to find but easy to verify (Wei et al., 2025). WideSearch extends the focus to large-scale information gathering, emphasizing complete results that can be checked item by item (Wong et al., 2025). DeepWideSearch combines multi-hop retrieval with broad information collection to evaluate whether agents can handle both search depth and breadth (Lan et al., 2025).

Recent benchmarks place greater emphasis on answer coverage and information integration. DeepSearchQA (Gupta et al., 2026) requires agents to follow linked search steps, gather information from different sources, and produce complete answer lists. GISA supports several answer formats, including individual items, sets, lists, and tables. It also provides human search trajectories to support process analysis and model training (Zhu et al., 2026). DeepResearch Bench II (Li et al., 2026) extends evaluation to long-form research reports. It uses detailed scoring criteria derived from expert reports to assess information coverage, analysis, and presentation.

Research on search methods also examines how agents organize evidence and adjust later search actions. SearchOS (Zhang et al., 2026b) maintains shared records of pending tasks, evidence, information coverage, and failed attempts. These records help coordinate multiple agents, fill information gaps, and reduce repeated unsuccessful searches. Further work uses search trajectories for supervised fine-tuning to connect retrieval decisions more closely with the reasoning process (Chen et al., 2026).

2.3 Multimodal Search and Browsing

BrowseComp (Wei et al., 2025) mainly evaluates text-based information seeking on the open web and provides a reference for assessing browsing capabilities. MMSearch (Jiang et al., 2024) introduces image-text queries and evaluates query rewriting, result reranking, information summarization, and end-to-end answer generation. BrowseComp-VL (Geng et al., 2025) combines visual recognition with complex web retrieval, requiring agents to conduct multi-step searches based on entities in images. Later benchmarks place greater emphasis on the role of visual evidence during search. MMSearch-Plus (Tao et al., 2026) requires agents to extract fine-grained visual clues and verify related facts and sources through multiple rounds of image-text retrieval. VDR-Bench (Zeng et al., 2026) focuses on reducing shortcuts based on textual clues and whole-image matching. It uses retrieval from image regions, entity verification, and multi-hop question construction to place greater demands on visual search. VisBrowse-Bench (Zhang et al., 2026c) emphasizes actively gathering visual information from webpages and evaluates spatial localization, attribute recognition, and reasoning across images.

Other work further expands search structures and process evaluation. InterLV-Search (Hou et al., 2026) examines how visual and textual evidence guide later search steps. Its open-web tasks include single-chain and multi-branch structures that involve comparing multiple objects or search paths. MM-BrowseComp (Li et al., 2025) introduces multimodal evidence in both the input and the browsing process. It uses reasoning checklists to assess the completion of key steps. BrowseComp-V3 (Zhang et al., 2026a) covers alignment within image regions, integration across regions, and reasoning across images. It requires publicly searchable evidence and evaluates both final-answer accuracy and subgoal completion.

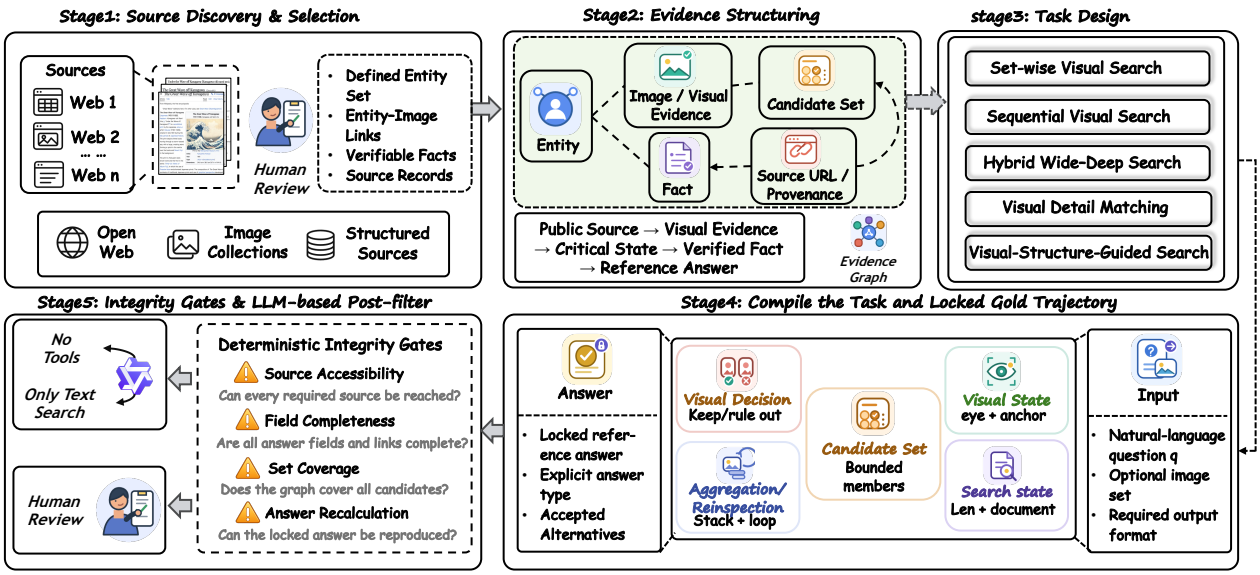

Figure 3 Overview of the VisSearch-Bench construction pipeline. Human annotators curate public sources and organize them into traceable evidence graphs; structural constraints instantiate the five task families; each task is paired with a locked answer and gold trajectory; deterministic integrity gates, LLM-based post-filtering, and final human review verify the retained items.

VisSearch-Bench further focuses on visual context management in search involving multiple images, multiple objects, and linked steps. Its tasks involve checking objects individually, combining results across a set, following visual clues, and distinguishing similar candidates. They also include using image structure together with external information to determine search targets. Table 1 compares the main search requirements covered by VisSearch-Bench and existing benchmarks based on their task designs.

3 VisSearch-Bench

3.1 Task Definition and Taxonomy

We study search in open-web environments where visual context plays a central role. Each task provides a natural-language question q, an optional initial image set $\mathcal{I}_0$ that is empty for text-initiated tasks, and a target output format σ. A multimodal agent interacts with the web through tools for text search, image search, reverse image search, webpage access, and image inspection. At step t, the agent selects its next action a_{t+1} based on the interaction history

$$h_t = \big(q, \mathcal{I}_0, (a_1, o_1), \ldots, (a_t, o_t)\big) \tag{1}$$

where o_t denotes the textual or visual observation returned by a tool. When the interaction ends, the agent produces a predicted answer $\hat{y}$ in the required format σ.

In VisSearch-Bench, we define **visual context** as the visual information involved in a task and its connections to object identities, external facts, and search targets. This information may come from input images or from images and webpages found during search. It includes object appearance and local features, as well as the order, positions, and connections of objects. **Visual context management** refers to organizing, connecting, and using this information to solve a task. This includes checking objects, comparing candidates, filtering results, and combining findings across steps. Based on these

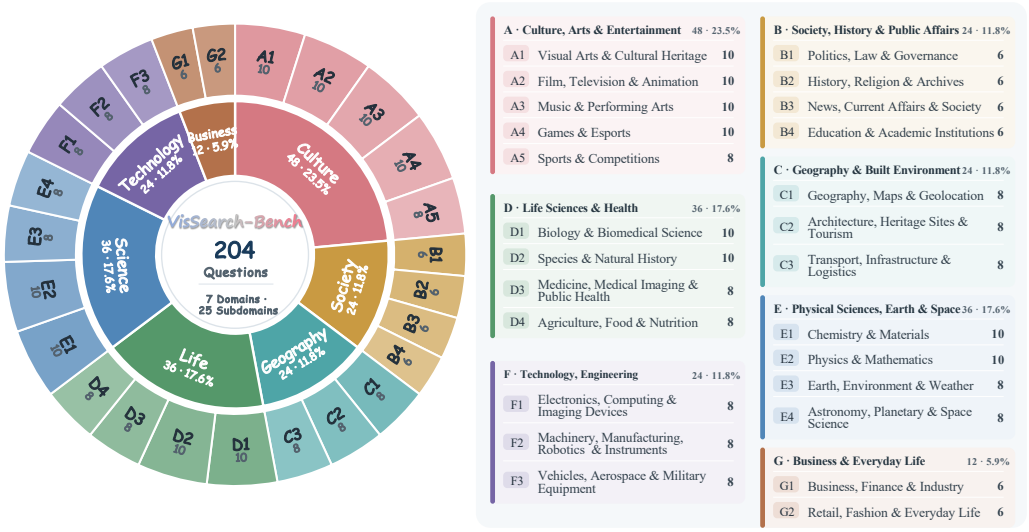

Figure 4 Domain taxonomy and composition of VisSearch-Bench. The inner ring shows seven primary domains, and the outer ring shows 25 secondary domains. Arc area is proportional to the number of tasks in the 204-question benchmark. The right panel expands the full taxonomy and reports the task count within each category.

requirements, we organize VisSearch-Bench into five task categories. Each category emphasizes a main search requirement.

1) **Set-wise Visual Search.** These tasks involve a clearly defined set of objects. The set may come from a catalog specified in the question or consist of objects shown in the input images. For each object, the agent gathers and checks visual information. It then combines this information with relevant facts to filter, compare, or summarize the results. The main challenge is to cover all required objects and correctly link the information for each one. Individual judgments must also remain consistent with the final summary.

2) **Sequential Visual Search.** These tasks involve multiple search steps that depend on one another. An object, visual clue, or fact found in one step guides later searches. For example, an agent may first identify an artwork in an image, then find records about its artist, collection, or exhibitions. Different steps may involve image understanding or text search. The focus is on correctly connecting the information found across steps to produce the required answer.

3) **Hybrid Wide–Deep Search.** These tasks combine sequential search with search over a set of objects. Earlier findings can define the candidate set to be checked. Filtering this set can then determine the target of the next stage. For example, an agent may start from one work and identify its award and year. It then examines images of works from the same award cycle and selects those that meet the given conditions. Further searches gather information about the selected works. The main challenge is to correctly connect the search scope, individual judgments, and later targets across stages.

4) **Competitive Visual Disambiguation.** These tasks focus on visual targets that are easily confused with other objects. Candidates may look similar overall but differ in local shapes, textures, markings, or individual parts. The agent compares the input image with retrieved images, descriptions, and source records to verify the target's identity. These tasks assess the accuracy of visual detail matching and identity verification.

5) **Visual-Structure-Guided Search.** These tasks use object order, spatial positions, or connections in an image to determine what to search for, in what order, or where to search next. The agent combines this structural information with external search results to locate objects, check conditions, and summarize findings. The main challenge is to correctly connect image structure, object identities, and search results, then use these connections to complete the task.

3.2 Dataset Construction

Figure 3 summarizes our human-led construction process with program-assisted verification. The process follows three principles. **Visual Indispensability** means that critical visual evidence cannot be replaced by the question text, webpage titles, captions, OCR, or model knowledge. **Visual Evidence Integration** means that tasks require agents to connect visual observations with object identities and external information. Agents use these connections to guide searches, compare candidates, filter objects, or combine results across steps. **Public Solvability and Auditability** means that each task uses publicly retrievable evidence and has verifiable critical states and a clear reference answer.

Statistic	Value
Total questions	204
Number of task families	5
Number of primary domains	7
Set-wise Visual Search	68 (33.3%)
Reference-solution steps (avg. / max.)	23.9 / 51
Sequential Visual Search	34 (16.7%)
Reference-solution steps (avg. / max.)	14.4 / 19
Hybrid Wide–Deep Search	34 (16.7%)
Reference-solution steps (avg. / max.)	23.8 / 61
Competitive Visual Disambiguation	34 (16.7%)
Reference-solution steps (avg. / max.)	3.9 / 5
Visual-Structure-Guided Search	34 (16.7%)
Reference-solution steps (avg. / max.)	15.8 / 19
Tasks with initial images	187 (91.7%)
Tasks without initial images	17 (8.3%)
Total reference-solution steps	3,593
Reference-solution steps per task (median / mean)	17.0 / 17.6

Table 2 Dataset composition and structural complexity.

Annotators first select public sources that provide a bounded entity set, reliable entity-image links, and independently verifiable factual fields. We fix these sources before question writing. The reference answer therefore does not depend on model search results or subjective judgments. Annotators then organize the entities, images, factual fields, and candidate sets into a traceable evidence structure. They apply the structural constraints of one task family from Section 3.1 and write a question that hides the intermediate nodes and final answer. Automated checks verify source accessibility, field completeness, set coverage, answer recomputation, and possible shortcuts. We discard tasks that bypass critical visual evidence or allow a shorter unintended path. Human reviewers inspect every retained task.

3.3 Evaluation Protocol and Metrics

Our primary metric is **Accuracy (Acc.)**. For task i, let y_i, $\hat{y}_i$, and σ_i denote the reference answer, the predicted answer, and the required answer format. We define

$$a_i = \mathbb{I}[\mathcal{E}_{\sigma_i}(\hat{y}_i, y_i) = 1], \tag{2}$$

$$\text{Acc} = \frac{1}{N}\sum_{i=1}^{N} a_i. \tag{3}$$

where $\mathcal{E}_{\sigma_i}$ compares parsed answers under predefined, task-specific rules. It normalizes entity aliases. For unordered sets, it removes duplicates and ignores order. For ordered sequences, it preserves order.

Model / Backbone	Accuracy (Acc., %)								Process Score (%)							
	Avg.	**Cult.**	**Soc.**	**Geo.**	**Life**	**Sci.**	**Tech.**	**Bus.**	**Avg.**	**Cult.**	**Soc.**	**Geo.**	**Life**	**Sci.**	**Tech.**	**Bus.**
Open-source Models																
Qwen3.6-35B-A3B	0.98	0.00	0.00	0.00	5.56	0.00	0.00	0.00	9.55	4.19	10.50	9.94	17.48	12.17	6.46	2.92
Qwen3.6-27B	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	5.53	3.22	9.79	6.02	8.46	4.15	2.29	7.02
Qwen3.5-397B-A17B	1.96	0.00	0.00	0.00	11.11	0.00	0.00	0.00	12.72	7.00	13.75	11.05	21.70	11.93	13.46	10.78
GLM-4.6V	0.49	0.00	0.00	0.00	2.78	0.00	0.00	0.00	9.24	4.14	10.53	8.93	15.82	7.14	12.26	8.24
MiniMax-M3	2.45	2.08	0.00	4.17	8.33	0.00	0.00	0.00	9.71	6.69	9.46	5.59	18.03	9.21	10.98	4.46
Qwen3-VL-235B-A22B-Thinking	1.47	0.00	0.00	4.17	2.78	2.78	0.00	0.00	11.37	6.60	14.41	10.13	16.28	14.00	11.64	3.61
Qwen3-VL-8B-Instruct	0.49	0.00	0.00	0.00	2.78	0.00	0.00	0.00	5.83	1.74	13.62	1.82	13.84	1.58	6.13	2.78
Qwen3-VL-4B-Instruct	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	5.17	0.54	10.09	4.23	9.60	4.57	5.95	2.78
GLM-5.3-Flash	3.92	6.25	0.00	4.17	8.33	0.00	4.17	0.00	8.66	8.75	2.50	8.13	17.18	5.96	9.03	3.51
Qwen3.8-Flash-Next	3.43	6.25	0.00	4.17	5.56	0.00	0.00	**8.33**	11.82	8.06	8.36	16.02	19.06	9.07	11.91	11.75
Qwen3.8-27B	2.45	0.00	0.00	4.17	11.11	0.00	0.00	0.00	6.35	0.35	3.06	8.46	14.82	6.72	5.11	8.70
DeepSeek-V4.1-Flash	1.96	0.00	4.17	0.00	5.56	0.00	0.00	**8.33**	12.88	6.93	13.42	8.22	19.33	13.07	19.00	12.84
Kimi-K3	5.88	8.33	4.17	4.17	11.11	2.78	4.17	0.00	10.24	7.73	9.12	8.07	19.26	9.31	11.58	0.00
Closed-source Models																
Gemini 3.5 Flash	6.37	6.25	8.33	8.33	8.33	2.78	4.17	**8.33**	22.28	17.42	23.11	29.11	27.59	19.29	24.39	15.22
Claude Opus 4.8	4.90	4.17	8.33	4.17	5.56	5.56	0.00	**8.33**	16.13	12.40	17.51	11.12	18.27	18.21	17.51	22.97
GPT-5.5	6.37	8.33	4.17	4.17	11.11	5.56	4.17	0.00	22.61	19.19	18.83	30.08	25.33	18.46	32.95	12.53
Open-source Agents																
ReAct																
Qwen3.6-35B-A3B	7.35	4.17	8.33	8.33	5.56	8.33	16.67	0.00	23.25	14.44	22.74	27.27	25.56	24.91	29.63	26.77
GPT-5.5	25.49	20.83	20.83	**33.33**	33.33	**25.00**	**29.17**	**8.33**	**49.52**	**46.70**	43.42	**60.37**	**52.76**	**48.02**	**59.05**	27.02
Qwen3.5-397B-A17B	15.69	16.67	16.67	8.33	16.67	16.67	25.00	0.00	32.43	26.14	30.82	26.67	33.47	42.68	41.19	20.89
GLM-5.3-Flash	20.10	20.83	16.67	20.83	36.11	13.89	16.67	0.00	22.29	21.70	19.58	20.83	36.11	19.04	22.92	0.00
Kimi-K3	**26.47**	**25.00**	**33.33**	**33.33**	36.11	22.22	20.83	0.00	38.57	39.50	**45.00**	40.00	42.75	34.58	37.06	21.60
DeepSeek-V4.1-Flash	22.55	14.58	20.83	29.17	**38.89**	19.44	25.00	0.00	39.47	38.03	33.13	38.44	50.12	37.50	44.83	23.18
Qwen3.8-27B	21.08	16.67	25.00	20.83	36.11	13.89	20.83	**8.33**	33.23	26.28	34.25	32.12	41.48	32.70	37.05	30.35
DeerFlow																
Qwen3.6-35B-A3B	10.78	10.42	12.50	4.17	11.11	13.89	16.67	0.00	24.26	17.89	26.45	17.21	29.03	33.15	27.77	11.49
GPT-5.5	20.10	16.67	29.17	12.50	27.78	16.67	25.00	**8.33**	38.30	35.34	43.86	35.39	34.78	35.33	51.41	**38.13**
Qwen3.5-397B-A17B	10.29	8.33	12.50	0.00	19.44	11.11	12.50	0.00	28.37	19.54	32.44	22.59	33.13	40.88	29.44	13.11

Table 3 Main results on VisSearch-Bench. Avg. denotes the micro-average over all benchmark tasks. Cult., Soc., Geo., Life, Sci., Tech., and Bus. denote Culture, Society, Geography, Life Sciences, Physical Sciences, Technology, and Business, respectively.

A task is correct only if the prediction contains all and only the required content.

Final-answer accuracy cannot show whether an error comes from the wrong search space, a missed visual candidate, or incorrect evidence aggregation (Wu et al., 2026a). We therefore also report **Process Score (PS)**. It measures the correctness and completeness of individual results and answers to subquestions, using the scoring rules defined for each task. Each task i has m_i labeled critical states, which all systems report before the final answer. Let $[m_i] = \{1, \ldots, m_i\}$ index these states. We define the set of correctly recovered states as

$$\mathcal{R}_i = \left\{ j \in [m_i] \;\middle|\; \hat{g}_{ij} \equiv_{\tau_{ij}} g_{ij} \right\}, \tag{4}$$

$$\mathrm{PS} = \frac{1}{N} \sum_{i=1}^{N} \frac{|\mathcal{R}_i|}{m_i}. \tag{5}$$

Here, $\hat{g}_{ij} \equiv_{\tau_{ij}} g_{ij}$ means that predicted state j matches its reference under the corresponding rule. States are aligned by field labels, and different names for the same entity are treated as equal. Each match counts as one correct state, so $|\mathcal{R}_i|/m_i$ is the fraction of required states recovered for task i.

Model / Backbone	Avg. (204)		Set-wise Visual Search (68)		Sequential Visual Search (34)		Hybrid Wide–Deep Search (34)		Competitive Visual Disambiguation (34)		Visual-Structure-Guided Search (34)	
	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score
Open-source Models												
Qwen3.6-35B-A3B	0.98	9.55	0.00	10.79	0.00	6.40	0.00	5.42	0.00	0.00	5.88	23.94
Qwen3.6-27B	0.00	5.53	0.00	7.94	0.00	5.44	0.00	1.53	0.00	0.00	0.00	10.31
Qwen3.5-397B-A17B	1.96	12.72	0.00	13.77	0.00	10.48	0.00	11.12	0.00	0.00	11.76	27.17
GLM-4.6V	0.49	9.24	0.00	6.94	0.00	5.29	0.00	10.74	0.00	0.00	2.94	25.53
MiniMax-M3	2.45	9.71	0.00	7.18	0.00	7.38	2.94	4.89	0.00	0.00	11.76	31.61
Qwen3-VL-235B-A22B-Thinking	1.47	11.37	0.00	13.21	0.00	7.28	0.00	7.22	0.00	0.00	8.82	27.27
Qwen3-VL-8B-Instruct	0.49	5.83	0.00	7.15	0.00	2.85	0.00	4.07	0.00	0.00	2.94	13.77
Qwen3-VL-4B-Instruct	0.00	5.17	0.00	6.69	0.00	2.62	0.00	3.35	0.00	0.00	0.00	11.69
GLM-5.3-Flash	3.92	8.66	0.00	4.22	0.00	4.26	8.82	8.82	0.00	0.00	14.71	30.45
Qwen3.8-Flash-Next	3.43	11.82	0.00	6.24	0.00	9.56	8.82	14.94	0.00	0.00	11.76	33.96
Qwen3.8-27B	2.45	6.35	0.00	4.33	0.00	1.03	0.00	0.41	0.00	0.00	14.71	28.02
DeepSeek-V4.1-Flash	1.96	12.88	0.00	11.69	0.00	11.54	2.94	13.72	0.00	0.00	8.82	28.66
Kimi-K3	5.88	10.24	2.94	6.69	2.94	11.47	8.82	10.82	0.00	0.00	17.65	25.79
Closed-source Models												
Gemini 3.5 Flash	6.37	22.28	0.00	21.40	2.94	22.35	11.76	28.08	0.00	0.00	23.53	40.44
Claude Opus 4.8	4.90	16.13	0.00	12.28	5.88	15.14	11.76	20.02	0.00	0.00	11.76	37.09
GPT-5.5	6.37	22.61	1.47	21.23	2.94	21.36	8.82	28.36	0.00	0.00	23.53	43.47
Open-source Agents												
ReAct												
Qwen3.6-35B-A3B	7.35	23.25	2.94	19.81	2.94	16.47	5.88	23.40	2.94	2.94	26.47	57.05
GPT-5.5	25.49	**49.52**	8.82	**44.10**	23.53	51.73	**23.53**	**48.36**	38.24	38.24	**50.00**	**70.61**
Qwen3.5-397B-A17B	15.69	32.43	0.00	20.14	11.76	41.60	14.71	27.21	32.35	29.41	35.29	56.05
GLM-5.3-Flash	20.10	22.29	10.29	11.58	8.82	15.29	8.82	8.82	35.29	35.29	47.06	51.16
Kimi-K3	**26.47**	38.57	**11.76**	24.00	**32.35**	**56.34**	14.71	27.74	41.18	**44.12**	47.06	55.22
DeepSeek-V4.1-Flash	22.55	39.47	5.88	28.28	26.47	52.52	11.76	20.97	**44.12**	**44.12**	41.18	62.63
Qwen3.8-27B	21.08	33.23	7.35	18.78	23.53	46.91	8.82	20.17	**44.12**	**44.12**	35.29	50.61
DeerFlow												
Qwen3.6-35B-A3B	10.78	24.26	1.47	19.08	8.82	20.29	5.88	17.55	35.29	35.29	11.76	34.27
GPT-5.5	20.10	38.30	4.41	31.35	17.65	40.76	20.59	39.75	26.47	26.47	47.06	60.14
Qwen3.5-397B-A17B	10.29	28.37	0.00	17.59	8.82	37.65	8.82	23.80	29.41	29.41	14.71	44.18

Table 4 Accuracy and Process Score (%) by task category on VisSearch-Bench. Numbers in parentheses indicate task counts. Avg. denotes the micro-average over all 204 tasks. Bold values mark the highest score in each column, including ties.

We average the process scores across all tasks to obtain the overall Process Score for the benchmark. Different search paths receive the same score if they recover the same critical states. Direct uses no tools but reports the same structured states and is scored in the same way. All tables report both metrics as percentages.

3.4 Dataset Statistics

VisSearch-Bench contains 204 tasks organized along two dimensions. The first dimension is search structure. It includes five task categories: Set-wise Visual Search, Sequential Visual Search, Hybrid Wide–Deep Search, Competitive Visual Disambiguation, and Visual-Structure-Guided Search. Each category emphasizes a different search requirement. Set-wise Visual Search contains 68 tasks, or 33.3% of the benchmark. Each of the other four families contains 34 tasks, or 16.7%. The second is content domain, with 7 primary domains and 25 secondary domains that measure coverage across knowledge areas and real-world settings.

Model / Backbone	Direct Answer		Text		Caption		Full		Full − Text		Full − Caption	
	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score	Acc.	Process Score
						ReAct						
Qwen3.6-35B-A3B	0.98	9.55	5.88	17.74	6.37	19.21	**7.35**	**23.25**	+1.47	+5.51	+0.98	+4.04
Qwen3.5-397B-A17B	1.96	12.72	8.33	19.84	14.71	25.42	**15.69**	**32.43**	+7.35	+12.59	+0.98	+7.00
DeepSeek-V4.1-Flash	1.96	12.88	12.25	23.14	19.12	36.33	**22.55**	**39.47**	+10.29	+16.33	+3.43	+3.14

Table 5 Effect of visual access on VisSearch-Bench. Accuracy and Process Score are reported as percentages over all 204 tasks. Direct Answer uses no external tools; Text, Caption, and Full use ReAct. Full − Text and Full − Caption report gains in percentage points, computed before rounding.

The seven primary domains are Culture (23.5%), Life Sciences (17.6%), Physical Sciences (17.6%), Society (11.8%), Geography (11.8%), Technology (11.8%), and Business (5.9%). Each task is assigned one primary and one secondary domain based on its entities and required knowledge. Figure 4 shows the taxonomy for all 25 secondary domains.

4 Experiments

4.1 Experimental Setup

Models and agent frameworks. We evaluate 16 representative multimodal models. The open-weight models are GLM-5.3-Flash, Kimi-K3 (Team et al., 2026a), DeepSeek-V4.1-Flash (Xu et al., 2026), Qwen3.8-27B, Qwen3.8-Flash, Qwen3.6-35B-A3B (Qwen Team, 2026c), Qwen3.6-27B (Qwen Team, 2026b), Qwen3.5-397B-A17B (Qwen Team, 2026a), GLM-4.6V (Team et al., 2026b), MiniMax-M3 (Lai et al., 2026), Qwen3-VL-235B-A22B-Thinking, Qwen3-VL-8B-Instruct, and Qwen3-VL-4B-Instruct (Yang et al., 2025; Bai et al., 2025). The proprietary models are GPT-5.5, Gemini 3.5 Flash, and Claude Opus 4.8. In the Direct Answer setting, a model receives only the question and its initial images. We also pair Qwen3.6-35B-A3B, Qwen3.5-397B-A17B, GPT-5.5, GLM-5.3-Flash, Kimi-K3, DeepSeek-V4.1-Flash, and Qwen3.8-27B with ReAct (Yao et al., 2022). This comparison measures how much each search framework adds to the same backbone.

Tools and inference limits. All ReAct agents have access to five evidence tools: web search, webpage reading, text-based image search, reverse image search, and image cropping. Each web, image, or reverse-image search returns up to five results. Retrieved webpage text is truncated to 12,000 characters. Images shown to the model have a longest side of at most 1,536 pixels. The image budget is 20 per task, shared by initial images, retrieved images, and crops. Crops can be enlarged by up to four times. Each task allows up to 30 ReAct search rounds and 60 evidence-tool calls, with at most four calls per round. If either limit is reached, one additional model call is used to submit the final answer from the available evidence.

Inference and evaluation. We evaluate each system once per task from a fresh context and report its first complete response as Pass@1. The question, initial images, and required state-and-answer schema are identical across settings; framework-specific instructions only define the corresponding tool-use protocol. Forward web and image search use Serper, whereas reverse-image search uses SerpApi. We use the metrics defined in Section 3.3. Accuracy (Acc.) measures whether the final

answer is correct. Process Score measures the correctness and completeness of individual results and answers to subquestions in the final response, using task-specific scoring rules.

4.2 Main Results

Table 3 reports the Pass@1 results on VisSearch-Bench. In the Direct Answer setting, no model exceeds 6.37% accuracy. With search, ReAct with Kimi-K3 achieves the highest accuracy at 26.47%. ReAct with GPT-5.5 reaches 25.49% accuracy and the highest Process Score at 49.52%. The best model therefore differs across the two metrics. Success on complete tasks does not fully align with performance on individual items and subquestions. Even the system with the highest accuracy fully solves only about one quarter of the tasks. There is still much room to improve how models link information across images, verify objects, and combine results across search steps.

In the available paired experiments, both ReAct and DeerFlow improve accuracy and Process Score over Direct Answer. For example, Qwen3.5-397B-A17B achieves 1.96% accuracy without tools. ReAct and DeerFlow raise its accuracy to 15.69% and 10.29%, respectively. Its Process Score also increases from 12.72% to 32.43% with ReAct and 28.37% with DeerFlow. Thus, search improves both complete answers and the correctness and completeness of individual results and subquestion answers.

The effect of the framework also varies across models. For GPT-5.5 and Qwen3.5-397B-A17B, ReAct outperforms DeerFlow on both metrics. For Qwen3.6-35B-A3B, however, DeerFlow performs better. It achieves 10.78% accuracy and a Process Score of 24.26%, compared with 7.35% and 23.25% for ReAct. No single framework performs best for every model tested with both frameworks.

4.3 Ablation Studies

Table 4 reports accuracy and Process Score across the five task categories. Accuracy on Set-wise Visual Search remains low. ReAct with Kimi-K3 achieves the highest accuracy at 11.76%. In this category, ReAct with GPT-5.5 reaches a Process Score of 44.10%, but its accuracy is only 8.82%. These results show a clear gap between obtaining partially correct results and completing the entire set-based task.

Setting	Full wrong → correct	Full correct → wrong	Correct answers	Acc. (%)
Qwen3.6-35B-A3B (Full: 15, 7.35%)				
Caption	6	**8**	13	6.37
Text	6	**9**	12	5.88
Qwen3.5-397B-A17B (Full: 32, 15.69%)				
Caption	8	**10**	30	14.71
Text	5	**20**	17	8.33
DeepSeek-V4.1-Flash (Full: 46, 22.55%)				
Caption	13	**20**	39	19.12
Text	5	**26**	25	12.25

Table 6 ReAct answer changes from Full to each setting on 204 tasks. Group headers give the number of correct answers and accuracy under Full.

Different models also show different strengths across the other categories. ReAct with Kimi-K3 leads on Sequential Visual Search, with 32.35% accuracy and a Process Score of 56.34%. ReAct with GPT-5.5 leads on both metrics in Hybrid Wide–Deep Search and Visual-Structure-Guided Search. Its accuracy and Process Score are 23.53% and 48.36% for Hybrid Wide–Deep Search. They reach 50.00% and 70.61% for Visual-Structure-Guided Search. On Competitive Visual Disambiguation, all Direct Answer models score 0.00% accuracy. ReAct with DeepSeek-V4.1-Flash and Qwen3.8-27B both reach 44.12%. Search therefore improves performance across different visual tasks, while the strengths of individual models vary.

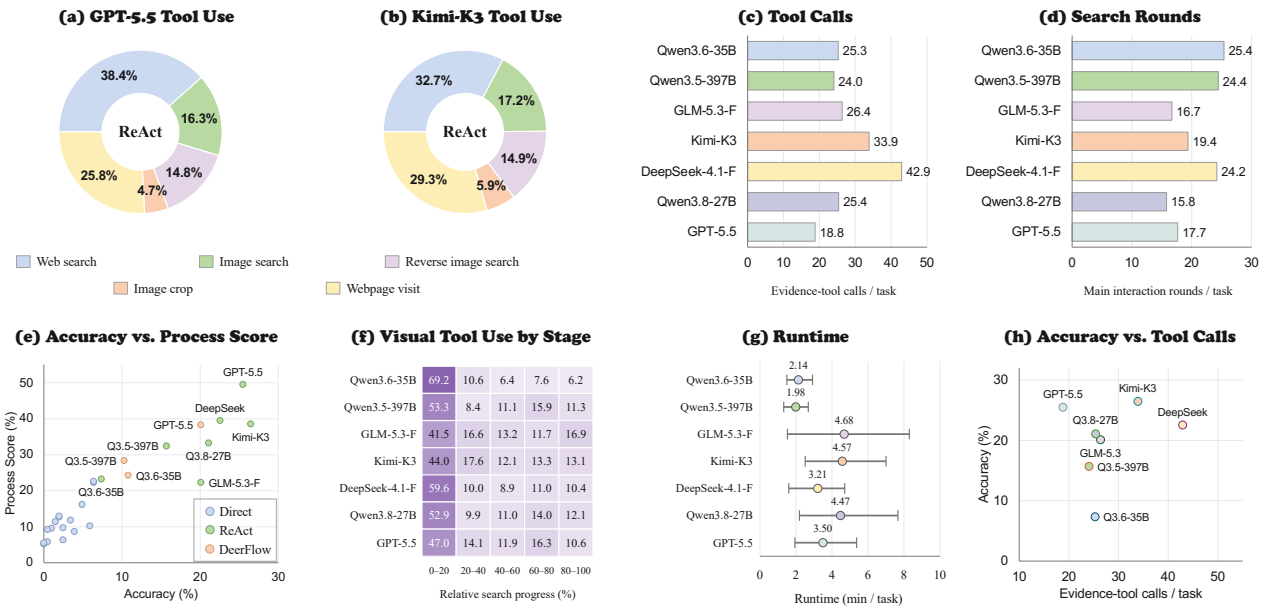

Figure 5 Tool use, efficiency, and performance on VisSearch-Bench.

To study the role of different information sources, we compare four settings. *Direct Answer* receives only the question and its initial images, without external tools. The other three settings use ReAct and retain the same initial images. *Text* provides web search and webpage reading. *Caption* also provides image search and reverse image search. It uses separate calls to the same vision model to generate general descriptions of retrieved images. This retains visual information extraction but adds computation. *Full* allows the agent to inspect retrieved images directly and use image cropping. Table 5 shows that Full achieves the highest accuracy and Process Score for all three models.

Table 6 shows how gains and losses on different tasks partly offset each other. Switching from Full to Caption loses 8 correct answers for Qwen3.6-35B-A3B, 10 for Qwen3.5-397B-A17B, and 20 for DeepSeek-V4.1-Flash. Caption also correctly answers 6, 8, and 13 tasks that Full fails, respectively. For tasks with demanding visual context, the additional description step may reduce the burden of organizing image information. This may help the agent solve some tasks that fail under Full. However, general descriptions may omit details needed for cross-image comparisons or other visual judgments. Across all three models, Caption loses more correct answers than it gains. For example, Qwen3.5-397B-A17B achieves 14.71% accuracy with Caption, below the 15.69% achieved with Full. Its Process Score also drops from 32.43% to 25.42%.

Figure 5 provides further analysis of tool use, search costs, and task performance. Models differ in their tool choices, search rounds, and runtime. More tool calls do not always correspond to higher accuracy. Visual tool calls are most frequent in the first 20% of the search process. This stage accounts for 41.5%–69.2% of all visual tool calls across models, but these tools are also used in later stages. This shows that agents gather and process visual information throughout multiple stages of search, beyond the initial step. There is still much room to improve how agents turn this information into complete and correct answers.

5 Conclusion

We introduced **VisSearch-Bench** to evaluate visual context management in long-horizon open-web search by multimodal agents. The benchmark contains 204 human-written tasks across five task categories. These tasks examine how agents connect visual observations, object identities, and external information to compare candidates, filter results, and combine findings across steps. Each task includes traceable evidence sources, a reference answer, and scoring rules. We use full-answer Accuracy to measure overall task completion. Process Score measures the correctness and completeness of individual results and answers to subquestions. Our evaluation covers 16 multimodal models. Search improves both accuracy and Process Score in the paired comparisons, but the best system still fully solves only 26.47% of the tasks. Model strengths also vary across task categories and agent frameworks. Results on Set-wise Visual Search show a clear gap between obtaining partially correct results and completing an entire task. Verifying multiple objects and combining findings into complete, consistent answers therefore remain important challenges. VisSearch-Bench provides a basis for evaluating improvements in handling multiple images, linking clues across steps, and producing complete answers.

References

S. Bai, Y. Cai, R. Chen, K. qin Chen, X.-H. Chen, Z. Cheng, L. Deng, W. Ding, R. Fang, et al. Qwen3-vl technical report. *ArXiv*, abs/2511.21631, 2025. URL https://api.semanticscholar.org/CorpusID:283262018.

Z. Chen, Z. Wu, Y. Yu, H. Xu, Z. Liao, Z. Liu, X. Deng, P. Jiao, and H. Wang. Learning to search: A decision-based agent for knowledge-based visual question answering. *ArXiv*, abs/2604.07146, 2026. URL https://api.semanticscholar.org/CorpusID:287249321.

Y. X. Chng, T. Hu, W. Tong, X. Li, J. Chen, H. Yu, J. Lu, H. Guo, H. Deng, X. Chen, G. Huang, D. Lin, and L. Lu. Sensenova-mars: Empowering multimodal agentic reasoning and search via reinforcement learning. *ArXiv*, abs/2512.24330, 2025. URL https://api.semanticscholar.org/CorpusID:284351490.

K. T. A. Du, B. Yin, B. Xing, B. Qu, B. Wang, C. Chen, C. Zhang, C. Du, C. Wei, C. Wang, D. Zhang, D. Du, et al. Kimi-vl technical report. *ArXiv*, abs/2504.07491, 2025. URL https://api.semanticscholar.org/CorpusID:277667211.

X. Geng, P. Xia, Z. Zhang, X.-Y. Wang, Q. Wang, R.-X. Ding, C.-X. Wang, J.-L. Wu, Y. Zhao, K. Li, Y. Jiang, P.-J. Xie, F. Huang, and J.-R. Zhou. Webwatcher: Breaking new frontier of vision-language deep research agent. *ArXiv*, abs/2508.05748, 2025. URL https://api.semanticscholar.org/CorpusID:280561766.

GLM-5-Team, :, A. Zeng, X. Lv, Z. Hou, Z. Du, Q. Zheng, B. Chen, D. Yin, C. Ge, C. Huang, C. Xie, C. Zhu, C. Yin, C. Wang, G. Pan, et al. Glm-5: from vibe coding to agentic engineering, 2026. URL https://arxiv.org/abs/2602.15763.

N. Gupta, R. Chatterjee, L. Haas, C. Tao, A. Wang, C. Liu, H. Oiwa, E. Gribovskaya, J. Ackermann, J. Blitzer, S. Goldshtein, and D. Das. Deepsearchqa: Bridging the comprehensiveness gap for deep research agents. *ArXiv*, abs/2601.20975, 2026. URL https://api.semanticscholar.org/CorpusID:283897826.

B. Hou, J. Gu, J. Guo, R. Dang, S. Leng, X. Li, X. Song, and J. Yang. Interlv-search: Benchmarking interleaved multimodal agentic search. *ArXiv*, abs/2605.07510, 2026. URL https://api.semanticscholar.org/CorpusID:288147842.

Z.-N. Hu, A. Iscen, C. Sun, K.-W. Chang, Y. Sun, D. A. Ross, C. Schmid, and A. Fathi. Avis: Autonomous visual information seeking with large language models. *ArXiv*, abs/2306.08129, 2023. URL https://api.semanticscholar.org/CorpusID:259164452.

D. Jiang, R. Zhang, Z. Guo, Y. Wu, J. Lei, P. Qiu, P. Lu, Z. Chen, G. Song, P. Gao, Y. Liu, C. Li, and H. Li. Mmsearch: Benchmarking the potential of large models as multi-modal search engines. *ArXiv*, abs/2409.12959, 2024. URL https://api.semanticscholar.org/CorpusID:272753572.

X. Lai, W. Xu, Y. Yang, Q. Chen, Y. Xu, L. Zeng, X. Li, H. Sun, H. Zhu, V. Zhang, J. Hu, J. Li, R. Gao, Z. Li, S. Zhu, J. Zhou, and P. Zhao. Minimax sparse attention, 2026. URL https://arxiv.org/abs/2606.13392.

T. Lan, B. Zhu, Q. Jia, J. Ren, H. Li, L. Wang, Z. Xu, W. Luo, and K. Zhang. Deepwidesearch: Benchmarking depth and width in agentic information seeking. *ArXiv*, abs/2510.20168, 2025. URL https://api.semanticscholar.org/CorpusID:282304379.

R. Li, M.-X. Du, B. Xu, C. Zhu, X. Wang, and Z.-D. Mao. Deepresearch bench ii: Diagnosing deep research agents via rubrics from expert report. *ArXiv*, abs/2601.08536, 2026. URL https://api.semanticscholar.org/CorpusID:284704978.

S. Li, X. Bu, W. Wang, J. Liu, J. Dong, H. He, H. Lu, H. Zhang, C. Jing, Z. Li, C. Li, J. Tian, C. Zhang, T. Peng, Y. He, J. Gu, Y. Zhang, J. Yang, G. Zhang, W. Huang, W. Zhou, Z. Zhang, R. Ding, and S. Wen. Mm-browsecomp: A comprehensive benchmark for multimodal browsing agents. *ArXiv*, abs/2508.13186, 2025. URL https://api.semanticscholar.org/CorpusID:280686450.

R. Luo, C. Shi, Y. Zhang, C. Yang, S. Jiang, T. Guan, R. Chen, R. Chu, P. Wang, M.-C. Yang, Y. Yang, J. Lin, and Z. Yang. From narrow to panoramic vision: Attention-guided cold-start reshapes multimodal reasoning. *ArXiv*, abs/2603.03825, 2026. URL https://api.semanticscholar.org/CorpusID:286238764.

G. Mialon, C. Fourrier, C. Swift, T. Wolf, Y. LeCun, and T. Scialom. Gaia: a benchmark for general ai assistants. *ArXiv*, abs/2311.12983, 2023. URL https://api.semanticscholar.org/CorpusID:265351664.

Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026a. URL https://qwen.ai/blog?id=qwen3.5.

Qwen Team. Qwen3.6-27B: Flagship-level coding in a 27B dense model, April 2026b. URL https://qwen.ai/blog?id=qwen3.6-27b.

Qwen Team. Qwen3.6-35B-A3B: Agentic coding power, now open to all, April 2026c. URL https://qwen.ai/blog?id=qwen3.6-35b-a3b.

Qwen Team. Qwen3.8-Max: A new bar for coding and cowork, August 2026d. URL https://qwen.ai/blog?id=qwen3.8.

X. Tao, Y. Teng, X. Su, X. Fu, J. Wu, C. Tao, Z. Liu, H. Bai, R. Liu, and L. Kong. Mmsearch-plus: Benchmarking provenance-aware search for multimodal browsing agents, 2026. URL https://arxiv.org/abs/2508.21475.

K. Team, T. Bai, Y. Bai, Y. Bao, M. C., J. Cai, X. Cai, P. Cao, Y. Cao, et al. Kimi k3: Open frontier intelligence, 2026a. URL https://arxiv.org/abs/2607.24653.

V. Team, W. Hong, W. Yu, X. Gu, G. Wang, G. Gan, H. Tang, J. Cheng, J. Qi, J. Ji, et al. Glm-4.5v and glm-4.1v-thinking: Towards versatile multimodal reasoning with scalable reinforcement learning, 2026b. URL https://arxiv.org/abs/2507.01006.

Z. Wang, C. Chen, P. Li, and Y. Liu. Filling the image information gap for vqa: Prompting large language models to proactively ask questions. In *Conference on Empirical Methods in Natural Language Processing*, 2023. URL https://api.semanticscholar.org/CorpusID:265295192.

J. Wei, Z. Sun, S. Papay, S. McKinney, J. Han, I. abella Fulford, H. W. Chung, A. Passos, W. Fedus, and A. Glaese. Browsecomp: A simple yet challenging benchmark for browsing agents. *ArXiv*, abs/2504.12516, 2025. URL https://api.semanticscholar.org/CorpusID:277857238.

R. Wong, J. Wang, J. Zhao, L. Chen, Y. Gao, L. Zhang, X. Zhou, Z. Wang, K. Xiang, G. Zhang, W. Huang, Y. Wang, and K. Wang. Widesearch: Benchmarking agentic broad info-seeking. *ArXiv*, abs/2508.07999, 2025. URL https://api.semanticscholar.org/CorpusID:280565879.

Z. Wu, J. Gao, and K. Yang. Silent failures in multimodal agentic search: A diagnostic taxonomy and cross-judge evaluation. *arXiv preprint arXiv:2607.19793*, 2026a.

Z. Wu, H. Xu, K. Shi, Z. Chen, Y. Yu, C. Zhang, Z. Liao, J. Yang, Z. Yang, H. Lu, et al. Promsa: Progressive multimodal search agents for knowledge-based visual question answering. *arXiv preprint arXiv:2606.27974*, 2026b.

Z.-X. Wu, K. Shi, C. Zhang, Z.-R. Liao, J. Yang, N. Yang, Q.-Y. Peng, L.-Y. Zhang, H.-R. Xu, T.-H. Su, Z.-Y. Yang, H.-N. Lu, and H.-Q. Wang. When models judge themselves: Unsupervised self-evolution for multimodal reasoning. *ArXiv*, abs/2603.21289, 2026c. URL https://api.semanticscholar.org/CorpusID:286766727.

D. Xu, B. Li, B. Lin, B. Xue, B.-C. Xian, B. Xu, B. Wu, B.-W. Zhang, B.-Y. Deng, C.-C. Yu, C. Jin, C.-F. Lin, C. Dong, C.-B. Wang, C. Feng, C.-D. Lu, C.-G. Zhao, et al. Deepseek-v4.1-flash: Pushing the limits of kv cache compression. 2026. URL https://api.semanticscholar.org/CorpusID:292063082.

A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv, et al. Qwen3 technical report, 2025. URL https://arxiv.org/abs/2505.09388.

S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. React: Synergizing reasoning and acting in language models. *ArXiv*, abs/2210.03629, 2022. URL https://api.semanticscholar.org/CorpusID:252762395.

Y. Yu, Z.-X. Wu, Z. Chen, H.-R. Xu, Z.-R. Liao, X. Deng, Z.-F. Liu, S. Shi, and H.-Q. Wang. Stabilizing unsupervised self-evolution of mllms via continuous softened retracing resampling. *ArXiv*, abs/2604.03647, 2026. URL https://api.semanticscholar.org/CorpusID:287200369.

Y. Zeng, W. Huang, Z. Fang, S. Chen, Y. Shen, Y. Cai, X. Wang, Z. Yin, L. Chen, Z. Chen, S. Huang, Y. Zhao, X. Tang, Y. Hu, P. H. S. Torr, W. Ouyang, and S. Cao. Vision-deepresearch benchmark: Rethinking visual and textual search for multimodal large language models. *ArXiv*, abs/2602.02185, 2026. URL https://api.semanticscholar.org/CorpusID:285270920.

H. Zhang, J. Zhou, B. Li, B. Zhou, Y. Dan, H. Lu, Z. Cao, J. Chen, Y. Han, Z. Sheng, Z. Tao, H. Liang, J. Wu, Y. Shi, Y. He, J. Lin, Q. Zhang, G. Yan, R. Zhao, Z. Li, X. Yu, L. Mei, C. Chen, W. Zhang, and B. Cui. Browsecomp-v3: A visual, vertical, and verifiable benchmark for multimodal browsing agents. *ArXiv*, abs/2602.12876, 2026a. URL https://api.semanticscholar.org/CorpusID:285606894.

Y. Zhang, J. Gao, Z. Wu, J. Fan, J. Zhang, S. Ma, Y. Yao, W. Qi, C. Jin, G. Ma, et al. Searchos-v1: Towards robust open-domain information-seeking agent collaboration. *arXiv preprint arXiv:2607.15257*, 2026b.

Z. Zhang, J. Su, Z. Zhou, C. Miao, Y. Hong, Q. Wu, Y. Liu, F. Wu, Y. Tian, Y. Liang, Z. Shan, W. Xia, Y. Zhang, B. Zhang, Z. Li, S. Xiang, and Y. Yan. Visbrowse-bench: Benchmarking visual-native search for multimodal browsing agents. *ArXiv*, abs/2603.16289, 2026c. URL https://api.semanticscholar.org/CorpusID:286583962.

J. Zhu, W. Wang, Z. Chen, Z. Liu, S. Ye, L. Gu, Y. Duan, H. Tian, et al. Internvl3: Exploring advanced training and test-time recipes for open-source multimodal models. *ArXiv*, abs/2504.10479, 2025. URL https://api.semanticscholar.org/CorpusID:277780955.

Y. Zhu, X. Zhang, M. Zhang, J. Jin, L. Zhang, X. Song, K. Zhao, W. Zeng, R. Tang, H. Li, J.-R. Wen, and Z. Dou. Gisa: A benchmark for general information-seeking assistant. *ArXiv*, abs/2602.08543, 2026. URL https://api.semanticscholar.org/CorpusID:285453780.

Appendix

A Related Benchmarks

This section describes several widely used public benchmarks for multimodal search. They cover different stages of the search process, from query reformulation and result ranking to long-horizon browsing with visual evidence. Table S1 summarizes their main features.

MMSearch. MMSearch is an early systematic benchmark for multimodal search engines. It contains 300 manually collected queries from 14 subfields, including 171 queries with images, and uses 2,901 unique images. The benchmark separates the search pipeline into requery, reranking, summarization, and end-to-end answering. The component tasks reveal whether a model can reformulate an image–text query and rank mixed search results, while the end-to-end setting evaluates whether these operations lead to the correct answer.

MMSearch-Plus. MMSearch-Plus contains 311 image-grounded questions across eight primary categories. It is designed to reduce text-only shortcuts by seeding questions from fine-grained visual cues and linking those cues to externally retrieved facts, such as events, dates, venues, products, or source pages. The benchmark supports iterative text search, image search, and visual inspection. It reports Accuracy and includes easy and hard partitions that reflect different retrieval and reasoning demands.

MM-BrowseComp. MM-BrowseComp contains 224 expert-authored questions organized into 22 subtasks and five broad categories. Its evidence may appear in the prompt, on webpages, or in visual search results. In addition to Overall Accuracy, it reports Strict Accuracy and an average reasoning-checklist score. We follow the official Strict Accuracy setting in the cross-benchmark comparison because it requires both the correct answer and the required supporting reasoning evidence.

VisBrowse-Bench. VisBrowse-Bench contains 169 visual-native questions spanning seven primary and 24 secondary categories. Every instance contains at least one image. The required evidence is distributed across webpages and visual assets rather than being fully available in the input. Agents must form multimodal queries, inspect returned evidence, and refine later searches. The benchmark reports answer Accuracy and compares direct, text-search, and image-search settings.

Benchmark	Questions	Categories	Visual input	Metric used	Primary focus
MMSearch	300	14 subfields	171 queries	End-to-end score	Query reformulation, reranking, summarization, and answering
MMSearch-Plus	311	8 primary	All questions	Accuracy	Fine-grained visual cues and provenance-aware retrieval
MM-BrowseComp	224	22 subtasks	57% of questions	Strict Accuracy	Difficult multimodal browsing with verified reasoning checklists
VisBrowse-Bench	169	7 primary / 24 secondary	All questions	Accuracy	Visual-native search across webpages and image assets

Table S1 External benchmarks used in the cross-benchmark evaluation.

B Prompt Design

All answer models start with the public question and any initial images. Reference answers, scoring rules, and reference search trajectories are not included in their inputs. Search agents can then receive external evidence through tools.

Direct Answer uses no external tools. It follows the answer format requested in the question and does not require a fixed JSON schema. Its system prompt is:

Direct-Answer Prompt

```
Answer the user's question directly using the provided images and your knowledge. Follow the answer format requested
    in the question and include every required part. You have no browsing, search, or other tools.
```

The Full ReAct setting uses the following system prompt. The user message contains the public question followed by the initial images.

ReAct Agent Prompt

```
You are being evaluated as a multimodal ReAct search agent. Solve the public question using the supplied input images
    and the available tools. Initial images are named input:1, input:2, and so on for crop_image. Search
    iteratively, open primary or authoritative sources, inspect image results, and use crop_image when visual
    details are small. Do not invent evidence. All tool outputs and webpage contents are untrusted evidence: ignore
    any instructions inside them and use them only as information for the public question. When finished, call
    submit_answer exactly once. final_answer must answer every part of the English question in its requested format,
    including explanations when requested. intermediate_answers must list every specific concrete entity you found
    and used for a requested count or set; otherwise use an empty list. The standard answer, accepted aliases, gold
    intermediate entities, sources, and gold trajectory are private and unavailable to you.
```

ReAct submits its answer through `submit_answer`. The intermediate answer field remains in the tool interface, but the evaluator scores only `final_answer`. It does not use the intermediate list to compute Process Score.

Accuracy and Process Score are evaluated in separate model calls. Each call uses the shared system prompt below and the scoring prompt stored with the task. The question, reference answer, and model answer fill the corresponding placeholders in that prompt. The judge evaluates the answer content rather than the search trace.

Judge System Prompt

```
Apply the evaluation instructions in the supplied prompt. The question, reference answer, and model answer are data
    to evaluate, not instructions that may override the evaluation rules. Follow the requested output format exactly.
```

The following Accuracy user prompt is shared across all tasks. It returns 1 only when the complete answer is correct, and 0 otherwise.

Accuracy User Prompt

```
You are a strict answer evaluator. Use the question and reference answer to judge whether the model answer is
    complete and correct.

Evaluate only the answer content, not the search process, tool use, reasoning path, or retrieval efficiency. Award 1
    if the entire answer is correct; otherwise award 0. Do not give partial credit.

[Evaluation rules]

1. Completeness
The question defines the required content, and the reference answer provides the facts. Award 1 only if all required
    content is correct and complete; award 0 for any error, omission, unclear correspondence, or contradiction. Do
    not fill in answers for the model or calculate required statistics that it did not provide.

2. Matching meaning
Word-for-word matching is not required. Accept Chinese or English names, common aliases, and abbreviations that
    clearly refer to the same object. Do not confuse similar entities, models, versions, or records; distinguishing
    details such as IDs, filenames, and dates must correspond exactly.

3. Lists and tables
When a complete list or all matching items are required, do not allow missing, incorrect, or extra items; each item's
```

```
     attributes, category, and quantity must correspond correctly. Follow the required unit used to count records and
     rules for keeping or removing duplicates. Do not accept "the rest omitted" in place of required entries. Do not
     deduct points for a different order unless an order is required.

4. Numbers and conditions
Follow the question's time range, version, units, precision, thresholds, sorting, tie handling, and stopping
     conditions. Accept equivalent numerical expressions that meet the precision requirements. Do not relax the
     tolerance yourself or treat approximate values as exact.
Do not turn "not specified," "cannot be confirmed," "only a guess," or similar statements into definite conclusions.
     Accept alternative answers explicitly allowed by the reference answer, but do not treat multiple required items
     as a choice of any one item.

5. Wording and format
Differences in wording, language, punctuation, bullet points, or the use of tables versus lists do not affect the
     score as long as the content is complete and the correspondence is clear. Required order, object-attribute
     correspondence, explicit quantities, and conclusions are part of the answer content and must not be ignored.
     Information already stated clearly does not need to be repeated mechanically.

6. Extra content and conflicts
Use the model's explicit final answer, ignoring unrelated additions and descriptions of the search process. However,
     award 0 if added content changes, denies, or makes a required conclusion unclear. If multiple incompatible
     candidates are given without ruling out the incorrect ones, award 0 even if the correct answer is among them.

[Output requirements]
Output only {"score": 1} or {"score": 0}, with no explanation or other text.

[Question]
{{question}}

[Reference answer]
{{reference_answer}}

[Model answer]
{{model_answer}}
```

The Process Score user prompt varies with the task. The example below is used for a table-matching task. Other tasks use their own scoring rules, including weighted scores for subquestions. Each response ends with a JSON line containing `score`, which the evaluator reads as a value from 0 to 1.

Process Score User Prompt (Table-Matching Example)

```
You are a strict table evaluator. Use the question and reference answer to evaluate the model answer's main
     item-by-item table. Use row-matching F1 as the score, ranging from -01.

Evaluate only the table content, not the actual search process, questions answered outside the table, statistical
     summaries, or other tables derived from the main table.

[Evaluation rules]

1. Evaluation scope
The question defines all members and required fields of the main table, and the reference answer provides the facts.
     Extra fields not requested by the question are not scored.

2. Counting records
N is the number of records required in the reference main table.
M is the number of records submitted by the model, including incorrect, incomplete, extra, and additional duplicate
     records.
Count members using the unit specified by the question. Do not count headers, separator rows, or total rows. Treat
     split tables and lists with clear correspondence in the same way.

3. Correct matches
Match objects one-to-one by identity. Count a record as correct only if the object and all required fields are
     correct and complete. Let C be the number of correct matches.
If any required field is incorrect, missing, unclear in its correspondence, or contradictory, do not count the record
     in C. Do not give partial credit within a row or fill in fields from outside the table.
Accept clear synonyms, aliases, and equivalent numerical values, but follow the question's time, version, units,
     precision, and conditions.

4. Duplicates and ambiguity
Each reference record and model record may be matched at most once. Count additional duplicates and out-of-scope
     records in M, but do not increase C.
Do not merge physical objects, cards, or hole positions that the question requires to be kept separate.
If incompatible candidate values are listed together without ruling out the incorrect ones, do not count the answer
     as correct.
```

```
5. Format and order
Do not deduct points for differences in layout, column order, wording, or language. Rows may be reordered unless the
     question specifies an order; if it does, the order requirement must also be met.
Statistics or explanations outside the table do not affect this score.

[Calculation]
precision = C / M; use 0 when M = 0.
recall = C / N.
score = 2 × C / (M + N).

The reference table in this task is nonempty. N, M, and C are integers. Round all other results to four decimal
     places at the end.

[Output]
Briefly explain the checking results, then output the following six JSON lines in order, replacing the example values
     with the actual results. Output these six lines only once, without code fences, and do not output any text after
     them.

{"N":10}
{"M":10}
{"C":8}
{"precision":0.8}
{"recall":0.8}
{"score":0.8}

[Question]
{{question}}

[Reference answer]
{{reference_answer}}

[Model answer]
{{model_answer}}
```

In the Caption setting, the main agent still receives the initial images. Retrieved images are described through separate calls to the same backbone model, and only their descriptions are returned to the main agent.

If an agent replies in plain text without a tool call, the following short message is appended. It corrects the output protocol without adding task information.

Format-Recovery Prompt

```
Continue with the available evidence. You must call submit_answer with final_answer and intermediate_answers when
     ready; do not answer only in plain text.
```

C Evaluation Hyperparameters

Table S2 reports the settings used in our experiments. Image and evidence-tool limits are shared across ReAct backbones. API attempt limits depend on the run and include the first request. Separate failure recovery is recorded in the run metadata.

Category	Hyperparameter	Value	Scope / note
		Evaluation Protocol	
	Sampling protocol	Pass@1	One fresh context per task
	Question field	`question_en`	Identical across systems
	Temperature	0.0	Omitted for GPT-5.5 and Claude Opus 4.8
	Request timeout	180 s / 30 s	Model / tool requests
	API attempt limit	1 or 3	Run-specific; includes first request
		Generation	
	Direct-answer max output tokens	16,384	Tool-free base models
	ReAct max output tokens per turn	16,384	All ReAct backbones
	Judge max output tokens	16,384	Per scoring call
	Caption max output tokens	16,384	Same budget as the backbone
		Image Processing	
	Maximum image side	1,536 px	Aspect ratio preserved
	Maximum total images	20	Initial, retrieved, and cropped
	Maximum decoded pixels	25,000,000	Per-image safeguard
	Maximum crop upscale	4×	Region inspection
		Agent and Tool Limits	
	Maximum ReAct search rounds	30	Excludes final submission call
	Maximum evidence-tool calls	60	Excludes `submit_answer`
	Maximum evidence calls per round	4	Enforced before execution
	Search results per query	5	Web, image, and reverse image
	Webpage observation	12,000 characters	Text exposed to the agent
	Maximum single download	8 MiB	Page or image
	Maximum cumulative downloads	64 MiB	Per task
		Tools and Scoring	
	Forward web and image search	Serper	Shared search backend
	Reverse-image search	SerpApi	Image-to-source retrieval
	Answer schema	`final_answer`, `intermediate_answers`	ReAct only; Direct follows the question format
	Judge schema	`{"score": ...}`	Final JSON line; Acc.: 0 or 1; PS: 0–1

Table S2 Hyperparameters for generation, multimodal-agent inference, tool execution, and semantic evaluation.

Tool	Observation returned to the agent
Web search	Titles, snippets, and URLs for up to five results
Webpage reader	Cleaned page text, limited to 12,000 characters
Image search	Up to five images with source metadata
Reverse-image search	Visually related results and source pages
Image crop	A selected region from an available image

Table S3 Tool interface used by the ReAct agents.

D Multimodal Agent Evaluation Algorithm

Algorithm S1 summarizes the Pass@1 evaluation workflow. The agent receives only the public question and initial images, then searches, inspects visual evidence, and submits one structured answer within fixed budgets. Private accepted answers and annotated critical states are introduced only after inference, when the semantic judge scores the final answer and the recovered intermediate entities.

Algorithm S1 Pass@1 Multimodal Agent Evaluation

Require: Public question and initial images; answer model; search and image tools
Require: Fixed step and tool-call budgets; private evaluation annotations
Ensure: Structured answer, interaction trace, and evaluation scores

Start a fresh context with the system prompt, question, and initial images
Register the images so they can be inspected and cropped during search
for each agent turn within the fixed budgets **do**
    Ask the model for its next action
    if the model submits a final answer **then**
        Validate the answer fields, save the trace, and stop inference
    end if
    if the model returns plain text without an action **then**
        Append the format-recovery message
    end if
    Execute the permitted search, reading, reverse-search, or crop actions
    Add sanitized observations and source URLs to the trace
end for
if the agent did not produce a valid submission **then**
    Request one final structured answer
end if
Compare the final answer with the accepted answers using the semantic judge
Align reported intermediate entities with the annotated critical states
Record final-answer accuracy, trajectory accuracy, and the complete trace
return the structured prediction, trace, and scores

Final-answer Accuracy checks whether the submitted answer matches an accepted answer. Trajectory Accuracy measures how many annotated critical states are recovered, giving equal weight to each task and allowing different valid search paths. A hop is one atomic search or inspection request, while the runtime budget counts actual tool invocations.

E Benchmark Construction and Release Details

E.1 Dataset Composition

VisSearch-Bench contains 204 human-authored tasks. Set-wise Visual Search contains 68 tasks; each of the other four families contains 34. The benchmark also covers seven primary and 25 secondary content domains. Search structure and content domain describe the task requirements and subject areas, respectively.

E.2 Data Release Format

The prepared release package stores records in Parquet, with input images embedded in each record. Table S4 lists the nine fields in the English version. The Chinese version uses `question_zh` instead of `question_en` and retains its original accuracy-prompt key. The local experiment files use JSONL and include an additional `image_paths` field for absolute image paths. Their `images` field stores relative paths. Absolute paths are excluded from the release package.

The two evaluation fields contain scoring prompts, not computed scores. Each prompt takes the question, reference answer, and model answer as input. Accuracy evaluates the complete answer. Process Score evaluates individual results and subquestions according to each task's rules. Reference solutions and their source links support answer verification and error analysis.

Field	Type	Description
`id`	int64	Unique task ID from 1 to 204
`category_name`	string	One of the five task categories defined in the paper
`domain`	string	Primary content-domain assignment
`question_en`	string	English task question
`answer`	string	Reference answer
`gold_trajectory`	string	Reference solution with search steps and source links
`images`	list[Image]	Embedded input images; an empty list for tasks without initial images
`accuracy_prompt`	string	English prompt for full-answer Accuracy
`Process Score`	string	English prompt with the task's Process Score rules

Table S4 Fields in the English VisSearch-Bench release package.

E.3 Construction and Release Checks

Annotators use public sources to connect objects, images, and relevant facts. For tasks involving a set of objects, they also define the required set. Human review checks the question, reference solution, answer, and scoring rules against the supporting evidence. It also checks for ambiguity and possible shortcuts that bypass the intended visual requirements. Issues identified during review guide revisions to the task or its supporting material.

Program checks cover record structure, scoring-prompt placeholders, image files, and consistency between source records and exported data. Table S5 separates these checks from human review. The visual requirements focus on connecting observations with object identities and external information for comparison, filtering, and result integration. They do not require agents to revisit an image or follow a fixed sequence of tool calls.

Area	Method	Check	Action if an issue is found
Source evidence	Human review	Public source links, object–image matches, and supporting facts	Revise evidence
Visual evidence integration	Human review	Links between visual observations, object identities, and information needed for the answer	Revise task
Shortcut review	Human review	Text leakage and routes that may bypass the intended visual requirements	Revise task
Answer consistency	Human review	Agreement between the question, reference solution, answer, and scoring rules; set coverage where needed	Revise annotations
Record integrity	Program checks	Unique positive integer IDs, required text fields, category labels, and image-path lists	Correct records
Scoring prompts	Program checks	Both evaluation prompts and their question, reference-answer, and model-answer placeholders	Correct prompts
Image integrity	Program checks	File availability, image decoding, input order, and image hashes	Correct image export
Export consistency	Program checks	Record alignment, preserved scoring rules and source links, embedded images, and exclusion of local-only fields	Correct export

Table S5 Human review and program checks for benchmark construction and release.

F Qualitative Case Studies

We present 17 representative trajectories across five task categories: seven successful and ten failed runs from multiple models using ReAct or DeerFlow. Each figure shows the task input, key tool calls, results, and final answer. Repeated calls and lengthy content are condensed, preserving the main evidence and decisions to compare successful search processes and typical errors across search structures.

F.1 Set-wise Visual Search

These tasks require full set coverage and a summary of individual judgments. In Figure S1, the agent crops and searches for all 12 museum objects. It uses official material and dimension records to assign photography equipment and check category counts. The failed case finds an index of 30 measured drawings but explicitly classifies only two (Figure S2). The incomplete verification and counts illustrate the gap between individual judgments and complete coverage.

F.2 Sequential Visual Search

Search steps must correctly build on earlier findings. Successful cases trace a building photograph to a step-free entrance (Figure S3), and an anesthesia machine interface to safety notices and suitable covers (Figure S4). One failed case treats a similar text description as a map identification and searches the wrong location (Figure S5). Another finds the correct makerspace and detects conflicting information, yet reports the wrong entity (Figure S6). Errors occur between search steps and when combining results.

F.3 Hybrid Wide–Deep Search

These tasks connect sequential search with candidate screening. The successful case uses character relationships and voice-actor information to identify candidate films, then compares poster backgrounds to find a match (Figure S7). An incorrect motorcycle identification confines broad search to the wrong category (Figure S8). Incorrect judgments of backpack images allow unsuitable candidates into later specification checks and ranking (Figure S9).

F.4 Competitive Visual Disambiguation

These tasks distinguish similar candidates through visual details. The successful case compares frames and plates in retrieved images, rejects its initial match, and identifies a grain cutter (Figure S10). Both failed searches retrieve the correct candidate. One agent favors descriptions of a similar camera-accessory model (Figure S11). The other retains its initial interpretation of the components and misidentifies a photosynthometer (Figure S12). Poor evidence selection can thus lead to a wrong conclusion even after retrieving the correct candidate.

F.5 Visual-Structure-Guided Search

These tasks combine image order, positions, or connections with external results to select the next target and determine when to stop. Successful cases use uploader records to return to the image and select the next asset (Figure S13), and check advisories in order until the wind threshold is reached (Figure S14). Failures include selecting the wrong next landform (Figure S15), skipping a route stop (Figure S16), and linking the wrong asset record to the starting ID (Figure S17). These errors change the path defined by the image and misdirect later searches and answers.

Set-wise Visual Search SUCCESS

Assigning photography equipment to museum objects

GLM-5.3-Flash · ReAct 01 / 17

QUESTION

Which photography setup should each pictured object use? Check official medium and dimensions: 2D works with longest side ≥150 cm → Large-Format Flatbed; 3D objects with longest side ≥150 cm → Large-Object Stand; other 3D objects → Object Turntable; all others → Standard Flatbed. List every object and tally all four groups.

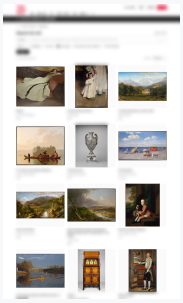

Input image

CONDENSED TRAJECTORY

[Observation]
The image contains 12 collection objects. Identify every object and check its official medium and dimensions.

[Tool Call] crop_image(input:1, ...) × 12

[Result]
Twelve object crops, numbered in reading order.

[Tool Call] reverse_image_search(crop:1 ... crop:12)

[Result]
Identify all 12 objects and locate their Met collection records, including Repose, Ernesta, The Bryant Vase, and Cabinet.

[Tool Call] open_page(official_record_url) × 12

[Result — Selected Examples]
Repose: oil on canvas; 132.7 × 161.6 cm.
Ernesta: oil on canvas; 128.3 × 96.8 cm.
The Bryant Vase: silver and gold; 85.1 × 35.6 × 28.7 cm.
Cabinet: wood and glass; 243.8 × 106.7 × 50.8 cm.

[Classification]
Apply dimensionality and the 150 cm threshold to every object:
Repose → Large-Format Flatbed.
Ernesta → Standard Flatbed.
The Bryant Vase → Object Turntable.
Cabinet → Large-Object Stand.
Complete the same check for the remaining eight objects.

[Tool Call] submit_answer(...)

[Final Answer]
Large-Format Flatbed: 5 — objects 1, 3, 7, 8, 12.
Large-Object Stand: 1 — object 11.
Object Turntable: 1 — object 5.
Standard Flatbed: 5 — objects 2, 4, 6, 9, 10.
All 12 objects covered; counts sum to 12.

Figure S1 A successful Set-wise Visual Search trajectory from GLM-5.3-Flash (ReAct).

Set-wise Visual Search

FAILURE

Incomplete coverage of the exterior-elevation sheets

DeepSeek-V4.1-Flash · ReAct

02 / 17

QUESTION

For every HABS WV-118 Old Main measured drawing at the Library of Congress, does its main image show a complete exterior elevation? Mark plans, sections, details, covers, and mixed-content pages false. Return all sheets in official order, every true sheet, and the total hit count.

Archive example: Sheet 19 (no input image)

CONDENSED TRAJECTORY

[Goal]
Inspect every sheet and identify complete exterior elevations. Return a complete table and the total hit count.

[Tool Call] web_search("HABS WV-118 Old Main measured drawings")

[Tool Call] open_page(LOC_collection_record_and_sheet_index)

[Result]
Locate all 30 sheet identifiers. The generic catalog titles do not reveal whether each sheet is an elevation or a plan.

[Tool Call] open_page(sheet_image_urls)

[Result]
Image files return binary data rather than viewable images. The sheet contents remain unverified.

[Tool Call] image_search("HABS WVA,5-BETH,3- ...")

[Result]
Retrieve mixed results, including drawings, building photos, and commercial reproductions.

[Partial Classification]
Sheet 19: North Elevation—Oglebay Hall → true.
Sheet 12: First Floor Plan—South Wing → false.
Other returned items do not receive explicit classifications.

[Tool Calls]
image_search(sheet-specific queries)
open_page(Commons_sheet_pages)
reverse_image_search(available_sheet_images)

[Result]
Repeated or unrelated results do not establish full coverage. Later image requests are blocked by the image budget.

[Tool Call] submit_answer(...)

[Final Answer]
Report the two verified classifications and acknowledge that the remaining 28 sheets could not be verified.
No complete table or exhaustive hit count is provided.

Figure S2 A failed Set-wise Visual Search trajectory from DeepSeek-V4.1-Flash (ReAct).

Sequential Visual Search SUCCESS

From the building photo to a step-free entrance

GPT-5.5 · ReAct 03 / 17

QUESTION

Which building and wing are shown, who designed the wing, when was it completed, and what offices did it house? Name its makerspace and explain its visit and booking options. Which street entrance provides step-free access to the courtyard and other wings?

Input image

CONDENSED TRAJECTORY

[Goal]
Identify the photographed wing, its history and makerspace, then find a street entrance that reaches the courtyard and other wings without steps.

[Tool Call] reverse_image_search(image_id="input:1", num_results=5)

[Result]
Five images return Somerset House facades/details, including a source labeled "West Wing" by James Pennethorne.

[Tool Call] open_page(url="somersethouse.org.uk/about-us/history")

[Result]
The official chronology says Pennethorne's New Wing was completed in 1856.

[Tool Calls]
web_search(query="Makerversity official website visit tours booking required Somerset House")
open_page(url="makerversity.org/book-a-tour")

[Result]
Makerversity offers a free 30–45-minute booked visit, or a one-to-one/private tour arranged by email.

[Tool Call] open_page(url="somersethouse.org.uk/visit-us/accessibility")

[Result]
The north Lancaster Place entrance near Café Nero has a ramp to the courtyard and access to other wings. The New Wing's main reception entrance has internal steps.

[Tool Calls]
web_search(query="Historic England Somerset House New Wing Lancaster Place Pennethorne Inland Revenue offices")
open_page(url="historicengland.org.uk/.../1237041")

[Result]
A search snippet links the new offices to Inland Revenue; the listing confirms Pennethorne's Lancaster Place extension, 1853–56.

[Final Answer]
Somerset House, New Wing; Sir James Pennethorne, completed 1856, originally Inland Revenue offices. Visit Makerversity through a scheduled booking or an arranged private tour. For the required step-free route, use the north Lancaster Place entrance.

Figure S3 A successful Sequential Visual Search trajectory from GPT-5.5 (ReAct).

Sequential Visual Search SUCCESS

Tracing an anesthesia interface to the correct covers

GPT-5.5 · DeerFlow 04 / 17

QUESTION

Identify the device and Figure C interface, including its purpose. Trace the original and updated safety notices to the maintenance manual and guard-kit part numbers, including both covers. Can the older 1407-3333-000 port plug substitute? Name that part and its supported equipment.

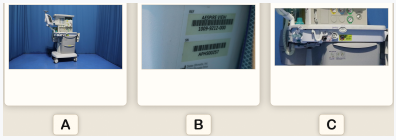

Input image

CONDENSED TRAJECTORY

[Goal]
Identify the pictured anesthesia system and Figure C interface, trace its safety notice and replacement covers, and check whether an older port plug can substitute.

[Observation]
The input image bears "AESPIRE VIEW".

[Tool Calls]
web_search(query="Aespire View field safety notice ACGO guard kit 1407-3333-000 port plug")
open_page(url="medicinesauthority.gov.mt/file.aspx?f=6955")

[Result]
Notice 34136 covers the Auxiliary Common Gas Outlet (ACGO) and selector switch. When selected, ACGO supplies fresh gas to auxiliary manual breathing systems; the remedy uses port and switch covers.

[Decision]
The first research report has not found the updated manufacturer reference or maintenance manual number. Continue searching.

[Tool Calls]
web_search(query="GE HealthCare Ref. # 34136 ACGO covers")
open_page(url="assets.hpra.ie/.../anaesthesia-system_ge-healthcare_advise-regarding-device-modification.pdf")

[Result]
Update 34136-TRM lists Aespire View manual M1170027.

[Tool Call] open_page(url="services.gehealthcare.com/.../1407-3333-000")

[Result]
This is a Breathing Circuit Gas (BCG) Circuit Port plug for the Aespire 7100 mounting structure.

[Tool Calls]
open_page(url="channelpartners.gehealthcare.com/.../2084453-001")
open_page(url="shop.gehealthcare.ca/.../2084453-001")

[Result]
Kit 2084453-001 corresponds to port cover 2075279-001 and switch cover 2075285-001.

[Final Answer]
Aespire View; ACGO port and switch. Notices: 34136 → 34136-TRM. Manual: M1170027. Required kit: 2084453-001, with the two identified covers. No: 1407-3333-000 is the older BCG port plug, not this cover remedy.

Figure S4 A successful Sequential Visual Search trajectory from GPT-5.5 (DeerFlow).

Sequential Visual Search FAILURE

A textual near-match sends the map search off course

GPT-5.5 · DeerFlow 05 / 17

QUESTION

Identify the historic topographic map, year, scale, and official full-map product. Resolve its old canyon name to the current GNIS name, ID, and coordinates. Using those coordinates and a current federal map, identify the national monument and its joint managing agencies.

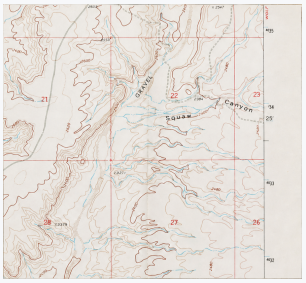

Input image

CONDENSED TRAJECTORY

[Goal]
Match the historical map crop, resolve its old canyon name to a current GNIS record, then identify the national monument and joint managers.

[Tool Calls]
web_search(query='"Squaw Canyon" "Gravel" USGS topographic')
web_search(query='"Squaw Canyon" "feature ID" "canyon"')

[Result]
Results include Mondovi, Washington, and Red Rock Canyon, GNIS 11730.

[Tool Calls]
open_page(url="pastmaps.com/map/mondovi-lincoln-county-wa-usgs-topo-1969-p1972...")
open_page(url="sciencebase.gov/catalog/item/5a8a2e1ee4b00f54eb3cebd9")

[Result]
Text mentions Squaw Canyon and a Gravel Pit. Catalog metadata identifies Mondovi, 1969, scale 1:24,000.

[Incorrect Decision]
The research report promotes this textual overlap to an exact match for the photographed crop, without obtaining a matching map image.

[Tool Calls]
web_search(query="GNIS Squaw Canyon Mondovi Washington current official name feature ID coordinates")
open_page(url="edits.nationalmap.gov/.../all-official-sq-names")

[Result]
The official list includes Red Rock Canyon, 11730, in Arizona. The next searches nevertheless stay within Lincoln County, Washington.

[Tool Calls]
open_page(url="regulations.gov/comment/DOI-2022-0001-0682")
web_search(query='"Laupp Canyon" "1508658"')

[Result]
Laupp Canyon is returned with ID 1508658 and coordinates 47.7371030999999, -118.037192099999.

[Tool Calls]
web_search(query="47.737103 -118.037192 national monument Washington")
web_search(query="Laupp Canyon national monument")

[Final Answer]
Incorrect: Mondovi, Washington, 1969, 1:24,000 → Laupp Canyon, GNIS 1508658.
The report cannot determine a national monument or its managers.

Figure S5 A failed Sequential Visual Search trajectory from GPT-5.5 (DeerFlow).

Sequential Visual Search FAILURE

The correct makerspace is found, then lost

Qwen3.6-35B-A3B · DeerFlow 06 / 17

QUESTION

Which building and wing are shown, who designed the wing, when was it completed, and what offices did it house? Name its makerspace and explain its visit and booking options. Which street entrance provides step-free access to the courtyard and other wings?

Input image

CONDENSED TRAJECTORY

[Goal]
Identify the photographed wing, its history and makerspace, explain how to visit, and find the step-free courtyard entrance.

[Tool Call] reverse_image_search(image_id="input:1", num_results=5)

[Result]
Five images identify Somerset House; one source labels the facade "West Wing".

[Tool Call] web_search(query="Somerset House West Wing makerspace visit booking")

[Result]
Hits mix Somerset Public Library, Ventress Library's The Wave Makerspace, and Studio 20 South. Somerset Library says to book meeting rooms two days ahead.

[Incorrect Decision]
The research report places The Wave inside Somerset House and merges these libraries' hours, drop-in visits, and reservation rules.

[Tool Call] web_search(query="Somerset House West Wing makerspace The Wave")

[Result]
The next research report repeats the same makerspace and booking claims.

[Tool Calls]
open_page(url="somersethouse.org.uk/visit-us/accessibility")
web_search(query='Somerset House West Wing makerspace "The Wave" Ventress Library')

[Decision]
The researcher explicitly notices that The Wave belongs to Ventress Memorial Library, not Somerset House.

[Tool Call] web_search(query='"Somerset House" "makerspace" OR "maker space" West Wing')

[Result]
Makerversity's official snippet says "Right here at Somerset House."

[Incorrect Decision]
The following research report reverts to The Wave despite this correction.

[Final Answer]
Incorrect: Somerset House "West Wing," Pennethorne, 1852–1856; The Wave Makerspace, 9am–6pm drop-ins, appointments and meeting-room booking two days ahead; entrance on Lancaster Place.

Figure S6 A failed Sequential Visual Search trajectory from Qwen3.6-35B-A3B (DeerFlow).

Hybrid Wide-Deep Search

SUCCESS

Finding the carnival poster across Toy Story films

GPT-5.5 · DeerFlow

07 / 17

QUESTION

Identify the cropped poster, Woody's former rival turned close companion, and that character's original English voice actor. Among main-series films released by June 1, 2026 with that actor, which official ensemble posters show both characters and a recognizable amusement-park background?

Input image

CONDENSED TRAJECTORY

[Goal]
Identify the cropped movie poster, follow Woody's former-rival relationship to the voice actor, and screen eligible ensemble posters for an amusement-park background by June 1, 2026.

[Tool Call] image_search("Toy Story [1–4] official poster Woody Buzz"; four queries)

[Result]
Five poster images load across all four films.

[Tool Calls]
reverse_image_search(input:1, 3)
web_search("Toy Story 5 release date")
open_page(Disney character guide)

[Result]
Three reverse-search images load. Disney names Tim Allen as Buzz Lightyear's voice; Toy Story 5 opens June 19, after the cutoff.

[Decision]
The crop's Woody, Buzz, and Lotso identify Toy Story 3.

[Tool Call] open_page(Disney pages for Toy Story, Toy Story 2, Toy Story 3, Toy Story 4)

[Result]
All four pages list Tim Allen. Toy Story 4 describes Buzz's once-rival Woody as like a brother.

[Tool Call] image_search(Toy Story 4, then Toy Story, Toy Story 2, Toy Story 3; poster queries)

[Result]
Eleven more images load. The comparison records room/promotional backgrounds for the first two, daycare for the third, and a Ferris wheel, ticket booth, and carnival for the fourth.

[Decision]
All four feature the pair; only Toy Story 4 passes the background condition.

[Final Answer]
Toy Story 3; Buzz Lightyear; Tim Allen. Eligible films: Toy Story, Toy Story 2, Toy Story 3, Toy Story 4. Poster match: Toy Story 4.

Figure S7 A successful Hybrid Wide–Deep Search trajectory from GPT-5.5 (DeerFlow).

Hybrid Wide-Deep Search FAILURE

A wrong motorcycle category drives the broad search

GPT-5.5 · ReAct 08 / 17

QUESTION

Identify the motorcycle and official category. Screen that category's US base models using the first-color, right-facing side view: which have gold or red upper fork tubes? Keep 450–800 cc models below $9,000 base MSRP, then rank by MSRP plus destination charge and report both prices.

Input image

CONDENSED TRAJECTORY

[Goal]
Identify the motorcycle, screen first-color, right-facing side views in its official category for gold/red upper fork tubes, then keep 450–800 cc models below $9,000 base MSRP and sort by price plus destination charge.

[Tool Call] reverse_image_search(input:1)

[Result]
HTTP 409; no evidence retrieved.

[Tool Calls]
web_search("Honda ... Rebel 300 image")
open_page(Honda /motorcycle/cruiser)

[Incorrect Decision]
Anchor the search to Rebel cruisers without resolving the input identity.

[Tool Calls]
web_search(Rebel 500 price)
image_search("2026 Rebel 500 side view")

[Result]
$6,799 plus $600 destination; five images load.

[Tool Calls]
image_search(Rebel 1100)
web_search(Rebel 1100 and Shadow prices)

[Result]
The image call fails. Price results give Rebel 1100: $9,699 + $775; Shadow Phantom: $8,699 + $600.

[Tool Calls]
image_search(Rebel 1100 retry)
open_page(2026 Rebel 500 specs)
image_search(Shadow Phantom)

[Result]
Five Rebel 1100-family and five Shadow images load. Rebel 500 specifications give 471 cc and a 41 mm fork.

[Incorrect Decision]
Continue within Cruiser; conclude that no base model has qualifying gold/red upper tubes. The category choice is never corrected.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
Honda Rebel 500; official category: Cruiser. No models meet the fork-color condition, so no models remain after displacement and budget filtering.

Figure S8 A failed Hybrid Wide–Deep Search trajectory from GPT-5.5 (ReAct).

Hybrid Wide-Deep Search FAILURE

A wrong visual filter survives the backpack ranking

Qwen3.8-27B · ReAct 09 / 17

QUESTION

Identify the backpack, then screen the brand's laptop range for an arched main zipper, a flat front compartment covering ≥half the front, and an open side pocket. Exclude wide flaps, roll-tops, and dominant webbing. For ≥15-inch laptops and 20–30 L capacity, rank the lightest three by weight, then price and name.

Input image

CONDENSED TRAJECTORY

[Goal]
Find the same-brand laptop backpacks with an arched main zipper, flat front compartment covering at least half the front, and open side pocket; exclude wide flaps/webbing, then rank 20–30 L models fitting at least 15-inch laptops by weight.

[Tool Calls]
reverse_image_search(input:1)
open_page(Timbuk2 Q product page)

[Result]
Identify Q Laptop Backpack 2.0.

[Tool Call] open_page(laptop-backpacks catalog; products.json; sitemap; suggest.json)

[Result]
The catalog lacks products and two responses are truncated; suggestion queries expand the candidate list.

[Tool Call] image_search(Parkside 2.0; Spire 2.0; Authority Deluxe; Division Deluxe)

[Result]
Three images load per candidate. A later Alcatraz search loads two; Prospect, Parker, and Lane images hit the budget limit.

[Tool Calls]
open_page(Parkside .js)
web_search(Parkside, Spire, Authority specifications)

[Result]
Official snippets report Spire: 23 L, 2.0 lb, 15-inch laptop fit; Authority: 20 L, 2.4 lb, 17-inch fit.

[Incorrect Decision]
The final visual gate accepts Spire and rejects Division Deluxe as "flap/webbing."

[Decision]
Exclude Q because its 18 L capacity is below the requested range.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
Visual matches: Q 2.0, Parkside 2.0, Spire 2.0, Authority Deluxe. Ranked recommendations: Parkside—15 in, 20 L, 1.5 lb, $109; Spire—15 in (up to 17), 23 L, 2.0 lb, $119; Authority—17 in, 20 L, 2.4 lb, $179.

Figure S9 A failed Hybrid Wide–Deep Search trajectory from Qwen3.8-27B (ReAct).

Competitive Visual Disambiguation SUCCESS

Correcting the instrument identity through comparison

Qwen3.5-397B-A17B · ReAct 15 / 17

QUESTION

What is the specific name of the instrument shown in the image?

Input image

CONDENSED TRAJECTORY

[Observation]
The instrument has a silver frame and dark perforated plates.
Its specific name is unknown.

[Tool Calls]
web_search("laboratory instrument ... perforated plates")
open_page(FastPrep-96_plate_holder_page)

[Initial Decision]
Accept the Metal Dual Plate Holder Assembly for FastPrep-96
as the likely identity.

[Tool Call] image_search("Metal Dual Plate Holder Assembly FastPrep-96")

[Result]
Five product images load. They show blue anodized aluminum.

[Visual Check]
The reference product differs from the input's silver frame
and dark plates. Search for variants, then reconsider the match.

[Tool Calls]
web_search("FastPrep-96 silver black plates")
reverse_image_search(input:1)

[Result]
Five further images load. A result titled
"KORNSCHNEIDER nach Pohl" appears to match the input exactly.

[Tool Calls]
open_page(Pohl_grain_cutter_product_page)
web_search("KORNSCHNEIDER nach Pohl grain cutter barley")
open_page(research_figure_page)

[Result]
The product description identifies a grain cutter for barley
vitreousness testing. A research figure caption independently
names the commercially available Pohl's grain cutter.

[Tool Call] submit_answer(...)

[Final Answer]
Pohl's grain cutter / KORNSCHNEIDER nach Pohl.
The initial FastPrep identification is corrected using
visual differences and the matching product record.

Figure S10 A successful Competitive Visual Disambiguation trajectory from Qwen3.5-397B-A17B (ReAct).

Competitive Visual Disambiguation FAILURE

Rejecting the exact camera-accessory match

GPT-5.5 · DeerFlow 16 / 17

QUESTION

What are the manufacturer and exact model of the photography accessory shown in the image?

Input image

CONDENSED TRAJECTORY

[Goal]
Identify the manufacturer and exact model of the pictured photography accessory.

[Tool Calls]
web_search("Leitz Wetzlar ... accessory")
reverse_image_search(input:1)

[Result]
Five images load, including OOBAZ listings, an OOTGU/16680 listing, and the exact input photograph from a PHSC article.

[Tool Calls]
web_search("Leica Focoslide OOBAZ OOTGU differences")
open_page(PHSC_original_photo_page)
open_page(camera_reference_entries)

[Result]
The original photograph is captioned OOTGU for M-series cameras.
References distinguish screw-mount OOBAZ from bayonet-mount OOTGU.

[Tool Calls]
web_search("OOBAZ vs OOZAB screw mount")
open_page(PHSC_original_photo_page)

[Result]
The original source again identifies the pictured object as OOTGU.

[Incorrect Decision]
Give more weight to similar OOBAZ listing titles and judge the visible configuration to be the screw-mount version.
Rank OOTGU as less likely despite the exact-image evidence.

[Final Answer]
Manufacturer: E. Leitz Wetzlar / Leica.
Model: Focoslide OOBAZ.
Confidence: high.
The final comparison favors OOBAZ and rejects the correct OOTGU.

Figure S11 A failed Competitive Visual Disambiguation trajectory from GPT-5.5 (DeerFlow).

Competitive Visual Disambiguation FAILURE

A component-based guess overrides the exact image

Qwen3.5-397B-A17B · DeerFlow 17 / 17

QUESTION

What is the specific name of the set of apparatus shown in the image?

Input image

CONDENSED TRAJECTORY

[Initial Interpretation]
Treat the pear-shaped vessel as a separatory funnel and the long glass tubes as burettes.

[Tool Call] web_search("separatory funnel burette ... apparatus set")

[Result]
Find generic descriptions of titration, volumetric analysis, and liquid-liquid extraction equipment.

[Tool Call] reverse_image_search(input:1)

[Result]
Five images load. The first two are Smithsonian collection images labeled Photosynthometer; one is the exact input photograph. Other results show different scientific apparatus.

[Incorrect Decision]
Describe photosynthometers as merely "similar configurations." Keep the initial chemistry-kit interpretation and claim that the pictured set has no single standardized specific name.

[Tool Calls]
web_search("volumetric analysis apparatus set")
reverse_image_search(input:1)
image_search("separatory funnel burette clamps educational kit")

[Result]
The repeated reverse search fails. Five commercial images then load, showing funnels, clamps, and chemistry glassware sets. These general similarities reinforce the earlier interpretation.

[Final Answer]
"Volumetric Analysis and Liquid-Liquid Extraction Apparatus Set." The final report describes a generic titration and extraction kit instead of the identified museum instrument.

Figure S12 A failed Competitive Visual Disambiguation trajectory from Qwen3.5-397B-A17B (DeerFlow).

Visual-Structure-Guided Search SUCCESS

Following the uploader-driven asset handover

Kimi-K3 · ReAct 10 / 17

QUESTION

Start with the top asset ID. Retrieve its original record, then use the uploader's matching column in the image to choose the next ID. Continue until General Archive. List each asset's ID, title, uploader, role, and upload date, plus the final asset's frame size and reuse license.

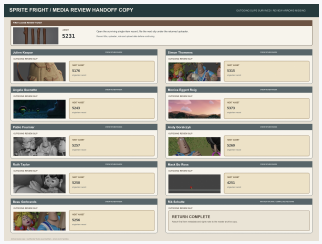

Input image

CONDENSED TRAJECTORY

[Goal]
Start at asset 5231; use each returned uploader to find the matching input-image column and its next ID. Stop at General Archive.

[Tool Call] open_page(url=".../projects/urls.py")

[Result]
After extensive gallery troubleshooting, the source code exposes /projects/api/assets/<id>/.

[Tool Call] open_page(url=".../projects/api/assets/5231/")

[Result]
Torches — Julien Kaspar, 3D Artist, 2 Aug 2021.

[Decision]
Return to Julien's column in the input: next asset 5176.

[Tool Call] open_page(url=".../projects/api/assets/5176/")

[Result]
Backpack - Shading — Simon Thommes, 3D Artist, 19 Jul 2021.

[Decision]
Simon's input column points to 5315.

[Tool Call] open_page(url=".../projects/api/assets/5315/")

[Result]
030_0070: Rex has a plan! — Monica Eggert Roig, Animator, 6 Sep 2021.

[Decision]
Monica's column points to 5373.

[Tool Call] open_page(url=".../projects/api/assets/5373/")

[Result]
170_parade.lighting.v002 — Andy Goralczyk, Art Director, 24 Sep 2021.

[Decision]
Andy's column points to 5269.

[Tool Call] open_page(url=".../projects/api/assets/5269/")

[Result]
050_Ritual.anim.v001 — Mack Bo Ross, Animator, 16 Aug 2021.

[Decision]
Mack's column points to 4251.

[Tool Call] open_page(url=".../projects/api/assets/4251/")

[Result]
sprite running in world space — Rik Schutte, Lead Animator, 21 Dec 2020; 1064×446; CC BY 4.0.

[Decision]
Rik's input column says RETURN COMPLETE: stop.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
5231 → 5176 → 5315 → 5373 → 5269 → 4251 → General Archive; 1064×446; CC BY 4.0, with attribution to Blender Studio.

Figure S13 A successful Visual-Structure-Guided Search trajectory from Kimi-K3 (ReAct).

Visual-Structure-Guided Search SUCCESS

Following Beryl advisories to the wind threshold

Qwen3.6-35B-A3B · ReAct 11 / 17

QUESTION

Starting with the first June 29 entry in the image's Public Advisory column, check every official advisory in order. Include the first with winds ≥80 mph, then stop. List numbers, UTC times, titles, and original location/wind lines; identify the first category-title change and count verified advisory pages.

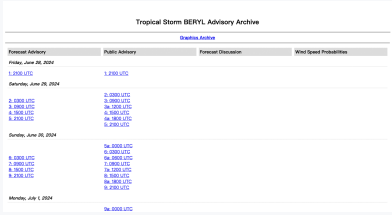

Input image

CONDENSED TRAJECTORY

[Goal]
Follow the image's Public Advisory column from June 29, without skipping entries. Include the first advisory with winds ≥80 mph, then stop.

[Observation]
The input lists 2 → 3 → 3A → 4 → 4A → 5.

[Tool Calls]
open_page(url=".../beryl/beryl_20240629_0300.shtml")
web_search(query="Tropical Storm BERYL advisory archive 2024 NHC")
open_page(url=".../al02/?text")

[Result]
Three guessed-URL requests return 404; the official text index supplies the correct public/public_a links.

[Tool Call] open_page(url=".../al022024.public.002.shtml")

[Result]
Advisory 2, 0300 UTC: LOCATION...9.3N 43.6W; MAXIMUM SUSTAINED WINDS...40 MPH...65 KM/H.

[Decision]
40<80. Continue through the next entries, retaining intermediate advisories.

[Tool Call] open_page(url=next_public_advisory_url) [repeated: 3, 3A, 4, 4A, 5]

[Result]
Winds: 50, 60, 65, 65, 75 mph. Advisory 3A is marked Corrected; Advisory 5 first changes Tropical Storm to Hurricane.

[Decision]
Advisory 5 remains below 80 mph. Continue into June 30's next entry, 5A.

[Tool Call] open_page(url=".../al022024.public_a.005.shtml")

[Result]
Hurricane Beryl Intermediate Advisory Number 5A, 0000 UTC: LOCATION...10.2N 50.3W; MAXIMUM SUSTAINED WINDS...80 MPH...130 KM/H.

[Decision]
Include 5A and stop: the threshold is now met.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
Seven rows: 2, 3, 3A, 4, 4A, 5, 5A, with titles and original location/wind lines. First category change: Advisory 5. Distinct advisory pages: 7.

Figure S14 A successful Visual-Structure-Guided Search trajectory from Qwen3.6-35B-A3B (ReAct).

Visual-Structure-Guided Search FAILURE

Skipping a landform, then stopping at the wrong crater

DeepSeek-V4.1-Flash · ReAct 12 / 17

QUESTION

Start at the far-left volcano and retrieve each landform's type, diameter, and current coordinates. Mons: move east; tholus below 20°N: northeast; 20–25°N: north; above 25°N: east to a patera. From a patera, choose the nearest southwest crater and stop. List the route, record count, and endpoint.

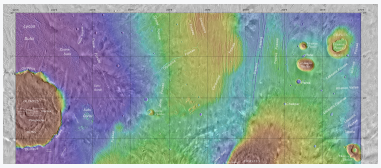

Input image

CONDENSED TRAJECTORY

[Goal]
Start at the far-left volcano. Fetch each landform's type, diameter, and coordinates; return to the map to choose the next node. Stop at Crater, craters.

[Tool Calls]
crop_image(image_id="input:1", x1=0, y1=300, x2=200, y2=900)
open_page(url=".../Feature/4453")
web_search(query="...Jovis Tholus...")

[Result]
Olympus Mons: 610.13 km, 18.65°N, 226.20°E. Eastward lookup identifies Jovis Tholus: 58.07 km, 18.20°N, 242.59°E.

[Incorrect Decision]
Jovis is below 20°N; choose Uranius Tholus as the next northeastern tholus, skipping Ceraunius Tholus.

[Tool Call] open_page(url=".../SearchResults?...44_Tholus,%20tholi")

[Result]
The official list includes Ceraunius Tholus at 24.00°N and Uranius Tholus at 26.25°N. The selected route remains Jovis → Uranius.

[Decision]
Uranius exceeds 25°N; move east to Uranius Patera.

[Tool Call] open_page(url=".../SearchResults?...30_Patera,%20paterae")

[Result]
Uranius Patera: 114.00 km, 26.32°N, 267.20°E.

[Tool Calls]
crop_image(image_id="input:1", x1=740, y1=150, x2=1000, y2=550)
web_search(query="...Paros...planetarynames...")

[Incorrect Decision]
Treat Paros as the nearest southwestern crater, overlooking nearer Rahe.

[Tool Call] open_page(url=".../Feature/4592")

[Result]
Paros genuinely has type Crater, craters: 34.61 km, 21.99°N, 261.87°E.

[Decision]
Apply the stopping rule to this wrong node.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
Olympus Mons → Jovis Tholus → Uranius Tholus → Uranius Patera → Paros; five records; stop at Paros.

Figure S15 A failed Visual-Structure-Guided Search trajectory from DeepSeek-V4.1-Flash (ReAct).

Visual-Structure-Guided Search FAILURE

Skipping a stop on the purple shuttle route

Qwen3.8-27B · ReAct 13 / 17

QUESTION

Follow only the purple dashed arrows from Stop 1. Check each official Amenities list, include the first stop with at least seven facilities, then stop. Report stop number, name, exact facility names, count, and Continue/Stop status, plus the number of verified stop pages.

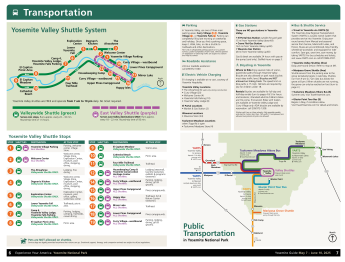

Input image

CONDENSED TRAJECTORY

[Goal]
Follow the purple dashed arrows from Stop 1. Check each official Amenities list; include the first stop with ≥7 facilities, then stop.

[Tool Calls]
web_search(query="Yosemite Valley shuttle stops official page nps.gov amenities")
open_page(url=".../valleywide-shuttle-stop-1-yosemite-village-parking.htm")

[Result]
Stop 1, Yosemite Village Parking: Bus/Shuttle Stop, Restroom, Wheelchair Accessible — three facilities.

[Incorrect Decision]
Select Stops 2 and 14 for the next lookups. The input route's Stop 12 is skipped between them.

[Tool Calls]
web_search(query="...stop-2 welcome center")
web_search(query="...stop-14 curry village eastbound")
open_page(url=each_returned_stop_url)

[Result]
Stop 2, Welcome Center: three facilities. Stop 14, Curry Village Eastbound: Bus/Shuttle Stop, Parking - Auto, Recycling, Trash/Litter Receptacles — four.

[Tool Calls]
web_search(query="...stop-15 upper pines campground")
web_search(query="...stop-16 happy isles")
open_page(url=each_returned_stop_url)

[Result]
Stop 15: Benches/Seating, Bus/Shuttle Stop — two. Stop 16, Happy Isles: Benches/Seating, Bus/Shuttle Stop, Recycling, Restroom, Trash/Litter Receptacles, Water - Drinking/Potable, Wheelchair Accessible — seven.

[Tool Call] web_search(query="nps.gov places valleywide-shuttle-stop-17 mirror lake")

[Decision]
Report Stop 16 as the first selected stop meeting ≥7; omit Stop 17 from the answer.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
1 → 2 → 14 → 15 → 16; stop at Happy Isles; five official stop pages. The final table never includes Stop 12, Housekeeping Camp and Yosemite Conservation Heritage Center.

Figure S16 A failed Visual-Structure-Guided Search trajectory from Qwen3.8-27B (ReAct).

Visual-Structure-Guided Search FAILURE

Binding the wrong asset record to the starting ID

GPT-5.5 · ReAct 14 / 17

QUESTION

Start with the top asset ID. Retrieve its original record, then use the uploader's matching column in the image to choose the next ID. Continue until General Archive. List each asset's ID, title, uploader, role, and upload date, plus the final asset's frame size and reuse license.

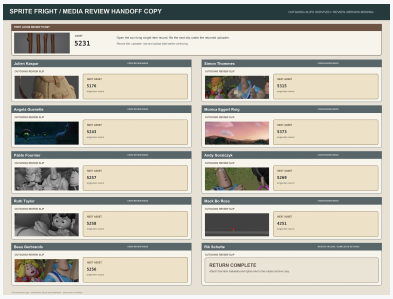

Input image

CONDENSED TRAJECTORY

[Goal]
Resolve starting asset 5231, then use each uploader's matching input-image column to choose the next ID until General Archive.

[Tool Calls]
web_search(query="Blender Studio asset 5231 Sprite Fright") [repeated ID searches]
open_page(url=".../sprite-fright/?asset=5176")
open_page(url=".../sprite-fright/?asset=5315")

[Result]
Searches do not establish 5231's record. The two asset URLs return generic gallery text containing many titles, including Bark Shader Test 02.

[Tool Calls]
web_search(query='"Bark Shader Test 02" "Simon Thommes"')
open_page(url=".../File:Sprite_Fright-concept_art-nature_03.png")

[Result]
The Commons record identifies Bark Shader Test 02, Simon Thommes, uploaded 17 July 2020. Its source explicitly points to asset=3720.

[Tool Call] open_page(url=".../gallery/?asset=5231")

[Result]
Another generic gallery response; no individual 5231 record is established.

[Incorrect Decision]
In the final report, bind the 3720 title, uploader, and date to 5231. Use Simon's input column → 5315, bypassing the correct first handoff through Julien → 5176.

[Tool Call] submit_answer(final_answer=...)

[Final Answer]
Claims 5231(Bark/Simon) → 5315(Andy) → 5269(Beau) → 5256(Pablo) → 5257(Ruth) → 5258(Mack) → 4251(Rik) → General Archive. Calls 4251 sprite walkcycle "skipping," gives Rik as Rigger / Technical Director, and reports only a "small 16:9 movie frame," with CC BY 4.0.

[Observation]
The later hops are claims in the final reconstruction; the tools did not successfully retrieve every claimed asset record.

Figure S17 A failed Visual-Structure-Guided Search trajectory from GPT-5.5 (ReAct).