arXiv:2611.06781v1 [cs.CV] 4 Nov 2026

Semantic Residual Vector Quantization for Sign Language Generation

Chang Liu[1] Ke Han[1,*] Niccolò Biondi[1] Gloria Isotton[1]
chang.liu-2@unitn.it ke.han.aca@gmail.com niccolo.biondi@unitn.it gloria.isotton@unitn.it
Elisa Ricci[1,2] Davide Talon[2] Nicu Sebe[1]
e.ricci@unitn.it dtalon@fbk.eu niculae.sebe@unitn.it
[1]University of Trento [2]Fondazione Bruno Kessler

Abstract

Sign language generation (SLG) maps written sentences into signing motion while preserving both lexical content and fine-grained articulation. This requires a model to decide what signs should be expressed and to realize them through precise handshape, trajectory, facial expression and body coordination. Existing pose-based systems and reconstruction-only motion tokenizers often entangle these two decisions, leaving the generator without an explicit text-aligned intermediate target for what to sign before refining how to sign it. To address this issue, we introduce Semantic Residual Vector Quantization (SRVQ), a framework for SLG that frames generation as a two-level process moving from semantics to kinematics. The newly introduced SRV Quantizer learns a hierarchy in which the first residual level encodes sign content and subsequent levels refine pose, handshape, trajectory, and body coordination, all within a single quantizer and without separate semantic representations or external planning stages. Through word-aware conditioning for hierarchical code generation, a masked semantic generator predicts semantic codes and a kinematic generator completes the residual codes for full SMPL-X motion synthesis. We evaluate using retrieval and FID in a shared text-motion embedding space, alongside standard geometric metrics. On CSL-Daily and Phoenix-2014T, SRVQ achieves R@1 of 82.3 and 70.3, outperforming the previous SOTA by 30 and 46 points, respectively.

1. Introduction

Sign language is the primary language of many deaf and hard-of-hearing people, yet today's digital society remains overwhelmingly built around spoken and written language [12, 27]. As a result, deaf users are often pushed toward text-based interaction rather than being supported in their native sign language. Sign language generation (SLG) aims to bridge this access gap by mapping written-language sentences into continuous signing motion, enabling more inclusive and accessible interaction.

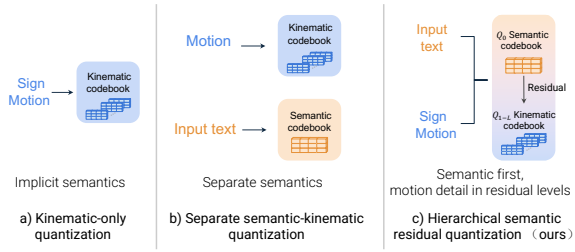

Figure 1. Comparison of quantization paradigms for SLG. (a) Kinematic-only quantization learns motion codes from motion alone, leaving semantics implicit. (b) Separate semantic-kinematic quantization models semantics and motion separately using additional semantic representations. (c) Our hierarchical semantic residual quantization embeds semantics and kinematics within a single residual quantizer, where the first residual level captures semantics and the remaining levels capture kinematic details.

SLG poses two coupled challenges: generating a signing sentence requires determining the linguistic content to express, *i.e.*, which signs should be produced and in what order, while realizing each sign through fine-grained articulatory patterns, including handshapes, movement, facial expression, and body coordination. This makes SLG a two-level generation problem, entangling semantic sign selection with precise motion realization. Yet, current models do not explicitly account for this two-level structure. They either generate signing motion end-to-end, mapping text directly to continuous poses [32–34], leverage discrete motion tokens optimized for reconstruction-only [39, 44, 51], or assume fixed intermediate sign representations that constrain the generation [50]. As a result, semantic sign selection and motion realization are modeled within a shared representation, making it difficult to learn representations tailored to their distinct roles.

*Corresponding author.

Latent image generation [5, 41] and text-to-motion models [14, 19, 47] suggest a broader principle: complex generation benefits from intermediate token spaces, where high-level content can be organized before low-level details are synthesized. Recent models in these domains introduce semantic structure by learning reasoning tokens on top of pretrained motion tokenizers [30], separating semantic and pixel branches in unified image tokenizers [31], or first learning a semantic codebook and then adding low-level refinement codebooks through staged training [7]. These designs show the value of separating high-level semantics from low-level realization, but they usually require separate semantic spaces, staged training, or additional codebooks that are challenging to coordinate with reconstruction.

To address these gaps, we propose *Semantic Residual Vector Quantization (SRVQ)*, a framework that formulates SMPL-X-based SLG in a *semantic-to-kinematic* fashion, explicitly reflecting the two-level structure of signing, as illustrated in Fig. 1. At the core of SRVQ, the *SRV Quantizer* employs a residual quantization hierarchy where the first residual codebook is semantically aligned through text and gloss supervision to capture what should be signed, while subsequent residual codebooks progressively encode the fine-grained kinematic details needed to realize it. All within the same quantizer hierarchy. Generation follows the same decomposition: a *semantic generator* predicts what to sign using word-aware conditioning that captures both global sentence and token-level lexical information, while a *kinematic generator* completes how the signs are articulated. We evaluate SRVQ on CSL-Daily [49] and Phoenix-2014T [21] using retrieval in a shared text-motion space, together with standard geometric distances. SRVQ achieves 82.3 R@1 on CSL-Daily and 70.3 on Phoenix-2014T, outperforming SOKE [51] by 30 and 46 points. Our contributions are summarized as follows:

- We introduce SRVQ, a semantic-to-kinematic generation framework that decomposes SLG into a two-level process that first predicts semantic information and then complement it with residual kinematics.
- We present the SRV Quantizer, which encodes motion through a single residual hierarchy learned with both reconstruction and text-alignment objectives. The first-level codebook captures high-level semantics, while later residual levels encode fine-grained kinematic details.
- We introduce word-aware conditioning, in which both generators cross-attend to word-token embeddings augmented with a global sentence token, giving motion tokens direct access to local lexical cues while retaining sentence-level context.
- SRVQ achieves state-of-the-art performance in SLG as measured with both retrieval and geometric metrics.

2. Related Work

Sign Language Generation. SLG aims to synthesize continuous signing motion from text input, and prior work has explored several ways to bridge this linguistic-to-kinematic gap. Progressive Transformers [33] directly regress skeletal motion from text, while adversarial multi-channel training [32] and motion-primitive modeling [34] improve continuous pose generation. Other methods synthesize human appearance through 3D avatars or meshes [2, 38], while related controllable image-generation methods introduce explicit pose or appearance conditions [35, 36]. Diffusion-based methods model signing motion directly [10]. Recent discrete approaches introduce a motion vocabulary before generation. T2S-GPT [44] uses dynamic vector quantization for autoregressive sign generation, SOKE [51] combines discrete sign tokens with retrieval-enhanced conditioning, and M3T [39] introduces discrete multi-modal motion tokens. Several of these systems additionally use gloss annotations as intermediate supervision [33, 43], where glosses provide word-like labels for individual sign units and serve as sign-level anchors between spoken-language text and signing motion. These methods improve sequence modeling, but the tokenizer is usually trained mainly for motion reconstruction. As a result, the discrete code space is not explicitly organized to separate sign content from kinematic realization. This motivates us to look beyond SLG to tokenized text-to-image and motion generation, where discrete or latent representations bridge high-level content and low-level synthesis.

Semantic Organization in Tokenized Generation. Vector quantization provides a discrete interface between continuous signals and sequence models [41]. In text-to-image generation, masked and autoregressive token models predict visual tokens before decoding them into pixels [5, 15, 22], suggesting that an intermediate token space can separate high-level content prediction from low-level synthesis. Text-to-motion generation adopts a similar tokenized formulation for temporally extended human motion. T2M-GPT [47] and MotionGPT [19] model motion as discrete token sequences, while later methods further explore semantic-to-kinematic separation through global-local attention [48], text-aligned motion representations [26], explicit planning and posture stages [24, 37], or reasoning and execution latents [30]. These works indicate that generation benefits from exposing an intermediate semantic structure before producing detailed motion. However, this semantic-reconstruction balance is often introduced through additional objectives, separate representations, or multi-stage generator designs [24, 26, 30, 37, 48]. Residual quantization [8, 14, 22, 46] offers a natural hierarchy of codebooks, where later levels progressively reduce the residual error left by earlier ones. In existing image, audio, and text-to-motion tokenizers, this hierarchy is mainly used to improve recon-

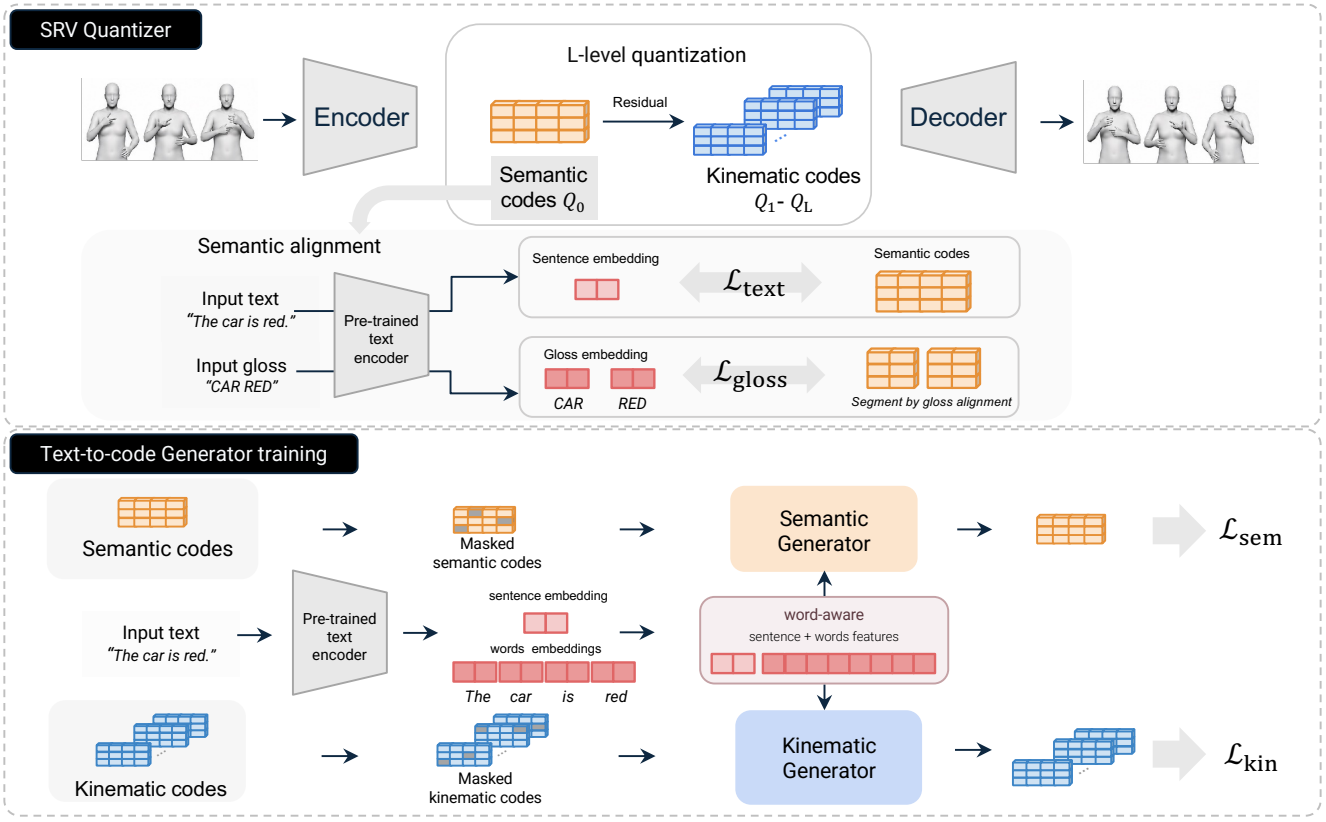

Figure 2. **Overview of the SRVQ framework.** *Top:* The SRV Quantizer learns a hierarchical discrete representation of each SMPL-X motion. The first residual level (Q_0) learns a text-aligned semantic codebook, while the remaining residual levels Q_1–Q_L progressively capture finer kinematic details. *Bottom:* During generation, word embeddings together with a global sentence token form a word-aware textual representation. A Semantic Generator predicts Q_0 from this representation, and a kinematic generator predicts Q_1-Q_L before the part decoders reconstruct the continuous SMPL-X motion.

struction fidelity or rate-distortion trade-offs. SRVQ instead uses the hierarchy as a semantic-to-kinematic interface for SLG, assigning the first residual level to text-aligned sign content and the later levels to kinematic refinement.

Text Conditioning for Sign Generation. Beyond the tokenizer, the text representation used by the generator also matters. Many SLG systems inherit the text-to-motion practice of conditioning generation on a pooled sentence embedding or a small set of global text tokens [23, 40, 44, 51]. This design gives the generator sentence-level context, but it can hide the local correspondences between words and sign segments. The limitation is especially important for SLG because individual lexical units often determine handshape, spatial location, handedness, and local motion timing. Our generator instead predicts semantic motion codes by cross-attending to word-token embeddings augmented with a global sentence token. This exposes local lexical cues to every generated motion token while retaining sentence-level context for the whole sequence.

3. Method

We introduce SRVQ, a semantic-to-kinematic framework for SLG, as shown in Fig. 2. It first learns a *SRV Quantizer* (Fig. 2-top, Sec. 3.1) that decomposes signing motion into a text-aligned semantic level and residual kinematic levels. Later, in the *Word-Aware Conditioning for Hierarchical Code Generation* (Fig. 2-bottom, Sec. 3.2), we train a semantic generator to predict semantic codes (Q_0) from text, while a kinematic generator completes the residual codes (Q_1–Q_L) required for full motion synthesis.

Problem Formulation. Given an input sentence $\mathbf{c} = (c_1, \ldots, c_M)$, where c_i denotes the i-th word, SLG aims to generate a continuous SMPL-X signing motion sequence $\mathbf{S} \in \mathbb{R}^{T \times 133}$ that preserves the sentence meaning. Following [51], we represent $\mathbf{S}$ using a 133-dimensional SMPL-X parameterization and decompose it into body (including facial features), left-hand, and right-hand streams, denoted by $\mathbf{S}^B$, $\mathbf{S}^L$, and $\mathbf{S}^R$. We first describe how the SRV Quantizer learns a motion vocabulary whose first code level is aligned with sign content, and then introduce the hierarchical gen-

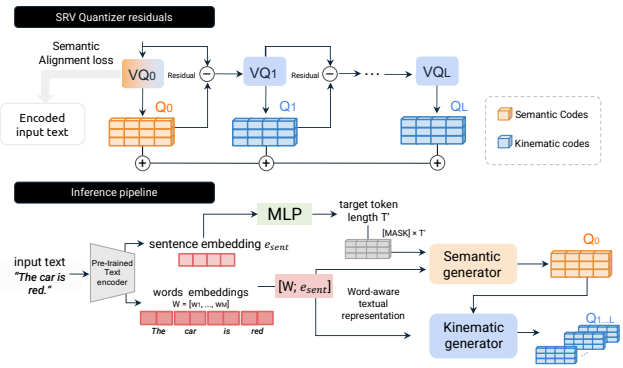

Figure 3. *Top:* Each quantization level encodes the residual from the previous level, with semantic supervision applied only to Q_0. The final representation consists of the semantic code and residual kinematic codes. *Bottom:* During inference, the Semantic Generator predicts Q_0 from the word-aware textual representation, followed by the Kinematic Generator predicting $Q_1, \ldots, Q_L$.

erators that predict semantic and residual codes from text.

3.1. SRV Quantizer

Residual Vector Quantization (RVQ) [8, 46] provides a natural mechanism for learning hierarchical discrete representations. We exploit this structure to separate semantics, *i.e.*, what to sign, from kinematic refinement, *i.e.*, how to sign it, within a single end-to-end quantizer, without imposing explicit constraints on how modeling capacity is allocated across levels. Specifically, the first residual level is trained as a text-aligned semantic codebook, capturing the meaning of the sign content to be produced, while the subsequent levels encode residual motion details that specify how the sign is articulated. This design keeps semantic code prediction and motion reconstruction within the same hierarchy, providing the generator with an explicit semantic target while subsequent levels capture fine-grained motion details.

The SRV Quantizer processes $\mathbf{S}^B$, $\mathbf{S}^L$, and $\mathbf{S}^R$ independently, with each body-part using its own encoder/decoder and codebooks. For clarity, we describe the quantizer for a generic stream $\mathbf{X} = \mathbf{S}^p, p \in \{B, L, R\}$ and omit the part superscript. For each body part, a 1D-CNN encoder maps the motion $\mathbf{X}$ to a latent sequence $\mathbf{F} \in \mathbb{R}^{T' \times d}$, with temporal downsampling factor $u = 4$. We apply an $L+1$-level residual quantizer to $\mathbf{F}$. Starting from the initial residual sequence $\mathbf{r}^0 = \mathbf{F}$, the quantizer recursively approximates the residual at each level and passes the remaining residual to the next level. At level ℓ, each residual token is assigned to its nearest entry in a learned codebook $\mathbf{Z}^\ell = \{\mathbf{z}_n^\ell\}_{n=1}^{N^\ell}$:

$$
\begin{aligned}
i_t^\ell &= \arg\min_n \left\| \mathbf{r}_t^\ell - \mathbf{z}_n^\ell \right\|_2^2, \qquad \hat{\mathbf{z}}_t^\ell = \mathbf{z}_{i_t^\ell}^\ell, \\
\hat{\mathbf{Z}}^\ell &= \left(\hat{\mathbf{z}}_1^\ell, \ldots, \hat{\mathbf{z}}_{T'}^\ell\right), \qquad \mathbf{r}^{\ell+1} = \mathbf{r}^\ell - \hat{\mathbf{Z}}^\ell.
\end{aligned}
\tag{1}
$$

Here, $\mathbf{r}^\ell \in \mathbb{R}^{T' \times d}$ denotes the residual sequence at level ℓ, $\mathbf{r}_t^\ell \in \mathbb{R}^d$ denotes its token at temporal position t, and $\hat{\mathbf{Z}}^\ell \in \mathbb{R}^{T' \times d}$ is the quantized sequence selected at level ℓ. We denote the level-ℓ discrete index sequence by $Q_\ell = (i_1^\ell, \ldots, i_{T'}^\ell)$; looking up these indices in the codebook $\mathbf{Z}^\ell$ yields the quantized embeddings $\hat{\mathbf{Z}}^\ell$.

After the final residual level, the decoder $\mathcal{D}$ reconstructs the motion stream from the sum of the selected code embeddings: $\hat{\mathbf{X}} = \mathcal{D}\left(\sum_{\ell=0}^{L} \hat{\mathbf{Z}}^\ell\right)$.

We encourage level-0 codes to model semantic information, while leaving the remaining residual levels free to capture fine-grained kinematic details (Fig. 3-top). To this end, we concatenate the embeddings from the three part branches $\hat{\mathbf{Z}}^{p,0}$ and project them into a text-motion d-dimensional contrastive space:

$$
\mathbf{H}^0 = [\hat{\mathbf{Z}}^{B,0}; \hat{\mathbf{Z}}^{L,0}; \hat{\mathbf{Z}}^{R,0}], \quad \mathbf{M} = \mathrm{norm}(\phi(\mathbf{H}^0)) \in \mathbb{R}^{T' \times d},
\tag{2}
$$

with ϕ denoting a two-layer MLP. A pretrained text encoder [11] produces word-token embeddings for the input sentence $\mathbf{c}$, which are projected to the same d-dimensional contrastive space as $\mathbf{M}$, yielding $\mathbf{W} = (\mathbf{w}_1, \ldots, \mathbf{w}_M)$. Here, $\mathbf{M}$ is a continuous feature for semantic alignment and is distinct from the discrete index sequence Q_0. Let $\mathrm{sim}(\cdot, \cdot)$ denote cosine similarity. We define a bidirectional token-level matching score between motion features $\mathbf{M}$ and word tokens $\mathbf{W}$ as:

$$
\begin{aligned}
\mathrm{match}(\mathbf{M}, \mathbf{W}) = \frac{1}{2}\Bigg(& \frac{1}{T'} \sum_{t=1}^{T'} \max_{1 \le i \le M} \mathrm{sim}(\mathbf{m}_t, \mathbf{w}_i) \\
& + \frac{1}{M} \sum_{i=1}^{M} \max_{1 \le t \le T'} \mathrm{sim}(\mathbf{m}_t, \mathbf{w}_i) \Bigg),
\end{aligned}
\tag{3}
$$

where $\mathbf{m}_t$ and $\mathbf{w}_i$ denote the t-th row of $\mathbf{M}$ and the i-th row of $\mathbf{W}$, respectively, and the first term measures how well each motion feature is explained by its most similar word token, while the second term measures how well each word token is covered by the motion sequence. We use this matching score in an InfoNCE loss [28] over the samples in each mini-batch, encouraging the level-0 motion features to align with their paired text while contrasting against others:

$$
\mathcal{L}_{\text{text}} = -\frac{1}{|\mathcal{B}|} \sum_{j \in \mathcal{B}} \log \frac{\exp(\mathrm{match}(\mathbf{M}_j, \mathbf{W}_j)/\tau)}{\sum_{k \in \mathcal{B}} \exp(\mathrm{match}(\mathbf{M}_j, \mathbf{W}_k)/\tau)},
\tag{4}
$$

with $\mathcal{B}$ the mini-batch and τ the contrastive temperature.

Similarly, we use frame-level gloss annotations to provide local semantic supervision for the level-0 motion features $\mathbf{M}$. Each gloss label is associated with a start and end frame in the motion sequence, which we use to crop the corresponding segment of $\mathbf{M}$. Let $(\mathbf{M}_i, \mathbf{G}_i)$ denote the cropped motion-feature segment and its associated gloss text features in a mini-batch, where $\mathbf{G}_i$ is ob-

tained by encoding the gloss label with the same text encoder [11]. We define $\mathcal{L}_{\text{gloss}}$ as the same InfoNCE objective as Eq. 4, replacing full motion-text pairs $(\mathbf{M}_j, \mathbf{W}_j)$ with gloss cropped-motion ones.

For motion reconstruction, we use a Smooth-L1 loss $\mathcal{L}_{\text{rec}}$ between reconstructed and ground-truth motion streams. Following prior motion tokenization practice [14], we add the first-order velocity loss $\mathcal{L}_{\text{vel}}$ on $\Delta\mathbf{X}_t = \mathbf{X}_t - \mathbf{X}_{t-1}$ to encourage temporal smoothness. We also use the standard VQ commitment loss $\mathcal{L}_{\text{commit}}$ [8, 41, 46], which keeps encoder residuals close to their assigned code embeddings. The final objective for the SRV Quantizer is:

$$\mathcal{L} = \mathcal{L}_{\text{rec}} + \lambda_{\text{text}}\mathcal{L}_{\text{text}} + \lambda_{\text{gloss}}\mathcal{L}_{\text{gloss}} + \lambda_{\text{vel}}\mathcal{L}_{\text{vel}} + \lambda_{\text{commit}}\mathcal{L}_{\text{commit}}, \tag{5}$$

with $\lambda_{\text{text}}, \lambda_{\text{gloss}}, \lambda_{\text{vel}}$ and λ_{commit} hyperparameters regulating the strengths of the respective losses.

3.2. Word-Aware Conditioning for Hierarchical Motion Generation

In the second stage, we freeze the SRV Quantizer and train generators to predict its discrete code hierarchy from text. The quantizer converts each motion sequence into part-level code streams: Q_0 codes represent text-aligned sign content, while higher residual levels represent kinematic detail. Generation follows the same hierarchy. A semantic generator first predicts the Q_0 codes, and a kinematic generator then completes the residual levels for reconstruction.

Word-Aware Conditioning. To condition generation on both local lexical information and sentence-level context, we encode the input sentence $\mathbf{c}$ with a pretrained text encoder [11]. This produces word-token embeddings $\mathbf{W}$ and a global sentence embedding $\mathbf{e}_{\text{sent}}$. We append the sentence embedding to the word-token sequence and use the resulting word-aware textual representation for conditioning generation via cross-attention. This lets each motion token attend to local lexical cues while retaining the sentence-level context of the full sentence.

Semantic Generator. The semantic generator predicts Q_0 with a bidirectional masked transformer [5, 14]. During training, we randomly mask a subset of Q_0 tokens with a learnable `[MASK]` token and train the model to recover the masked code indices by cross-attending to the word-aware textual representation. We rely on a cross-entropy loss over masked positions:

$$\mathcal{L}_{\text{sem}} = -\frac{1}{|\mathcal{I}|}\sum_{j\in\mathcal{I}} \log p\left(i_j^0 \mid Q_{0,\setminus\mathcal{I}}, \mathbf{W}, \mathbf{e}_{\text{sent}}\right), \tag{6}$$

where i_j^0 is the masked level-0 code index at position j and $Q_{0,\setminus\mathcal{I}}$ denotes the visible level-0 indices. Since each body part is quantized with a separate codebook, we predict all body-part codes jointly using separate part heads.

Kinematic Generator. The Kinematic Generator follows the same masked prediction formulation, but operates on the residual levels rather than the Q_0 semantic codes. We denote the summed cross-entropy loss over residual levels 1 to L as $\mathcal{L}_{\text{kin}}$. Training leverages a teacher-forcing strategy where lower-level codes are obtained by quantization of the ground-truth motion.

Inference. As shown in Fig. 3-bottom, at inference time, a lightweight MLP first predicts the target token length T' from the global sentence embedding $\mathbf{e}_{\text{sent}}$. The length predictor is trained with cross-entropy over discretized length bins, using the length of the ground-truth motion as supervision. Given the predicted length, the semantic generator starts from a fully masked sequence of length T', to progressively decode semantic codes using confidence-based unmasking [5]. To improve text adherence, we apply classifier-free guidance (CFG) [17] during decoding. The resulting semantic codes are then passed to the kinematic generator, which predicts the residual codes needed to recover fine-grained articulatory details.

4. Experiments

We evaluate SRVQ on standard sign language generation benchmarks. We first describe the evaluation protocol, including implementation details, datasets, and metrics. We then compare against prior work through both quantitative and qualitative evaluations, followed by ablation studies that isolate the contribution of the key design choices.

4.1. Experimental Setup

Implementation Details. SRVQ is trained jointly on CSL-Daily and Phoenix-2014T. It uses 6 quantization levels ($L = 5$) with code dimension 512, body codebook size $N_B^\ell = 96$, and hand codebook sizes $N_L^\ell = N_R^\ell = 256$ at each level. The quantizer loss of Eq. 5 uses $\lambda_{\text{text}} = 0.1$, $\lambda_{\text{gloss}} = 0.1$, $\lambda_{\text{vel}} = 0.5$ and $\lambda_{\text{commit}} = 0.02$. The contrastive temperature of Eq. 4 is $\tau = 0.07$, $d = 512$. Both the semantic and kinematic generators are 8-layer masked Transformers. We train all modules with AdamW [25] using learning rate 2×10^{-4} and batch size 64 per GPU on 4 GPUs. The quantizer is trained for 200 epochs, while the two generators are trained for 300. Evaluation is conducted on CSL-Daily and Phoenix-2014T test sets. At inference, we use CFG scale $\gamma = 2.5$, and $K = 10$ decoding steps. Unless otherwise stated, we use inference seed 3407, following the default configuration of MoMask [14].

Datasets. We evaluate on CSL-Daily [49] (18K training / 1176 test sentences, Chinese sign language) and Phoenix-2014T [21] (∼7K training / 642 test sentences, German sign language). Following prior work [51], we use the released 133-dimensional SMPL-X pose annotations sampled at 25 fps. For gloss supervision, we rely on the frame-level annotations from Spoken2Sign [50], which provide gloss labels

Table 1. **Comparison with state-of-the-art methods on CSL-Daily and Phoenix-2014T.** The best and second-best results are highlighted in **bold** and underline, respectively.

Method	CSL-Daily								Phoenix-2014T							
	Retrieval / FID				DTW-JPE		DTW-PA		Retrieval / FID				DTW-JPE		DTW-PA	
	R@1↑	R@5↑	R@10↑	FID↓	Body↓	Hands↓	Body↓	Hands↓	R@1↑	R@5↑	R@10↑	FID↓	Body↓	Hands↓	Body↓	Hands↓
Prog. Trans. [33]	–	–	–	–	16.30	32.63	15.98	12.91	–	–	–	–	15.01	31.77	13.67	11.95
Text2Mesh [38]	–	–	–	–	13.76	30.37	13.47	12.10	–	–	–	–	14.04	31.64	13.48	12.06
T2S-GPT [44]	–	–	–	–	12.32	15.43	11.94	5.93	–	–	–	–	11.65	19.09	10.38	6.47
S-MotionGPT [19]	–	–	–	–	11.58	11.31	10.81	3.78	–	–	–	–	10.42	9.08	9.45	3.41
SOKE [51]	51.96	69.30	75.77	0.062	6.541	8.524	7.407	**1.503**	24.14	45.64	54.67	0.151	5.914	7.316	6.464	1.368
SRVQ (semantic-only)	77.81	89.97	93.28	**0.057**	**6.296**	**8.376**	**7.113**	1.508	57.48	78.97	85.83	0.083	**4.700**	**6.410**	**5.315**	**1.262**
SRVQ (sem+kinematic)	**82.31**	**94.13**	**96.34**	0.058	6.787	8.535	7.614	1.544	**70.25**	**87.54**	**92.52**	**0.059**	5.330	6.920	6.003	1.354

together with their start and end frames.

Evaluation Metrics. Following prior works [33, 51], we report DTW-JPE and DTW-PA. DTW first aligns generated and reference motions over time with dynamic time warping, which reduces the effect of length and speed mismatch. DTW-JPE then computes the average joint position error over the aligned frame pairs. DTW-PA additionally applies Procrustes alignment before computing the same joint error, so it discounts global translation, rotation, and scale differences. Lower values indicate a smaller geometric distance to the reference motion and, therefore, better fidelity. However, geometric metrics alone are insufficient for assessing semantic correctness. A sentence can be signed with valid variations in timing, coarticulation, and style, and imperfect spatiotemporal alignment further complicates reference-based evaluation. As a result, geometric distance to a single reference cannot fully capture whether the generated motion preserves the intended meaning. Recent SLG evaluation studies therefore advocate embedding-based semantic evaluation as a complement to pose-distance metrics [18, 20]. Retrieval-based semantic matching is also widely used in text-to-motion generation [14, 16, 29, 47] and has recently been adopted in sign-language generation [3, 45]. We use retrieval metrics as our primary semantic indicator. To this end, a pretrained text-motion retrieval model [1] embeds generated motions and text into a shared space. For each benchmark, this evaluator is trained on real SMPL-X motion-text pairs from the corresponding training split and is used for all generated and real motions. We report motion-to-text $R@1$, $R@5$, $R@10$, and FID. We refer to the Supplementary A for additional validation using Meaningful Pose and back-translation evaluators adapted to SMPL-X motion.

4.2. Main Results

Baselines. We compare SRVQ against representative sign language generation baselines reported by SOKE [51]. These include Progressive Transformer (Prog. Trans.) [33], which directly regresses continuous signing poses; Text2Mesh [38], which leverages mesh-based generation; and discrete motion-token generation methods, namely T2S-GPT [44] and S-MotionGPT [19] (MotionGPT [19] adapted for sign language as evaluated by SOKE). To ensure comparability, we report geometric metrics using the numbers provided by SOKE [51]. Since prior work does not report retrieval-based metrics, we additionally evaluate SOKE under the same retrieval protocol used for SRVQ.

Quantitative Results. Tab. 1 compares SRVQ with state-of-the-art SLG models. We report two variants: *Semantic-only*, which decodes motion after predicting only the level-0 semantic codes Q_0, and *sem+kinematic*, which additionally predicts the residual kinematic codes before decoding.

SRVQ outperforms prior methods across nearly all considered metrics. Sem+kinematic achieves the strongest semantic retrieval performance on both datasets, while semantic-only performs best on most geometric metrics.

Intuitively, semantic-only produces a coarse, temporally smooth, low-frequency representation that is close to the average of many valid articulations of the same sign, which is well aligned with what DTW-based metrics reward. We revisit this observation in Sec. 4.4. In contrast, sem+kinematic adds articulation details, improving the model's ability to capture fine-grained distinctions between samples. On CSL-Daily, sem+kinematic improves R@1 from 51.96 to 82.31, yielding a gain of 30.35 percentage points over the strongest prior baseline. On Phoenix-2014T, it improves R@1 from 24.14 to 70.25, a gain of 46.11 points, and also obtains the best FID. Notably, most of the semantic retrieval gain already comes from explicit semantic modeling, supporting the proposed hierarchical semantic modelling. Retrieval improvements are consistent across R@5 and R@10, indicating that the generated motions are not only ranked closer to the correct sentence at top-1, but also occupy a better neighborhood in the embedding space. For reference, the retrieval R@1 of real test motions under the same protocol is 87.0 on CSL-Daily and 71.0 on Phoenix-2014T, showing that Sem+kinematic substantially closes the gap to ground-truth retrievability.

Qualitative Results. Fig. 4 visualizes representative gen-

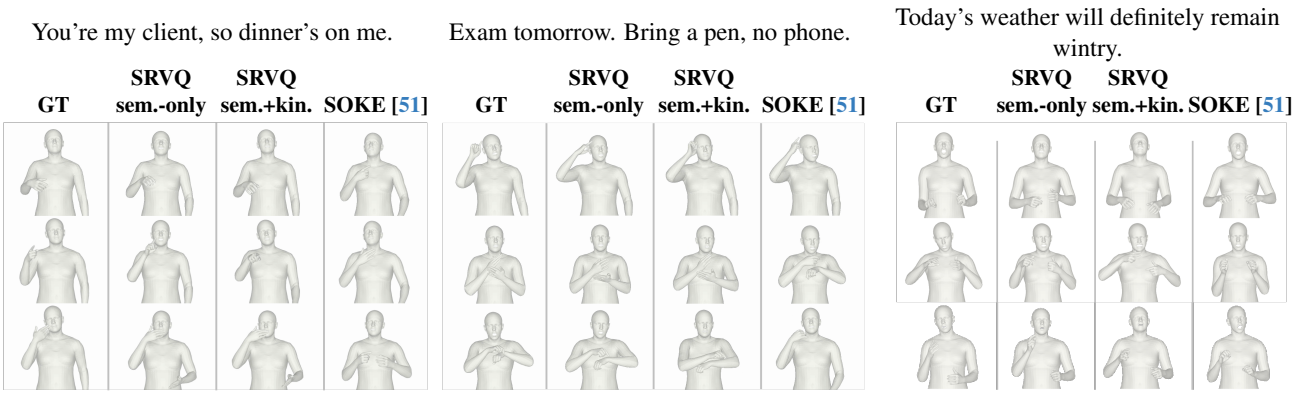

Figure 4. **Qualitative comparison of generated sign sequences.** We show representative frames from ground truth (GT), SRVQ semantic-only generation, SRVQ sem+kinematic generation, and SOKE. The text above each example is the corresponding input sentence. The first two examples are from CSL-Daily, and the last example is from Phoenix-2014T.

erated sequences alongside the ground-truth signing motion. Semantic-only generation already recovers sentence-dependent signing patterns, showing that the semantic codes provide a meaningful text-aligned target for generation. Sem+kinematic preserves this semantic structure while adding residual hand and body details, producing motions that are locally coherent over the sampled sequence. Compared with SOKE, SRVQ yields sign sequences that better follow the intended lexical content, which is consistent with the retrieval gains in Tab. 1.

4.3. Ablation Studies

Semantic Alignment. Tab. 2 examines the effect of semantic supervision on the quantizer. Results show that removing both text and gloss alignment ($\lambda_{\text{text}} = 0, \lambda_{\text{gloss}} = 0$) yields the worst performance across most metrics, highlighting the importance of injecting semantic information into the learned motion codes. Removing gloss supervision (*No-gloss*) still provides a moderate improvement over the unaligned variant, reaching sem+kinematic R@1 scores of 75.60 on CSL-Daily and 58.88 on Phoenix-2014T. The full objective further improves performance,($\lambda_{\text{text}} = \lambda_{\text{gloss}} = 0.1$) raising generation to 82.31 on CSL-Daily and 70.25 on Phoenix-2014T. Increasing to $\lambda = 0.3$ does not consistently improve generation. It lowers CSL-Daily R@1 to 80.78 while improving Phoenix-2014T R@1 to 74.92. We use $\lambda = 0.1$ by default, which provides a balanced trade-off between semantic alignment and motion reconstruction.

Lexical Organization of Semantic Codes. We further examine whether the learned semantic codes organize motions according to lexical content, comparing their latent structure with the *reconstruction-only* codes from SOKE [51]. To this end, we select ten frequent glosses from CSL-Daily and extract the token features corresponding to their annotated gloss spans. We visualize the resulting representations with t-SNE [42], color-coding samples by their

Table 2. **Semantic alignment ablation. Bold** indicates the best result and underline the second best.

Method	CSL-Daily			Phoenix-2014T		
	R@1↑	R@5↑	R@10↑	R@1↑	R@5↑	R@10↑
$\lambda_{\text{text}} = 0, \lambda_{\text{gloss}} = 0$	72.96	87.24	91.75	52.96	84.27	91.12
$\lambda_{\text{text}} = 0.3, \lambda_{\text{gloss}} = 0.3$	80.78	91.16	93.96	**74.92**	**92.83**	**95.79**
No-gloss ($\lambda_{\text{gloss}} = 0.0$)	75.60	89.80	93.62	58.88	84.27	90.50
SRVQ (sem+kinematic)	**82.31**	**94.13**	**96.34**	70.25	87.54	92.52

Table 3. **Quantitative lexical organization of semantic codes.**

Representation	NN↑	P@5↑	P@10↑	Sil.↑
Reconstruction-only [51]	56.2	52.8	49.6	0.032
SRVQ (Q_0)	**73.8**	**73.4**	**73.2**	**0.260**

associated gloss. Fig. 5 shows that our text-aligned tokenization form clearer lexical neighborhoods than the reconstruction-only VQ-VAE from SOKE [51]. This suggests that reconstruction-only tokenization does not explicitly encourage lexical organization, whereas our semantic alignment objective induces a motion code space that better reflects sign-level content. The quantitative neighborhood statistics in Tab. 3 confirm the same trend. Our representation increases the nearest-neighbor same-keyword rate (NN) from 56.2% to 73.8%, improves both precision-at-K metrics (P@K), and raises the silhouette score (Sil.) from 0.032 to 0.260. This confirms that semantic supervision shapes level-0 codes toward text-aligned generation.

Text Conditioning Design. Tab. 4 compares conditioning strategies for the full generation pipeline, namely the use of a global sentence embedding *vs.* the word-aware textual representation. Cross-attending both the global textual representation ($\mathbf{e}_{\text{sent}}$) and the word ($\mathbf{W}$) consistently improves over global sentence conditioning alone, *e.g.*, 82.31 *vs.* 75.34 R@1 in CSL-Daily and 92.52 *vs.* 79.28 R@10 for Phoenix-2014T, showing the benefit of exposing lexi-

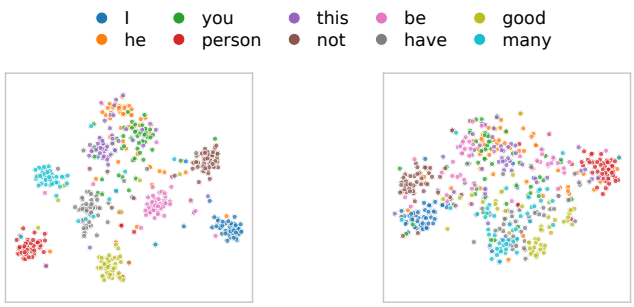

(a) SRVQ level-0 (b) Reconstruction-only [51]

Figure 5. **Lexical organization of motion codes.** t-SNE visualization of gloss-span motion features on CSL-Daily. Colors denote gloss identities. SRVQ level-0 forms clearer lexical neighborhoods than reconstruction-only tokenization.

Table 4. **Text conditioning ablation.**

$\mathbf{e}_{\text{sent}}$	$\mathbf{W}$	CSL-Daily			Phoenix-2014T		
		R@1↑	R@5↑	R@10↑	R@1↑	R@5↑	R@10↑
✓	✗	75.34	90.14	93.62	41.28	71.81	79.28
✓	✓	**82.31**	**94.13**	**96.34**	**70.25**	**87.54**	**92.52**

cal information to the generators. Notably, the gain is particularly large on Phoenix-2014T, with an approximately 29-point improvement in R@1. We hypothesize that this is due to the Phoenix-2014T's narrow weather-forecast domain, where sentences share coarse semantics but differ in key lexical cues such as cold, rain, or winter. Thus, global sentence embeddings may fail to capture the exact signs required. In contrast, CSL-Daily's broader everyday vocabulary offers greater sentence-level variation, where word-level lexical information have diminishing returns.

Residual Levels. Tab. 5 investigates how many residual levels are needed for sign-motion quantization. We report retrieval metrics as the main comparison, since geometric metrics are less sensitive to the fine-grained kinematic details that are primarily encoded from level 1 onward. SRVQ remains robust when using fewer residual levels, with the number of residuals having a smaller impact on performance than the SRVQ semantic alignment and word-aware conditioning design. Training fewer residual generators does not consistently improve over the six-level model. We refer to the Supplementary B for additional ablations.

4.4. Analysis of the DTW–Retrieval Discrepancy

We conduct a diagnostic prefix-truncation experiment on the same generated motions on full test set: for a kept prefix ratio r, we evaluate only the first r fraction of generated frames while leaving the reference motion and input text unchanged. This truncation removes the latter part of the predicted signing sequence, and should therefore reduce generation quality because some sign content no longer exists. As shown in Tab. 6, retrieval and FID follow this expectation. In contrast, DTW scores improve as the evaluated prefix becomes shorter. This indicates that DTW reflects geometric alignment under temporal warping, which may not always track sentence-level semantic preservation.

Table 5. **Number of residuals ablation.** † One-level VQ uses semantic-only decoding.

Depth	CSL-Daily			Phoenix-2014T		
	R@1↑	R@5↑	R@10↑	R@1↑	R@5↑	R@10↑
1†	75.77	89.54	93.28	61.21	83.49	88.32
2	80.53	93.11	95.24	65.89	87.54	92.37
4	**82.48**	92.18	94.47	68.85	87.54	**92.68**
6	82.31	**94.13**	**96.34**	**70.25**	87.54	92.52

Table 6. **Length sensitivity of DTW.**

Dataset / Ratio	DTW-B↓	DTW-H↓	FID↓	R@1↑
CSL-Daily				
1.0	6.787	8.535	**0.058**	**82.31**
0.9	6.763	8.532	0.059	80.53
0.8	6.751	8.530	0.062	74.15
0.7	6.613	8.408	0.080	64.37
0.6	6.561	8.365	0.105	55.02
0.5	**6.534**	**8.351**	0.140	43.45
Phoenix-2014T				
1.0	5.330	6.920	**0.059**	**70.25**
0.9	5.146	6.524	0.068	65.58
0.8	5.031	6.394	0.089	53.12
0.7	4.956	6.335	0.121	41.12
0.6	4.896	6.278	0.162	28.50
0.5	**4.852**	**6.274**	0.217	16.36

5. Conclusion

We presented SRVQ, an SLG framework built around a semantic residual quantization hierarchy in which the coarsest level is jointly trained with a semantic alignment objective. This design creates a discrete semantic intermediate representation that bridges text and motion, enabling a two-level generation process within a single unified quantizer hierarchy, without separate semantic and reconstruction architectures. Paired with word-aware conditioning for hierarchical code generation, SRVQ surpasses the state-of-the-art on CSL-Daily and Phoenix-2014T by large margins in retrieval quality while the semantic-only variant improves upon prior baselines on most geometric metrics.

Limitations. SRVQ still relies on high-quality SMPL-X motion annotations and frame-level gloss supervision during quantizer training, although no gloss annotations are required at inference. Future work will study more cross-lingual settings, explicit modeling of non-manual markers, and will incorporate a Deaf-participant human study.

Acknowledgments

We acknowledge ISCRA for awarding this project access to the LEONARDO supercomputer, owned by the EuroHPC Joint Undertaking, hosted by CINECA (Italy). GI acknowledges financial support from VEASYT. NB acknowledges the support of the EU Horizon Europe projects ELIAS (No. 101120237) and ELLIOT (No. 101214398), and of the FIS project GUIDANCE (No. FIS2023-03251). This work was carried out in the Vision and Learning joint laboratory of FBK and UniTN.

References

[1] Anonymous. Anonymous paper, provided in supplementary, 2026. Accepted; venue omitted for anonymity.

[2] Vasileios Baltatzis, Rolandos Alexandros Potamias, Evangelos Ververas, Guanxiong Sun, Jiankang Deng, and Stefanos Zafeiriou. Neural sign actors: A diffusion model for 3d sign language production from text. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.

[3] Léore Bensabath, Mathis Petrovich, and Gül Varol. Text-driven 3d hand motion generation from sign language data. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026.

[4] Yiyang Cao, Yunze Deng, Ziyu Lin, Bin Feng, Xinggang Wang, Wenyu Liu, Dandan Zheng, and Jingdong Chen. Tric-motion: Tri-domain causal modeling grounded text-to-motion generation. In *International Conference on Learning Representations*, 2026.

[5] Huiwen Chang, Han Zhang, Lu Jiang, Ce Liu, and William T Freeman. Maskgit: Masked generative image transformer. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2022.

[6] Yutong Chen, Ronglai Zuo, Fangyun Wei, Yu Wu, Shujie Liu, and Brian Mak. Two-stream network for sign language recognition and translation. In *Advances in Neural Information Processing Systems*, 2022.

[7] Zisheng Chen, Chunwei Wang, Runhui Huang, Hongbin Xu, Xiuwei Chen, Jun Zhou, Jianhua Han, Hang Xu, and Xiaodan Liang. Semhitok: A unified image tokenizer via semantic-guided hierarchical codebook for multimodal understanding and generation. *arXiv preprint arXiv:2503.06764*, 2025.

[8] Alexandre Défossez, Jade Copet, Gabriel Synnaeve, and Yossi Adi. High fidelity neural audio compression. *arXiv preprint arXiv:2210.13438*, 2022.

[9] Pengcheng Fang, Tengjiao Sun, Xiaoyu Zhan, Xiaohao Cai, and Dongjie Fu. Mogeflow: Flowing through motion codebook geometry for text-to-motion generation. *arXiv preprint arXiv:2606.11656*, 2026.

[10] Sen Fang, Chunyu Sui, Xuedong Zhang, and Yapeng Tian. Signdiff: Learning diffusion models for american sign language production. *arXiv preprint arXiv:2308.16082*, 2023.

[11] Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. Language-agnostic bert sentence embedding. In *Proceedings of the 60th annual meeting of the association for computational linguistics (volume 1: Long papers)*, 2022.

[12] Abraham Glasser, Kesavan Kushalnagar, and Raja Kushalnagar. Deaf, hard of hearing, and hearing perspectives on using automatic speech recognition in conversation. In *Proceedings of the 19th International ACM SIGACCESS Conference on Computers and Accessibility*, 2017.

[13] Mo Guan, Yan Wang, Guangkun Ma, Jiarui Liu, and Mingzu Sun. Mska: Multi-stream keypoint attention network for sign language recognition and translation. *Pattern Recognition*, 165:111602, 2025.

[14] Chuan Guo, Yuxuan Mu, Muhammad Gohar Javed, Sen Wang, and Li Cheng. Momask: Generative masked modeling of 3d human motions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.

[15] Ziyao Guo, Kaipeng Zhang, and Michael Qizhe Shieh. Improving autoregressive image generation through coarse-to-fine token prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026.

[16] KunHo Heo, SuYeon Kim, Yonghyun Gwon, Youngbin Kim, and MyeongAh Cho. Party: Part-guidance for expressive text-to-motion synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026.

[17] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.

[18] Saki Imai, Mert Inan, Anthony B Sicilia, and Malihe Alikhani. Silverscore: Semantically-aware embeddings for sign language generation evaluation. In *Proceedings of the 15th International Conference on Recent Advances in Natural Language Processing-Natural Language Processing in the Generative AI Era*, 2025.

[19] Biao Jiang, Xin Chen, Wen Liu, Jingyi Yu, Gang Yu, and Tao Chen. Motiongpt: Human motion as a foreign language. *Advances in Neural Information Processing Systems*, 2023.

[20] Zifan Jiang, Colin Leong, Amit Moryossef, Oliver Cory, Maksym Ivashechkin, Neha Tarigopula, Biao Zhang, Anne Göhring, Annette Rios Gonzales, Rico Sennrich, et al. Meaningful pose-based sign language evaluation. In *Proceedings of the Tenth Conference on Machine Translation*, 2025.

[21] Oscar Koller, Jens Forster, and Hermann Ney. Continuous sign language recognition: Towards large vocabulary statistical recognition systems handling multiple signers. *Computer vision and image understanding*, 2015.

[22] Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, and Wook-Shin Han. Autoregressive image generation using residual quantization. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2022.

[23] Chang Liu, Mengyi Zhao, Bin Ren, Mengyuan Liu, Nicu Sebe, et al. Spatio-temporal graph diffusion for text-driven human motion generation. In *Proceedings of the 34th British Machine Vision Conference (BMVC)*, 2023.

[24] Jinpeng Liu, Wenxun Dai, Chunyu Wang, Yiji Cheng, Yansong Tang, and Xin Tong. Plan, posture and go: To-

wards open-world text-to-motion generation. *arXiv preprint arXiv:2312.14828*, 2023.

[25] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.

[26] Shunlin Lu, Ling-Hao Chen, Ailing Zeng, Jing Lin, Ruimao Zhang, Lei Zhang, and Heung-Yeung Shum. Humantomato: Text-aligned whole-body motion generation. *arXiv preprint arXiv:2310.12978*, 2023.

[27] Kelly Mack, Danielle Bragg, Meredith Ringel Morris, Maarten W Bos, Isabelle Albi, and Andrés Monroy-Hernández. Social app accessibility for deaf signers. *Proceedings of the ACM on Human-Computer Interaction*, 2020.

[28] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.

[29] Mathis Petrovich, Michael J Black, and Gül Varol. Tmr: Text-to-motion retrieval using contrastive 3d human motion synthesis. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9488–9497, 2023.

[30] Yijie Qian, Juncheng Wang, Yuxiang Feng, Chao Xu, Wang Lu, Yang Liu, Baigui Sun, Yiqiang Chen, Yong Liu, and Shujun Wang. Think before you move: Latent motion reasoning for text-to-motion generation. *arXiv preprint arXiv:2512.24100*, 2025.

[31] Liao Qu, Huichao Zhang, Yiheng Liu, Xu Wang, Yi Jiang, Yiming Gao, Hu Ye, Daniel K Du, Zehuan Yuan, and Xinglong Wu. Tokenflow: Unified image tokenizer for multimodal understanding and generation. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, 2025.

[32] Ben Saunders, Necati Cihan Camgoz, and Richard Bowden. Adversarial training for multi-channel sign language production. *arXiv preprint arXiv:2008.12405*, 2020.

[33] Ben Saunders, Necati Cihan Camgoz, and Richard Bowden. Progressive transformers for end-to-end sign language production. In *European Conference on Computer Vision*. Springer, 2020.

[34] Ben Saunders, Necati Cihan Camgoz, and Richard Bowden. Mixed signals: Sign language production via a mixture of motion primitives. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021.

[35] Fei Shen and Jinhui Tang. Imagpose: A unified conditional framework for pose-guided person generation. In *Advances in Neural Information Processing Systems*, 2024.

[36] Fei Shen, Xin Jiang, Xin He, Hu Ye, Cong Wang, Xiaoyu Du, Zechao Li, and Jinhui Tang. Imagdressing-v1: Customizable virtual dressing. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025.

[37] Nikolay Shvetsov, Maksim Bobrin, Nazar Buzun, and Dmitry V Dylov. Plan, don't pose: Long composite motion generation with text-aligned bfm. *arXiv preprint arXiv:2605.29906*, 2026.

[38] Stephanie Stoll, Armin Mustafa, and Jean-Yves Guillemaut. There and back again: 3d sign language generation from text using back-translation. In *2022 International Conference on 3D Vision (3DV)*. IEEE, 2022.

[39] Alexandre Symeonidis-Herzig, Jianhe Low, Ozge Mercanoglu Sincan, and Richard Bowden. M3t: Discrete multimodal motion tokens for sign language production. *arXiv preprint arXiv:2603.23617*, 2026.

[40] Guy Tevet, Sigal Raab, Brian Gordon, Yonatan Shafir, Daniel Cohen-Or, and Amit H Bermano. Human motion diffusion model. *arXiv preprint arXiv:2209.14916*, 2022.

[41] Aaron Van Den Oord, Oriol Vinyals, et al. Neural discrete representation learning. *Advances in neural information processing systems*, 2017.

[42] Laurens Van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine learning research*, 2008.

[43] Pan Xie, Qipeng Zhang, Peng Taiying, Hao Tang, Yao Du, and Zexian Li. G2p-ddm: Generating sign pose sequence from gloss sequence with discrete diffusion model. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.

[44] Aoxiong Yin, Haoyuan Li, Kai Shen, Siliang Tang, and Yueting Zhuang. T2s-gpt: Dynamic vector quantization for autoregressive sign language production from text. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024.

[45] Zhengdi Yu, Shaoli Huang, Yongkang Cheng, and Tolga Birdal. Signavatars: A large-scale 3d sign language holistic motion dataset and benchmark. In *European Conference on Computer Vision*, 2024.

[46] Neil Zeghidour, Alejandro Luebs, Ahmed Omran, Jan Skoglund, and Marco Tagliasacchi. Soundstream: An end-to-end neural audio codec. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 2021.

[47] Jianrong Zhang, Yangsong Zhang, Xiaodong Cun, Yong Zhang, Hongwei Zhao, Hongtao Lu, Xi Shen, and Ying Shan. Generating human motion from textual descriptions with discrete representations. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2023.

[48] Chongyang Zhong, Lei Hu, Zihao Zhang, and Shihong Xia. Attt2m: Text-driven human motion generation with multi-perspective attention mechanism. In *Proceedings of the IEEE/CVF international conference on computer vision*, 2023.

[49] Hao Zhou, Wengang Zhou, Weizhen Qi, Junfu Pu, and Houqiang Li. Improving sign language translation with monolingual data by sign back-translation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2021.

[50] Ronglai Zuo, Fangyun Wei, Zenggui Chen, Brian Mak, Jiaolong Yang, and Xin Tong. A simple baseline for spoken language to sign language translation with 3d avatars. In *European Conference on Computer Vision*. Springer, 2024.

[51] Ronglai Zuo, Rolandos Alexandros Potamias, Evangelos Ververas, Jiankang Deng, and Stefanos Zafeiriou. Signs as tokens: A retrieval-enhanced multilingual sign language generator. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2025.

Supplementary Material

This supplement provides additional analyses supporting the main paper. Specifically, we first report additional validation results based on back-translation and meaningful pose (Sec. A). We then present ablations (Sec. B) , including text conditioning, tokenizer reconstruction and residual analysis, gloss-annotation robustness, codebook utilization, and seed sensitivity. Finally, we conclude with a study on the effect of motion length (Sec. D).

A. Additional Evaluators

This section complements the main paper retrieval validation with additional evaluators relying on back-translation (Sec. A.1) and pose retrieval and reconstruction (Sec. A.2).

A.1. Back-translation with SLT Evaluators Adapted to SMPL-X

To obtain an independent semantic check, we adapt two keypoint-based sign-language translation systems to accept SMPL-X motion as input. The change targets the input interface: the original keypoint input is replaced by the corresponding SMPL-X representation. Table 7 reports metrics under the Two-Stream SLT system [6], while Table 8 relies on MSKA [13]. Both evaluators rank SRVQ above SOKE on every reported BLEU and ROUGE metric, providing a cross-evaluator semantic check in addition to retrieval.

Table 7. **Back-translation with the Two-Stream SLT evaluator.** The Two-Stream Network [6] is adapted from keypoint input to SMPL-X input.

Dataset / Method	BLEU-1	BLEU-2	BLEU-4	ROUGE
Phoenix-2014T				
SOKE	24.75	13.65	7.05	25.14
SRVQ (sem+kinematic)	**35.92**	**23.56**	**13.47**	**36.05**
CSL-Daily				
SOKE	21.60	11.62	4.33	22.58
SRVQ (sem+kinematic)	**25.81**	**14.99**	**5.97**	**26.93**

Table 8. **Back-translation with the MSKA SLT evaluator.** MSKA [13] is adapted from keypoint input to SMPL-X input.

Dataset / Method	BLEU-1	BLEU-2	BLEU-4	ROUGE
Phoenix-2014T				
SOKE	24.44	14.28	7.90	25.80
SRVQ (sem+kinematic)	**39.36**	**27.40**	**16.39**	**40.96**
CSL-Daily				
SOKE	25.42	14.78	5.50	25.99
SRVQ (sem+kinematic)	**29.49**	**18.02**	**7.42**	**30.38**

A.2. Meaningful Pose Evaluation

We additionally evaluate our method using the Meaningful Pose toolkit [20], which provides complementary pose-retrieval and reconstruction metrics over all joints, body, hands and face. Since the original toolkit does not directly support SMPL-X motion, we adapt it by converting SMPL-X sequences to OpenPose keypoints, resampling them to a fixed length, and performing geometry-only pose retrieval using L2 distance. Unlike the main text–motion retrieval, this is motion-to-motion retrieval: each generated pose sequence is directly compared with a gallery of ground-truth pose sequences using only resampled keypoint geometry. The low absolute R@1 values are therefore expected because retrieving the exact paired motion requires matching signer-, timing-, and articulation-specific geometry; they should not be numerically compared with the text–motion R@1, but only compared across methods under this same protocol. Table 9 reports results for the sem+kinematic and sem-only variants against SOKE, confirming the trends observed with the text–motion retrieval and geometric metrics in the main paper. Specifically, sem+kinematic achieves the strongest pose-retrieval performance on both datasets for all-joint and hand queries, while sem-only generally yields the best reconstruction. Sem+kinematic also improves over SOKE on Phoenix-2014T and achieves comparable or slightly lower reconstruction performance on CSL-Daily. No learned text–motion, back-translation model is involved in this evaluation.

B. Additional Ablations

B.1. Additional Text-Conditioning Analysis

Tab. 10 extends the main text-conditioning comparisons by evaluating the final word-aware model against a word-token-only variant under the same generation-validation protocol. The word-aware model outperforms the word-token-only variant across all retrieval metrics (+1.44 R@1, +3.06 R@5, +1.19 R@10) on CSL-Daily, consistent with the main-paper conditioning ablation (Tab. 4), where removing word tokens caused a larger drop (-6.97 R@1) than removing the global sentence embedding here (-1.44), confirming that lexical word-level information is the stronger conditioning signal.

Table 10. **Word-token-only comparison on CSL-Daily.** Both rows use the final generation-validation protocol. DTW-B: DTW-JPE body; DTW-H: DTW-JPE averaged over left and right hands.

Conditioning	R@1↑	R@5↑	R@10↑	FID↓	DTW-B↓	DTW-H↓
Word-aware (ours)	82.31	94.13	96.34	0.058	6.79	8.54
Word tokens only	80.87	91.07	95.15	0.057	6.83	8.63

Table 9. **Meaningful Pose evaluation** on CSL-Daily and Phoenix-2014T. SMPL-X motion sequences are converted to OpenPose keypoints. The official DTW and Aggregated Power Distance metrics reported by the evaluation toolkit correspond to the "All" DTW and MJE columns, respectively. Best results on each dataset are shown in **bold**.

(a) Pose Retrieval

	All joints					Hands				
Dataset / Method	R@1↑	R@5↑	R@10↑	MRR↑	MedR↓	R@1↑	R@5↑	R@10↑	MRR↑	MedR↓
CSL-Daily										
SOKE [51]	1.956	5.697	9.014	0.0472	216.0	2.126	5.527	8.844	0.0474	216.0
SRVQ (sem+kinematic)	**2.806**	**7.823**	**11.650**	**0.0607**	**192.0**	**2.976**	**7.908**	**11.310**	**0.0603**	**195.5**
SRVQ (sem-only)	2.551	6.718	10.714	0.0548	202.0	2.381	6.378	9.779	0.0533	206.5
Phoenix-2014T										
SOKE [51]	0.467	3.583	6.075	0.0287	174.0	0.156	3.427	5.919	0.0250	164.5
SRVQ (sem+kinematic)	**3.115**	**9.502**	**14.019**	**0.0719**	**111.0**	**2.025**	**8.723**	**14.486**	**0.0634**	**104.5**
SRVQ (sem-only)	1.558	6.386	10.125	0.0526	118.5	1.246	6.698	9.502	0.0489	125.5

(b) Pose Reconstruction Error

Dataset / Metric	**Method**	All	Body	Hands	L-Hand	R-Hand	Face
CSL-Daily							
DTW↓	SOKE [51]	0.7420	0.2428	1.8177	2.0795	1.5560	0.2611
	SRVQ (sem+kinematic)	0.7435	0.2422	1.8120	2.0755	1.5485	0.2678
	SRVQ (sem-only)	**0.7014**	**0.2260**	**1.7340**	**1.9755**	**1.4925**	**0.2384**
MJE↓	SOKE [51]	0.0836	0.0274	0.2065	**0.2469**	0.1660	0.0284
	SRVQ (sem+kinematic)	0.0864	0.0284	0.2124	0.2552	0.1697	0.0298
	SRVQ (sem-only)	**0.0824**	**0.0268**	**0.2056**	0.2479	**0.1634**	**0.0268**
MAE↓	SOKE [51]	0.1246	0.0410	0.3062	**0.3643**	0.2481	0.0431
	SRVQ (sem+kinematic)	0.1287	0.0424	0.3154	0.3766	0.2542	0.0452
	SRVQ (sem-only)	**0.1226**	**0.0400**	**0.3046**	0.3644	**0.2448**	**0.0407**
Phoenix-2014T							
DTW↓	SOKE [51]	0.7089	0.2410	1.6583	1.8333	1.4832	0.2944
	SRVQ (sem+kinematic)	0.6446	0.2172	1.4936	1.6659	1.3212	0.2773
	SRVQ (sem-only)	**0.5978**	**0.2004**	**1.3924**	**1.4844**	**1.3003**	**0.2532**
MJE↓	SOKE [51]	0.0838	0.0284	0.1978	0.2219	0.1738	0.0338
	SRVQ (sem+kinematic)	0.0778	0.0262	0.1814	0.2020	0.1607	0.0329
	SRVQ (sem-only)	**0.0700**	**0.0236**	**0.1638**	**0.1750**	**0.1526**	**0.0291**
MAE↓	SOKE [51]	0.1217	0.0421	0.2844	0.3143	0.2545	0.0505
	SRVQ (sem+kinematic)	0.1143	0.0390	0.2649	0.2914	0.2384	0.0490
	SRVQ (sem-only)	**0.1026**	**0.0349**	**0.2386**	**0.2522**	**0.2249**	**0.0435**

B.2. Robustness to Gloss Perturbations

We evaluate robustness to temporal gloss-span noise and gloss dropout under the training protocol. Each row in Table 11 corresponds to a separately trained variant with the same architecture and optimization settings; only the gloss supervision is perturbed during training. All variants are evaluated on the unmodified test sets, and no gloss annotations are provided at inference. The final model obtains 82.31 R@1 on CSL-Daily and 70.25 R@1 on Phoenix-2014T. A ± 4-frame span shift causes a modest decrease on CSL-Daily (82.31→80.61), whereas 20% gloss dropout has a larger effect (82.31→77.47). Phoenix-2014T R@1 remains nearly unchanged (70.25, 70.40, and 69.78), indicating that the learned representation is relatively tolerant to moderate gloss-boundary noise. The isolated R@5/R@10 increases under perturbation are not consistent across datasets and do not change this R@1 trend.

B.3. Tokenizer Analysis

Tab. 12 reports reconstruction analysis for the SRV Quantizer. Full reconstruction encodes the ground-truth motion with all SRVQ levels and decodes it back to motion, while Q_0 reconstruction decodes only the semantic code level.

Table 11. Robustness to gloss perturbations. We report sem+kinematic generation for the final model and perturbation settings.

Setting	CSL-Daily				Phoenix-2014T			
	R@1↑	R@5↑	R@10↑	FID↓	R@1↑	R@5↑	R@10↑	FID↓
SRVQ (sem+kinematic)	**82.31**	**94.13**	**96.34**	0.0580	70.25	87.54	92.52	0.0590
Gloss span shift ±4 frames	80.61	92.09	94.81	**0.0488**	**70.40**	90.19	**94.86**	**0.0562**
Gloss dropout 20%	77.47	91.24	94.22	0.0495	69.78	**91.12**	94.70	0.0567

Table 13. Controlled residual contribution analysis. $\hat{Q}_\ell$ denotes predicted codes at level ℓ and Q_ℓ^{GT} denotes codes from the ground-truth motion. DTW-B, DTW-H, and PA-H are lower-is-better; hand metrics average the left and right hands.

Dataset / Code combination	R@1	R@5	DTW-B	DTW-H	PA-H
CSL-Daily					
$\hat{Q}_0$+zero residual	77.81	89.97	6.296	8.376	1.508
$\hat{Q}_0 + \hat{Q}_{1:L}$	82.31	94.13	6.787	8.535	1.544
$\hat{Q}_0 + Q^{\mathrm{GT}}_{1:L}$	77.98	88.95	5.140	7.519	1.513
$Q^{\mathrm{GT}}_0 + \hat{Q}_{1:L}$	74.15	86.90	4.945	6.430	0.845
$Q^{\mathrm{GT}}_0 + Q^{\mathrm{GT}}_{1:L}$	85.97	93.45	1.779	2.727	0.390
Phoenix-2014T					
$\hat{Q}_0$+zero residual	57.48	78.97	4.700	6.410	1.262
$\hat{Q}_0 + \hat{Q}_{1:L}$	70.25	87.54	5.330	6.920	1.354
$\hat{Q}_0 + Q^{\mathrm{GT}}_{1:L}$	60.90	82.71	3.047	4.651	1.168
$Q^{\mathrm{GT}}_0 + \hat{Q}_{1:L}$	53.43	78.50	4.411	6.193	0.882
$Q^{\mathrm{GT}}_0 + Q^{\mathrm{GT}}_{1:L}$	64.49	85.05	1.564	2.889	0.406

Table 12. **Tokenizer reconstruction analysis.** DTW-B: DTW-JPE body; DTW-H: DTW-JPE averaged over left and right hands.

Dataset / Method	R@1↑	FID↓	DTW-B↓	DTW-H↓
CSL-Daily				
SRVQ full reconstruction	85.97	0.002	1.78	2.73
SRVQ Q_0 reconstruction	74.49	0.018	3.33	4.81
Phoenix-2014T				
SRVQ full reconstruction	64.49	0.008	1.56	2.89
SRVQ Q_0 reconstruction	40.03	0.067	3.15	5.14

B.4. Residual Code Contribution

The controlled code-combination analysis in Table 13 isolates the contribution of the semantic code and the residual levels without retraining. Each row reports a decoding condition applied to the same test samples and checkpoint. The standard sem+kinematic variant is $\hat{Q}_0 + \hat{Q}_{1:L}$; the remaining conditions replace predicted semantic or residual codes with codes extracted from the ground-truth motion and serve as diagnostic oracles. Relative to semantic-only decoding ($\hat{Q}_0$+zero residual), predicted residuals improve R@1 by 4.50 and 12.77 points on CSL-Daily and Phoenix-2014T, respectively, while increasing DTW-B by 0.491 and 0.630. Replacing predicted residuals with ground-truth residuals substantially reduces DTW-B (6.787→5.140 and 5.330→3.047) but lowers retrieval. Thus, residuals improve text-aligned generation, while their predicted realization need not match one signer-specific reference trajectory in timing, coarticulation, or style. Notably, adding ground-truth residuals $Q^{\mathrm{GT}}_{1:L}$ improves geometric metrics showcasing how reference-specific kinematics allow better reconstruction of the target.

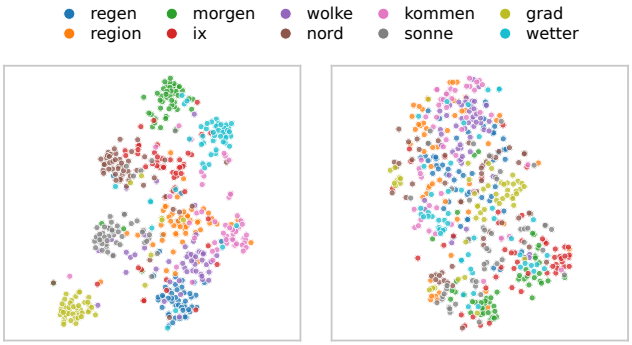

(a) SRVQ level-0 (b) Reconstruction-only [51]

Figure 6. **Lexical organization of motion codes.** t-SNE visualization of gloss-span motion features on Phoenix-2014T. Colors denote gloss identities. SRVQ level-0 forms clearer lexical neighborhoods than reconstruction-only tokenization.

B.5. Lexical Organization on Phoenix-2014T

The main paper visualizes lexical organization on CSL-Daily. We additionally show the similar analysis on Phoenix-2014T in Fig. 6. The selected glosses are frequent Phoenix-2014T weather-domain signs. Despite the limited semantic diversity of the dataset, SRVQ Q_0 still forms clear lexical neighborhoods, whereas the reconstruction-only tokenizer exhibits little separation between gloss groups. The quantitative neighborhood statistics in Tab. 16 confirm the same trend: SRVQ Q_0 improves the nearest-neighbor same-keyword rate from 52.4% to 69.4% and the silhouette score from -0.083 to 0.199.

Table 14. Per-level codebook utilization of the final SRVQ tokenizer. Active codes and normalized perplexity are higher-is-better, while Top-1 usage is lower-is-better.

Metric	Dataset	Part	Q0	Q1	Q2	Q3	Q4	Q5
Active codes (%) ↑	CSL-Daily	Body	79.17	96.88	95.83	97.92	97.92	100.00
		Left hand	99.61	100.00	100.00	100.00	100.00	100.00
		Right hand	100.00	100.00	99.61	99.61	97.66	84.38
	Phoenix-2014T	Body	22.92	91.67	94.79	94.79	100.00	97.92
		Left hand	98.83	99.22	98.83	99.61	100.00	99.61
		Right hand	99.61	99.61	100.00	100.00	98.05	91.41
Norm. perplexity (%) ↑	CSL-Daily	Body	72.01	89.30	92.01	93.86	94.07	92.65
		Left hand	68.97	68.59	68.96	67.88	69.88	67.60
		Right hand	79.48	86.66	92.03	95.16	92.63	72.16
	Phoenix-2014T	Body	22.47	87.38	87.18	86.37	84.79	84.35
		Left hand	75.63	88.37	89.48	89.49	90.57	88.39
		Right hand	76.00	92.61	95.67	96.67	93.21	73.83
Top-1 usage (%) ↓	CSL-Daily	Body	2.32	2.15	1.59	1.70	1.51	2.03
		Left hand	2.39	2.78	2.85	3.33	3.46	4.61
		Right hand	1.33	1.40	0.77	0.73	0.60	0.82
	Phoenix-2014T	Body	6.67	2.24	2.60	4.07	3.91	3.94
		Left hand	1.73	1.20	1.32	1.50	2.05	1.91
		Right hand	1.56	0.84	0.76	0.68	0.76	1.17

C. Qualitative Sample Metadata

For completeness, Table 15 lists the metadata of the qualitative examples used in the paper and supplementary visualizations, including the English display text, dataset, sample ID, and original annotated sentence.

Table 17 separates cumulative representations from individual residual levels. Cumulative codes preserve strong same-gloss lexical neighborhoods as residual levels are appended: on Phoenix-2014T, NN increases from 0.6940 to 0.7380 and P@10 from 0.6520 to 0.6964, while on CSL-Daily P@5 reaches 0.7592 for Q_0–Q_4. In contrast, each individual residual level Q_1–Q_5 has much lower NN/P@K scores and a negative silhouette score on both datasets. These results indicate that Q_0 is the primary lexical anchor: residual levels preserve the semantic organization when combined with Q_0, but do not independently encode lexical identity.

C.1. Codebook Utilization

On each test set, we report three utilization statistics for each part and RVQ level. Active is the fraction of codebook entries used by at least one valid token, Norm. perplexity is the code usage perplexity normalized by the codebook size, and Top-1 is the token fraction assigned to the most frequent code. As shown in Tab. 14, the hand codebooks are almost fully activated across all levels on both datasets, with high normalized perplexity and low Top-1 usage. This indicates that SRVQ does not collapse to a few dominant codes. The only low-activation case is body Q0 on Phoenix-2014T, which is expected for a constrained broadcast domain with limited body pose. All higher body residual levels remain highly active, showing that the residual hierarchy is still used for kinematic detail.

C.2. Seed Stability

Tab. 18 evaluates the final sem+kinematic generator on CSL-Daily across three inference seeds. Retrieval and geometric metrics show only minor variation, indicating that the reported performance is stable and not dependent on a particular random seed.

Table 18. **Seed stability on CSL-Daily.** All rows use the same sem+kinematic checkpoint. DTW-B: DTW-JPE body; DTW-H: DTW-JPE averaged over left and right hands.

Seed	R@1↑	R@5↑	R@10↑	FID↓	DTW-B↓	DTW-H↓
0	82.23	93.28	96.00	0.050	6.62	8.40
1	82.91	93.96	95.92	0.050	6.60	8.41
2	83.08	94.05	96.09	0.050	6.58	8.41
Mean ± std	82.74±0.45	93.76±0.42	96.00±0.09	0.050±0.000	6.60±0.02	8.41±0.01

D. Effect of Motion Length

Tab. 19 reports performance across different ground-truth motion lengths. For retrieval, only the queries are grouped by length, while the gallery remains the full test set, ensuring that retrieval difficulty is unchanged across groups.

Table 15. **Qualitative sample metadata.** The English display text, dataset, sample ID, and original annotated sentence for the qualitative examples.

English display text	Dataset	Sample ID	Original annotated sentence
You're my client, so dinner's on me.	CSL-Daily	S001423_P0004_T00	你是我的客户，我应该请你吃饭。
Exam tomorrow. Bring a pen, no phone.	CSL-Daily	S001372_P0008_T00	明天考试，带上笔，不要带手机。
Today's weather will definitely remain wintry.	PHOENIX-2014T	26January_2010_Tuesday_heute-105	in den nächsten tagen geht es auf jeden fall winterlich weiter immer wieder mal mit etwas schnee .

Table 16. **Quantitative evaluation of lexical organization on Phoenix-2014T.**

Representation	NN↑	P@5↑	P@10↑	Sil.↑
Reconstruction-only [51]	52.4	43.7	37.4	-0.083
SRVQ Q_0	**69.4**	**66.9**	**65.2**	**0.199**

Table 17. **Lexical organization of cumulative and individual SRVQ codes.** NN, P@5, P@10, and Sil. measure same-gloss neighborhood quality, with higher values being better. "Cumulative" concatenates the codes from Q_0 through the indicated residual level, while "Individual" evaluates each residual level independently.

Dataset / Type	Codes	NN↑	P@5↑	P@10↑	Sil.↑
Phoenix-2014T					
Cumulative	Q_0	0.6940	0.6692	0.6520	0.1995
	Q_0–Q_1	0.7120	0.6888	0.6748	0.1942
	Q_0–Q_2	0.7160	0.7008	0.6896	0.1984
	Q_0–Q_3	0.7220	0.7064	0.6880	0.1988
	Q_0–Q_4	0.7280	**0.7132**	0.6900	0.1999
	Q_0–Q_5	**0.7380**	0.7116	**0.6964**	**0.2007**
Individual	Q_1	0.2160	0.1848	0.1754	-0.0265
	Q_2	0.1860	0.1692	0.1684	-0.0216
	Q_3	0.1180	0.1124	0.1054	-0.0335
	Q_4	0.1620	0.1428	0.1304	-0.0386
	Q_5	0.1620	0.1608	0.1556	-0.0285
CSL-Daily					
Cumulative	Q_0	0.7380	0.7344	0.7316	**0.2596**
	Q_0–Q_1	0.7360	0.7420	0.7354	0.2365
	Q_0–Q_2	0.7500	0.7432	0.7318	0.2233
	Q_0–Q_3	0.7580	0.7576	0.7394	0.2148
	Q_0–Q_4	**0.7660**	**0.7592**	**0.7430**	0.2114
	Q_0–Q_5	0.7580	0.7560	0.7422	0.2099
Individual	Q_1	0.1300	0.1068	0.1150	-0.0351
	Q_2	0.1320	0.1248	0.1196	-0.0341
	Q_3	0.1260	0.1076	0.1056	-0.0352
	Q_4	0.1020	0.1256	0.1216	-0.0350
	Q_5	0.1220	0.1124	0.1188	-0.0349

DTW increases steadily with motion length, whereas retrieval performance does not exhibit a monotonic trend. On Phoenix-2014T, the shortest sequences achieve the lowest DTW scores but the lowest R@1, suggesting that short weather-domain sentences are geometrically easier to align yet semantically more difficult to distinguish.

Table 19. **Performance by ground-truth motion length.** Retrieval is evaluated against the full test gallery. DTW-B: DTW-JPE body; DTW-H: DTW-JPE averaged over left and right hands.

Dataset	Length	N	R@1↑	R@5↑	R@10↑	DTW-B↓	DTW-H↓
CSL-Daily	≤100	422	82.94	94.31	96.45	4.61	5.78
	101–140	367	80.11	94.28	96.46	6.62	8.32
	>140	387	83.72	93.80	96.12	9.31	11.75
Phoenix-2014T	≤76	219	63.01	82.65	89.50	3.14	4.20
	77–112	216	75.00	89.35	95.37	5.14	6.70
	>112	207	72.95	90.82	92.75	7.85	10.03